

Python code for Artificial Intelligence: Foundations of Computational Agents

David L. Poole and Alan K. Mackworth

Version 0.9.0 of June 22, 2021.

<http://aipython.org> <http://artint.info>

©David L Poole and Alan K Mackworth 2017-2021.

All code is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. See: http://creativecommons.org/licenses/by-nc-sa/4.0/deed.en_US

This document and all the code can be downloaded from
<http://artint.info/AIPython/> or from <http://aipython.org>

The authors and publisher of this book have used their best efforts in preparing this book. These efforts include the development, research and testing of the theories and programs to determine their effectiveness. The authors and publisher make no warranty of any kind, expressed or implied, with regard to these programs or the documentation contained in this book. The author and publisher shall not be liable in any event for incidental or consequential damages in connection with, or arising out of, the furnishing, performance, or use of these programs.

Contents

Contents	3
1 Python for Artificial Intelligence	9
1.1 Why Python?	9
1.2 Getting Python	9
1.3 Running Python	10
1.4 Pitfalls	11
1.5 Features of Python	11
1.5.1 Lists, Tuples, Sets, Dictionaries and Comprehensions . .	11
1.5.2 Functions as first-class objects	12
1.5.3 Generators and Coroutines	14
1.6 Useful Libraries	15
1.6.1 Timing Code	15
1.6.2 Plotting: Matplotlib	16
1.7 Utilities	16
1.7.1 Display	16
1.7.2 Argmax	18
1.7.3 Probability	19
1.7.4 Dictionary Union	19
1.8 Testing Code	19
2 Agents and Control	21
2.1 Representing Agents and Environments	21
2.2 Paper buying agent and environment	22
2.2.1 The Environment	22
2.2.2 The Agent	24

2.2.3	Plotting	24
2.3	Hierarchical Controller	25
2.3.1	Environment	25
2.3.2	Body	26
2.3.3	Middle Layer	28
2.3.4	Top Layer	29
2.3.5	Plotting	30
3	Searching for Solutions	33
3.1	Representing Search Problems	33
3.1.1	Explicit Representation of Search Graph	34
3.1.2	Paths	36
3.1.3	Example Search Problems	37
3.2	Generic Searcher and Variants	41
3.2.1	Searcher	41
3.2.2	Frontier as a Priority Queue	42
3.2.3	A^* Search	44
3.2.4	Multiple Path Pruning	45
3.3	Branch-and-bound Search	46
4	Reasoning with Constraints	51
4.1	Constraint Satisfaction Problems	51
4.1.1	Constraints	51
4.1.2	CSPs	52
4.1.3	Examples	53
4.2	Solving a CSP using Search	59
4.3	Consistency Algorithms	61
4.3.1	Direct Implementation of Domain Splitting	63
4.3.2	Domain Splitting as an interface to graph searching	65
4.4	Solving CSPs using Stochastic Local Search	67
4.4.1	Any-conflict	69
4.4.2	Two-Stage Choice	70
4.4.3	Updatable Priority Queues	72
4.4.4	Plotting Runtime Distributions	74
4.4.5	Testing	74
5	Propositions and Inference	77
5.1	Representing Knowledge Bases	77
5.2	Bottom-up Proofs	79
5.3	Top-down Proofs	81
5.4	Assumables	82
6	Planning with Certainty	85
6.1	Representing Actions and Planning Problems	85
6.1.1	Robot Delivery Domain	86

6.1.2	Blocks World	88
6.2	Forward Planning	90
6.2.1	Defining Heuristics for a Planner	92
6.3	Regression Planning	94
6.3.1	Defining Heuristics for a Regression Planner	97
6.4	Planning as a CSP	98
6.5	Partial-Order Planning	101
7	Supervised Machine Learning	107
7.1	Representations of Data and Predictions	108
7.1.1	Evaluating Predictions	109
7.1.2	Creating Test and Training Sets	111
7.1.3	Importing Data From File	111
7.1.4	Creating Binary Features	113
7.1.5	Augmented Features	116
7.2	Generic Learner Interface	118
7.3	Learning With No Input Features	118
7.3.1	Evaluation	120
7.4	Decision Tree Learning	121
7.5	Cross Validation and Parameter Tuning	125
7.6	Linear Regression and Classification	127
7.6.1	Batched Stochastic Gradient Descent	133
7.7	Deep Neural Network Learning	133
7.8	Boosting	138
8	Reasoning Under Uncertainty	141
8.1	Representing Probabilistic Models	141
8.2	Representing Factors	142
8.3	Conditional Probability Distributions	143
8.3.1	Logistic Regression	144
8.3.2	Noisy-or	144
8.3.3	Tabular Factors	145
8.4	Graphical Models	146
8.4.1	Example Belief Networks	148
8.5	Inference Methods	152
8.6	Recursive Conditioning	153
8.7	Variable Elimination	156
8.8	Stochastic Simulation	160
8.8.1	Sampling from a discrete distribution	160
8.8.2	Sampling Methods for Belief Network Inference	162
8.8.3	Rejection Sampling	162
8.8.4	Likelihood Weighting	163
8.8.5	Particle Filtering	164
8.8.6	Examples	165
8.8.7	Gibbs Sampling	166

8.8.8	Plotting Behaviour of Stochastic Simulators	168
8.9	Hidden Markov Models	170
8.9.1	Exact Filtering for HMMs	171
8.9.2	Particle Filtering for HMMs	173
8.9.3	Generating Examples	175
8.10	Dynamic Belief Networks	175
8.10.1	Representing Dynamic Belief Networks	176
8.10.2	Unrolling DBNs	179
8.10.3	DBN Filtering	180
9	Planning with Uncertainty	183
9.1	Decision Networks	183
9.1.1	Example Decision Networks	185
9.1.2	Recursive Conditioning for decision networks	190
9.1.3	Variable elimination for decision networks	193
9.2	Markov Decision Processes	195
9.2.1	Value Iteration	198
9.2.2	Showing Grid MDPs	199
9.2.3	Asynchronous Value Iteration	201
10	Learning with Uncertainty	205
10.1	K-means	205
10.2	EM	209
11	Multiagent Systems	215
11.1	Minimax	215
11.1.1	Creating a two-player game	215
11.1.2	Minimax and α - β Pruning	219
12	Reinforcement Learning	223
12.1	Representing Agents and Environments	223
12.1.1	Simulating an environment from an MDP	224
12.1.2	Simple Game	225
12.1.3	Evaluation and Plotting	227
12.2	Q Learning	229
12.2.1	Testing Q-learning	231
12.3	Model-based Reinforcement Learner	232
12.4	Reinforcement Learning with Features	234
12.4.1	Representing Features	235
12.4.2	Feature-based RL learner	237
12.5	Multiagent Learning	240
13	Relational Learning	247
13.1	Collaborative Filtering	247
13.1.1	Alternative Formulation	249

<i>Contents</i>	7
13.1.2 Plotting	249
13.1.3 Creating Rating Sets	251
14 Version History	255
Bibliography	257
Index	259

Python for Artificial Intelligence

1.1 Why Python?

We use Python because Python programs can be close to pseudo-code. It is designed for humans to read.

Python is reasonably efficient. Efficiency is usually not a problem for small examples. If your Python code is not efficient enough, a general procedure to improve it is to find out what is taking most the time, and implement just that part more efficiently in some lower-level language. Most of these lower-level languages interoperate with Python nicely. This will result in much less programming and more efficient code (because you will have more time to optimize) than writing everything in a low-level language. You will not have to do that for the code here if you are using it for course projects.

1.2 Getting Python

You need Python 3 (<http://python.org/>) and matplotlib (<http://matplotlib.org/>) that runs with Python 3. This code is *not* compatible with Python 2 (e.g., with Python 2.7).

Download and install the latest Python 3 release from <http://python.org/>. This should also install *pip3*. You can install matplotlib using

```
pip3 install matplotlib
```

in a terminal shell (not in Python). That should “just work”. If not, try using *pip* instead of *pip3*.

The command `python` or `python3` should then start the interactive python shell. You can quit Python with a control-D or with `quit()`.

To upgrade matplotlib to the latest version (which you should do if you install a new version of Python) do:

```
pip3 install --upgrade matplotlib
```

We recommend using the enhanced interactive python **ipython** (<http://ipython.org/>). To install ipython after you have installed python do:

```
pip3 install ipython
```

1.3 Running Python

We assume that everything is done with an interactive Python shell. You can either do this with an IDE, such as IDLE that comes with standard Python distributions, or just running ipython3 (or perhaps just ipython) from a shell.

Here we describe the most simple version that uses no IDE. If you download the zip file, and cd to the “aipython” folder where the .py files are, you should be able to do the following, with user input following : . The first ipython3 command is in the operating system shell (note that the -i is important to enter interactive mode), with user input in bold:

```
$ ipython -i searchGeneric.py
```

```
Python 3.6.5 (v3.6.5:f59c0932b4, Mar 28 2018, 05:52:31)
```

```
Type 'copyright', 'credits' or 'license' for more information
```

```
IPython 6.2.1 -- An enhanced Interactive Python. Type '?' for help.
```

```
Testing problem 1:
```

```
7 paths have been expanded and 4 paths remain in the frontier
```

```
Path found: a --> b --> c --> d --> g
```

```
Passed unit test
```

```
In [1]: searcher2 = AStarSearcher(searchProblem.acyclic_delivery_problem) #A*
```

```
In [2]: searcher2.search() # find first path
```

```
16 paths have been expanded and 5 paths remain in the frontier
```

```
Out[2]: o103 --> o109 --> o119 --> o123 --> r123
```

```
In [3]: searcher2.search() # find next path
```

```
21 paths have been expanded and 6 paths remain in the frontier
```

```
Out[3]: o103 --> b3 --> b4 --> o109 --> o119 --> o123 --> r123
```

```
In [4]: searcher2.search() # find next path
```

```
28 paths have been expanded and 5 paths remain in the frontier
```

```
Out[4]: o103 --> b3 --> b1 --> b2 --> b4 --> o109 --> o119 --> o123 --> r123
```

```
In [5]: searcher2.search() # find next path
```

```
No (more) solutions. Total of 33 paths expanded.
```

In [6]:

You can then interact at the last prompt.

There are many textbooks for Python. The best source of information about python is <https://www.python.org/>. We will be using Python 3; please download the latest release. The documentation is at <https://docs.python.org/3/>.

The rest of this chapter is about what is special about the code for AI tools. We will only use the Standard Python Library and matplotlib. All of the exercises can be done (and should be done) without using other libraries; the aim is for you to spend your time thinking about how to solve the problem rather than searching for pre-existing solutions.

1.4 Pitfalls

It is important to know when side effects occur. Often AI programs consider what would happen or what may have happened. In many such cases, we don't want side effects. When an agent acts in the world, side effects are appropriate.

In Python, you need to be careful to understand side effects. For example, the inexpensive function to add an element to a list, namely *append*, changes the list. In a functional language like Haskell or Lisp, adding a new element to a list, without changing the original list, is a cheap operation. For example if x is a list containing n elements, adding an extra element to the list in Python (using *append*) is fast, but it has the side effect of changing the list x . To construct a new list that contains the elements of x plus a new element, without changing the value of x , entails copying the list, or using a different representation for lists. In the searching code, we will use a different representation for lists for this reason.

1.5 Features of Python

1.5.1 Lists, Tuples, Sets, Dictionaries and Comprehensions

We make extensive uses of lists, tuples, sets and dictionaries (dicts). See <https://docs.python.org/3/library/stdtypes.html>

One of the nice features of Python is the use of list comprehensions (and also tuple, set and dictionary comprehensions).

$(fe \text{ for } e \text{ in } iter \text{ if } cond)$

enumerates the values fe for each e in $iter$ for which $cond$ is true. The “if $cond$ ” part is optional, but the “for” and “in” are not optional. Here e has to be a variable, $iter$ is an iterator, which can generate a stream of data, such as a list, a set, a range object (to enumerate integers between ranges) or a file. $cond$

is an expression that evaluates to either True or False for each e , and fe is an expression that will be evaluated for each value of e for which $cond$ returns True.

The result can go in a list or used in another iteration, or can be called directly using *next*. The procedure *next* takes an iterator returns the next element (advancing the iterator) and raises a StopIteration exception if there is no next element. The following shows a simple example, where user input is prepended with >>>

```
>>> [e*e for e in range(20) if e%2==0]
[0, 4, 16, 36, 64, 100, 144, 196, 256, 324]
>>> a = (e*e for e in range(20) if e%2==0)
>>> next(a)
0
>>> next(a)
4
>>> next(a)
16
>>> list(a)
[36, 64, 100, 144, 196, 256, 324]
>>> next(a)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
StopIteration
```

Notice how *list(a)* continued on the enumeration, and got to the end of it.

Comprehensions can also be used for dictionaries. The following code creates an index for list *a*:

```
>>> a = ["a","f","bar","b","a","aaaaa"]
>>> ind = {a[i]:i for i in range(len(a))}
>>> ind
{'a': 4, 'f': 1, 'bar': 2, 'b': 3, 'aaaaa': 5}
>>> ind['b']
3
```

which means that 'b' is the 3rd element of the list.

The assignment of *ind* could have also be written as:

```
>>> ind = {val:i for (i,val) in enumerate(a)}
```

where *enumerate* returns an iterator of (*index,value*) pairs.

1.5.2 Functions as first-class objects

Python can create lists and other data structures that contain functions. There is an issue that tricks many newcomers to Python. For a local variable in a function, the function uses the last value of the variable when the function is

called, not the value of the variable when the function was defined (this is called “late binding”). This means if you want to use the value a variable has when the function is created, you need to save the current value of that variable. Whereas Python uses “late binding” by default, the alternative that newcomers often expect is “early binding”, where a function uses the value a variable had when the function was defined, can be easily implemented.

Consider the following programs designed to create a list of 5 functions, where the i th function in the list is meant to add i to its argument:¹

```
pythonDemo.py — Some tricky examples
11 fun_list1 = []
12 for i in range(5):
13     def fun1(e):
14         return e+i
15     fun_list1.append(fun1)
16
17 fun_list2 = []
18 for i in range(5):
19     def fun2(e,iv=i):
20         return e+iv
21     fun_list2.append(fun2)
22
23 fun_list3 = [lambda e: e+i for i in range(5)]
24
25 fun_list4 = [lambda e,iv=i: e+iv for i in range(5)]
26
27 i=56
```

Try to predict, and then test to see the output, of the output of the following calls, remembering that the function uses the latest value of any variable that is not bound in the function call:

```
pythonDemo.py — (continued)
29 # in Shell do
30 ## ipython -i pythonDemo.py
31 # Try these (copy text after the comment symbol and paste in the Python prompt):
32 # print([f(10) for f in fun_list1])
33 # print([f(10) for f in fun_list2])
34 # print([f(10) for f in fun_list3])
35 # print([f(10) for f in fun_list4])
```

In the first for-loop, the function *fun* uses *i*, whose value is the last value it was assigned. In the second loop, the function *fun2* uses *iv*. There is a separate *iv* variable for each function, and its value is the value of *i* when the function was defined. Thus *fun1* uses late binding, and *fun2* uses early binding. *fun_list3* and *fun_list4* are equivalent to the first two (except *fun_list4* uses a different *i* variable).

¹Numbered lines are Python code available in the code-directory, aipython. The name of the file is given in the gray text above the listing. The numbers correspond to the line numbers in that file.

One of the advantages of using the embedded definitions (as in *fun1* and *fun2* above) over the *lambda* is that it is possible to add a `__doc__` string, which is the standard for documenting functions in Python, to the embedded definitions.

1.5.3 Generators and Coroutines

Python has generators which can be used for a form of coroutines.

The *yield* command returns a value that is obtained with *next*. It is typically used to enumerate the values for a *for* loop or in generators.

A version of the built-in *range*, with 2 or 3 arguments (and positive steps) can be implemented as:

```
pythonDemo.py — (continued)
37 def myrange(start, stop, step=1):
38     """enumerates the values from start in steps of size step that are
39     less than stop.
40     """
41     assert step>0, "only positive steps implemented in myrange"
42     i = start
43     while i<stop:
44         yield i
45         i += step
46
47 print("myrange(2,30,3):", list(myrange(2,30,3)))
```

Note that the built-in *range* is unconventional in how it handles a single argument, as the single argument acts as the second argument of the function. Note also that the built-in *range* also allows for indexing (e.g., *range(2,30,3)[2]* returns 8), which the above implementation does not. However *myrange* also works for floats, which the built-in *range* does not.

Exercise 1.1 Implement a version of *myrange* that acts like the built-in version when there is a single argument. (Hint: make the second argument have a default value that can be recognized in the function.)

Yield can be used to generate the same sequence of values as in the example of Section 1.5.1:

```
pythonDemo.py — (continued)
49 def ga(n):
50     """generates square of even nonnegative integers less than n"""
51     for e in range(n):
52         if e%2==0:
53             yield e*e
54 a = ga(20)
```

The sequence of *next(a)*, and *list(a)* gives exactly the same results as the comprehension in Section 1.5.1.

It is straightforward to write a version of the built-in *enumerate*. Let's call it *myenumerate*:

```
pythonDemo.py — (continued)
56 def myenumerate(enum):
57     for i in range(len(enum)):
58         yield i,enum[i]
```

Exercise 1.2 Write a version of *enumerate* where the only iteration is “for val in enum”. Hint: keep track of the index.

1.6 Useful Libraries

1.6.1 Timing Code

In order to compare algorithms, we often want to compute how long a program takes; this is called the **runtime** of the program. The most straightforward way to compute runtime is to use *time.perf_counter()*, as in:

```
import time
start_time = time.perf_counter()
compute_for_a_while()
end_time = time.perf_counter()
print("Time:", end_time - start_time, "seconds")
```

Note that *time.perf_counter()* measures clock time; so this should be done without user interaction between the calls. On the console, you should do:

```
start_time = time.perf_counter(); compute_for_a_while(); end_time = time.perf_counter()
```

If this time is very small (say less than 0.2 second), it is probably very inaccurate, and it may be better to run your code many times to get a more accurate count. For this you can use *timeit* (<https://docs.python.org/3/library/timeit.html>). To use *timeit* to time the call to *foo.bar(aaa)* use:

```
import timeit
time = timeit.timeit("foo.bar(aaa)",
                    setup="from __main__ import foo,aaa", number=100)
```

The setup is needed so that Python can find the meaning of the names in the string that is called. This returns the number of seconds to execute *foo.bar(aaa)* 100 times. The variable *number* should be set so that the runtime is at least 0.2 seconds.

You should not trust a single measurement as that can be confounded by interference from other processes. *timeit.repeat* can be used for running *timeit* a few (say 3) times. Usually the minimum time is the one to report, but you should be explicit and explain what you are reporting.

1.6.2 Plotting: Matplotlib

The standard plotting for Python is matplotlib (<http://matplotlib.org/>). We will use the most basic plotting using the pyplot interface.

Here is a simple example that uses everything we will use.

```
pythonDemo.py — (continued)
60 import matplotlib.pyplot as plt
61
62 def myplot(min,max,step,fun1,fun2):
63     plt.ion() # make it interactive
64     plt.xlabel("The x axis")
65     plt.ylabel("The y axis")
66     plt.xscale('linear') # Makes a 'log' or 'linear' scale
67     xvalues = range(min,max,step)
68     plt.plot(xvalues,[fun1(x) for x in xvalues],
69             label="The first fun")
70     plt.plot(xvalues,[fun2(x) for x in xvalues], linestyle='--',color='k',
71             label=fun2.__doc__) # use the doc string of the function
72     plt.legend(loc="upper right") # display the legend
73
74 def slin(x):
75     """y=2x+7"""
76     return 2*x+7
77 def sqfun(x):
78     """y=(x-40)^2/10-20"""
79     return (x-40)**2/10-20
80
81 # Try the following:
82 # from pythonDemo import myplot, slin, sqfun
83 # import matplotlib.pyplot as plt
84 # myplot(0,100,1,slin,sqfun)
85 # plt.legend(loc="best")
86 # import math
87 # plt.plot([41+40*math.cos(th/10) for th in range(50)],
88 #          [100+100*math.sin(th/10) for th in range(50)])
89 # plt.text(40,100,"ellipse?")
90 # plt.xscale('log')
```

At the end of the code are some commented-out commands you should try in interactive mode. Cut from the file and paste into Python (and remember to remove the comments symbol and leading space).

1.7 Utilities

1.7.1 Display

In this distribution, to keep things simple and to only use standard Python, we use a text-oriented tracing of the code. A graphical depiction of the code could

override the definition of *display* (but we leave it as a project).

The method *self.display* is used to trace the program. Any call

```
self.display(level,to_print...)
```

where the *level* is less than or equal to the value for *max_display_level* will be printed. The *to_print...* can be anything that is accepted by the built-in *print* (including any keyword arguments).

The definition of *display* is:

```

display.py — A simple way to trace the intermediate steps of algorithms.
11 class Displayable(object):
12     """Class that uses 'display'.
13     The amount of detail is controlled by max_display_level
14     """
15     max_display_level = 1 # can be overridden in subclasses
16
17     def display(self,level,*args,**nargs):
18         """print the arguments if level is less than or equal to the
19         current max_display_level.
20         level is an integer.
21         the other arguments are whatever arguments print can take.
22         """
23         if level <= self.max_display_level:
24             print(*args, **nargs) ##if error you are using Python2 not Python3

```

Note that *args* gets a tuple of the positional arguments, and *nargs* gets a dictionary of the keyword arguments). This will not work in Python 2, and will give an error.

Any class that wants to use *display* can be made a subclass of *Displayable*.

To change the maximum display level to say 3, for a class do:

```
Classname.max_display_level = 3
```

which will make calls to *display* in that class print when the value of *level* is less than-or-equal to 3. The default display level is 1. It can also be changed for individual objects (the object value overrides the class value).

The value of *max_display_level* by convention is:

- 0 display nothing
- 1 display solutions (nothing that happens repeatedly)
- 2 also display the values as they change (little detail through a loop)
- 3 also display more details
- 4 and above even more detail

In order to implement more sophisticated visualizations of the algorithm, we add a **visualize** “decorator” to the methods to be visualized. The following code ignores the decorator:

```

display.py — (continued)
26 def visualize(func):
27     """A decorator for algorithms that do interactive visualization.
28     Ignored here.
29     """
30     return func

```

1.7.2 Argmax

Python has a built-in *max* function that takes a generator (or a list or set) and returns the maximum value. The *argmax* method returns the index of an element that has the maximum value. If there are multiple elements with the maximum value, one of the indexes to that value is returned at random. *argmaxe* assumes an enumeration; a generator of (*element*, *value*) pairs, as for example is generated by the built-in *enumerate(list)* for lists or *dict.items()* for dicts.

```

utilities.py — AIPython useful utilities
11 import random
12 import math
13
14 def argmaxall(gen):
15     """gen is a generator of (element,value) pairs, where value is a real.
16     argmaxall returns a list of all of the elements with maximal value.
17     """
18     maxv = -math.inf    # negative infinity
19     maxvals = []       # list of maximal elements
20     for (e,v) in gen:
21         if v>maxv:
22             maxvals,maxv = [e], v
23         elif v==maxv:
24             maxvals.append(e)
25     return maxvals
26
27 def argmaxe(gen):
28     """gen is a generator of (element,value) pairs, where value is a real.
29     argmaxe returns an element with maximal value.
30     If there are multiple elements with the max value, one is returned at random.
31     """
32     return random.choice(argmaxall(gen))
33
34 def argmax(lst):
35     """returns maximum index in a list"""
36     return argmaxe(enumerate(lst))
37 # Try:
38 # argmax([1,6,3,77,3,55,23])
39
40 def argmaxd(dct):
41     """returns the arg max of a dictionary dct"""
42     return argmaxe(dct.items())

```

```

43 | # Try:
44 | # arxmaxd({2:5,5:9,7:7})

```

Exercise 1.3 Change `argmax` to have an optional argument that specifies whether you want the “first”, “last” or a “random” index of the maximum value returned. If you want the first or the last, you don’t need to keep a list of the maximum elements.

1.7.3 Probability

For many of the simulations, we want to make a variable `True` with some probability. `flip(p)` returns `True` with probability *p*, and otherwise returns `False`.

```

_____utilities.py — (continued) _____
45 | def flip(prob):
46 |     """return true with probability prob"""
47 |     return random.random() < prob

```

1.7.4 Dictionary Union

This is now | in Python 3.9, so will be replaced.

The function `dict_union(d1,d2)` returns the union of dictionaries *d1* and *d2*. If the values for the keys conflict, the values in *d2* are used. This is similar to `dict(d1,**d2)`, but that only works when the keys of *d2* are strings.

```

_____utilities.py — (continued) _____
49 | def dict_union(d1,d2):
50 |     """returns a dictionary that contains the keys of d1 and d2.
51 |     The value for each key that is in d2 is the value from d2,
52 |     otherwise it is the value from d1.
53 |     This does not have side effects.
54 |     """
55 |     d = dict(d1) # copy d1
56 |     d.update(d2)
57 |     return d

```

1.8 Testing Code

It is important to test code early and test it often. We include a simple form of **unit test**. The value of the current module is in `__name__` and if the module is run at the top-level, it’s value is `“__main__”`. See https://docs.python.org/3/library/__main__.html.

The following code tests `argmax` and `dict_union`, but only when if `utilities` is loaded in the top-level. If it is loaded in a module the test code is not run.

In your code you should do more substantial testing than we do here, in particular testing the boundary cases.

```
utilities.py — (continued) —
59 def test():
60     """Test part of utilities"""
61     assert argmax(enumerate([1,6,55,3,55,23])) in [2,4]
62     assert dict_union({1:4, 2:5, 3:4},{5:7, 2:9}) == {1:4, 2:9, 3:4, 5:7}
63     print("Passed unit test in utilities")
64
65 if __name__ == "__main__":
66     test()
```

Agents and Control

This implements the controllers described in Chapter 2.

In this version the higher-levels call the lower-levels. A more sophisticated version may have them run concurrently (either as coroutines or in parallel). The higher-levels calling the lower-level works in simulated environments when there is a single agent, and where the lower-level are written to make sure they return (and don't go on forever), and the higher level doesn't take too long (as the lower-levels will wait until called again).

2.1 Representing Agents and Environments

An agent observes the world, and carries out actions in the environment, it also has an internal state that it updates. The environment takes in actions of the agents, updates its internal state and returns the percepts.

In this implementation, the state of the agent and the state of the environment are represented using standard Python variables, which are updated as the state changes. The percepts and the actions are represented as variable-value dictionaries.

An agent implements the $go(n)$ method, where n is an integer. This means that the agent should run for n time steps.

In the following code `raise NotImplementedError()` is a way to specify an abstract method that needs to be overridden in any implemented agent or environment.

```
agents.py — Agent and Controllers
11 import random
12
13 class Agent(object):
14     def __init__(self, env):
```

```

15     """set up the agent"""
16     self.env=env
17
18     def go(self,n):
19         """acts for n time steps"""
20         raise NotImplementedError("go") # abstract method

```

The environment implements a *do(action)* method where *action* is a variable-value dictionary. This returns a percept, which is also a variable-value dictionary. The use of dictionaries allows for structured actions and percepts.

Note that *Environment* is a subclass of *Displayable* so that it can use the *display* method described in Section 1.7.1.

```

agents.py — (continued)
22 from display import Displayable
23 class Environment(Displayable):
24     def initial_percepts(self):
25         """returns the initial percepts for the agent"""
26         raise NotImplementedError("initial_percepts") # abstract method
27
28     def do(self,action):
29         """does the action in the environment
30         returns the next percept """
31         raise NotImplementedError("do") # abstract method

```

2.2 Paper buying agent and environment

To run the demo, in folder "aipython", load "agents.py", using e.g., `ipython -i agents.py`, and copy and paste the commented-out commands at the bottom of that file. This requires Python 3 with matplotlib.

This is an implementation of the paper buying example.

2.2.1 The Environment

The environment state is given in terms of the *time* and the amount of paper in *stock*. It also remembers the in-stock history and the price history. The percepts are the price and the amount of paper in stock. The action of the agent is the number to buy.

Here we assume that the prices are obtained from the *prices* list plus a random integer in range $[0, \text{max_price_addon})$ plus a linear "inflation". The agent cannot access the price model; it just observes the prices and the amount in stock.

```

agents.py — (continued)
33 class TP_env(Environment):

```

```

34 prices = [234, 234, 234, 234, 255, 255, 275, 275, 211, 211, 211,
35           234, 234, 234, 234, 199, 199, 275, 275, 234, 234, 234, 234, 255,
36           255, 260, 260, 265, 265, 265, 265, 270, 270, 255, 255, 260, 260,
37           265, 265, 150, 150, 265, 265, 270, 270, 255, 255, 260, 260, 265,
38           265, 265, 265, 270, 270, 211, 211, 255, 255, 260, 260, 265, 265,
39           260, 265, 270, 270, 205, 255, 255, 260, 260, 265, 265, 265, 265,
40           270, 270]
41 max_price_addon = 20 # maximum of random value added to get price
42
43 def __init__(self):
44     """paper buying agent"""
45     self.time=0
46     self.stock=20
47     self.stock_history = [] # memory of the stock history
48     self.price_history = [] # memory of the price history
49
50 def initial_percepts(self):
51     """return initial percepts"""
52     self.stock_history.append(self.stock)
53     price = self.prices[0]+random.randrange(self.max_price_addon)
54     self.price_history.append(price)
55     return {'price': price,
56           'instock': self.stock}
57
58 def do(self, action):
59     """does action (buy) and returns percepts (price and instock)"""
60     used = pick_from_dist({6:0.1, 5:0.1, 4:0.2, 3:0.3, 2:0.2, 1:0.1})
61     bought = action['buy']
62     self.stock = self.stock+bought-used
63     self.stock_history.append(self.stock)
64     self.time += 1
65     price = (self.prices[self.time%len(self.prices)] # repeating pattern
66             +random.randrange(self.max_price_addon) # plus randomness
67             +self.time//2) # plus inflation
68     self.price_history.append(price)
69     return {'price': price,
70           'instock':self.stock}

```

The *pick_from_dist* method takes in a *item : probability* dictionary, and returns one of the items in proportion to its probability.

agents.py — (continued)

```

72 def pick_from_dist(item_prob_dist):
73     """ returns a value from a distribution.
74     item_prob_dist is an item:probability dictionary, where the
75     probabilities sum to 1.
76     returns an item chosen in proportion to its probability
77     """
78     ranreal = random.random()
79     for (it,prob) in item_prob_dist.items():
80         if ranreal < prob:

```

```

81         return it
82     else:
83         ranreal -= prob
84     raise RuntimeError(str(item_prob_dist)+" is not a probability distribution")

```

2.2.2 The Agent

The agent does not have access to the price model but can only observe the current price and the amount in stock. It has to decide how much to buy.

The belief state of the agent is an estimate of the average price of the paper, and the total amount of money the agent has spent.

```

agents.py — (continued)
86 class TP_agent(Agent):
87     def __init__(self, env):
88         self.env = env
89         self.spent = 0
90         percepts = env.initial_percepts()
91         self.ave = self.last_price = percepts['price']
92         self.instock = percepts['instock']
93
94     def go(self, n):
95         """go for n time steps
96         """
97         for i in range(n):
98             if self.last_price < 0.9*self.ave and self.instock < 60:
99                 tobuy = 48
100             elif self.instock < 12:
101                 tobuy = 12
102             else:
103                 tobuy = 0
104             self.spent += tobuy*self.last_price
105             percepts = env.do({'buy': tobuy})
106             self.last_price = percepts['price']
107             self.ave = self.ave+(self.last_price-self.ave)*0.05
108             self.instock = percepts['instock']

```

Set up an environment and an agent. Uncomment the last lines to run the agent for 90 steps, and determine the average amount spent.

```

agents.py — (continued)
110 env = TP_env()
111 ag = TP_agent(env)
112 #ag.go(90)
113 #ag.spent/env.time ## average spent per time period

```

2.2.3 Plotting

The following plots the price and number in stock history:


```

agents.py — (continued)
115 import matplotlib.pyplot as plt
116
117 class Plot_prices(object):
118     """Set up the plot for history of price and number in stock"""
119     def __init__(self, ag, env):
120         self.ag = ag
121         self.env = env
122         plt.ion()
123         plt.xlabel("Time")
124         plt.ylabel("Number in stock.                                Price.")
125
126     def plot_run(self):
127         """plot history of price and instock"""
128         num = len(env.stock_history)
129         plt.plot(range(num), env.stock_history, label="In stock")
130         plt.plot(range(num), env.price_history, label="Price")
131         #plt.legend(loc="upper left")
132         plt.draw()
133
134 # pl = Plot_prices(ag, env)
135 # ag.go(90); pl.plot_run()

```

2.3 Hierarchical Controller

To run the hierarchical controller, in folder "aipython", load "agentTop.py", using e.g., `ipython -i agentTop.py`, and copy and paste the commands near the bottom of that file. This requires Python 3 with matplotlib.

In this implementation, each layer, including the top layer, implements the environment class, because each layer is seen as an environment from the layer above.

We arbitrarily divide the environment and the body, so that the environment just defines the walls, and the body includes everything to do with the agent. Note that the named locations are part of the (top-level of the) agent, not part of the environment, although they could have been.

2.3.1 Environment

The environment defines the walls.

```

agentEnv.py — Agent environment
11 import math
12 from agents import Environment
13
14 class Rob_env(Environment):

```

```

15     def __init__(self, walls = {}):
16         """walls is a set of line segments
17             where each line segment is of the form ((x0,y0),(x1,y1))
18         """
19         self.walls = walls

```

2.3.2 Body

The body defines everything about the agent body.

```

agentEnv.py — (continued)
21 import math
22 from agents import Environment
23 import matplotlib.pyplot as plt
24 import time
25
26 class Rob_body(Environment):
27     def __init__(self, env, init_pos=(0,0,90)):
28         """ env is the current environment
29             init_pos is a triple of (x-position, y-position, direction)
30             direction is in degrees; 0 is to right, 90 is straight-up, etc
31         """
32         self.env = env
33         self.rob_x, self.rob_y, self.rob_dir = init_pos
34         self.turning_angle = 18 # degrees that a left makes
35         self.whisker_length = 6 # length of the whisker
36         self.whisker_angle = 30 # angle of whisker relative to robot
37         self.crashed = False
38         # The following control how it is plotted
39         self.plotting = True # whether the trace is being plotted
40         self.sleep_time = 0.05 # time between actions (for real-time plotting)
41         # The following are data structures maintained:
42         self.history = [(self.rob_x, self.rob_y)] # history of (x,y) positions
43         self.wall_history = [] # history of hitting the wall
44
45     def percepts(self):
46         return {'rob_x_pos':self.rob_x, 'rob_y_pos':self.rob_y,
47             'rob_dir':self.rob_dir, 'whisker':self.whisker() , 'crashed':self.crashed}
48     initial_percepts = percepts # use percept function for initial percepts too
49
50     def do(self, action):
51         """ action is {'steer':direction}
52             direction is 'left', 'right' or 'straight'
53         """
54         if self.crashed:
55             return self.percepts()
56         direction = action['steer']
57         compass_deriv = {'left':1, 'straight':0, 'right':-1}[direction]*self.turning_angle
58         self.rob_dir = (self.rob_dir + compass_deriv + 360)%360 # make in range [0,360)
59         rob_x_new = self.rob_x + math.cos(self.rob_dir*math.pi/180)

```

```

60     rob_y_new = self.rob_y + math.sin(self.rob_dir*math.pi/180)
61     path = ((self.rob_x,self.rob_y),(rob_x_new,rob_y_new))
62     if any(line_segments_intersect(path,wall) for wall in self.env.walls):
63         self.crashed = True
64         if self.plotting:
65             plt.plot([self.rob_x],[self.rob_y],"r*",markersize=20.0)
66             plt.draw()
67     self.rob_x, self.rob_y = rob_x_new, rob_y_new
68     self.history.append((self.rob_x, self.rob_y))
69     if self.plotting and not self.crashed:
70         plt.plot([self.rob_x],[self.rob_y],"go")
71         plt.draw()
72         plt.pause(self.sleep_time)
73     return self.percepts()

```

This detects if the whisker and the wall intersect. It's value is returned as a percept.

agentEnv.py — (continued)

```

75     def whisker(self):
76         """returns true whenever the whisker sensor intersects with a wall
77         """
78         whisk_ang_world = (self.rob_dir-self.whisker_angle)*math.pi/180
79         # angle in radians in world coordinates
80         wx = self.rob_x + self.whisker_length * math.cos(whisk_ang_world)
81         wy = self.rob_y + self.whisker_length * math.sin(whisk_ang_world)
82         whisker_line = ((self.rob_x,self.rob_y),(wx,wy))
83         hit = any(line_segments_intersect(whisker_line,wall)
84                 for wall in self.env.walls)
85         if hit:
86             self.wall_history.append((self.rob_x, self.rob_y))
87             if self.plotting:
88                 plt.plot([self.rob_x],[self.rob_y],"ro")
89                 plt.draw()
90         return hit
91
92     def line_segments_intersect(linea,lineb):
93         """returns true if the line segments, linea and lineb intersect.
94         A line segment is represented as a pair of points.
95         A point is represented as a (x,y) pair.
96         """
97         ((x0a,y0a),(x1a,y1a)) = linea
98         ((x0b,y0b),(x1b,y1b)) = lineb
99         da, db = x1a-x0a, x1b-x0b
100        ea, eb = y1a-y0a, y1b-y0b
101        denom = db*ea-eb*da
102        if denom==0: # line segments are parallel
103            return False
104        cb = (da*(y0b-y0a)-ea*(x0b-x0a))/denom # position along line b
105        if cb<0 or cb>1:
106            return False

```

```

107     ca = (db*(y0b-y0a)-eb*(x0b-x0a))/denom # position along line a
108     return 0<=ca<=1
109
110 # Test cases:
111 # assert line_segments_intersect(((0,0),(1,1)),((1,0),(0,1)))
112 # assert not line_segments_intersect(((0,0),(1,1)),((1,0),(0.6,0.4)))
113 # assert line_segments_intersect(((0,0),(1,1)),((1,0),(0.4,0.6)))

```

2.3.3 Middle Layer

The middle layer acts like both a controller (for the environment layer) and an environment for the upper layer. It has to tell the environment how to steer. Thus it calls *env.do(·)*. It also is told the position to go to and the timeout. Thus it also has to implement *do(·)*.

```

agentMiddle.py — Middle Layer
11 from agents import Environment
12 import math
13
14 class Rob_middle_layer(Environment):
15     def __init__(self,env):
16         self.env=env
17         self.percepts = env.initial_percepts()
18         self.straight_angle = 11 # angle that is close enough to straight ahead
19         self.close_threshold = 2 # distance that is close enough to arrived
20         self.close_threshold_squared = self.close_threshold**2 # just compute it once
21
22     def initial_percepts(self):
23         return {}
24
25     def do(self, action):
26         """action is {'go_to':target_pos,'timeout':timeout}
27         target_pos is (x,y) pair
28         timeout is the number of steps to try
29         returns {'arrived':True} when arrived is true
30         or {'arrived':False} if it reached the timeout
31         """
32         if 'timeout' in action:
33             remaining = action['timeout']
34         else:
35             remaining = -1 # will never reach 0
36         target_pos = action['go_to']
37         arrived = self.close_enough(target_pos)
38         while not arrived and remaining != 0:
39             self.percepts = self.env.do({"steer":self.steer(target_pos)})
40             remaining -= 1
41             arrived = self.close_enough(target_pos)
42         return {'arrived':arrived}

```

This determines how to steer depending on whether the goal is to the right or the left of where the robot is facing.

```

agentMiddle.py — (continued)
44 def steer(self, target_pos):
45     if self.percepts['whisker']:
46         self.display(3, 'whisker on', self.percepts)
47         return "left"
48     else:
49         gx, gy = target_pos
50         rx, ry = self.percepts['rob_x_pos'], self.percepts['rob_y_pos']
51         goal_dir = math.acos((gx-rx)/math.sqrt((gx-rx)*(gx-rx)
52                                     +(gy-ry)*(gy-ry)))*180/math.pi
53         if ry>gy:
54             goal_dir = -goal_dir
55         goal_from_rob = (goal_dir - self.percepts['rob_dir']+540)%360-180
56         assert -180 < goal_from_rob <= 180
57         if goal_from_rob > self.straight_angle:
58             return "left"
59         elif goal_from_rob < -self.straight_angle:
60             return "right"
61         else:
62             return "straight"
63
64 def close_enough(self, target_pos):
65     gx, gy = target_pos
66     rx, ry = self.percepts['rob_x_pos'], self.percepts['rob_y_pos']
67     return (gx-rx)**2 + (gy-ry)**2 <= self.close_threshold_squared

```

2.3.4 Top Layer

The top layer treats the middle layer as its environment. Note that the top layer is an environment for us to tell it what to visit.

```

agentTop.py — Top Layer
11 from agentMiddle import Rob_middle_layer
12 from agents import Environment
13
14 class Rob_top_layer(Environment):
15     def __init__(self, middle, timeout=200, locations = {'mail':(-5,10),
16                                                         'o103':(50,10), 'o109':(100,10), 'storage':(101,51)} ):
17         """middle is the middle layer
18         timeout is the number of steps the middle layer goes before giving up
19         locations is a loc:pos dictionary
20         where loc is a named location, and pos is an (x,y) position.
21         """
22         self.middle = middle
23         self.timeout = timeout # number of steps before the middle layer should give up
24         self.locations = locations
25

```

```

26     def do(self, plan):
27         """carry out actions.
28         actions is of the form {'visit':list_of_locations}
29         It visits the locations in turn.
30         """
31         to_do = plan['visit']
32         for loc in to_do:
33             position = self.locations[loc]
34             arrived = self.middle.do({'go_to':position, 'timeout':self.timeout})
35             self.display(1, "Arrived at", loc, arrived)

```

2.3.5 Plotting

The following is used to plot the locations, the walls and (eventually) the movement of the robot. It can either plot the movement if the robot as it is going (with the default `env.plotting = True`), or not plot it as it is going (setting `env.plotting = False`; in this case the trace can be plotted using `pl.plot_run()`).

```

agentTop.py — (continued)
37 import matplotlib.pyplot as plt
38
39 class Plot_env(object):
40     def __init__(self, body, top):
41         """sets up the plot
42         """
43         self.body = body
44         plt.ion()
45         plt.clf()
46         plt.axes().set_aspect('equal')
47         for wall in body.env.walls:
48             ((x0,y0),(x1,y1)) = wall
49             plt.plot([x0,x1],[y0,y1], "-k", linewidth=3)
50         for loc in top.locations:
51             (x,y) = top.locations[loc]
52             plt.plot([x],[y], "k<")
53             plt.text(x+1.0,y+0.5,loc) # print the label above and to the right
54         plt.plot([body.rob_x],[body.rob_y], "go")
55         plt.draw()
56
57     def plot_run(self):
58         """plots the history after the agent has finished.
59         This is typically only used if body.plotting==False
60         """
61         xs,ys = zip(*self.body.history)
62         plt.plot(xs,ys, "go")
63         wxs,wys = zip(*self.body.wall_history)
64         plt.plot(wxs,wys, "ro")
65         #plt.draw()

```

The following code plots the agent as it acts in the world:

```

agentTop.py — (continued)
67 from agentEnv import Rob_body, Rob_env
68
69 env = Rob_env({((20,0),(30,20)), ((70,-5),(70,25))})
70 body = Rob_body(env)
71 middle = Rob_middle_layer(body)
72 top = Rob_top_layer(middle)
73
74 # try:
75 # pl=Plot_env(body,top)
76 # top.do({'visit':['o109','storage','o109','o103']})
77 # You can directly control the middle layer:
78 # middle.do({'go_to':(30,-10), 'timeout':200})
79 # Can you make it crash?

```

Exercise 2.1 The following code implements a robot trap. Write a controller that can escape the “trap” and get to the goal. See textbook for hints.

```

agentTop.py — (continued)
81 # Robot Trap for which the current controller cannot escape:
82 trap_env = Rob_env({((10,-21),(10,0)), ((10,10),(10,31)), ((30,-10),(30,0)),
83                     ((30,10),(30,20)), ((50,-21),(50,31)), ((10,-21),(50,-21)),
84                     ((10,0),(30,0)), ((10,10),(30,10)), ((10,31),(50,31))})
85 trap_body = Rob_body(trap_env,init_pos=(-1,0,90))
86 trap_middle = Rob_middle_layer(trap_body)
87 trap_top = Rob_top_layer(trap_middle,locations={'goal':(71,0)})
88
89 # Robot trap exercise:
90 # pl=Plot_env(trap_body,trap_top)
91 # trap_top.do({'visit':['goal']})

```


Searching for Solutions

3.1 Representing Search Problems

A search problem consists of:

- a start node
- a neighbors function that given a node, returns an enumeration of the arcs from the node
- a specification of a goal in terms of a Boolean function that takes a node and returns true if the node is a goal
- a (optional) heuristic function that, given a node, returns a non-negative real number. The heuristic function defaults to zero.

As far as the searcher is concerned a node can be anything. If multiple-path pruning is used, a node must be hashable. In the simple examples, it is a string, but in more complicated examples (in later chapters) it can be a tuple, a frozen set, or a Python object.

In the following code raise `NotImplementedError()` is a way to specify that this is an abstract method that needs to be overridden to define an actual search problem.

```
searchProblem.py — representations of search problems
11 class Search_problem(object):
12     """A search problem consists of:
13     * a start node
14     * a neighbors function that gives the neighbors of a node
15     * a specification of a goal
16     * a (optional) heuristic function.
```

```

17     The methods must be overridden to define a search problem."""
18
19     def start_node(self):
20         """returns start node"""
21         raise NotImplementedError("start_node") # abstract method
22
23     def is_goal(self,node):
24         """is True if node is a goal"""
25         raise NotImplementedError("is_goal") # abstract method
26
27     def neighbors(self,node):
28         """returns a list of the arcs for the neighbors of node"""
29         raise NotImplementedError("neighbors") # abstract method
30
31     def heuristic(self,n):
32         """Gives the heuristic value of node n.
33         Returns 0 if not overridden."""
34         return 0

```

The neighbors is a list of arcs. A (directed) arc consists of a *from_node* node and a *to_node* node. The arc is the pair $\langle from_node, to_node \rangle$, but can also contain a non-negative *cost* (which defaults to 1) and can be labeled with an *action*.

searchProblem.py — (continued)

```

36 class Arc(object):
37     """An arc has a from_node and a to_node node and a (non-negative) cost"""
38     def __init__(self, from_node, to_node, cost=1, action=None):
39         assert cost >= 0, ("Cost cannot be negative for"+
40                             str(from_node)+"->"+str(to_node)+"", cost: "+str(cost))
41         self.from_node = from_node
42         self.to_node = to_node
43         self.action = action
44         self.cost=cost
45
46     def __repr__(self):
47         """string representation of an arc"""
48         if self.action:
49             return str(self.from_node)+" --"+str(self.action)+"--> "+str(self.to_node)
50         else:
51             return str(self.from_node)+" --> "+str(self.to_node)

```

3.1.1 Explicit Representation of Search Graph

The first representation of a search problem is from an explicit graph (as opposed to one that is generated as needed).

An **explicit graph** consists of

- a list or set of nodes
- a list or set of arcs

- a start node
- a list or set of goal nodes
- (optionally) a dictionary that maps a node to a heuristic value for that node

To define a search problem, we need to define the start node, the goal predicate, the neighbors function and the heuristic function.

```

53 class Search_problem_from_explicit_graph(Search_problem):
54     """A search problem consists of:
55     * a list or set of nodes
56     * a list or set of arcs
57     * a start node
58     * a list or set of goal nodes
59     * a dictionary that maps each node into its heuristic value.
60     * a dictionary that maps each node into its (x,y) position
61     """
62
63     def __init__(self, nodes, arcs, start=None, goals=set(), hmap={}, positions={}):
64         self.neighs = {}
65         self.nodes = nodes
66         for node in nodes:
67             self.neighs[node]=[]
68         self.arcs = arcs
69         for arc in arcs:
70             self.neighs[arc.from_node].append(arc)
71         self.start = start
72         self.goals = goals
73         self.hmap = hmap
74         self.positions = positions
75
76     def start_node(self):
77         """returns start node"""
78         return self.start
79
80     def is_goal(self,node):
81         """is True if node is a goal"""
82         return node in self.goals
83
84     def neighbors(self,node):
85         """returns the neighbors of node"""
86         return self.neighs[node]
87
88     def heuristic(self,node):
89         """Gives the heuristic value of node n.
90         Returns 0 if not overridden in the hmap."""
91         if node in self.hmap:
92             return self.hmap[node]

```

```

93         else:
94             return 0
95
96     def __repr__(self):
97         """returns a string representation of the search problem"""
98         res=""
99         for arc in self.arcs:
100             res += str(arc)+" "
101         return res

```

The following is used for the depth-first search implementation below.

```

searchProblem.py — (continued)
103     def neighbor_nodes(self,node):
104         """returns an iterator over the neighbors of node"""
105         return (path.to_node for path in self.neighs[node])

```

3.1.2 Paths

A searcher will return a path from the start node to a goal node. A Python list is not a suitable representation for a path, as many search algorithms consider multiple paths at once, and these paths should share initial parts of the path. If we wanted to do this with Python lists, we would need to keep copying the list, which can be expensive if the list is long. An alternative representation is used here in terms of a recursive data structure that can share subparts.

A path is either:

- a node (representing a path of length 0) or
- a path, *initial* and an arc, where the *from_node* of the arc is the node at the end of *initial*.

These cases are distinguished in the following code by having *arc = None* if the path has length 0, in which case *initial* is the node of the path.

```

searchProblem.py — (continued)
107 class Path(object):
108     """A path is either a node or a path followed by an arc"""
109
110     def __init__(self,initial,arc=None):
111         """initial is either a node (in which case arc is None) or
112         a path (in which case arc is an object of type Arc)"""
113         self.initial = initial
114         self.arc=arc
115         if arc is None:
116             self.cost=0
117         else:
118             self.cost = initial.cost+arc.cost
119
120     def end(self):

```

```

121     """returns the node at the end of the path"""
122     if self.arc is None:
123         return self.initial
124     else:
125         return self.arc.to_node
126
127     def nodes(self):
128         """enumerates the nodes for the path.
129         This starts at the end and enumerates nodes in the path backwards."""
130         current = self
131         while current.arc is not None:
132             yield current.arc.to_node
133             current = current.initial
134         yield current.initial
135
136     def initial_nodes(self):
137         """enumerates the nodes for the path before the end node.
138         This starts at the end and enumerates nodes in the path backwards."""
139         if self.arc is not None:
140             for nd in self.initial.nodes(): yield nd # could be "yield from"
141
142     def __repr__(self):
143         """returns a string representation of a path"""
144         if self.arc is None:
145             return str(self.initial)
146         elif self.arc.action:
147             return (str(self.initial)+"\n --"+str(self.arc.action)
148                 +"--> "+str(self.arc.to_node))
149         else:
150             return str(self.initial)+" --> "+str(self.arc.to_node)

```

3.1.3 Example Search Problems

The first search problem is one with 5 nodes where the least-cost path is one with many arcs. See Figure 3.1. Note that this example is used for the unit tests, so the test (in `searchGeneric`) will need to be changed if this is changed.

```

searchProblem.py — (continued)
152 problem1 = Search_problem_from_explicit_graph(
153     {'a','b','c','d','g'},
154     [Arc('a','c',1), Arc('a','b',3), Arc('c','d',3), Arc('c','b',1),
155       Arc('b','d',1), Arc('b','g',3), Arc('d','g',1)],
156     start = 'a',
157     goals = {'g'},
158     positions={'a': (0, 0), 'b': (1, 1), 'c': (0,1), 'd': (1,2), 'g': (2,2)})

```

The second search problem is one with 8 nodes where many paths do not lead to the goal. See Figure 3.2.

```

searchProblem.py — (continued)

```

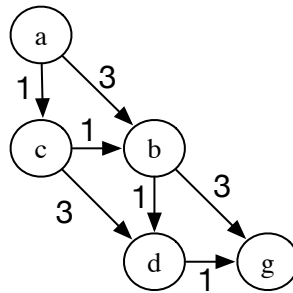


Figure 3.1: problem1

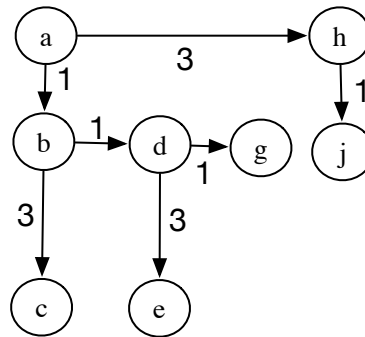


Figure 3.2: problem2

```

159 | problem2 = Search_problem_from_explicit_graph(
160 |     {'a','b','c','d','e','g','h','j'},
161 |     [Arc('a','b',1), Arc('b','c',3), Arc('b','d',1), Arc('d','e',3),
162 |       Arc('d','g',1), Arc('a','h',3), Arc('h','j',1)],
163 |     start = 'a',
164 |     goals = {'g'},
165 |     positions={'a': (0, 0), 'b': (0, 1), 'c': (0,4), 'd': (1,1), 'e': (1,4),
166 |               'g': (2,1), 'h': (3,0), 'j': (3,1)})

```

The third search problem is a disconnected graph (contains no arcs), where the start node is a goal node. This is a boundary case to make sure that weird cases work.

searchProblem.py — (continued) —

```

168 | problem3 = Search_problem_from_explicit_graph(
169 |     {'a','b','c','d','e','g','h','j'},
170 |     [],
171 |     start = 'g',
172 |     goals = {'k','g'})

```

The `acyclic_delivery_problem` is the delivery problem described in Example 3.4 and shown in Figure 3.2 of the textbook.

```

searchProblem.py — (continued)
174 acyclic_delivery_problem = Search_problem_from_explicit_graph(
175     {'mail', 'ts', 'o103', 'o109', 'o111', 'b1', 'b2', 'b3', 'b4', 'c1', 'c2', 'c3',
176     'o125', 'o123', 'o119', 'r123', 'storage'},
177     [Arc('ts', 'mail', 6),
178       Arc('o103', 'ts', 8),
179       Arc('o103', 'b3', 4),
180       Arc('o103', 'o109', 12),
181       Arc('o109', 'o119', 16),
182       Arc('o109', 'o111', 4),
183       Arc('b1', 'c2', 3),
184       Arc('b1', 'b2', 6),
185       Arc('b2', 'b4', 3),
186       Arc('b3', 'b1', 4),
187       Arc('b3', 'b4', 7),
188       Arc('b4', 'o109', 7),
189       Arc('c1', 'c3', 8),
190       Arc('c2', 'c3', 6),
191       Arc('c2', 'c1', 4),
192       Arc('o123', 'o125', 4),
193       Arc('o123', 'r123', 4),
194       Arc('o119', 'o123', 9),
195       Arc('o119', 'storage', 7)],
196     start = 'o103',
197     goals = {'r123'},
198     hmap = {
199         'mail' : 26,
200         'ts' : 23,
201         'o103' : 21,
202         'o109' : 24,
203         'o111' : 27,
204         'o119' : 11,
205         'o123' : 4,
206         'o125' : 6,
207         'r123' : 0,
208         'b1' : 13,
209         'b2' : 15,
210         'b3' : 17,
211         'b4' : 18,
212         'c1' : 6,
213         'c2' : 10,
214         'c3' : 12,
215         'storage' : 12
216     }
217 )

```

The cyclic_delivery_problem is the delivery problem described in Example 3.8 and shown in Figure 3.6 of the textbook. This is the same as acyclic_delivery_problem, but almost every arc also has its inverse.

```

219 cyclic_delivery_problem = Search_problem_from_explicit_graph(
220     {'mail', 'ts', 'o103', 'o109', 'o111', 'b1', 'b2', 'b3', 'b4', 'c1', 'c2', 'c3',
221      'o125', 'o123', 'o119', 'r123', 'storage'},
222     [ Arc('ts', 'mail', 6), Arc('mail', 'ts', 6),
223       Arc('o103', 'ts', 8), Arc('ts', 'o103', 8),
224       Arc('o103', 'b3', 4),
225       Arc('o103', 'o109', 12), Arc('o109', 'o103', 12),
226       Arc('o109', 'o119', 16), Arc('o119', 'o109', 16),
227       Arc('o109', 'o111', 4), Arc('o111', 'o109', 4),
228       Arc('b1', 'c2', 3),
229       Arc('b1', 'b2', 6), Arc('b2', 'b1', 6),
230       Arc('b2', 'b4', 3), Arc('b4', 'b2', 3),
231       Arc('b3', 'b1', 4), Arc('b1', 'b3', 4),
232       Arc('b3', 'b4', 7), Arc('b4', 'b3', 7),
233       Arc('b4', 'o109', 7),
234       Arc('c1', 'c3', 8), Arc('c3', 'c1', 8),
235       Arc('c2', 'c3', 6), Arc('c3', 'c2', 6),
236       Arc('c2', 'c1', 4), Arc('c1', 'c2', 4),
237       Arc('o123', 'o125', 4), Arc('o125', 'o123', 4),
238       Arc('o123', 'r123', 4), Arc('r123', 'o123', 4),
239       Arc('o119', 'o123', 9), Arc('o123', 'o119', 9),
240       Arc('o119', 'storage', 7), Arc('storage', 'o119', 7)],
241     start = 'o103',
242     goals = {'r123'},
243     hmap = {
244         'mail' : 26,
245         'ts' : 23,
246         'o103' : 21,
247         'o109' : 24,
248         'o111' : 27,
249         'o119' : 11,
250         'o123' : 4,
251         'o125' : 6,
252         'r123' : 0,
253         'b1' : 13,
254         'b2' : 15,
255         'b3' : 17,
256         'b4' : 18,
257         'c1' : 6,
258         'c2' : 10,
259         'c3' : 12,
260         'storage' : 12
261     }
262 )

```


3.2 Generic Searcher and Variants

To run the search demos, in folder “aipython”, load “searchGeneric.py” , using e.g., `ipython -i searchGeneric.py`, and copy and paste the example queries at the bottom of that file. This requires Python 3.

3.2.1 Searcher

A *Searcher* for a problem can be asked repeatedly for the next path. To solve a problem, we can construct a *Searcher* object for the problem and then repeatedly ask for the next path using *search*. If there are no more paths, *None* is returned.

```

searchGeneric.py — Generic Searcher, including depth-first and A*
11 from display import Displayable, visualize
12
13 class Searcher(Displayable):
14     """returns a searcher for a problem.
15     Paths can be found by repeatedly calling search().
16     This does depth-first search unless overridden
17     """
18     def __init__(self, problem):
19         """creates a searcher from a problem
20         """
21         self.problem = problem
22         self.initialize_frontier()
23         self.num_expanded = 0
24         self.add_to_frontier(Path(problem.start_node()))
25         super().__init__()
26
27     def initialize_frontier(self):
28         self.frontier = []
29
30     def empty_frontier(self):
31         return self.frontier == []
32
33     def add_to_frontier(self, path):
34         self.frontier.append(path)
35
36     @visualize
37     def search(self):
38         """returns (next) path from the problem's start node
39         to a goal node.
40         Returns None if no path exists.
41         """
42         while not self.empty_frontier():
43             path = self.frontier.pop()
44             self.display(2, "Expanding:", path, "(cost:", path.cost, ")")
45             self.num_expanded += 1

```

```

46         if self.problem.is_goal(path.end()): # solution found
47             self.display(1, self.num_expanded, "paths have been expanded and",
48                         len(self.frontier), "paths remain in the frontier")
49             self.solution = path # store the solution found
50             return path
51         else:
52             neighs = self.problem.neighbors(path.end())
53             self.display(3, "Neighbors are", neighs)
54             for arc in reversed(list(neighs)):
55                 self.add_to_frontier(Path(path, arc))
56             self.display(3, "Frontier:", self.frontier)
57     self.display(1, "No (more) solutions. Total of",
58                 self.num_expanded, "paths expanded.")

```

Note that this reverses the neighbours so that it implements depth-first search in an intuitive manner (expanding the first neighbor first), and *list* is needed if the neighbours are generated. Reversing the neighbours might not be required for other methods. The calls to *reversed* and *list* can be removed, and the algorithm still implements depth-first search.

Exercise 3.1 When it returns a path, the algorithm can be used to find another path by calling *search()* again. However, it does not find other paths that go through one goal node to another. Explain why, and change the code so that it can find such paths when *search()* is called again.

3.2.2 Frontier as a Priority Queue

In many of the search algorithms, such as A^* and other best-first searchers, the frontier is implemented as a priority queue. Here we use the Python's built-in priority queue implementations, *heapq*.

Following the lead of the Python documentation, <http://docs.python.org/3.3/library/heapq.html>, a frontier is a list of triples. The first element of each triple is the value to be minimized. The second element is a unique index which specifies the order when the first elements are the same, and the third element is the path that is on the queue. The use of the unique index ensures that the priority queue implementation does not compare paths; whether one path is less than another is not defined. It also lets us control what sort of search (e.g., depth-first or breadth-first) occurs when the value to be minimized does not give a unique next path.

The variable *frontier_index* is the total number of elements of the frontier that have been created. As well as being used as a unique index, it is useful for statistics, particularly in conjunction with the current size of the frontier.

```

searchGeneric.py — (continued)
60 import heapq          # part of the Python standard library
61 from searchProblem import Path
62
63 class FrontierPQ(object):

```

```

64     """A frontier consists of a priority queue (heap), frontierpq, of
65         (value, index, path) triples, where
66         * value is the value we want to minimize (e.g., path cost + h).
67         * index is a unique index for each element
68         * path is the path on the queue
69         Note that the priority queue always returns the smallest element.
70     """
71
72     def __init__(self):
73         """constructs the frontier, initially an empty priority queue
74         """
75         self.frontier_index = 0 # the number of items ever added to the frontier
76         self.frontierpq = [] # the frontier priority queue
77
78     def empty(self):
79         """is True if the priority queue is empty"""
80         return self.frontierpq == []
81
82     def add(self, path, value):
83         """add a path to the priority queue
84         value is the value to be minimized"""
85         self.frontier_index += 1 # get a new unique index
86         heapq.heappush(self.frontierpq, (value, -self.frontier_index, path))
87
88     def pop(self):
89         """returns and removes the path of the frontier with minimum value.
90         """
91         (_,_,path) = heapq.heappop(self.frontierpq)
92         return path

```

The following methods are used for finding and printing information about the frontier.

```

searchGeneric.py — (continued)
94     def count(self, val):
95         """returns the number of elements of the frontier with value=val"""
96         return sum(1 for e in self.frontierpq if e[0]==val)
97
98     def __repr__(self):
99         """string representation of the frontier"""
100        return str([(n,c,str(p)) for (n,c,p) in self.frontierpq])
101
102    def __len__(self):
103        """length of the frontier"""
104        return len(self.frontierpq)
105
106    def __iter__(self):
107        """iterate through the paths in the frontier"""
108        for (_,_,path) in self.frontierpq:
109            yield path

```

3.2.3 A* Search

For an A* **Search** the frontier is implemented using the FrontierPQ class.

```

searchGeneric.py — (continued)
111 class AStarSearcher(Searcher):
112     """returns a searcher for a problem.
113     Paths can be found by repeatedly calling search().
114     """
115
116     def __init__(self, problem):
117         super().__init__(problem)
118
119     def initialize_frontier(self):
120         self.frontier = FrontierPQ()
121
122     def empty_frontier(self):
123         return self.frontier.empty()
124
125     def add_to_frontier(self, path):
126         """add path to the frontier with the appropriate cost"""
127         value = path.cost + self.problem.heuristic(path.end())
128         self.frontier.add(path, value)

```

Code should always be tested. The following provides a simple **unit test**, using problem1 as the the default problem.

```

searchGeneric.py — (continued)
130 import searchProblem as searchProblem
131
132 def test(SearchClass, problem=searchProblem.problem1, solutions=[['g','d','b','c','a']]):
133     """Unit test for aipython searching algorithms.
134     SearchClass is a class that takes a problem and implements search()
135     problem is a search problem
136     solutions is a list of optimal solutions
137     """
138     print("Testing problem 1:")
139     schr1 = SearchClass(problem)
140     path1 = schr1.search()
141     print("Path found:", path1)
142     assert path1 is not None, "No path is found in problem1"
143     assert list(path1.nodes()) in solutions, "Shortest path not found in problem1"
144     print("Passed unit test")
145
146 if __name__ == "__main__":
147     #test(Searcher)
148     test(AStarSearcher)
149
150 # example queries:
151 # searcher1 = Searcher(searchProblem.acyclic_delivery_problem) # DFS
152 # searcher1.search() # find first path

```

```

153 # searcher1.search() # find next path
154 # searcher2 = AStarSearcher(searchProblem.acyclic_delivery_problem) # A*
155 # searcher2.search() # find first path
156 # searcher2.search() # find next path
157 # searcher3 = Searcher(searchProblem.cyclic_delivery_problem) # DFS
158 # searcher3.search() # find first path with DFS. What do you expect to happen?
159 # searcher4 = AStarSearcher(searchProblem.cyclic_delivery_problem) # A*
160 # searcher4.search() # find first path

```

Exercise 3.2 Change the code so that it implements (i) best-first search and (ii) lowest-cost-first search. For each of these methods compare it to A^* in terms of the number of paths expanded, and the path found.

Exercise 3.3 In the *add* method in *FrontierPQ* what does the “-” in front of *frontier_index* do? When there are multiple paths with the same f -value, which search method does this act like? What happens if the “-” is removed? When there are multiple paths with the same value, which search method does this act like? Does it work better with or without the “-”? What evidence did you base your conclusion on?

Exercise 3.4 The searcher acts like a Python iterator, in that it returns one value (here a path) and then returns other values (paths) on demand, but does not implement the iterator interface. Change the code so it implements the iterator interface. What does this enable us to do?

3.2.4 Multiple Path Pruning

To run the multiple-path pruning demo, in folder “aipython”, load “searchMPP.py”, using e.g., `ipython -i searchMPP.py`, and copy and paste the example queries at the bottom of that file.

The following implements A^* with multiple-path pruning. It overrides *search()* in *Searcher*.

```

searchMPP.py — Searcher with multiple-path pruning
11 from searchGeneric import AStarSearcher, visualize
12 from searchProblem import Path
13
14 class SearcherMPP(AStarSearcher):
15     """returns a searcher for a problem.
16     Paths can be found by repeatedly calling search().
17     """
18     def __init__(self, problem):
19         super().__init__(problem)
20         self.explored = set()
21
22     @visualize
23     def search(self):
24         """returns next path from an element of problem's start nodes
25         to a goal node.
26         Returns None if no path exists.

```

```

27 """
28 while not self.empty_frontier():
29     path = self.frontier.pop()
30     if path.end() not in self.explored:
31         self.display(2, "Expanding:", path, "(cost:", path.cost, ")")
32         self.explored.add(path.end())
33         self.num_expanded += 1
34         if self.problem.is_goal(path.end()):
35             self.display(1, self.num_expanded, "paths have been expanded and",
36                         len(self.frontier), "paths remain in the frontier")
37             self.solution = path # store the solution found
38             return path
39         else:
40             neighs = self.problem.neighbors(path.end())
41             self.display(3, "Neighbors are", neighs)
42             for arc in neighs:
43                 self.add_to_frontier(Path(path, arc))
44             self.display(3, "Frontier:", self.frontier)
45     self.display(1, "No (more) solutions. Total of",
46                 self.num_expanded, "paths expanded.")
47
48 from searchGeneric import test
49 if __name__ == "__main__":
50     test(SearcherMPP)
51
52 import searchProblem
53 # searcherMPPcdp = SearcherMPP(searchProblem.cyclic_delivery_problem)
54 # print(searcherMPPcdp.search()) # find first path

```

Exercise 3.5 Implement a searcher that implements cycle pruning instead of multiple-path pruning. You need to decide whether to check for cycles when paths are added to the frontier or when they are removed. (Hint: either method can be implemented by only changing one or two lines in SearcherMPP. Hint: there is a cycle if `path.end()` in `path.initial_nodes()`) Compare no pruning, multiple path pruning and cycle pruning for the cyclic delivery problem. Which works better in terms of number of paths expanded, computational time or space?

3.3 Branch-and-bound Search

To run the demo, in folder "aipython", load "searchBranchAndBound.py", and copy and paste the example queries at the bottom of that file.

Depth-first search methods do not need an a priority queue, but can use a list as a stack. In this implementation of branch-and-bound search, we call *search* to find an optimal solution with cost less than bound. This uses depth-first search to find a path to a goal that extends *path* with cost less than the

bound. Once a path to a goal has been found, that path is remembered as the *best_path*, the bound is reduced, and the search continues.

```

searchBranchAndBound.py — Branch and Bound Search
11 from searchProblem import Path
12 from searchGeneric import Searcher
13 from display import Displayable, visualize
14
15 class DF_branch_and_bound(Searcher):
16     """returns a branch and bound searcher for a problem.
17     An optimal path with cost less than bound can be found by calling search()
18     """
19     def __init__(self, problem, bound=float("inf")):
20         """creates a searcher than can be used with search() to find an optimal path.
21         bound gives the initial bound. By default this is infinite - meaning there
22         is no initial pruning due to depth bound
23         """
24         super().__init__(problem)
25         self.best_path = None
26         self.bound = bound
27
28     @visualize
29     def search(self):
30         """returns an optimal solution to a problem with cost less than bound.
31         returns None if there is no solution with cost less than bound."""
32         self.frontier = [Path(self.problem.start_node())]
33         self.num_expanded = 0
34         while self.frontier:
35             path = self.frontier.pop()
36             if path.cost+self.problem.heuristic(path.end()) < self.bound:
37                 # if path.end() not in path.initial_nodes(): # for cycle pruning
38                 self.display(3,"Expanding:",path,"cost:",path.cost)
39                 self.num_expanded += 1
40                 if self.problem.is_goal(path.end()):
41                     self.best_path = path
42                     self.bound = path.cost
43                     self.display(2,"New best path:",path," cost:",path.cost)
44                 else:
45                     neighs = self.problem.neighbors(path.end())
46                     self.display(3,"Neighbors are", neighs)
47                     for arc in reversed(list(neighs)):
48                         self.add_to_frontier(Path(path, arc))
49         self.display(1,"Number of paths expanded:",self.num_expanded,
50                     "(optimal" if self.best_path else "(no", "solution found)")
51         self.solution = self.best_path
52         return self.best_path

```

Note that this code used *reversed* in order to expand the neighbors of a node in the left-to-right order one might expect. It does this because *pop()* removes the rightmost element of the list. The call to *list* is there because *reversed* only works on lists and tuples, but the neighbours can be generated.

Here is a unit test and some queries:

```

searchBranchAndBound.py — (continued)
54 from searchGeneric import test
55 if __name__ == "__main__":
56     test(DF_branch_and_bound)
57
58 # Example queries:
59 import searchProblem
60 # searcherb1 = DF_branch_and_bound(searchProblem.acyclic_delivery_problem)
61 # print(searcherb1.search()) # find optimal path
62 # searcherb2 = DF_branch_and_bound(searchProblem.cyclic_delivery_problem, bound=100)
63 # print(searcherb2.search()) # find optimal path

```

Exercise 3.6 Implement a branch-and-bound search uses recursion. Hint: you don't need an explicit frontier, but can do a recursive call for the children.

Exercise 3.7 After the branch-and-bound search found a solution, Sam ran search again, and noticed a different count. Sam hypothesized that this count was related to the number of nodes that an A^* search would use (either expand or be added to the frontier). Or maybe, Sam thought, the count for a number of nodes when the bound is slightly above the optimal path case is related to how A^* would work. Is there relationship between these counts? Are there different things that it could count so they are related? Try to find the most specific statement that is true, and explain why it is true.

To test the hypothesis, Sam wrote the following code, but isn't sure it is helpful:

```

searchTest.py — code that may be useful to compare A* and branch-and-bound
11 from searchGeneric import Searcher, AStarSearcher
12 from searchBranchAndBound import DF_branch_and_bound
13 from searchMPP import SearcherMPP
14
15 DF_branch_and_bound.max_display_level = 1
16 Searcher.max_display_level = 1
17
18 def run(problem,name):
19     print("\n\n*****",name)
20
21     print("\nA*:")
22     asearcher = AStarSearcher(problem)
23     print("Path found:",asearcher.search()," cost=",asearcher.solution.cost)
24     print("there are",asearcher.frontier.count(asearcher.solution.cost),
25           "elements remaining on the queue with f-value=",asearcher.solution.cost)
26
27     print("\nA* with MPP:"),
28     msearcher = SearcherMPP(problem)
29     print("Path found:",msearcher.search()," cost=",msearcher.solution.cost)
30     print("there are",msearcher.frontier.count(msearcher.solution.cost),
31           "elements remaining on the queue with f-value=",msearcher.solution.cost)
32

```



```

33     bound = asearcher.solution.cost+0.01
34     print("\nBranch and bound (with too-good initial bound of", bound,")")
35     tbb = DF_branch_and_bound(problem,bound) # cheating!!!
36     print("Path found:",tbb.search()," cost=",tbb.solution.cost)
37     print("Rerunning B&B")
38     print("Path found:",tbb.search())
39
40     bbound = asearcher.solution.cost*2+10
41     print("\nBranch and bound (with not-very-good initial bound of", bbound, ")")
42     tbb2 = DF_branch_and_bound(problem,bbound) # cheating!!!
43     print("Path found:",tbb2.search()," cost=",tbb2.solution.cost)
44     print("Rerunning B&B")
45     print("Path found:",tbb2.search())
46
47     print("\nDepth-first search: (Use ^C if it goes on forever)")
48     tsearcher = Searcher(problem)
49     print("Path found:",tsearcher.search()," cost=",tsearcher.solution.cost)
50
51
52 import searchProblem
53 from searchTest import run
54 if __name__ == "__main__":
55     run(searchProblem.problem1,"Problem 1")
56     # run(searchProblem.acyclic_delivery_problem,"Acyclic Delivery")
57     # run(searchProblem.cyclic_delivery_problem,"Cyclic Delivery")
58 # also test some graphs with cycles, and some with multiple least-cost paths

```


Reasoning with Constraints

4.1 Constraint Satisfaction Problems

4.1.1 Constraints

A **variable** is a string or any value that is printable and can be the key of a Python dictionary.

A **constraint** consists of a list (or tuple) of variables and a condition.

- The tuple (or list) of variables is called the **scope**.
- The **condition** is a Boolean function that takes the same number of arguments as there are variables in the scope. The condition must have a `__name__` property that gives a printable name of the function; built-in functions and functions that are defined using *def* have such a property; for other functions you may need to define this property.

```
cspProblem.py — Representations of a Constraint Satisfaction Problem
11 class Constraint(object):
12     """A Constraint consists of
13     * scope: a tuple of variables
14     * condition: a function that can applied to a tuple of values
15     * string: a string for printing the constraints. All of the strings must be unique.
16     for the variables
17     """
18     def __init__(self, scope, condition, string=None):
19         self.scope = scope
20         self.condition = condition
21         if string is None:
22             self.string = self.condition.__name__ + str(self.scope)
23         else:
```

```

24         self.string = string
25
26     def __repr__(self):
27         return self.string

```

An **assignment** is a *variable:value* dictionary.

If *con* is a constraint, *con.holds(assignment)* returns True or False depending on whether the condition is true or false for that assignment. The assignment *assignment* must assigns a value to every variable in the scope of the constraint *con* (and could also assign values other variables); *con.holds* gives an error if not all variables in the scope of *con* are assigned in the assignment. It ignores variables in *assignment* that are not in the scope of the constraint.

In Python, the *** notation is used for unpacking a tuple. For example, *F(*(1,2,3))* is the same as *F(1,2,3)*. So if *t* has value (1,2,3), then *F(*t)* is the same as *F(1,2,3)*.

```

_____cspProblem.py — (continued)_____
29     def holds(self, assignment):
30         """returns the value of Constraint con evaluated in assignment.
31
32         precondition: all variables are assigned in assignment
33         """
34         return self.condition(*tuple(assignment[v] for v in self.scope))

```

4.1.2 CSPs

A constraint satisfaction problem (CSP) requires:

- *domains*: a dictionary that maps variables to the set of possible values. Thus *domains[var]* is the domain of variable *var*.
- *constraints*: a set or list of constraints.

Other properties are inferred from these:

- *variables* is the set of variables. The variables can be enumerated by using “for var in domains” because iterating over a dictionary gives the keys, which in this case are the variables.
- *var_to_const* is a mapping from variables to set of constraints, such that *var_to_const[var]* is the set of constraints with *var* in the scope.

```

_____cspProblem.py — (continued)_____
36 class CSP(object):
37     """A CSP consists of
38     * domains, a dictionary that maps each variable to its domain
39     * constraints, a list of constraints
40     * variables, a set of variables

```

```

41 * var_to_const, a variable to set of constraints dictionary
42 """
43 def __init__(self, domains, constraints, positions={}):
44     """domains is a variable:domain dictionary
45     constraints is a list of constraints
46     """
47     self.variables = set(domains)
48     self.domains = domains
49     self.constraints = constraints
50     self.positions = positions
51     self.var_to_const = {var:set() for var in self.variables}
52     for con in constraints:
53         for var in con.scope:
54             self.var_to_const[var].add(con)
55
56 def __str__(self):
57     """string representation of CSP"""
58     return str(self.domains)
59
60 def __repr__(self):
61     """more detailed string representation of CSP"""
62     return "CSP("+str(self.domains)+", "+str([str(c) for c in self.constraints])+")"

```

`csp.consistent(assignment)` returns true if the assignment is consistent with each of the constraints in `csp` (i.e., all of the constraints that can be evaluated evaluate to true). Note that this is a local consistency with each constraint; it does *not* imply the CSP is consistent or has a solution.

```

cspProblem.py — (continued)
64 def consistent(self, assignment):
65     """assignment is a variable:value dictionary
66     returns True if all of the constraints that can be evaluated
67     evaluate to True given assignment.
68     """
69     return all(con.holds(assignment)
70               for con in self.constraints
71               if all(v in assignment for v in con.scope))

```

4.1.3 Examples

In the following code `ne_`, when given a number, returns a function that is true when its argument is not that number. For example, if $f = ne_ (3)$, then $f(2)$ is True and $f(3)$ is False. That is, $ne_ (x)(y)$ is true when $x \neq y$. Allowing a function of multiple arguments to use its arguments one at a time is called **currying**, after the logician Haskell Curry. Functions used as conditions in constraints require names (so they can be printed).

```

cspExamples.py — Example CSPs
11 from cspProblem import CSP, Constraint

```

```

12 from operator import lt,ne,eq,gt
13
14 def ne_(val):
15     """not equal value"""
16     # nev = lambda x: x != val # alternative definition
17     # nev = partial(neq,val) # another alternative definition
18     def nev(x):
19         return val != x
20     nev.__name__ = str(val)+"!=" # name of the function
21     return nev

```

Similarly $is_x(y)$ is true when $x = y$.

```

cspExamples.py — (continued)
23 def is_(val):
24     """is a value"""
25     # isv = lambda x: x == val # alternative definition
26     # isv = partial(eq,val) # another alternative definition
27     def isv(x):
28         return val == x
29     isv.__name__ = str(val)+"=="
30     return isv

```

The CSP, *csp0* has variables X , Y and Z , each with domain $\{1,2,3\}$. The constraints are $X < Y$ and $Y < Z$.

```

cspExamples.py — (continued)
32 csp0 = CSP({'X':{1,2,3}, 'Y':{1,2,3}, 'Z':{1,2,3}},
33           [ Constraint(['X','Y'],lt),
34             Constraint(['Y','Z'],lt)])

```


The CSP, *csp1* has variables A , B and C , each with domain $\{1,2,3,4\}$. The constraints are $A < B$, $B \neq 2$ and $B < C$. This is slightly more interesting than *csp0* as it has more solutions. This example is used in the unit tests, and so if it is changed, the unit tests need to be changed.

```

cspExamples.py — (continued)
36 C0 = Constraint(['A','B'], lt, "A < B")
37 C1 = Constraint(['B'], ne_(2), "B != 2")
38 C2 = Constraint(['B','C'], lt, "B < C")
39 csp1 = CSP({'A':{1,2,3,4}, 'B':{1,2,3,4}, 'C':{1,2,3,4}},
40           [C0, C1, C2],
41           positions={"A": (1, 0),
42                     "B": (3, 0),
43                     "C": (5, 0),
44                     "A < B": (2, 1),
45                     "B < C": (4, 1),
46                     "B != 2": (3, 2)})

```

The next CSP, *csp2* is Example 4.9 of the textbook; the domain consistent network (after applying the unary constraints) is shown in Figure 4.1.

Figure 4.1: Domain-consistent constraint network (*csp2*).

cspExamples.py — (continued)

```

48 csp2 = CSP({'A':{1,2,3,4}, 'B':{1,2,3,4}, 'C':{1,2,3,4},
49           'D':{1,2,3,4}, 'E':{1,2,3,4}},
50           [ Constraint(['B'], ne_(3), "B != 3"),
51             Constraint(['C'], ne_(2), "C != 2"),
52             Constraint(['A', 'B'], ne, "A != B"),
53             Constraint(['B', 'C'], ne, "A != C"),
54             Constraint(['C', 'D'], lt, "C < D"),
55             Constraint(['A', 'D'], eq, "A = D"),
56             Constraint(['A', 'E'], gt, "A > E"),
57             Constraint(['B', 'E'], gt, "B > E"),
58             Constraint(['C', 'E'], gt, "C > E"),
59             Constraint(['D', 'E'], gt, "D > E"),
60             Constraint(['B', 'D'], ne, "B != D")])

```


The following example is another scheduling problem (but with multiple answers). This is the same a scheduling 2 in the original AIspace.org consistency app.

cspExamples.py — (continued)

```

62 csp3 = CSP({'A':{1,2,3,4}, 'B':{1,2,3,4}, 'C':{1,2,3,4},
63           'D':{1,2,3,4}, 'E':{1,2,3,4}},
64           [Constraint(['A', 'B'], ne, "A != B"),
65             Constraint(['A', 'D'], lt, "A < D"),

```

**Words:**

ant, big, bus, car, has,
book, buys, hold, lane,
year, ginger, search,
symbol, syntax.

Figure 4.2: A crossword puzzle to be solved

```

66     Constraint(['A','E'], lambda a,e: (a-e)%2 == 1, "A-E is odd"), # A-E is odd
67     Constraint(['B','E'], lt, "B < E"),
68     Constraint(['D','C'], lt, "D < C"),
69     Constraint(['C','E'], ne, "C != E"),
70     Constraint(['D','E'], ne, "D != E"))

```

The following example is another abstract scheduling problem. What are the solutions?

```

cspExamples.py — (continued)
72 def adjacent(x,y):
73     """True when x and y are adjacent numbers"""
74     return abs(x-y) == 1
75
76 csp4 = CSP({'A':{1,2,3,4,5}, 'B':{1,2,3,4,5}, 'C':{1,2,3,4,5},
77           'D':{1,2,3,4,5}, 'E':{1,2,3,4,5}},
78           [Constraint(['A','B'], adjacent, "adjacent(A,B)"),
79             Constraint(['B','C'], adjacent, "adjacent(B,C)"),
80             Constraint(['C','D'], adjacent, "adjacent(C,D)"),
81             Constraint(['D','E'], adjacent, "adjacent(D,E)"),
82             Constraint(['A','C'], ne, "A != C"),
83             Constraint(['B','D'], ne, "A != D"),
84             Constraint(['C','E'], ne, "C != E")])

```

The following examples represent the crossword shown in Figure 4.2.

In the first representation, the variables represent words. The constraint imposed by the crossword is that where two words intersect, the letter at the intersection must be the same. The method `meet_at` is used to test whether two words intersect with the same letter. For example, the constraint `meet_at(2,0)` means that the third letter (at position 2) of the first argument is the same as the first letter of the second argument.

cspExamples.py — (continued)


```

86 def meet_at(p1,p2):
87     """returns a function of two words that is true when the words intersect at postions p1, p2.
88     The positions are relative to the words; starting at position 0.
89     meet_at(p1,p2)(w1,w2) is true if the same letter is at position p1 of word w1
90     and at position p2 of word w2.
91     """
92     def meets(w1,w2):
93         return w1[p1] == w2[p2]
94     meets.__name__ = "meet_at("+str(p1)+','+str(p2)+')'
95     return meets
96
97 crossword1 = CSP({'one_across':{'ant', 'big', 'bus', 'car', 'has'},
98                 'one_down':{'book', 'buys', 'hold', 'lane', 'year'},
99                 'two_down':{'ginger', 'search', 'symbol', 'syntax'},
100                 'three_across':{'book', 'buys', 'hold', 'land', 'year'},
101                 'four_across':{'ant', 'big', 'bus', 'car', 'has'}},
102                 [Constraint(['one_across','one_down'], meet_at(0,0)),
103                 Constraint(['one_across','two_down'], meet_at(2,0)),
104                 Constraint(['three_across','two_down'], meet_at(2,2)),
105                 Constraint(['three_across','one_down'], meet_at(0,2)),
106                 Constraint(['four_across','two_down'], meet_at(0,4))])

```

In an alternative representation of a crossword (the “dual” representation), the variables represent letters, and the constraints are that adjacent sequences of letters form words.

```

cspExamples.py — (continued)
108 words = {'ant', 'big', 'bus', 'car', 'has', 'book', 'buys', 'hold',
109          'lane', 'year', 'ginger', 'search', 'symbol', 'syntax'}
110
111 def is_word(*letters, words=words):
112     """is true if the letters concatenated form a word in words"""
113     return "".join(letters) in words
114
115 letters = ["a", "b", "c", "d", "e", "f", "g", "h", "i", "j", "k", "l",
116           "m", "n", "o", "p", "q", "r", "s", "t", "u", "v", "w", "x", "y",
117           "z"]
118
119 crossword1d = CSP({'p00':letters, 'p10':letters, 'p20':letters, # first row
120                  'p01':letters, 'p21':letters, # second row
121                  'p02':letters, 'p12':letters, 'p22':letters, 'p32':letters, # third row
122                  'p03':letters, 'p23':letters, #fourth row
123                  'p24':letters, 'p34':letters, 'p44':letters, # fifth row
124                  'p25':letters # sixth row
125                  },
126                  [Constraint(['p00', 'p10', 'p20'], is_word), #1-across
127                  Constraint(['p00', 'p01', 'p02', 'p03'], is_word), # 1-down
128                  Constraint(['p02', 'p12', 'p22', 'p32'], is_word), # 3-across
129                  Constraint(['p20', 'p21', 'p22', 'p23', 'p24', 'p25'], is_word), # 2-down
130                  Constraint(['p24', 'p34', 'p44'], is_word) # 4-across
131                  ])

```

The queens problem is a puzzle on a chess board, where the idea is to place a queen on each column so the queens cannot take each other! there are no two queens on the same row, column or diagonal. The **n-queens problem** is a generalization where the size of the board is an $n \times n$, and n queens have to be placed.

Here is a representation of the n-queens problem, where the variables are the columns and the values are the rows in which the queen is placed. The original queens problem on a standard (8×8) chess board is `n_queens(8)`

```

_____cspExamples.py — (continued)_____
133 def queens(ri,rj):
134     """ri and rj are different rows, return the condition that the queens cannot take each other"""
135     def no_take(ci,cj):
136         """is true if queen at (ri,ci) cannot take a queen at (rj,cj)"""
137         return ci != cj and abs(ri-ci) != abs(rj-cj)
138     return no_take
139
140 def n_queens(n):
141     """returns a CSP for n-queens"""
142     columns = list(range(n))
143     return CSP(
144         {'R'+str(i):columns for i in range(n)},
145         [Constraint(['R'+str(i),'R'+str(j)], queens(i,j)) for i in range(n) for j in range(n) if i != j])
146
147 # try the CSP n_queens(8) in one of the solvers.
148 # What is the smallest n for which there is a solution?

```

Exercise 4.1 How many constraints does this representation of the n-queens problem produce? Can it be done with fewer constraints? Either explain why it can't be done with fewer constraints, or give a solution using fewer constraints.

Unit tests

The following defines a **unit test** for solvers, by default using example `csp1`.

```

_____cspExamples.py — (continued)_____
149 def test(CSP_solver, csp=csp1,
150         solutions=[{'A': 1, 'B': 3, 'C': 4}, {'A': 2, 'B': 3, 'C': 4}]):
151     """CSP_solver is a solver that takes a csp and returns a solution
152     csp is a constraint satisfaction problem
153     solutions is the list of all solutions to csp
154     This tests whether the solution returned by CSP_solver is a solution.
155     """
156     print("Testing csp with",CSP_solver.__doc__)
157     sol0 = CSP_solver(csp)
158     print("Solution found:",sol0)
159     assert sol0 in solutions, "Solution not correct for "+str(csp)
160     print("Passed unit test")

```

Exercise 4.2 Modify *test* so that instead of taking in a list of solutions, it checks whether the returned solution actually is a solution.

Exercise 4.3 Propose a test that is appropriate for CSPs with no solutions. Assume that the test designer knows there are no solutions. Consider what a CSP solver should return if there are no solutions to the CSP.

Exercise 4.4 Write a unit test that checks whether all solutions (e.g., for the search algorithms that can return multiple solutions) are correct, and whether all solutions can be found.

4.2 Solving a CSP using Search

To run the demo, in folder "aipython", load "cspSearch.py", and copy and paste the example queries at the bottom of that file.

The first solver searches through the space of partial assignments. This takes in a CSP problem and an optional variable ordering, which is a list of the variables in the CSP. It then constructs a search space that can be solved using the search methods of the previous chapter. In this search space:

- A node is a *variable : value* dictionary which does not violate any constraints (so that dictionaries that violate any constraints are not added).
- An arc corresponds to an assignment of a value to the next variable. This assumes a static ordering; the next variable chosen to split does not depend on the context. If no variable ordering is given, this makes no attempt to choose a good ordering.

```

cspSearch.py — Representations of a Search Problem from a CSP.
11 from cspProblem import CSP, Constraint
12 from searchProblem import Arc, Search_problem
13 from utilities import dict_union
14
15 class Search_from_CSP(Search_problem):
16     """A search problem directly from the CSP.
17
18     A node is a variable:value dictionary"""
19     def __init__(self, csp, variable_order=None):
20         self.csp=csp
21         if variable_order:
22             assert set(variable_order) == set(csp.variables)
23             assert len(variable_order) == len(csp.variables)
24             self.variables = variable_order
25         else:
26             self.variables = list(csp.variables)
27
28     def is_goal(self, node):

```

```

29     """returns whether the current node is a goal for the search
30     """
31     return len(node)==len(self.csp.variables)
32
33     def start_node(self):
34         """returns the start node for the search
35         """
36         return {}

```

The *neighbors*(*node*) method uses the fact that the length of the node, which is the number of variables already assigned, is the index of the next variable to split on. Note that we do not need to check whether there are no more variables to split on, as the nodes are all consistent, by construction, and so when there are no more variables we have a solution, and so don't need the neighbours.

```

cspSearch.py — (continued)
38     def neighbors(self, node):
39         """returns a list of the neighboring nodes of node.
40         """
41         var = self.variables[len(node)] # the next variable
42         res = []
43         for val in self.csp.domains[var]:
44             new_env = dict_union(node,{var:val}) #dictionary union
45             if self.csp.consistent(new_env):
46                 res.append(Arc(node,new_env))
47         return res

```

The unit tests relies on a solver. The following procedure creates a solver using search that can be tested.

```

cspSearch.py — (continued)
49 from cspExamples import csp1,csp2,test, crossword1, crossword1d
50 from searchGeneric import Searcher
51
52 def dfs_solver(csp):
53     """depth-first search solver"""
54     path = Searcher(Search_from_CSP(csp)).search()
55     if path is not None:
56         return path.end()
57     else:
58         return None
59
60 if __name__ == "__main__":
61     test(dfs_solver)
62
63 ## Test Solving CSPs with Search:
64 searcher1 = Searcher(Search_from_CSP(csp1))
65 #print(searcher1.search()) # get next solution
66 searcher2 = Searcher(Search_from_CSP(csp2))
67 #print(searcher2.search()) # get next solution
68 searcher3 = Searcher(Search_from_CSP(crossword1))

```

```

69 | #print(searcher3.search()) # get next solution
70 | searcher4 = Searcher(Search_from_CSP(crossword1d))
71 | #print(searcher4.search()) # get next solution (warning: slow)

```

Exercise 4.5 What would happen if we constructed the new assignment by assigning `node[var] = val` (with side effects) instead of using dictionary union? Give an example of where this could give a wrong answer. How could the algorithm be changed to work with side effects? (Hint: think about what information needs to be in a node).

Exercise 4.6 Change neighbors so that it returns an iterator of values rather than a list. (Hint: use *yield*.)

4.3 Consistency Algorithms

To run the demo, in folder "aipython", load "cspConsistency.py", and copy and paste the commented-out example queries at the bottom of that file.

A *Con_solver* is used to simplify a CSP using arc consistency.

```

_____cspConsistency.py — Arc Consistency and Domain splitting for solving a CSP_____
11 | from display import Displayable
12 |
13 | class Con_solver(Displayable):
14 |     """Solves a CSP with arc consistency and domain splitting
15 |     """
16 |     def __init__(self, csp, **kwargs):
17 |         """a CSP solver that uses arc consistency
18 |         * csp is the CSP to be solved
19 |         * kwargs is the keyword arguments for Displayable superclass
20 |         """
21 |         self.csp = csp
22 |         super().__init__(**kwargs) # Or Displayable.__init__(self,**kwargs)

```

The following implementation of arc consistency maintains the set *to_do* of (variable, constraint) pairs that are to be checked. It takes in a domain dictionary and returns a new domain dictionary. It needs to be careful to avoid side effects (by copying the *domains* dictionary and the *to_do* set).

```

_____cspConsistency.py — (continued) _____
24 | def make_arc_consistent(self, orig_domains=None, to_do=None):
25 |     """Makes this CSP arc-consistent using generalized arc consistency
26 |     orig_domains is the original domains
27 |     to_do is a set of (variable,constraint) pairs
28 |     returns the reduced domains (an arc-consistent variable:domain dictionary)
29 |     """
30 |     if orig_domains is None:
31 |         orig_domains = self.csp.domains

```

```

32     if to_do is None:
33         to_do = {(var, const) for const in self.csp.constraints
34                 for var in const.scope}
35     else:
36         to_do = to_do.copy() # use a copy of to_do
37     domains = orig_domains.copy()
38     self.display(2, "Performing AC with domains", domains)
39     while to_do:
40         var, const = self.select_arc(to_do)
41         self.display(3, "Processing arc (" + var + ", " + const + ")")
42         other_vars = [ov for ov in const.scope if ov != var]
43         new_domain = {val for val in domains[var]
44                       if self.any_holds(domains, const, {var: val}, other_vars)}
45         if new_domain != domains[var]:
46             self.display(4, "Arc: (" + var + ", " + const + ") is inconsistent")
47             self.display(3, "Domain pruned", "dom(" + var + ") =", new_domain,
48                           " due to " + const)
49             domains[var] = new_domain
50             add_to_do = self.new_to_do(var, const) - to_do
51             to_do |= add_to_do # set union
52             self.display(3, " adding", add_to_do if add_to_do else "nothing", "to to_do.")
53             self.display(4, "Arc: (" + var + ", " + const + ") now consistent")
54         self.display(2, "AC done. Reduced domains", domains)
55     return domains
56
57 def new_to_do(self, var, const):
58     """returns new elements to be added to to_do after assigning
59     variable var in constraint const.
60     """
61     return {(nvar, nconst) for nconst in self.csp.var_to_const[var]
62           if nconst != const
63           for nvar in nconst.scope
64           if nvar != var}

```

The following selects an arc. Any element of *to_do* can be selected. The selected element needs to be removed from *to_do*. The default implementation just selects which ever element *pop* method for sets returns. A user interface could allow the user to select an arc. Alternatively a more sophisticated selection could be employed (or just a stack or a queue).

```

_____cspConsistency.py — (continued) _____
66 def select_arc(self, to_do):
67     """Selects the arc to be taken from to_do .
68     * to_do is a set of arcs, where an arc is a (variable,constraint) pair
69     the element selected must be removed from to_do.
70     """
71     return to_do.pop()

```

The value of *new_domain* is the subset of the domain of *var* that is consistent with the assignment to the other variables. It might be easier to understand the following code, which treats unary (with no other variables in the constraint)

and binary (with one other variables in the constraint) constraints as special cases (this can replace the assignment to `new_domain` in the above code):

```

    if len(other_vars)==0:          # unary constraint
        new_domain = {val for val in domains[var]
                        if const.holds({var:val})}
    elif len(other_vars)==1:        # binary constraint
        other = other_vars[0]
        new_domain = {val for val in domains[var]
                        if any(const.holds({var: val, other: other_val})
                               for other_val in domains[other])}
    else:                            # general case
        new_domain = {val for val in domains[var]
                        if self.any_holds(domains, const, {var: val}, other_vars)}

```

any_holds is a recursive function that tries to find an assignment of values to the other variables (*other_vars*) that satisfies constraint *const* given the assignment in *env*. The integer variable *ind* specifies which index to *other_vars* needs to be checked next. As soon as one assignment returns *True*, the algorithm returns *True*. Note that it has side effects with respect to *env*; it changes the values of the variables in *other_vars*. It should only be called when the side effects have no ill effects.

```

_____cspConsistency.py — (continued)_____
73 def any_holds(self, domains, const, env, other_vars, ind=0):
74     """returns True if Constraint const holds for an assignment
75     that extends env with the variables in other_vars[ind:]
76     env is a dictionary
77     Warning: this has side effects and changes the elements of env
78     """
79     if ind == len(other_vars):
80         return const.holds(env)
81     else:
82         var = other_vars[ind]
83         for val in domains[var]:
84             # env = dict_union(env, {var:val}) # no side effects!
85             env[var] = val
86             if self.any_holds(domains, const, env, other_vars, ind + 1):
87                 return True
88     return False

```

4.3.1 Direct Implementation of Domain Splitting

The following is a direct implementation of domain splitting with arc consistency that uses recursion. It finds one solution if one exists or returns *False* if there are no solutions.

```
_____cspConsistency.py — (continued)_____
```

```

90 def solve_one(self, domains=None, to_do=None):
91     """return a solution to the current CSP or False if there are no solutions
92     to_do is the list of arcs to check
93     """
94     if domains is None:
95         domains = self.csp.domains
96     new_domains = self.make_arc_consistent(domains, to_do)
97     if any(len(new_domains[var]) == 0 for var in domains):
98         return False
99     elif all(len(new_domains[var]) == 1 for var in domains):
100         self.display(2, "solution:", {var: select(
101             new_domains[var] for var in new_domains})
102         return {var: select(new_domains[var] for var in domains}
103     else:
104         var = self.select_var(x for x in self.csp.variables if len(new_domains[x]) > 1)
105         if var:
106             dom1, dom2 = partition_domain(new_domains[var])
107             self.display(3, "...splitting", var, "into", dom1, "and", dom2)
108             new_doms1 = copy_with_assign(new_domains, var, dom1)
109             new_doms2 = copy_with_assign(new_domains, var, dom2)
110             to_do = self.new_to_do(var, None)
111             self.display(3, " adding", to_do if to_do else "nothing", "to to_do.")
112             return self.solve_one(new_doms1, to_do) or self.solve_one(new_doms2, to_do)
113
114 def select_var(self, iter_vars):
115     """return the next variable to split"""
116     return select(iter_vars)
117
118 def partition_domain(dom):
119     """partitions domain dom into two.
120     """
121     split = len(dom) // 2
122     dom1 = set(list(dom)[:split])
123     dom2 = dom - dom1
124     return dom1, dom2

```

The domains are implemented as a dictionary that maps each variables to its domain. Assigning a value in Python has side effects which we want to avoid. *copy_with_assign* takes a copy of the domains dictionary, perhaps allowing for a new domain for a variable. It creates a copy of the CSP with an (optional) assignment of a new domain to a variable. Only the domains are copied.

cspConsistency.py — (continued)

```

126 def copy_with_assign(domains, var=None, new_domain={True, False}):
127     """create a copy of the domains with an assignment var=new_domain
128     if var==None then it is just a copy.
129     """
130     newdoms = domains.copy()
131     if var is not None:
132         newdoms[var] = new_domain

```



```
133 |     return newdoms
```

cspConsistency.py — (continued)

```
135 | def select(iterable):
136 |     """select an element of iterable. Returns None if there is no such element.
137 |
138 |     This implementation just picks the first element.
139 |     For many of the uses, which element is selected does not affect correctness,
140 |     but may affect efficiency.
141 |     """
142 |     for e in iterable:
143 |         return e # returns first element found
```

Exercise 4.7 Implement *solve_all* that is like *solve_one* but returns the set of all solutions.

Exercise 4.8 Implement *solve_enum* that enumerates the solutions. It should use Python's *yield* (and perhaps *yield from*).

Unit test:

cspConsistency.py — (continued)

```
145 | from cspExamples import test
146 | def ac_solver(csp):
147 |     "arc consistency (solve_one)"
148 |     return Con_solver(csp).solve_one()
149 | if __name__ == "__main__":
150 |     test(ac_solver)
```

4.3.2 Domain Splitting as an interface to graph searching

An alternative implementation is to implement domain splitting in terms of the search abstraction of Chapter 3.

A node is domains dictionary.

cspConsistency.py — (continued)

```
152 | from searchProblem import Arc, Search_problem
153 |
154 | class Search_with_AC_from_CSP(Search_problem, Displayable):
155 |     """A search problem with arc consistency and domain splitting
156 |
157 |     A node is a CSP """
158 |     def __init__(self, csp):
159 |         self.cons = Con_solver(csp) #copy of the CSP
160 |         self.domains = self.cons.make_arc_consistent()
161 |
162 |     def is_goal(self, node):
163 |         """node is a goal if all domains have 1 element"""
164 |         return all(len(node[var])==1 for var in node)
165 |
```

```

166     def start_node(self):
167         return self.domains
168
169     def neighbors(self,node):
170         """returns the neighboring nodes of node.
171         """
172         neighs = []
173         var = select(x for x in node if len(node[x])>1)
174         if var:
175             dom1, dom2 = partition_domain(node[var])
176             self.display(2,"Splitting", var, "into", dom1, "and", dom2)
177             to_do = self.cons.new_to_do(var,None)
178             for dom in [dom1,dom2]:
179                 newdoms = copy_with_assign(node,var,dom)
180                 cons_doms = self.cons.make_arc_consistent(newdoms,to_do)
181                 if all(len(cons_doms[v])>0 for v in cons_doms):
182                     # all domains are non-empty
183                     neighs.append(Arc(node,cons_doms))
184             else:
185                 self.display(2,"...",var,"in",dom,"has no solution")
186         return neighs

```

Exercise 4.9 When splitting a domain, this code splits the domain into half, approximately in half (without any effort to make a sensible choice). Does it work better to split one element from a domain?

Unit test:

```

_____cspConsistency.py — (continued)_____
188 from cspExamples import test
189 from searchGeneric import Searcher
190
191 def ac_search_solver(csp):
192     """arc consistency (search interface)"""
193     sol = Searcher(Search_with_AC_from_CSP(csp)).search()
194     if sol:
195         return {v:select(d) for (v,d) in sol.end().items()}
196
197 if __name__ == "__main__":
198     test(ac_search_solver)

```

Testing:

```

_____cspConsistency.py — (continued)_____
200 from cspExamples import csp1, csp2, csp3, csp4, crossword1, crossword1d
201
202 ## Test Solving CSPs with Arc consistency and domain splitting:
203 #Con_solver.max_display_level = 4 # display details of AC (0 turns off)
204 #Con_solver(csp1).solve_one()
205 #searcher1d = Searcher(Search_with_AC_from_CSP(csp1))
206 #print(searcher1d.search())

```

```

207 #Searcher.max_display_level = 2 # display search trace (0 turns off)
208 #searcher2c = Searcher(Search_with_AC_from_CSP(csp2))
209 #print(searcher2c.search())
210 #searcher3c = Searcher(Search_with_AC_from_CSP(crossword1))
211 #print(searcher3c.search())
212 #searcher5c = Searcher(Search_with_AC_from_CSP(crossword1d))
213 #print(searcher5c.search())

```

4.4 Solving CSPs using Stochastic Local Search

To run the demo, in folder "aipython", load "cspSLS.py", and copy and paste the commented-out example queries at the bottom of that file. This assumes Python 3. Some of the queries require matplotlib.

This implements both the two-stage choice, the any-conflict algorithm and a random choice of variable (and a probabilistic mix of the three).

Given a CSP, the stochastic local searcher (*SLSearcher*) creates the data structures:

- *variables_to_select* is the set of all of the variables with domain-size greater than one. For a variable not in this set, we cannot pick another value from that variable.
- *var_to_constraints* maps from a variable into the set of constraints it is involved in. Note that the inverse mapping from constraints into variables is part of the definition of a constraint.

```

_____cspSLS.py — Stochastic Local Search for Solving CSPs_____
11 from cspProblem import CSP, Constraint
12 from searchProblem import Arc, Search_problem
13 from display import Displayable
14 import random
15 import heapq
16
17 class SLSearcher(Displayable):
18     """A search problem directly from the CSP..
19
20     A node is a variable:value dictionary"""
21     def __init__(self, csp):
22         self.csp = csp
23         self.variables_to_select = {var for var in self.csp.variables
24                                     if len(self.csp.domains[var]) > 1}
25         # Create assignment and conflicts set
26         self.current_assignment = None # this will trigger a random restart
27         self.number_of_steps = 0 #number of steps after the initialization

```

restart creates a new total assignment, and constructs the set of conflicts (the constraints that are false in this assignment).

```

cspSLS.py — (continued)
29 def restart(self):
30     """creates a new total assignment and the conflict set
31     """
32     self.current_assignment = {var:random_sample(dom) for
33                               (var,dom) in self.csp.domains.items()}
34     self.display(2,"Initial assignment",self.current_assignment)
35     self.conflicts = set()
36     for con in self.csp.constraints:
37         if not con.holds(self.current_assignment):
38             self.conflicts.add(con)
39     self.display(2,"Number of conflicts",len(self.conflicts))
40     self.variable_pq = None

```

The *search* method is the top-level searching algorithm. It can either be used to start the search or to continue searching. If there is no current assignment, it must create one. Note that, when counting steps, a restart is counted as one step.

This method selects one of two implementations. The argument *prob_best* is the probability of selecting a best variable (one involving the most conflicts). When the value of *prob_best* is positive, the algorithm needs to maintain a priority queue of variables and the number of conflicts (using *search_with_var_pq*). If the probability of selecting a best variable is zero, it does not need to maintain this priority queue (as implemented in *search_with_any_conflict*).

The argument *prob_anycon* is the probability that the any-conflict strategy is used (which selects a variable at random that is in a conflict), assuming that it is not picking a best variable. Note that for the probability parameters, any value less than zero acts like probability zero and any value greater than 1 acts like probability 1. This means that when *prob_anycon* = 1.0, a best variable is chosen with probability *prob_best*, otherwise a variable in any conflict is chosen. A variable is chosen at random with probability $1 - \text{prob_anycon} - \text{prob_best}$ as long as that is positive.

This returns the number of steps needed to find a solution, or *None* if no solution is found. If there is a solution, it is in *self.current_assignment*.

```

cspSLS.py — (continued)
42 def search(self,max_steps, prob_best=0, prob_anycon=1.0):
43     """
44     returns the number of steps or None if there is no solution.
45     If there is a solution, it can be found in self.current_assignment
46
47     max_steps is the maximum number of steps it will try before giving up
48     prob_best is the probability that a best variable (one in most conflict) is selected
49     prob_anycon is the probability that a variable in any conflict is selected
50     (otherwise a variable is chosen at random)
51     """
52     if self.current_assignment is None:
53         self.restart()

```

```

54         self.number_of_steps += 1
55         if not self.conflicts:
56             self.display(1,"Solution found:", self.current_assignment, "after restart")
57             return self.number_of_steps
58     if prob_best > 0: # we need to maintain a variable priority queue
59         return self.search_with_var_pq(max_steps, prob_best, prob_anycon)
60     else:
61         return self.search_with_any_conflict(max_steps, prob_anycon)

```

Exercise 4.10 This does an initial random assignment but does not do any random restarts. Implement a searcher that takes in the maximum number of walk steps (corresponding to existing *max_steps*) and the maximum number of restarts, and returns the total number of steps for the first solution found. (As in *search*, the solution found can be extracted from the variable *self.current_assignment*).

4.4.1 Any-conflict

If the probability of picking a best variable is zero, the implementation need to keeps track of which variables are in conflicts.

```

cspSLS.py — (continued)
63 def search_with_any_conflict(self, max_steps, prob_anycon=1.0):
64     """Searches with the any_conflict heuristic.
65     This relies on just maintaining the set of conflicts;
66     it does not maintain a priority queue
67     """
68     self.variable_pq = None # we are not maintaining the priority queue.
69                             # This ensures it is regenerated if
70                             # we call search_with_var_pq.
71     for i in range(max_steps):
72         self.number_of_steps +=1
73         if random.random() < prob_anycon:
74             con = random_sample(self.conflicts) # pick random conflict
75             var = random_sample(con.scope) # pick variable in conflict
76         else:
77             var = random_sample(self.variables_to_select)
78         if len(self.csp.domains[var]) > 1:
79             val = random_sample(self.csp.domains[var] -
80                               {self.current_assignment[var]})
81         self.display(2,self.number_of_steps,": Assigning",var,"=",val)
82         self.current_assignment[var]=val
83         for varcon in self.csp.var_to_const[var]:
84             if varcon.holds(self.current_assignment):
85                 if varcon in self.conflicts:
86                     self.conflicts.remove(varcon)
87             else:
88                 if varcon not in self.conflicts:
89                     self.conflicts.add(varcon)
90         self.display(2,"  Number of conflicts",len(self.conflicts))
91     if not self.conflicts:

```

```

92         self.display(1,"Solution found:", self.current_assignment,
93                     "in", self.number_of_steps,"steps")
94         return self.number_of_steps
95     self.display(1,"No solution in",self.number_of_steps,"steps",
96                 len(self.conflicts),"conflicts remain")
97     return None

```

Exercise 4.11 This makes no attempt to find the best alternative value for a variable. Modify the code so that after selecting a variable it selects a value that reduces the number of conflicts by the most. Have a parameter that specifies the probability that the best value is chosen.

4.4.2 Two-Stage Choice

This is the top-level searching algorithm that maintains a priority queue of variables ordered by (the negative of) the number of conflicts, so that the variable with the most conflicts is selected first. If there is no current priority queue of variables, one is created.

The main complexity here is to maintain the priority queue. This uses the dictionary *var_differential* which specifies how much the values of variables should change. This is used with the updatable queue (page 72) to find a variable with the most conflicts.

```

cspSLS.py — (continued)
99     def search_with_var_pq(self,max_steps, prob_best=1.0, prob_anycon=1.0):
100         """search with a priority queue of variables.
101         This is used to select a variable with the most conflicts.
102         """
103         if not self.variable_pq:
104             self.create_pq()
105         pick_best_or_con = prob_best + prob_anycon
106         for i in range(max_steps):
107             self.number_of_steps +=1
108             randnum = random.random()
109             ## Pick a variable
110             if randnum < prob_best: # pick best variable
111                 var,oldval = self.variable_pq.top()
112             elif randnum < pick_best_or_con: # pick a variable in a conflict
113                 con = random_sample(self.conflicts)
114                 var = random_sample(con.scope)
115             else: #pick any variable that can be selected
116                 var = random_sample(self.variables_to_select)
117             if len(self.csp.domains[var]) > 1: # var has other values
118                 ## Pick a value
119                 val = random_sample(self.csp.domains[var] -
120                                   {self.current_assignment[var]})
121                 self.display(2,"Assigning",var,val)
122                 ## Update the priority queue
123                 var_differential = {}

```

```

124         self.current_assignment[var]=val
125         for varcon in self.csp.var_to_const[var]:
126             self.display(3,"Checking",varcon)
127             if varcon.holds(self.current_assignment):
128                 if varcon in self.conflicts: #was incons, now consis
129                     self.display(3,"Became consistent",varcon)
130                     self.conflicts.remove(varcon)
131                     for v in varcon.scope: # v is in one fewer conflicts
132                         var_differential[v] = var_differential.get(v,0)-1
133             else:
134                 if varcon not in self.conflicts: # was consis, not now
135                     self.display(3,"Became inconsistent",varcon)
136                     self.conflicts.add(varcon)
137                     for v in varcon.scope: # v is in one more conflicts
138                         var_differential[v] = var_differential.get(v,0)+1
139             self.variable_pq.update_each_priority(var_differential)
140             self.display(2,"Number of conflicts",len(self.conflicts))
141             if not self.conflicts: # no conflicts, so solution found
142                 self.display(1,"Solution found:", self.current_assignment,"in",
143                             self.number_of_steps,"steps")
144                 return self.number_of_steps
145             self.display(1,"No solution in",self.number_of_steps,"steps",
146                         len(self.conflicts),"conflicts remain")
147         return None

```

create_pq creates an updatable priority queue of the variables, ordered by the number of conflicts they participate in. The priority queue only includes variables in conflicts and the value of a variable is the *negative* of the number of conflicts the variable is in. This ensures that the priority queue, which picks the minimum value, picks a variable with the most conflicts.

```

cspSLS.py — (continued)
149 def create_pq(self):
150     """Create the variable to number-of-conflicts priority queue.
151     This is needed to select the variable in the most conflicts.
152
153     The value of a variable in the priority queue is the negative of the
154     number of conflicts the variable appears in.
155     """
156     self.variable_pq = Updatable_priority_queue()
157     var_to_number_conflicts = {}
158     for con in self.conflicts:
159         for var in con.scope:
160             var_to_number_conflicts[var] = var_to_number_conflicts.get(var,0)+1
161     for var,num in var_to_number_conflicts.items():
162         if num>0:
163             self.variable_pq.add(var,-num)

cspSLS.py — (continued)
165 def random_sample(st):

```

```

166 | """selects a random element from set st"""
167 | return random.sample(st,1)[0]

```

Exercise 4.12 This makes no attempt to find the best alternative value for a variable. Modify the code so that after selecting a variable it selects a value that reduces the number of conflicts by the most. Have a parameter that specifies the probability that the best value is chosen.

Exercise 4.13 These implementations always select a value for the variable selected that is different from its current value (if that is possible). Change the code so that it does not have this restriction (so it can leave the value the same). Would you expect this code to be faster? Does it work worse (or better)?

4.4.3 Updatable Priority Queues

An **updatable priority queue** is a priority queue, where key-value pairs can be stored, and the pair with the smallest key can be found and removed quickly, and where the values can be updated. This implementation follows the idea of <http://docs.python.org/3.5/library/heapq.html>, where the updated elements are marked as removed. This means that the priority queue can be used unmodified. However, this might be expensive if changes are more common than popping (as might happen if the probability of choosing the best is close to zero).

In this implementation, the equal values are sorted randomly. This is achieved by having the elements of the heap being *[val, rand, elt]* triples, where the second element is a random number. Note that Python requires this to be a list, not a tuple, as the tuple cannot be modified.

```

cspSLS.py — (continued)
169 | class Updatable_priority_queue(object):
170 |     """A priority queue where the values can be updated.
171 |     Elements with the same value are ordered randomly.
172 |
173 |     This code is based on the ideas described in
174 |     http://docs.python.org/3.3/library/heapq.html
175 |     It could probably be done more efficiently by
176 |     shuffling the modified element in the heap.
177 |     """
178 |     def __init__(self):
179 |         self.pq = [] # priority queue of [val,rand,elt] triples
180 |         self.elt_map = {} # map from elt to [val,rand,elt] triple in pq
181 |         self.REMOVED = "*removed*" # a string that won't be a legal element
182 |         self.max_size=0
183 |
184 |     def add(self,elt,val):
185 |         """adds elt to the priority queue with priority=val.
186 |         """
187 |         assert val <= 0, val
188 |         assert elt not in self.elt_map, elt

```



```

189         new_triple = [val, random.random(),elt]
190         heapq.heappush(self.pq, new_triple)
191         self.elt_map[elt] = new_triple
192
193     def remove(self,elt):
194         """remove the element from the priority queue"""
195         if elt in self.elt_map:
196             self.elt_map[elt][2] = self.REMOVED
197             del self.elt_map[elt]
198
199     def update_each_priority(self,update_dict):
200         """update values in the priority queue by subtracting the values in
201         update_dict from the priority of those elements in priority queue.
202         """
203         for elt,incr in update_dict.items():
204             if incr != 0:
205                 newval = self.elt_map.get(elt,[0])[0] - incr
206                 assert newval <= 0, str(elt)+": "+str(newval+incr)+"-"+str(incr)
207                 self.remove(elt)
208                 if newval != 0:
209                     self.add(elt,newval)
210
211     def pop(self):
212         """Removes and returns the (elt,value) pair with minimal value.
213         If the priority queue is empty, IndexError is raised.
214         """
215         self.max_size = max(self.max_size, len(self.pq)) # keep statistics
216         triple = heapq.heappop(self.pq)
217         while triple[2] == self.REMOVED:
218             triple = heapq.heappop(self.pq)
219         del self.elt_map[triple[2]]
220         return triple[2], triple[0] # elt, value
221
222     def top(self):
223         """Returns the (elt,value) pair with minimal value, without removing it.
224         If the priority queue is empty, IndexError is raised.
225         """
226         self.max_size = max(self.max_size, len(self.pq)) # keep statistics
227         triple = self.pq[0]
228         while triple[2] == self.REMOVED:
229             heapq.heappop(self.pq)
230             triple = self.pq[0]
231         return triple[2], triple[0] # elt, value
232
233     def empty(self):
234         """returns True iff the priority queue is empty"""
235         return all(triple[2] == self.REMOVED for triple in self.pq)

```

4.4.4 Plotting Runtime Distributions

Runtime_distribution uses matplotlib to plot runtime distributions. Here the runtime is a misnomer as we are only plotting the number of steps, not the time. Computing the runtime is non-trivial as many of the runs have a very short runtime. To compute the time accurately would require running the same code, with the same random seed, multiple times to get a good estimate of the runtime. This is left as an exercise.

```

cspSLS.py — (continued)
237 import matplotlib.pyplot as plt
238
239 class Runtime_distribution(object):
240     def __init__(self, csp, xscale='log'):
241         """Sets up plotting for csp
242         xscale is either 'linear' or 'log'
243         """
244         self.csp = csp
245         plt.ion()
246         plt.xlabel("Number of Steps")
247         plt.ylabel("Cumulative Number of Runs")
248         plt.xscale(xscale) # Makes a 'log' or 'linear' scale
249
250     def plot_runs(self, num_runs=100, max_steps=1000, prob_best=1.0, prob_anycon=1.0):
251         """Plots num_runs of SLS for the given settings.
252         """
253         stats = []
254         SLSearcher.max_display_level, temp_mdl = 0, SLSearcher.max_display_level # no display
255         for i in range(num_runs):
256             searcher = SLSearcher(self.csp)
257             num_steps = searcher.search(max_steps, prob_best, prob_anycon)
258             if num_steps:
259                 stats.append(num_steps)
260         stats.sort()
261         if prob_best >= 1.0:
262             label = "P(best)=1.0"
263         else:
264             p_ac = min(prob_anycon, 1-prob_best)
265             label = "P(best)=%.2f, P(ac)=%.2f" % (prob_best, p_ac)
266         plt.plot(stats, range(len(stats)), label=label)
267         plt.legend(loc="upper left")
268         #plt.draw()
269         SLSearcher.max_display_level= temp_mdl #restore display

```

4.4.5 Testing

```

cspSLS.py — (continued)
271 from cspExamples import test
272 def sls_solver(csp, prob_best=0.7):

```

```

273     """stochastic local searcher (prob_best=0.7)"""
274     se0 = SLSearcher(csp)
275     se0.search(1000,prob_best)
276     return se0.current_assignment
277 def any_conflict_solver(csp):
278     """stochastic local searcher (any-conflict)"""
279     return sls_solver(csp,0)
280
281 if __name__ == "__main__":
282     test(sls_solver)
283     test(any_conflict_solver)
284
285 from cspExamples import csp1, csp2, crossword1
286
287 ## Test Solving CSPs with Search:
288 #se1 = SLSearcher(csp1); print(se1.search(100))
289 #se2 = SLSearcher(csp2); print(se2.search(1000,1.0)) # greedy
290 #se2 = SLSearcher(csp2); print(se2.search(1000,0)) # any_conflict
291 #se2 = SLSearcher(csp2); print(se2.search(1000,0.7)) # 70% greedy; 30% any_conflict
292 #SLSearcher.max_display_level=2 #more detailed display
293 #se3 = SLSearcher(crossword1); print(se3.search(100),0.7)
294 #p = Runtime_distribution(csp2)
295 #p.plot_runs(1000,1000,0) # any_conflict
296 #p.plot_runs(1000,1000,1.0) # greedy
297 #p.plot_runs(1000,1000,0.7) # 70% greedy; 30% any_conflict

```

Exercise 4.14 Modify this to plot the runtime, instead of the number of steps. To measure runtime use *timeit* (<https://docs.python.org/3.5/library/timeit.html>). Small runtimes are inaccurate, so *timeit* can run the same code multiple times. Stochastic local algorithms give different runtimes each time called. To make the timing meaningful, you need to make sure the random seed is the same for each repeated call (see *random.getstate* and *random.setstate* in <https://docs.python.org/3.5/library/random.html>). Because the runtime for different seeds can vary a great deal, for each seed, you should start with 1 iteration and multiplying it by, say 10, until the time is greater than 0.2 seconds. Make sure you plot the average time for each run. Before you start, try to estimate the total runtime, so you will be able to tell if there is a problem with the algorithm stopping.

Propositions and Inference

5.1 Representing Knowledge Bases

A clause consists of a head (an atom) and a body. A body is represented as a list of atoms. Atoms are represented as strings.

```
_____logicProblem.py — Representations Logics _____
11 class Clause(object):
12     """A definite clause"""
13
14     def __init__(self, head, body=[]):
15         """clause with atom head and lost of atoms body"""
16         self.head=head
17         self.body = body
18
19     def __str__(self):
20         """returns the string representation of a clause.
21         """
22         if self.body:
23             return self.head + " <- " + " & ".join(self.body) + "."
24         else:
25             return self.head + "."
```

An askable atom can be asked of the user. The user can respond in English or French or just with a “y”.

```
_____logicProblem.py — (continued) _____
27 class Askable(object):
28     """An askable atom"""
29
30     def __init__(self, atom):
31         """clause with atom head and lost of atoms body"""
```

```

32     self.atom=atom
33
34     def __str__(self):
35         """returns the string representation of a clause."""
36         return "askable " + self.atom + "."
37
38     def yes(ans):
39         """returns true if the answer is yes in some form"""
40         return ans.lower() in ['yes', 'yes.', 'oui', 'oui.', 'y', 'y.'] # bilingual

```

A knowledge base is a list of clauses and askables. In order to make top-down inference faster, this creates a dictionary that maps each atoms into the set of clauses with that atom in the head.

```

42 from display import Displayable
43
44 class KB(Displayable):
45     """A knowledge base consists of a set of clauses.
46     This also creates a dictionary to give fast access to the clauses with an atom in head.
47     """
48     def __init__(self, statements=[]):
49         self.statements = statements
50         self.clauses = [c for c in statements if isinstance(c, Clause)]
51         self.askables = [c.atom for c in statements if isinstance(c, Askable)]
52         self.atom_to_clauses = {} # dictionary giving clauses with atom as head
53         for c in self.clauses:
54             if c.head in self.atom_to_clauses:
55                 self.atom_to_clauses[c.head].add(c)
56             else:
57                 self.atom_to_clauses[c.head] = {c}
58
59     def clauses_for_atom(self,a):
60         """returns set of clauses with atom a as the head"""
61         if a in self.atom_to_clauses:
62             return self.atom_to_clauses[a]
63         else:
64             return set()
65
66     def __str__(self):
67         """returns a string representation of this knowledge base.
68         """
69         return '\n'.join([str(c) for c in self.statements])

```

Here is a trivial example (I think therefore I am) using in the unit tests:

```

71 triv_KB = KB([
72     Clause('i_am', ['i_think']),
73     Clause('i_think'),
74     Clause('i_smell', ['i_exist'])
75 ])

```

Here is a representation of the electrical domain of the textbook:

```

_____logicProblem.py — (continued)_____
77 elect = KB([
78     Clause('light_l1'),
79     Clause('light_l2'),
80     Clause('ok_l1'),
81     Clause('ok_l2'),
82     Clause('ok_cb1'),
83     Clause('ok_cb2'),
84     Clause('live_outside'),
85     Clause('live_l1', ['live_w0']),
86     Clause('live_w0', ['up_s2', 'live_w1']),
87     Clause('live_w0', ['down_s2', 'live_w2']),
88     Clause('live_w1', ['up_s1', 'live_w3']),
89     Clause('live_w2', ['down_s1', 'live_w3' ]),
90     Clause('live_l2', ['live_w4']),
91     Clause('live_w4', ['up_s3', 'live_w3' ]),
92     Clause('live_p_1', ['live_w3']),
93     Clause('live_w3', ['live_w5', 'ok_cb1']),
94     Clause('live_p_2', ['live_w6']),
95     Clause('live_w6', ['live_w5', 'ok_cb2']),
96     Clause('live_w5', ['live_outside']),
97     Clause('lit_l1', ['light_l1', 'live_l1', 'ok_l1']),
98     Clause('lit_l2', ['light_l2', 'live_l2', 'ok_l2']),
99     Askable('up_s1'),
100    Askable('down_s1'),
101    Askable('up_s2'),
102    Askable('down_s2'),
103    Askable('up_s3'),
104    Askable('down_s2')
105 ])
106
107 # print(kb)

```

5.2 Bottom-up Proofs

fixed_point computes the fixed point of the knowledge base *kb*.

```

_____logicBottomUp.py — Bottom-up Proof Procedure for Definite Clauses_____
11 from logicProblem import yes
12
13 def fixed_point(kb):
14     """Returns the fixed point of knowledge base kb.
15     """
16     fp = ask_askables(kb)
17     added = True
18     while added:
19         added = False # added is true when an atom was added to fp this iteration

```

```

20     for c in kb.clauses:
21         if c.head not in fp and all(b in fp for b in c.body):
22             fp.add(c.head)
23             added = True
24             kb.display(2,c.head,"added to fp due to clause",c)
25     return fp
26
27 def ask_askables(kb):
28     return {at for at in kb.askables if yes(input("Is "+at+" true? "))}

```

The following provides a trivial **unit test**, by default using the knowledge base `triv_KB`:

```

_____logicBottomUp.py — (continued)_____
30 from logicProblem import triv_KB
31 def test(kb=triv_KB, fixedpt = {'i_am','i_think'}):
32     fp = fixed_point(kb)
33     assert fp == fixedpt, "kb gave result "+str(fp)
34     print("Passed unit test")
35 if __name__ == "__main__":
36     test()
37
38 from logicProblem import elect
39 # elect.max_display_level=3 # give detailed trace
40 # fixed_point(elect)

```

Exercise 5.1 It is not very user-friendly to ask all of the askables up-front. Implement ask-the-user so that questions are only asked if useful, and are not re-asked. For example, if there is a clause $h \leftarrow a \wedge b \wedge c \wedge d \wedge e$, where c and e are askable, c and e only need to be asked if a, b, d are all in fp and they have not been asked before. Askable e only needs to be asked if the user says “yes” to c . Askable c doesn’t need to be asked if the user previously replied “no” to e .

This form of ask-the-user can ask a different set of questions than the top-down interpreter that asks questions when encountered. Give an example where they ask different questions (neither set of questions asked is a subset of the other).

Exercise 5.2 This algorithm runs in time $O(n^2)$, where n is the number of clauses, for a bounded number of elements in the body; each iteration goes through each of the clauses, and in the worst case, it will do an iteration for each clause. It is possible to implement this in time $O(n)$ time by creating an index that maps an atom to the set of clauses with that atom in the body. Implement this. What is its complexity as a function of n and b , the maximum number of atoms in the body of a clause?

Exercise 5.3 It is possible to be asymptotically more efficient (in terms of the number of elements in a body) than the method in the previous question by noticing that each element of the body of clause only needs to be checked once. For example, the clause $a \leftarrow b \wedge c \wedge d$, needs only be considered when b is added to fp . Once b is added to fp , if c is already in pf , we know that a can be added as soon as d is added. Implement this. What is its complexity as a function of n and b , the maximum number of atoms in the body of a clause?

5.3 Top-down Proofs

prove(kb, goal) is used to prove *goal* from a knowledge base, *kb*, where a *goal* is a list of atoms. It returns *True* if $kb \vdash goal$. The *indent* is used when displaying the code (and doesn't need to have a non-default value).

```

_____logicTopDown.py — Top-down Proof Procedure for Definite Clauses_____
11 from logicProblem import yes
12
13 def prove(kb, ans_body, indent=""):
14     """returns True if kb |- ans_body
15     ans_body is a list of atoms to be proved
16     """
17     kb.display(2, indent, 'yes <-', ' & '.join(ans_body))
18     if ans_body:
19         selected = ans_body[0] # select first atom from ans_body
20         if selected in kb.askables:
21             return (yes(input("Is "+selected+" true? "))
22                     and prove(kb, ans_body[1:], indent+" "))
23         else:
24             return any(prove(kb, cl.body+ans_body[1:], indent+" ")
25                         for cl in kb.clauses_for_atom(selected))
26     else:
27         return True # empty body is true

```

The following provides a simple **unit test** that is hard wired for *triv_KB*:

```

_____logicTopDown.py — (continued)_____
29 from logicProblem import triv_KB
30 def test():
31     a1 = prove(triv_KB, ['i_am'])
32     assert a1, "triv_KB proving i_am gave "+str(a1)
33     a2 = prove(triv_KB, ['i_smell'])
34     assert not a2, "triv_KB proving i_smell gave "+str(a2)
35     print("Passed unit tests")
36 if __name__ == "__main__":
37     test()
38 # try
39 from logicProblem import elect
40 # elect.max_display_level=3 # give detailed trace
41 # prove(elect, ['live_w6'])
42 # prove(elect, ['lit_l1'])

```

Exercise 5.4 This code can re-ask a question multiple times. Implement this code so that it only asks a question once and remembers the answer. Also implement a function to forget the answers.

Exercise 5.5 What search method is this using? Implement the search interface so that it can use A^* or other searching methods. Define an admissible heuristic that is not always 0.

5.4 Assumables

Atom a can be made assumable by including *Assumable*(a) in the knowledge base. A knowledge base that can include assumables is declared with *KBA*.

```

_____logicAssumables.py — Definite clauses with assumables_____
11 from logicProblem import Clause, Askable, KB, yes
12
13 class Assumable(object):
14     """An askable atom"""
15
16     def __init__(self, atom):
17         """clause with atom head and lost of atoms body"""
18         self.atom = atom
19
20     def __str__(self):
21         """returns the string representation of a clause.
22         """
23         return "assumable " + self.atom + "."
24
25 class KBA(KB):
26     """A knowledge base that can include assumables"""
27     def __init__(self, statements):
28         self.assumables = [c.atom for c in statements if isinstance(c, Assumable)]
29         KB.__init__(self, statements)

```

The top-down Horn clause interpreter, *prove_all_ass* returns a list of the sets of assumables that imply *ans_body*. This list will contain all of the minimal sets of assumables, but can also find non-minimal sets, and repeated sets, if they can be generated with separate proofs. The set *assumed* is the set of assumables already assumed.

```

_____logicAssumables.py — (continued)_____
31 def prove_all_ass(self, ans_body, assumed=set()):
32     """returns a list of sets of assumables that extends assumed
33     to imply ans_body from self.
34     ans_body is a list of atoms (it is the body of the answer clause).
35     assumed is a set of assumables already assumed
36     """
37     if ans_body:
38         selected = ans_body[0] # select first atom from ans_body
39         if selected in self.askables:
40             if yes(input("Is "+selected+" true? ")):
41                 return self.prove_all_ass(ans_body[1:], assumed)
42             else:
43                 return [] # no answers
44         elif selected in self.assumables:
45             return self.prove_all_ass(ans_body[1:], assumed|{selected})
46         else:
47             return [ass
48                     for cl in self.clauses_for_atom(selected)

```

```

49         for ass in self.prove_all_ass(cl.body+ans_body[1:],assumed)
50             ] # union of answers for each clause with head=selected
51     else:
52         # empty body
53         return [assumed] # one answer
54
55 def conflicts(self):
56     """returns a list of minimal conflicts"""
57     return minsets(self.prove_all_ass(['false']))

```

Given a list of sets, *minsets* returns a list of the minimal sets in the list. For example, *minsets*([{2, 3, 4}, {2, 3}, {6, 2, 3}, {2, 3}, {2, 4, 5}]) returns [{2, 3}, {2, 4, 5}].

```

-----logicAssumables.py --- (continued) -----
58 def minsets(ls):
59     """ls is a list of sets
60     returns a list of minimal sets in ls
61     """
62     ans = [] # elements known to be minimal
63     for c in ls:
64         if not any(c1<c for c1 in ls) and not any(c1 <= c for c1 in ans):
65             ans.append(c)
66     return ans
67
68 # minsets([{2, 3, 4}, {2, 3}, {6, 2, 3}, {2, 3}, {2, 4, 5}])

```

Warning: *minsets* works for a list of sets or for a set of (frozen) sets, but it does not work for a generator of sets. For example, try to predict and then test:

```
minsets(e for e in [{2, 3, 4}, {2, 3}, {6, 2, 3}, {2, 3}, {2, 4, 5}])
```

The diagnoses can be constructed from the (minimal) conflicts as follows. This also works if there are non-minimal conflicts, but is not as efficient.

```

-----logicAssumables.py --- (continued) -----
69 def diagnoses(cons):
70     """cons is a list of (minimal) conflicts.
71     returns a list of diagnoses."""
72     if cons == []:
73         return [set()]
74     else:
75         return minsets([(e|d) # | is set union
76                         for e in cons[0]
77                         for d in diagnoses(cons[1:])])

```

Test cases:

```

-----logicAssumables.py --- (continued) -----
80 electa = KBA([
81     Clause('light_l1'),
82     Clause('light_l2'),
83     Assumable('ok_l1'),
84     Assumable('ok_l2'),

```

```

85     Assumable('ok_s1'),
86     Assumable('ok_s2'),
87     Assumable('ok_s3'),
88     Assumable('ok_cb1'),
89     Assumable('ok_cb2'),
90     Assumable('live_outside'),
91     Clause('live_l1', ['live_w0']),
92     Clause('live_w0', ['up_s2', 'ok_s2', 'live_w1']),
93     Clause('live_w0', ['down_s2', 'ok_s2', 'live_w2']),
94     Clause('live_w1', ['up_s1', 'ok_s1', 'live_w3']),
95     Clause('live_w2', ['down_s1', 'ok_s1', 'live_w3']),
96     Clause('live_l2', ['live_w4']),
97     Clause('live_w4', ['up_s3', 'ok_s3', 'live_w3']),
98     Clause('live_p1', ['live_w3']),
99     Clause('live_w3', ['live_w5', 'ok_cb1']),
100    Clause('live_p2', ['live_w6']),
101    Clause('live_w6', ['live_w5', 'ok_cb2']),
102    Clause('live_w5', ['live_outside']),
103    Clause('lit_l1', ['light_l1', 'live_l1', 'ok_l1']),
104    Clause('lit_l2', ['light_l2', 'live_l2', 'ok_l2']),
105    Askable('up_s1'),
106    Askable('down_s1'),
107    Askable('up_s2'),
108    Askable('down_s2'),
109    Askable('up_s3'),
110    Askable('down_s2'),
111    Askable('dark_l1'),
112    Askable('dark_l2'),
113    Clause('false', ['dark_l1', 'lit_l1']),
114    Clause('false', ['dark_l2', 'lit_l2'])
115    ])
116 # electa.prove_all_ass(['false'])
117 # cs=electa.conflicts()
118 # print(cs)
119 # diagnoses(cs)          # diagnoses from conflicts

```

Exercise 5.6 To implement a version of *conflicts* that never generates non-minimal conflicts, modify *prove_all_ass* to implement iterative deepening on the number of assumables used in a proof, and prune any set of assumables that is a superset of a conflict.

Exercise 5.7 Implement *explanations(self, body)*, where *body* is a list of atoms, that returns the a list of the minimal explanations of the body. This does not require modification of *prove_all_ass*.

Exercise 5.8 Implement *explanations*, as in the previous question, so that it never generates non-minimal explanations. Hint: modify *prove_all_ass* to implement iterative deepening on the number of assumptions, generating conflicts and explanations together, and pruning as early as possible.

Planning with Certainty

6.1 Representing Actions and Planning Problems

The STRIPS representation of an action consists of:

- the name of the action
- preconditions: a dictionary of *feature:value* pairs that specifies that the feature must have this value for the action to be possible
- effects: a dictionary of *feature:value* pairs that are made true by this action. In particular, a feature in the dictionary has the corresponding value (and not its previous value) after the action, and a feature not in the dictionary keeps its old value.

```
stripsProblem.py — STRIPS Representations of Actions
11 class Strips(object):
12     def __init__(self, name, preconds, effects, cost=1):
13         """
14         defines the STRIPS representation for an action:
15         * name is the name of the action
16         * preconds, the preconditions, is feature:value dictionary that must hold
17         for the action to be carried out
18         * effects is a feature:value map that this action makes
19         true. The action changes the value of any feature specified
20         here, and leaves other features unchanged.
21         * cost is the cost of the action
22         """
23         self.name = name
```

```

24         self.preconds = preconds
25         self.effects = effects
26         self.cost = cost
27
28     def __repr__(self):
29         return self.name

```

A STRIPS domain consists of:

- A set of actions.
- A dictionary that maps each feature into a set of possible values for the feature.
- A list of the actions

```

_____stripsProblem.py — (continued)_____
31 class STRIPS_domain(object):
32     def __init__(self, feats_vals, actions):
33         """Problem domain
34         feats_vals is a feature:domain dictionary,
35         mapping each feature to its domain
36         actions
37         """
38         self.feats_vals = feats_vals
39         self.actions = actions

```

A planning problem consists of a planning domain, an initial state, and a goal. The goal does not need to fully specify the final state.

```

_____stripsProblem.py — (continued)_____
41 class Planning_problem(object):
42     def __init__(self, prob_domain, initial_state, goal):
43         """
44         a planning problem consists of
45         * a planning domain
46         * the initial state
47         * a goal
48         """
49         self.prob_domain = prob_domain
50         self.initial_state = initial_state
51         self.goal = goal

```

6.1.1 Robot Delivery Domain

The following specifies the robot delivery domain of Section 6.1, shown in Figure 6.1.

```

_____stripsProblem.py — (continued)_____
53 boolean = {True, False}

```

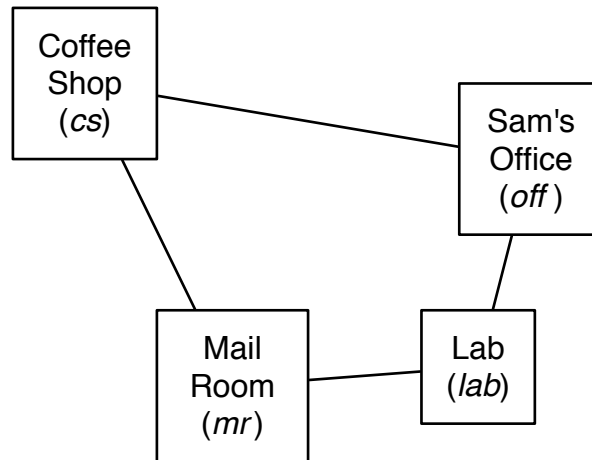
**Features to describe states***RLoc* – Rob's location*RHC* – Rob has coffee*SWC* – Sam wants coffee*MW* – Mail is waiting*RHM* – Rob has mail**Actions***mc* – move clockwise*mcc* – move counterclockwise*puc* – pickup coffee*dc* – deliver coffee*pum* – pickup mail*dm* – deliver mail

Figure 6.1: Robot Delivery Domain

```

54 delivery_domain = STRIPS_domain(
55     {'RLoc':{'cs', 'off', 'lab', 'mr'}, 'RHC':boolean, 'SWC':boolean,
56      'MW':boolean, 'RHM':boolean}, #feature:values dictionary
57     { Strips('mc_cs', {'RLoc':'cs'}, {'RLoc':'off'}),
58       Strips('mc_off', {'RLoc':'off'}, {'RLoc':'lab'}),
59       Strips('mc_lab', {'RLoc':'lab'}, {'RLoc':'mr'}),
60       Strips('mc_mr', {'RLoc':'mr'}, {'RLoc':'cs'}),
61       Strips('mcc_cs', {'RLoc':'cs'}, {'RLoc':'mr'}),
62       Strips('mcc_off', {'RLoc':'off'}, {'RLoc':'cs'}),
63       Strips('mcc_lab', {'RLoc':'lab'}, {'RLoc':'off'}),
64       Strips('mcc_mr', {'RLoc':'mr'}, {'RLoc':'lab'}),
65       Strips('puc', {'RLoc':'cs', 'RHC':False}, {'RHC':True}),
66       Strips('dc', {'RLoc':'off', 'RHC':True}, {'RHC':False, 'SWC':False}),
67       Strips('pum', {'RLoc':'mr', 'MW':True}, {'RHM':True, 'MW':False}),
68       Strips('dm', {'RLoc':'off', 'RHM':True}, {'RHM':False})
69     } )

```

stripsProblem.py — (continued)

```

71 problem0 = Planning_problem(delivery_domain,
72                             {'RLoc':'lab', 'MW':True, 'SWC':True, 'RHC':False,

```

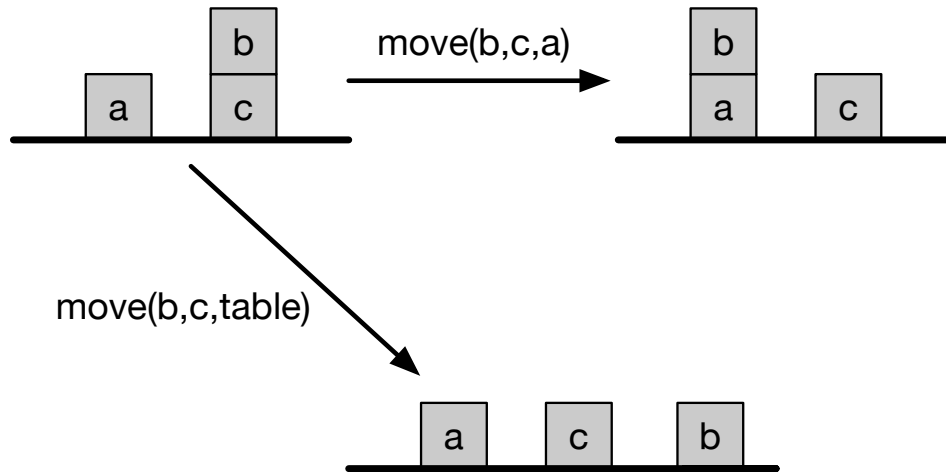


Figure 6.2: Blocks world with two actions

```

73         'RHM':False},
74         {'RLoc':'off'})
75 problem1 = Planning_problem(delivery_domain,
76                             {'RLoc':'lab', 'MW':True, 'SWC':True, 'RHC':False,
77                             'RHM':False},
78                             {'SWC':False})
79 problem2 = Planning_problem(delivery_domain,
80                             {'RLoc':'lab', 'MW':True, 'SWC':True, 'RHC':False,
81                             'RHM':False},
82                             {'SWC':False, 'MW':False, 'RHM':False})

```

6.1.2 Blocks World

The blocks world consist of blocks and a table. Each block can be on the table or on another block. A block can only have one other block on top of it. Figure 6.2 shows 3 states with some of the actions between them.

A state is defined by the two features:

- *on* where $on(x) = y$ when block x is on block or table y
- *clear* where $clear(x) = True$ when block x has nothing on it.

There is one parameterized action

- $move(x, y, z)$ move block x from y to z , where y and z could be a block or the table.

To handle parameterized actions (which depend on the blocks involved), the actions and the features are all strings, created for the all combinations of the

blocks. Note that we treat moving to a block separately from moving to the table, because the blocks needs to be clear, but the table always has room for another block.

```

stripsProblem.py — (continued)
84 ### blocks world
85 def move(x,y,z):
86     """string for the 'move' action"""
87     return 'move_'+x+'_from_'+y+'_to_'+z
88 def on(x):
89     """string for the 'on' feature"""
90     return x+'_is_on'
91 def clear(x):
92     """string for the 'clear' feature"""
93     return 'clear_'+x
94 def create_blocks_world(blocks = {'a','b','c','d'}):
95     blocks_and_table = blocks | {'table'}
96     stmap = {Strips(move(x,y,z),{on(x):y, clear(x):True, clear(z):True},
97                     {on(x):z, clear(y):True, clear(z):False})
98             for x in blocks
99             for y in blocks_and_table
100            for z in blocks
101            if x!=y and y!=z and z!=x}
102     stmap.update({Strips(move(x,y,'table'), {on(x):y, clear(x):True},
103                     {on(x):'table', clear(y):True})
104                 for x in blocks
105                 for y in blocks
106                 if x!=y})
107     feats_vals = {on(x):blocks_and_table-{x} for x in blocks}
108     feats_vals.update({clear(x):boolean for x in blocks_and_table})
109     return STRIPS_domain(feats_vals, stmap)

```

The problem *blocks1* is a classic example, with 3 blocks, and the goal consists of two conditions. See Figure 6.3. Note that this example is challenging because we can't achieve one of the goals and then the other; whichever one we achieve first has to be undone to achieve the second.

```

stripsProblem.py — (continued)
111 blocks1dom = create_blocks_world({'a','b','c'})
112 blocks1 = Planning_problem(blocks1dom,
113     {on('a'):'table', clear('a'):True,
114     on('b'):'c', clear('b'):True,
115     on('c'):'table', clear('c'):False}, # initial state
116     {on('a'):'b', on('c'):'a'}) #goal

```

The problem *blocks2* is one to invert a tower of size 4.

```

stripsProblem.py — (continued)
118 blocks2dom = create_blocks_world({'a','b','c','d'})
119 tower4 = {clear('a'):True, on('a'):'b',
120           clear('b'):False, on('b'):'c',

```

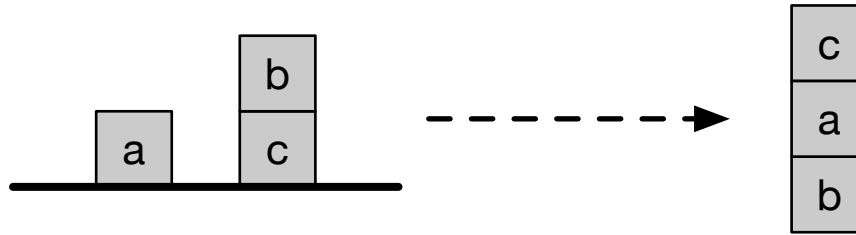


Figure 6.3: Blocks problem blocks1

```

121 |         clear('c'):False, on('c'):'d',
122 |         clear('d'):False, on('d'):'table'}
123 | blocks2 = Planning_problem(blocks2dom,
124 |     tower4, # initial state
125 |     {on('d'):'c',on('c'):'b',on('b'):'a'}) #goal

```

The problem *blocks3* is to move the bottom block to the top of a tower of size 4.

```

stripsProblem.py — (continued)
127 | blocks3 = Planning_problem(blocks2dom,
128 |     tower4, # initial state
129 |     {on('d'):'a', on('a'):'b', on('b'):'c'}) #goal

```

Exercise 6.1 Represent the problem of given a tower of 4 blocks (*a* on *b* on *c* on *d* on table), the goal is to have a tower with the previous top block on the bottom (*b* on *c* on *d* on *a*). Do not include the table in your goal (the goal does not care whether *a* is on the table). [Before you run the program, estimate how many steps it will take to solve this.] How many steps does an optimal planner take?

Exercise 6.2 Represent the domain so that *on*(*x*, *y*) is a Boolean feature that is True when *x* is on *y*. Does the representation of the state need to not include negative *on* facts? Why or why not? (Note that this may depend on the planner; write your answer with respect to particular planners.)

Exercise 6.3 It is possible to write the representation of the problem without using *clear*, where *clear*(*x*) means nothing is on *x*. Change the definition of the blocks world so that it does not use *clear* but uses *on* being false instead. Does this work better for any of the planners?

6.2 Forward Planning

To run the demo, in folder "aipython", load "stripsForwardPlanner.py", and copy and paste the commented-out example queries at the bottom of that file.

In a forward planner, a node is a state. A state consists of an assignment, which is a variable:value dictionary. In order to be able to do multiple-path pruning, we need to define a hash function, and equality between states.

```

stripsForwardPlanner.py — Forward Planner with STRIPS actions
11 from searchProblem import Arc, Search_problem
12 from stripsProblem import Strips, STRIPS_domain
13
14 class State(object):
15     def __init__(self, assignment):
16         self.assignment = assignment
17         self.hash_value = None
18     def __hash__(self):
19         if self.hash_value is None:
20             self.hash_value = hash(frozenset(self.assignment.items()))
21         return self.hash_value
22     def __eq__(self, st):
23         return self.assignment == st.assignment
24     def __str__(self):
25         return str(self.assignment)

```

In order to define a search problem (page 33), we need to define the goal condition, the start nodes, the neighbours, and (optionally) a heuristic function. Here *zero* is the default heuristic function.

```

stripsForwardPlanner.py — (continued)
27 def zero(*args,**nargs):
28     """always returns 0"""
29     return 0
30
31 class Forward_STRIPS(Search_problem):
32     """A search problem from a planning problem where:
33     * a node is a state object.
34     * the dynamics are specified by the STRIPS representation of actions
35     """
36     def __init__(self, planning_problem, heur=zero):
37         """creates a forward search space from a planning problem.
38         heur(state,goal) is a heuristic function,
39         an underestimate of the cost from state to goal, where
40         both state and goals are feature:value dictionaries.
41         """
42         self.prob_domain = planning_problem.prob_domain
43         self.initial_state = State(planning_problem.initial_state)
44         self.goal = planning_problem.goal
45         self.heur = heur
46
47     def is_goal(self, state):
48         """is True if node is a goal.
49
50         Every goal feature has the same value in the state and the goal."""
51         return all(state.assignment[prop]==self.goal[prop]

```

```

52         for prop in self.goal)
53
54     def start_node(self):
55         """returns start node"""
56         return self.initial_state
57
58     def neighbors(self, state):
59         """returns neighbors of state in this problem"""
60         return [ Arc(state, self.effect(act, state.assignment), act.cost, act)
61                 for act in self.prob_domain.actions
62                 if self.possible(act, state.assignment)]
63
64     def possible(self, act, state_asst):
65         """True if act is possible in state.
66         act is possible if all of its preconditions have the same value in the state"""
67         return all(state_asst[pre] == act.preconds[pre]
68                 for pre in act.preconds)
69
70     def effect(self, act, state_asst):
71         """returns the state that is the effect of doing act given state_asst
72         Python 3.9: return state_asst | act.effects"""
73         new_state_asst = state_asst.copy()
74         new_state_asst.update(act.effects)
75         return State(new_state_asst)
76
77     def heuristic(self, state):
78         """in the forward planner a node is a state.
79         the heuristic is an (under)estimate of the cost
80         of going from the state to the top-level goal.
81         """
82         return self.heur(state.assignment, self.goal)

```

Here are some test cases to try.

```

stripsForwardPlanner.py — (continued)
84 from searchBranchAndBound import DF_branch_and_bound
85 from searchMPP import SearcherMPP
86 from stripsProblem import problem0, problem1, problem2, blocks1, blocks2, blocks3
87
88 # SearcherMPP(Forward_STRIPS(problem1)).search() #A* with MPP
89 # DF_branch_and_bound(Forward_STRIPS(problem1), 10).search() #B&B
90 # To find more than one plan:
91 # s1 = SearcherMPP(Forward_STRIPS(problem1)) #A*
92 # s1.search() #find another plan

```

6.2.1 Defining Heuristics for a Planner

Each planning domain requires its own heuristics. If you change the actions, you will need to reconsider the heuristic function, as there might then be a lower-cost path, which might make the heuristic non-admissible.

Here is an example of defining a (not very good) heuristic for the coffee delivery planning domain.

First we define the distance between two locations, which is used for the heuristics.

```

stripsHeuristic.py — Planner with Heuristic Function
11 def dist(loc1, loc2):
12     """returns the distance from location loc1 to loc2
13     """
14     if loc1==loc2:
15         return 0
16     if {loc1,loc2} in [{'cs','lab'},{'mr','off'}]:
17         return 2
18     else:
19         return 1

```

Note that the current state is a complete description; there is a value for every feature. However the goal need not be complete; it does not need to define a value for every feature. Before checking the value for a feature in the goal, a heuristic needs to define whether the feature is defined in the goal.

```

stripsHeuristic.py — (continued)
21 def h1(state,goal):
22     """ the distance to the goal location, if there is one"""
23     if 'RLoc' in goal:
24         return dist(state['RLoc'], goal['RLoc'])
25     else:
26         return 0
27
28 def h2(state,goal):
29     """ the distance to the coffee shop plus getting coffee and delivering it
30     if the robot needs to get coffee
31     """
32     if ('SWC' in goal and goal['SWC']==False
33         and state['SWC']==True
34         and state['RHC']==False):
35         return dist(state['RLoc'],'cs')+3
36     else:
37         return 0

```

The maximum of the values of a set of admissible heuristics is also an admissible heuristic. The function `maxh` takes a number of heuristic functions as arguments, and returns a new heuristic function that takes the maximum of the values of the heuristics. For example, `h1` and `h2` are heuristic functions and so `maxh(h1,h2)` is also. `maxh` can take an arbitrary number of arguments.

```

stripsHeuristic.py — (continued)
39 def maxh(*heuristics):
40     """Returns a new heuristic function that is the maximum of the functions in heuristics.
41     heuristics is the list of arguments which must be heuristic functions.

```

```

42     """
43     # return lambda state,goal: max(h(state,goal) for h in heuristics)
44     def newh(state,goal):
45         return max(h(state,goal) for h in heuristics)
46     return newh

```

The following runs the example with and without the heuristic.

```

stripsHeuristic.py — (continued)
48 ##### Forward Planner #####
49 from searchMPP import SearcherMPP
50 from stripsForwardPlanner import Forward_STRIPS
51 from stripsProblem import problem0, problem1, problem2, blocks1, blocks2, blocks3
52
53 def test_forward_heuristic(thisproblem=problem1):
54     print("\n***** FORWARD NO HEURISTIC")
55     print(SearcherMPP(Forward_STRIPS(thisproblem)).search())
56
57     print("\n***** FORWARD WITH HEURISTIC h1")
58     print(SearcherMPP(Forward_STRIPS(thisproblem,h1)).search())
59
60     print("\n***** FORWARD WITH HEURISTIC h2")
61     print(SearcherMPP(Forward_STRIPS(thisproblem,h2)).search())
62
63     print("\n***** FORWARD WITH HEURISTICS h1 and h2")
64     print(SearcherMPP(Forward_STRIPS(thisproblem,maxh(h1,h2))).search())
65
66 if __name__ == "__main__":
67     test_forward_heuristic()

```

Exercise 6.4 Try the forward planner with a heuristic function of just h_1 , with just h_2 and with both. Explain how each one prunes or doesn't prune the search space.

Exercise 6.5 Create a better heuristic than $\max(h_1, h_2)$. Try it for a number of different problems. In particular, try and include the following costs:

- i) h_3 is like h_2 but also takes into account the case when R_{loc} is in goal.
- ii) h_4 uses the distance to the mail room plus getting mail and delivering it if the robot needs to get need to deliver mail.
- iii) h_5 is for getting mail when goal is for the robot to have mail, and then getting to the goal destination (if there is one).

Exercise 6.6 Create an admissible heuristic for the blocks world.

6.3 Regression Planning

To run the demo, in folder "aipython", load "stripsRegressionPlanner.py", and copy and paste the commented-out example queries at the bottom of that file.

In a regression planner a node is a subgoal that need to be achieved.

A *Subgoal* object consists of an assignment, which is *variable:value* dictionary. We make it hashable so that multiple path pruning can work. The hash is only computed when necessary (and only once).

```

stripsRegressionPlanner.py — Regression Planner with STRIPS actions
11 from searchProblem import Arc, Search_problem
12
13 class Subgoal(object):
14     def __init__(self, assignment):
15         self.assignment = assignment
16         self.hash_value = None
17     def __hash__(self):
18         if self.hash_value is None:
19             self.hash_value = hash(frozenset(self.assignment.items()))
20         return self.hash_value
21     def __eq__(self, st):
22         return self.assignment == st.assignment
23     def __str__(self):
24         return str(self.assignment)

```

A regression search has subgoals as nodes. The initial node is the top-level goal of the planner. The goal for the search (when the search can stop) is a subgoal that holds in the initial state.

```

stripsRegressionPlanner.py — (continued)
26 from stripsForwardPlanner import zero
27
28 class Regression_STRIPS(Search_problem):
29     """A search problem where:
30     * a node is a goal to be achieved, represented by a set of propositions.
31     * the dynamics are specified by the STRIPS representation of actions
32     """
33
34     def __init__(self, planning_problem, heur=zero):
35         """creates a regression search space from a planning problem.
36         heur(state,goal) is a heuristic function;
37         an underestimate of the cost from state to goal, where
38         both state and goals are feature:value dictionaries
39         """
40         self.prob_domain = planning_problem.prob_domain
41         self.top_goal = Subgoal(planning_problem.goal)
42         self.initial_state = planning_problem.initial_state
43         self.heur = heur
44
45     def is_goal(self, subgoal):
46         """if subgoal is true in the initial state, a path has been found"""
47         goal_asst = subgoal.assignment
48         return all(self.initial_state[g]==goal_asst[g]
49                   for g in goal_asst)
50

```

```

51 def start_node(self):
52     """the start node is the top-level goal"""
53     return self.top_goal
54
55 def neighbors(self, subgoal):
56     """returns a list of the arcs for the neighbors of subgoal in this problem"""
57     goal_asst = subgoal.assignment
58     return [ Arc(subgoal, self.weakest_precond(act, goal_asst), act.cost, act)
59             for act in self.prob_domain.actions
60             if self.possible(act, goal_asst)]
61
62 def possible(self, act, goal_asst):
63     """True if act is possible to achieve goal_asst.
64
65     the action achieves an element of the effects and
66     the action doesn't delete something that needs to be achieved and
67     the preconditions are consistent with other subgoals that need to be achieved
68     """
69     return ( any(goal_asst[prop] == act.effects[prop]
70                 for prop in act.effects if prop in goal_asst)
71             and all(goal_asst[prop] == act.effects[prop]
72                    for prop in act.effects if prop in goal_asst)
73             and all(goal_asst[prop] == act.preconds[prop]
74                    for prop in act.preconds if prop not in act.effects and prop in goal_asst)
75             )
76
77 def weakest_precond(self, act, goal_asst):
78     """returns the subgoal that must be true so goal_asst holds after act
79     should be: act.preconds | (goal_asst - act.effects)
80     """
81     new_asst = act.preconds.copy()
82     for g in goal_asst:
83         if g not in act.effects:
84             new_asst[g] = goal_asst[g]
85     return Subgoal(new_asst)
86
87 def heuristic(self, subgoal):
88     """in the regression planner a node is a subgoal.
89     the heuristic is an (under)estimate of the cost of going from the initial state to subgoal.
90     """
91     return self.heur(self.initial_state, subgoal.assignment)

```

stripsRegressionPlanner.py — (continued)

```

93 from searchBranchAndBound import DF_branch_and_bound
94 from searchMPP import SearcherMPP
95 from stripsProblem import problem0, problem1, problem2, blocks1, blocks2, blocks3
96
97 # SearcherMPP(Regression_STRIPS(problem1)).search() #A* with MPP
98 # DF_branch_and_bound(Regression_STRIPS(problem1), 10).search() #B&B

```


Exercise 6.7 Multiple path pruning could be used to prune more than the current code. In particular, if the current node contains more conditions than a previously visited node, it can be pruned. For example, if $\{a : \text{True}, b : \text{False}\}$ has been visited, then any node that is a superset, e.g., $\{a : \text{True}, b : \text{False}, d : \text{True}\}$, need not be expanded. If the simpler subgoal does not lead to a solution, the more complicated one won't either. Implement this more severe pruning. (Hint: This may require modifications to the searcher.)

Exercise 6.8 It is possible that, as knowledge of the domain, that some assignment of values to variables can never be achieved. For example, the robot cannot be holding mail when there is mail waiting (assuming it isn't holding mail initially). An assignment of values to (some of the) variables is incompatible if no possible (reachable) state can include that assignment. For example, $\{MW' : \text{True}, RHM' : \text{True}\}$ is an incompatible assignment. This information may be useful information for a planner; there is no point in trying to achieve these together. Define a subclass of *STRIPS.domain* that can accept a list of incompatible assignments. Modify the regression planner code to use such a list of incompatible assignments. Give an example where the search space is smaller.

Exercise 6.9 After completing the previous exercise, design incompatible assignments for the blocks world. (This should result in dramatic search improvements.)

6.3.1 Defining Heuristics for a Regression Planner

The regression planner can use the same heuristic function as the forward planner. However, just because a heuristic is useful for a forward planner does not mean it is useful for a regression planner, and vice versa. you should experiment with whether the same heuristic works well for both a regression planner and a forward planner.

The following runs the same example as the forward planner with and without the heuristic defined for the forward planner:

```

stripsHeuristic.py — (continued)
69 ##### Regression Planner
70 from stripsRegressionPlanner import Regression_STRIPS
71
72 def test_regression_heuristic(thisproblem=problem1):
73     print("\n***** REGRESSION NO HEURISTIC")
74     print(SearcherMPP(Regression_STRIPS(thisproblem)).search())
75
76     print("\n***** REGRESSION WITH HEURISTICS h1 and h2")
77     print(SearcherMPP(Regression_STRIPS(thisproblem,maxh(h1,h2))).search())
78
79 if __name__ == "__main__":
80     test_regression_heuristic()
```

Exercise 6.10 Try the regression planner with a heuristic function of just $h1$ and with just $h2$ (defined in Section 6.2.1). Explain how each one prunes or doesn't prune the search space.

Exercise 6.11 Create a better heuristic than *heuristic_fun* defined in Section 6.2.1.

6.4 Planning as a CSP

To run the demo, in folder "aipython", load "stripsCSPPlanner.py", and copy and paste the commented-out example queries at the bottom of that file. This assumes Python 3.

Here we implement the CSP planner assuming there is a single action at each step. This creates a CSP that can use any of the CSP algorithms to solve (e.g., stochastic local search or arc consistency with domain splitting).

This assumes the same action representation as before; we do not consider factored actions (action features), nor do we implement state constraints.

```

_____stripsCSPPlanner.py — CSP planner where actions are represented using STRIPS_____
11 from cspProblem import CSP, Constraint
12
13 class CSP_from_STRIPS(CSP):
14     """A CSP where:
15     * a CSP variable is constructed by st(var,stage).
16     * the dynamics are specified by the STRIPS representation of actions
17     """
18
19     def __init__(self, planning_problem, number_stages=2):
20         prob_domain = planning_problem.prob_domain
21         initial_state = planning_problem.initial_state
22         goal = planning_problem.goal
23         self.act_vars = [st('action',stage) for stage in range(number_stages)]
24         domains = {av:prob_domain.actions for av in self.act_vars}
25         domains.update({ st(var,stage):dom
26                         for (var,dom) in prob_domain.feats_vals.items()
27                         for stage in range(number_stages+1)})
28         # initial state constraints:
29         constraints = [Constraint((st(var,0),), is_(val))
30                         for (var,val) in initial_state.items()]
31         # goal constraints on the final state:
32         constraints += [Constraint((st(var,number_stages),),
33                                   is_(val))
34                         for (var,val) in goal.items()]
35         # precondition constraints:
36         constraints += [Constraint((st(var,stage), st('action',stage)),
37                                   if_(val,act)) # st(var,stage)==val if st('action',stage)=act
38                         for act in prob_domain.actions
39                         for var,val in act.preconds.items()
40                         for stage in range(number_stages)]
41         # effect constraints:
42         constraints += [Constraint((st(var,stage+1), st('action',stage)),
43                                   if_(val,act)) # st(var,stage+1)==val if st('action',stage)=act
44

```

```

44         for act in prob_domain.actions
45         for var,val in act.effects.items()
46         for stage in range(number_stages)]
47     # frame constraints:
48     constraints += [Constraint((st(var,stage), st('action',stage), st(var,stage+1)),
49                             eq_if_not_in_({act for act in prob_domain.actions
50                                     if var in act.effects}))
51         for var in prob_domain.feats_vals
52         for stage in range(number_stages) ]
53     CSP.__init__(self, domains, constraints)
54
55     def extract_plan(self,soln):
56         return [soln[a] for a in self.act_vars]
57
58     def st(var,stage):
59         """returns a string for the var-stage pair that can be used as a variable"""
60         return str(var)+"_"+str(stage)

```

The following methods return methods which can be applied to the particular environment.

For example, `is_(3)` returns a function that when applied to 3, returns True and when applied to any other value returns False. So `is_(3)(3)` returns *True* and `is_(3)(7)` returns *False*.

Note that the underscore ('_') is part of the name; here we use it as the convention that it is a function that returns a function. This uses two different styles to define `is_` and `if_`; returning a function defined by *lambda* is equivalent to returning the embedded function, except that the embedded function has a name. The embedded function can also be given a docstring.

```

stripsCSPPlanner.py — (continued)
62 def is_(val):
63     """returns a function that is true when it is it applied to val.
64     """
65     #return lambda x: x == val
66     def is_fun(x):
67         return x == val
68     is_fun.__name__ = "value_is_"+str(val)
69     return is_fun
70
71 def if_(v1,v2):
72     """if the second argument is v2, the first argument must be v1"""
73     #return lambda x1,x2: x1==v1 if x2==v2 else True
74     def if_fun(x1,x2):
75         return x1==v1 if x2==v2 else True
76     if_fun.__name__ = "if x2 is "+str(v2)+" then x1 is "+str(v1)
77     return if_fun
78
79 def eq_if_not_in_(actset):
80     """first and third arguments are equal if action is not in actset"""
81     # return lambda x1, a, x2: x1==x2 if a not in actset else True

```

```

82     def eq_if_not_fun(x1, a, x2):
83         return x1==x2 if a not in actset else True
84     eq_if_not_fun.__name__ = "first and third arguments are equal if action is not in "+str(actset)
85     return eq_if_not_fun

```

Putting it together, this returns a list of actions that solves the problem *prob* for a given horizon. If you want to do more than just return the list of actions, you might want to get it to return the solution. Or even enumerate the solutions (by using *Search_with_AC_from_CSP*).

```

_____stripsCSPPlanner.py — (continued) _____
87 def con_plan(prob,horizon):
88     """finds a plan for problem prob given horizon.
89     """
90     csp = CSP_from_STRIPS(prob, horizon)
91     sol = Con_solver(csp).solve_one()
92     return csp.extract_plan(sol) if sol else sol

```

The following are some example queries.

```

_____stripsCSPPlanner.py — (continued) _____
94 from searchGeneric import Searcher
95 from stripsProblem import delivery_domain
96 from cspConsistency import Search_with_AC_from_CSP, Con_solver
97 from stripsProblem import Planning_problem, problem0, problem1, problem2, blocks1, blocks2, blocks3
98
99 # Problem 0
100 # con_plan(problem0,1) # should it succeed?
101 # con_plan(problem0,2) # should it succeed?
102 # con_plan(problem0,3) # should it succeed?
103 # To use search to enumerate solutions
104 #searcher0a = Searcher(Search_with_AC_from_CSP(CSP_from_STRIPS(problem0, 1)))
105 #print(searcher0a.search())
106
107 ## Problem 1
108 # con_plan(problem1,5) # should it succeed?
109 # con_plan(problem1,4) # should it succeed?
110 ## To use search to enumerate solutions:
111 #searcher15a = Searcher(Search_with_AC_from_CSP(CSP_from_STRIPS(problem1, 5)))
112 #print(searcher15a.search())
113
114 ## Problem 2
115 #con_plan(problem2, 6) # should fail??
116 #con_plan(problem2, 7) # should succeed???
117
118 ## Example 6.13
119 problem3 = Planning_problem(delivery_domain,
120                             {'SWC':True, 'RHC':False}, {'SWC':False})
121 #con_plan(problem3,2) # Horizon of 2
122 #con_plan(problem3,3) # Horizon of 3
123

```

```

124 | problem4 = Planning_problem(delivery_domain,{'SWC':True},
125 |                             {'SWC':False, 'MW':False, 'RHM':False})
126 |
127 | # For the stochastic local search:
128 | #from cspSLS import SLSearcher, Runtime_distribution
129 | # cspplanning15 = CSP_from_STRIPS(problem1, 5) # should succeed
130 | #se0 = SLSearcher(cspplanning15); print(se0.search(100000,0.5))
131 | #p = Runtime_distribution(cspplanning15)
132 | #p.plot_run(1000,1000,0.7) # warning will take a few minutes

```

6.5 Partial-Order Planning

To run the demo, in folder "aipython", load "stripsPOP.py", and copy and paste the commented-out example queries at the bottom of that file.

A partial order planner maintains a partial order of action instances. An action instance consists of a name and an index. We need action instances because the same action could be carried out at different times.

```

_____stripsPOP.py — Partial-order Planner using STRIPS representation_____
11 | from searchProblem import Arc, Search_problem
12 | import random
13 |
14 | class Action_instance(object):
15 |     next_index = 0
16 |     def __init__(self,action,index=None):
17 |         if index is None:
18 |             index = Action_instance.next_index
19 |             Action_instance.next_index += 1
20 |         self.action = action
21 |         self.index = index
22 |
23 |     def __str__(self):
24 |         return str(self.action)+"#"+str(self.index)
25 |
26 |     __repr__ = __str__ # __repr__ function is the same as the __str__ function

```

A node (as in the abstraction of search space) in a partial-order planner consists of:

- *actions*: a set of action instances.
- *constraints*: a set of (a_1, a_2) pairs, where a_1 and a_2 are action instances, which represents that a_1 must come before a_2 in the partial order. There are a number of ways that this could be represented. Here we represent the set of pairs that are in transitive closure of the *before* relation. This lets us quickly determine whether some before relation is consistent with the current constraints.

- *agenda*: a list of (s, a) pairs, where s is a (var, val) pair and a is an action instance. This means that variable var must have value val before a can occur.
- *causal_links*: a set of (a_0, g, a_1) triples, where a_1 and a_2 are action instances and g is a (var, val) pair. This holds when action a_0 makes g true for action a_1 .

```

stripsPOP.py — (continued)
28 class POP_node(object):
29     """a (partial) partial-order plan. This is a node in the search space."""
30     def __init__(self, actions, constraints, agenda, causal_links):
31         """
32         * actions is a set of action instances
33         * constraints a set of (a0,a1) pairs, representing a0<a1,
34           closed under transitivity
35         * agenda list of (subgoal,action) pairs to be achieved, where
36           subgoal is a (variable,value) pair
37         * causal_links is a set of (a0,g,a1) triples,
38           where ai are action instances, and g is a (variable,value) pair
39         """
40         self.actions = actions # a set of action instances
41         self.constraints = constraints # a set of (a0,a1) pairs
42         self.agenda = agenda # list of (subgoal,action) pairs to be achieved
43         self.causal_links = causal_links # set of (a0,g,a1) triples
44
45     def __str__(self):
46         return ("actions: "+str({str(a) for a in self.actions})+
47             "\nconstraints: "+
48             str({(str(a1),str(a2)) for (a1,a2) in self.constraints})+
49             "\nagenda: "+
50             str([(str(s),str(a)) for (s,a) in self.agenda])+
51             "\ncausal_links:"+
52             str({(str(a0),str(g),str(a2)) for (a0,g,a2) in self.causal_links})
53         )

```

extract_plan constructs a total order of action instances that is consistent with the partial order.

```

stripsPOP.py — (continued)
54 def extract_plan(self):
55     """returns a total ordering of the action instances consistent
56     with the constraints.
57     raises IndexError if there is no choice.
58     """
59     sorted_acts = []
60     other_acts = set(self.actions)
61     while other_acts:
62         a = random.choice([a for a in other_acts if
63             all(((a1,a) not in self.constraints) for a1 in other_acts)])

```

```

64         sorted_acts.append(a)
65         other_acts.remove(a)
66     return sorted_acts

```

POP_search_from_STRIPS is an instance of a search problem. As such, we need to define the start nodes, the goal, and the neighbors of a node.

```

stripsPOP.py — (continued)
68 from display import Displayable
69
70 class POP_search_from_STRIPS(Search_problem, Displayable):
71     def __init__(self, planning_problem):
72         Search_problem.__init__(self)
73         self.planning_problem = planning_problem
74         self.start = Action_instance("start")
75         self.finish = Action_instance("finish")
76
77     def is_goal(self, node):
78         return node.agenda == []
79
80     def start_node(self):
81         constraints = {(self.start, self.finish)}
82         agenda = [(g, self.finish) for g in self.planning_problem.goal.items()]
83         return POP_node([self.start, self.finish], constraints, agenda, [])

```

The *neighbors* method is a coroutine that enumerates the neighbors of a given node.

```

stripsPOP.py — (continued)
85 def neighbors(self, node):
86     """enumerates the neighbors of node"""
87     self.display(3, "finding neighbors of\n", node)
88     if node.agenda:
89         subgoal, act1 = node.agenda[0]
90         self.display(2, "selecting", subgoal, "for", act1)
91         new_agenda = node.agenda[1:]
92         for act0 in node.actions:
93             if (self.achieves(act0, subgoal) and
94                 self.possible((act0, act1), node.constraints)):
95                 self.display(2, "reusing", act0)
96                 consts1 = self.add_constraint((act0, act1), node.constraints)
97                 new_clink = (act0, subgoal, act1)
98                 new_cls = node.causal_links + [new_clink]
99                 for consts2 in self.protect_cl_for_actions(node.actions, consts1, new_clink):
100                     yield Arc(node,
101                               POP_node(node.actions, consts2, new_agenda, new_cls),
102                               cost=0)
103         for a0 in self.planning_problem.prob_domain.actions: #a0 is an action
104             if self.achieves(a0, subgoal):
105                 #a0 achieves subgoal
106                 new_a = Action_instance(a0)

```

```

107         self.display(2, " using new action", new_a)
108         new_actions = node.actions + [new_a]
109         consts1 = self.add_constraint((self.start, new_a), node.constraints)
110         consts2 = self.add_constraint((new_a, act1), consts1)
111         new_agenda1 = new_agenda + [(pre, new_a) for pre in a0.preconds.items()]
112         new_clink = (new_a, subgoal, act1)
113         new_cls = node.causal_links + [new_clink]
114         for consts3 in self.protect_all_cls(node.causal_links, new_a, consts2):
115             for consts4 in self.protect_cl_for_actions(node.actions, consts3, new_clink):
116                 yield Arc(node,
117                           POP_node(new_actions, consts4, new_agenda1, new_cls),
118                           cost=1)

```

Given a casual link $(a0, subgoal, a1)$, the following method protects the causal link from each action in *actions*. Whenever an action deletes *subgoal*, the action needs to be before *a0* or after *a1*. This method enumerates all constraints that result from protecting the causal link from all actions.

```

stripsPOP.py — (continued)
120 def protect_cl_for_actions(self, actions, constrs, clink):
121     """yields constraints that extend constrs and
122     protect causal link (a0, subgoal, a1)
123     for each action in actions
124     """
125     if actions:
126         a = actions[0]
127         rem_actions = actions[1:]
128         a0, subgoal, a1 = clink
129         if a != a0 and a != a1 and self.deletes(a, subgoal):
130             if self.possible((a, a0), constrs):
131                 new_const = self.add_constraint((a, a0), constrs)
132                 for e in self.protect_cl_for_actions(rem_actions, new_const, clink): yield e
133             # could be "yield from"
134             if self.possible((a1, a), constrs):
135                 new_const = self.add_constraint((a1, a), constrs)
136                 for e in self.protect_cl_for_actions(rem_actions, new_const, clink): yield e
137             else:
138                 for e in self.protect_cl_for_actions(rem_actions, constrs, clink): yield e
139         else:
140             yield constrs

```

Given an action *act*, the following method protects all the causal links in *clinks* from *act*. Whenever *act* deletes *subgoal* from some causal link $(a0, subgoal, a1)$, the action *act* needs to be before *a0* or after *a1*. This method enumerates all constraints that result from protecting the causal links from *act*.

```

stripsPOP.py — (continued)
141 def protect_all_cls(self, clinks, act, constrs):
142     """yields constraints that protect all causal links from act"""
143     if clinks:
144         (a0, cond, a1) = clinks[0] # select a causal link

```



```

145         rem_clinks = clinks[1:] # remaining causal links
146         if act != a0 and act != a1 and self.deletes(act,cond):
147             if self.possible((act,a0),constrs):
148                 new_const = self.add_constraint((act,a0),constrs)
149                 for e in self.protect_all_cls(rem_clinks,act,new_const): yield e
150             if self.possible((a1,act),constrs):
151                 new_const = self.add_constraint((a1,act),constrs)
152                 for e in self.protect_all_cls(rem_clinks,act,new_const): yield e
153         else:
154             for e in self.protect_all_cls(rem_clinks,act,constrs): yield e
155     else:
156         yield constrs

```

The following methods check whether an action (or action instance) achieves or deletes some subgoal.

```

stripsPOP.py — (continued)
158 def achieves(self,action,subgoal):
159     var,val = subgoal
160     return var in self.effects(action) and self.effects(action)[var] == val
161
162 def deletes(self,action,subgoal):
163     var,val = subgoal
164     return var in self.effects(action) and self.effects(action)[var] != val
165
166 def effects(self,action):
167     """returns the variable:value dictionary of the effects of action.
168     works for both actions and action instances"""
169     if isinstance(action, Action_instance):
170         action = action.action
171     if action == "start":
172         return self.planning_problem.initial_state
173     elif action == "finish":
174         return {}
175     else:
176         return action.effects

```

The constraints are represented as a set of pairs closed under transitivity. Thus if (a,b) and (b,c) are the list, then (a,c) must also be in the list. This means that adding a new constraint means adding the implied pairs, but querying whether some order is consistent is quick.

```

stripsPOP.py — (continued)
178 def add_constraint(self, pair, const):
179     if pair in const:
180         return const
181     todo = [pair]
182     newconst = const.copy()
183     while todo:
184         x0,x1 = todo.pop()
185         newconst.add((x0,x1))

```

```

186         for x,y in newconst:
187             if x==x1 and (x0,y) not in newconst:
188                 todo.append((x0,y))
189             if y==x0 and (x,x1) not in newconst:
190                 todo.append((x,x1))
191         return newconst
192
193     def possible(self,pair,constraint):
194         (x,y) = pair
195         return (y,x) not in constraint

```

Some code for testing:

```

stripsPOP.py — (continued)
197 from searchBranchAndBound import DF_branch_and_bound
198 from searchMPP import SearcherMPP
199 from stripsProblem import problem0, problem1, problem2, blocks1, blocks2, blocks3
200
201 rplanning0 = POP_search_from_STRIPS(problem0)
202 rplanning1 = POP_search_from_STRIPS(problem1)
203 rplanning2 = POP_search_from_STRIPS(problem2)
204 searcher0 = DF_branch_and_bound(rplanning0,5)
205 searcher0a = SearcherMPP(rplanning0)
206 searcher1 = DF_branch_and_bound(rplanning1,10)
207 searcher1a = SearcherMPP(rplanning1)
208 searcher2 = DF_branch_and_bound(rplanning2,10)
209 searcher2a = SearcherMPP(rplanning2)
210 # Try one of the following searchers
211 # a = searcher0.search()
212 # a = searcher0a.search()
213 # a.end().extract_plan() # print a plan found
214 # a.end().constraints # print the constraints
215 # SearcherMPP.max_display_level = 0 # less detailed display
216 # DF_branch_and_bound.max_display_level = 0 # less detailed display
217 # a = searcher1.search()
218 # a = searcher1a.search()
219 # a = searcher2.search()
220 # a = searcher2a.search()

```

Supervised Machine Learning

This chapter is the first on machine learning. It covers the following topics:

- Data: how to load it, training and test sets
- Features: many of the features come directly from the data. Sometimes it is useful to construct features, e.g. $height > 1.9m$ might be a Boolean feature constructed from the real-values feature *height*. The next chapter is about how to learn these features; in this chapter we construct them by hand, in what is often known as **feature engineering**.
- Learning with no input features: this is the base case of many methods. What should we predict if we have no input features? This is basic knowledge that everyone doing machine learning should know.
- Decision tree learning: one of the classic and simplest learning algorithms, which is the basis of many other algorithms.
- Cross validations and parameter tuning: methods to prevent overfitting.
- Linear regression and classification: other classic and simple techniques that often work well (particularly combined with feature learning or engineering).
- Boosting: combining simpler learning methods to make even better learners.

A good source of classic datasets is the UCI machine Learning Repository [Lichman, 2013]. The SPECT and car datasets are from this repository.

7.1 Representations of Data and Predictions

The code uses the following definitions and conventions:

- A **data set** is an enumeration of examples.
- An **example** is a list (or tuple) of feature values. The feature values can be numbers or strings.
- A **feature** is a function from examples into the range of the feature. We assume each feature has a variable *frange* that gives the range of the feature.

A **Boolean feature** is a function from the examples into $\{False, True\}$. So, if f is a Boolean feature, $f.frange == [False, True]$, and if e is an example, $f(e)$ is either *True* or *False*.

The `__doc__` variable of the function contains the docstring, a string description of the function.

```

learnProblem.py — A Learning Problem
11 import math, random
12 import csv
13 from display import Displayable
14
15 boolean = [False, True]
```

When creating a data set, we partition the data into a training set (*train*) and a test set (*test*). The target feature is the feature that we are making a prediction of.

```

learnProblem.py — (continued)
17 class Data_set(Displayable):
18     """ A data set consists of a list of training data and a list of test data.
19     """
20     seed = None #123456 # make it None for a different test set each time
21
22     def __init__(self, train, test=None, prob_test=0.30, target_index=0, header=None):
23         """A dataset for learning.
24         train is a list of tuples representing the training examples
25         test is the list of tuples representing the test examples
26         if test is None, a test set is created by selecting each
27         example with probability prob_test
28         target_index is the index of the target. If negative, it counts from right.
29         If target_index is larger than the number of properties,
30         there is no target (for unsupervised learning)
31         header is a list of names for the features
32         """
33         if test is None:
34             train, test = partition_data(train, prob_test, seed=self.seed)
35             self.train = train
```

```

36         self.test = test
37         self.display(1,f"Training set has {len(train)} examples. Number of columns in ",{len(e) for e in train})
38         self.display(1,f"Test set has {len(test)} examples. Number of columns in ",{len(e) for e in test})
39         self.prob_test = prob_test
40         self.num_properties = len(self.train[0])
41         if target_index < 0: #allows for -1, -2, etc.
42             target_index = self.num_properties + target_index
43         self.target_index = target_index
44         self.header = header
45         self.create_features()
46         self.display(1,"There are",len(self.input_features),"input features")

```

Initially we assume that the properties can be mapped directly into features. If all values are 0 or 1 they can be used as Boolean features. This can be overridden to allow for more general features.

```

learnProblem.py — (continued)
48 def create_features(self):
49     """create the input features and target feature.
50     This assumes that the features all have range {0,1}.
51     This should be overridden if the features have a different range.
52     """
53     self.input_features = []
54     for i in range(self.num_properties):
55         def feat(e,index=i):
56             return e[index]
57         if self.header:
58             feat.__doc__ = self.header[i]
59         else:
60             feat.__doc__ = "e["+str(i)+"]"
61         feat.frange = [0,1]
62         if i == self.target_index:
63             self.target = feat
64         else:
65             self.input_features.append(feat)

```

7.1.1 Evaluating Predictions

A **predictor** is a function that takes an example and makes a prediction on the value of the target feature.

An error measure takes a prediction and the actual value and returns a non-negative real number, such that the error for a dataset is the mean of the errors for each example. We assume that a lower error is better.

The function *evaluate_dataset* returns the average error for each example, where the error for each example depends on the evaluation criteria. Here we consider three evaluation criteria, the squared error (average of the square of the difference between the actual and predicted values), absolute errors(average of the absolute difference between the actual and predicted values) and the logloss (the a average negative log-likelihood, which can be interpreted as the

number of bits to describe an example using a code based on the prediction treated as a probability).

```

learnProblem.py — (continued)
67 def evaluate_dataset(self, data, predictor, error_measure):
68     """Evaluates predictor on data according to the error_measure
69     predictor is a function that takes an example and returns a
70     prediction for the target feature.
71     error_measure(prediction,actual) -> non-negative reals
72     """
73     if data:
74         try:
75             error = mean(error_measure(predictor(e), self.target(e))
76                           for e in data)
77         except ValueError:
78             return float("inf") # infinity
79     return error

```

The following evaluation criteria are defined. (Please keep the `__doc__` strings a consistent length as they are used in tables.)

```

learnProblem.py — (continued)
81 def squared_error(prediction,actual):
82     "squared error "
83     return (prediction-actual)**2
84 def absolute_error(prediction, actual):
85     "absolute error"
86     return abs(prediction-actual)
87 def log_loss(prediction,actual):
88     "logloss      "
89     try:
90         if actual==0:
91             return -math.log2(1-prediction)
92         else:
93             return -math.log2(prediction)
94     except ValueError:
95         return float("inf") # infinity
96
97 evaluation_criteria = [squared_error, absolute_error, log_loss]

```

The following computes the mean of an enumeration, with an optional initial sum and initial count. This works for enumerations, even where `len()` is not defined, and only goes through the enumeration once. The obvious way to compute a mean: `sum(enum)/len(enum)` works for lists, but does not work for arbitrary enumerations.

```

learnProblem.py — (continued)
99 def mean(enum, isum=0, icount=0):
100     """returns the mean of enumeration enum,
101     isum is the initial sum, and icount is the initial count."""
102     for e in enum:

```

```

103         icount += 1
104         isum += e
105     return isum/icount

```

7.1.2 Creating Test and Training Sets

The following method partitions the data into a training set and a test set. Note that this does not guarantee that the test set will contain exactly a proportion of the data equal to *prob_test*.

[An alternative is to use *random.sample()* which can guarantee that the test set will contain exactly a particular proportion of the data. However this would require knowing how many elements are in the data set, which we may not know, as *data* may just be a generator of the data (e.g., when reading the data from a file).]

```

learnProblem.py — (continued)
107 def partition_data(data, prob_test=0.30, seed=None):
108     """partitions the data into a training set and a test set, where
109     prob_test is the probability of each example being in the test set.
110     """
111     train = []
112     test = []
113     if seed: # given seed makes the partition consistent from run-to-run
114         random.seed(seed)
115     for example in data:
116         if random.random() < prob_test:
117             test.append(example)
118         else:
119             train.append(example)
120     return train, test

```

7.1.3 Importing Data From File

A data set is typically loaded from a file. The default here is that it loaded from a CSV (comma separated values) file, although the default separator can be changed. This assumes that all lines that contain the separator are valid data (so we only include those data items that contain more than one element). This allows for blank lines and comment lines that do not contain the separator. However, it means that this method is not suitable for cases where there is only one feature.

Note that *data_all* and *data_tuples* are generators. *data_all* is a generator of a list of list of strings. This version assumes that CSV files are simple. The standard *csv* package, that allows quoted arguments, can be used by uncommenting the line for *data_all* and commenting out the following line. *data_tuples* contains only those lines that contain the delimiter (others lines are assumed to

be empty or comments), and tries to convert the elements to numbers whenever possible.

This allows for some of the columns to be included; specified by *include_only*. Note that if *include_only* is specified, the target index is the column in the remaining columns.

```

122 class Data_from_file(Data_set):
123     def __init__(self, file_name, separator=',', num_train=None, prob_test=0.3,
124                 has_header=False, target_index=0, boolean_features=True,
125                 categorical=[], include_only=None):
126         """create a dataset from a file
127         separator is the character that separates the attributes
128         num_train is a number n specifying the first n tuples are training, or None
129         prob_test is the probability an example should in the test set (if num_train is None)
130         has_header is True if the first line of file is a header
131         target_index specifies which feature is the target
132         boolean_features specifies whether we want to create Boolean features
133         (if False, it uses the original features).
134         categorical is a set (or list) of features that should be treated as categorical
135         include_only is a list or set of indexes of columns to include
136         """
137         self.boolean_features = boolean_features
138         with open(file_name,'r',newline='') as csvfile:
139             # data_all = csv.reader(csvfile,delimiter=separator) # for more complicated CSV files
140             data_all = (line.strip().split(separator) for line in csvfile)
141             if include_only is not None:
142                 data_all = ([v for (i,v) in enumerate(line) if i in include_only]
143                             for line in data_all)
144             if has_header:
145                 header = next(data_all)
146             else:
147                 header = None
148             data_tuples = (interpret_elements(d) for d in data_all if len(d)>1)
149             if num_train is not None:
150                 # training set is divided into training then test examples
151                 # the file is only read once, and the data is placed in appropriate list
152                 train = []
153                 for i in range(num_train): # will give an error if insufficient examples
154                     train.append(next(data_tuples))
155                 test = list(data_tuples)
156                 Data_set.__init__(self,train, test=test, target_index=target_index,header=header)
157             else: # randomly assign training and test examples
158                 Data_set.__init__(self,data_tuples, test=None, prob_test=prob_test,
159                                   target_index=target_index, header=header)
160
161     def __str__(self):
162         if self.train and len(self.train)>0:
163             return ("Data: "+str(len(self.train))+ " training examples, "
164                     +str(len(self.test))+ " test examples, ")

```



```

165         +str(len(self.train[0]))+" features.")
166     else:
167         return ("Data: "+str(len(self.train))+ " training examples, "
168               +str(len(self.test))+ " test examples.")

```

7.1.4 Creating Binary Features

Some of the algorithms require Boolean features or features with range $\{0, 1\}$. In order to be able to use these algorithms on datasets that allow for arbitrary ranges of input variables, we construct binary features from the attributes. This method overrides the method in *Data_set*.

There are 3 cases:

- When the range only has two values, we designate one to be the “true” value.
- When the values are all numeric, we assume they are ordered (as opposed to just being some classes that happen to be labelled with numbers) and construct Boolean features for splits of the data. That is, the feature is $e[ind] < cut$ for some value *cut*. We choose a number of *cut* values, up to a maximum number of cuts, given by *max_num_cuts*.
- When the values are not all numeric, we assume they are unordered, and create an indicator function for each value. An indicator function for a value returns true when that value is given and false otherwise. Note that we can’t create an indicator function for values that appear in the test set but not in the training set because we haven’t seen the test set. For the examples in the test set with a value that doesn’t appear in the training set for that feature, the indicator functions all return false.

```

learnProblem.py — (continued)
170 def create_features(self, max_num_cuts=8):
171     """creates boolean features from input features.
172     max_num_cuts is the maximum number of binary variables
173     to split a numerical feature into.
174     """
175     ranges = [set() for i in range(self.num_properties)]
176     for example in self.train:
177         for ind, val in enumerate(example):
178             ranges[ind].add(val)
179     if self.target_index <= self.num_properties:
180         # If target_index is larger than the number of properties,
181         # there is no target (for unsupervised learning)
182         def target(e, index=self.target_index):
183             return e[index]
184         if self.header:
185             target.__doc__ = self.header[ind]

```

```

186         else:
187             target.__doc__ = "e["+str(ind)+"]"
188             target.frange = ranges[self.target_index]
189             self.target = target
190     if self.boolean_features:
191         self.input_features = []
192         for ind,frange in enumerate(ranges):
193             if ind != self.target_index and len(frange)>1:
194                 if len(frange) == 2:
195                     # two values, the feature is equality to one of them.
196                     true_val = list(frange)[1] # choose one as true
197                     def feat(e, i=ind, tv=true_val):
198                         return e[i]==tv
199                     if self.header:
200                         feat.__doc__ = self.header[ind]+"==" +str(true_val)
201                     else:
202                         feat.__doc__ = "e["+str(ind)+"]==" +str(true_val)
203                     feat.frange = boolean
204                     self.input_features.append(feat)
205                 elif all(isinstance(val,(int,float)) for val in frange):
206                     # all numeric, create cuts of the data
207                     sorted_frange = sorted(frange)
208                     num_cuts = min(max_num_cuts,len(frange))
209                     cut_positions = [len(frange)*i//num_cuts for i in range(1,num_cuts)]
210                     for cut in cut_positions:
211                         cutat = sorted_frange[cut]
212                         def feat(e, ind_=ind, cutat=cutat):
213                             return e[ind_] < cutat
214
215                         if self.header:
216                             feat.__doc__ = self.header[ind]+"<" +str(cutat)
217                         else:
218                             feat.__doc__ = "e["+str(ind)+"]<" +str(cutat)
219                         feat.frange = boolean
220                         self.input_features.append(feat)
221                 else:
222                     # create an indicator function for every value
223                     for val in frange:
224                         def feat(e, ind_=ind, val_=val):
225                             return e[ind_] == val_
226                         if self.header:
227                             feat.__doc__ = self.header[ind]+"==" +str(val)
228                         else:
229                             feat.__doc__ = "e["+str(ind)+"]==" +str(val)
230                         feat.frange = boolean
231                         self.input_features.append(feat)
232     else: # boolean_features is off
233         self.input_features = []
234         for i in range(self.num_properties):
235             def feat(e,index=i):

```

```

236         return e[index]
237     if self.header:
238         feat.__doc__ = self.header[i]
239     else:
240         feat.__doc__ = "e["+str(i)+"]"
241     feat.frange = ranges[i]
242     if i == self.target_index:
243         self.target = feat
244     else:
245         self.input_features.append(feat)

```

Exercise 7.1 Change the code so that it splits using $e[ind] \leq cut$ instead of $e[ind] < cut$. Check boundary cases, such as 3 elements with 2 cuts. As a test case, make sure that when the range is the 30 integers from 100 to 129, and you want 2 cuts, the resulting Boolean features should be $e[ind] \leq 109$ and $e[ind] \leq 119$ to make sure that each of the resulting ranges is equal size.

Exercise 7.2 This splits on whether the feature is less than one of the values in the training set. Sam suggested it might be better to split between the values in the training set, and suggested using

$$cutat = (sorted_frange[cut] + sorted_frange[cut - 1])/2$$

Why might Sam have suggested this? Does this work better? (Try it on a few data sets).

When reading from a file all of the values are strings. This next method tries to convert each values into a number (an int or a float) or Boolean, if it is possible.

```

_____learnProblem.py — (continued)_____
246 def interpret_elements(str_list):
247     """make the elements of string list str_list numerical if possible.
248     Otherwise remove initial and trailing spaces.
249     """
250     res = []
251     for e in str_list:
252         try:
253             res.append(int(e))
254         except ValueError:
255             try:
256                 res.append(float(e))
257             except ValueError:
258                 se = e.strip()
259                 if se in ["True", "true", "TRUE"]:
260                     res.append(True)
261                 if se in ["False", "false", "FALSE"]:
262                     res.append(False)
263             else:
264                 res.append(e.strip())
265     return res

```

7.1.5 Augmented Features

Sometimes we want to augment the features with new features computed from the old features (eg. the product of features). Here we allow the creation of a new dataset from an old dataset but with new features. Note that these are sometimes called **kernels**; mapping the original feature space into a new space, from which we can use standard learning tools. For those interested in the mathematics, read about support vector machines, which have neat way to do learning in the augmented space (the “kernel trick”) that is beyond the scope of AIPython (currently).

A feature is a function of examples. A unary feature constructor takes a feature and returns a new feature. A binary feature combiner takes two features and returns a new feature.

```

learnProblem.py — (continued)
267 class Data_set_augmented(Data_set):
268     def __init__(self, dataset, unary_functions=[], binary_functions=[], include_orig=True):
269         """creates a dataset like dataset but with new features
270         unary_function is a list of unary feature constructors
271         binary_functions is a list of binary feature combiners.
272         include_orig specifies whether the original features should be included
273         """
274         self.orig_dataset = dataset
275         self.unary_functions = unary_functions
276         self.binary_functions = binary_functions
277         self.include_orig = include_orig
278         self.target = dataset.target
279         Data_set.__init__(self, dataset.train, test=dataset.test,
280                           target_index = dataset.target_index)
281
282     def create_features(self):
283         if self.include_orig:
284             self.input_features = self.orig_dataset.input_features.copy()
285         else:
286             self.input_features = []
287         for u in self.unary_functions:
288             for f in self.orig_dataset.input_features:
289                 self.input_features.append(u(f))
290         for b in self.binary_functions:
291             for f1 in self.orig_dataset.input_features:
292                 for f2 in self.orig_dataset.input_features:
293                     if f1 != f2:
294                         self.input_features.append(b(f1,f2))

```

The following are useful unary feature constructors and binary feature combiner.

```

learnProblem.py — (continued)
296 def square(f):
297     """a unary feature constructor to construct the square of a feature

```

```

298     """
299     def sq(e):
300         return f(e)**2
301     sq.__doc__ = f.__doc__+"**2"
302     return sq
303
304 def power_feat(n):
305     """given n returns a unary feature constructor to construct the nth power of a feature.
306     e.g., power_feat(2) is the same as square, defined above
307     """
308     def fn(f,n=n):
309         def pow(e,n=n):
310             return f(e)**n
311         pow.__doc__ = f.__doc__+"**"+str(n)
312         return pow
313     return fn
314
315 def prod_feat(f1,f2):
316     """a new feature that is the product of features f1 and f2
317     """
318     def feat(e):
319         return f1(e)*f2(e)
320     feat.__doc__ = f1.__doc__+"*" + f2.__doc__
321     return feat
322
323 def eq_feat(f1,f2):
324     """a new feature that is 1 if f1 and f2 give same value
325     """
326     def feat(e):
327         return 1 if f1(e)==f2(e) else 0
328     feat.__doc__ = f1.__doc__+"==" + f2.__doc__
329     return feat
330
331 def neq_feat(f1,f2):
332     """a new feature that is 1 if f1 and f2 give different values
333     """
334     def feat(e):
335         return 1 if f1(e)!=f2(e) else 0
336     feat.__doc__ = f1.__doc__+"!=" + f2.__doc__
337     return feat

```

Example:

```

learnProblem.py — (continued)
339 # from learnProblem import Data_set_augmented,prod_feat
340 # data = Data_from_file('data/holiday.csv', num_train=19, target_index=-1)
341 ## data = Data_from_file('data/SPECT.csv', prob_test=0.5, target_index=0)
342 # dataplus = Data_set_augmented(data,[],[prod_feat])
343 # dataplus = Data_set_augmented(data,[],[prod_feat,neq_feat])

```

Exercise 7.3 For symmetric properties, such as product, we don't need both

$f1 * f2$ as well as $f2 * f1$ as extra properties. Allow the user to be able to declare feature constructors as symmetric (by associating a Boolean feature with them). Change *construct_features* so that it does not create both versions for symmetric combiners.

7.2 Generic Learner Interface

A **learner** takes a dataset (and possibly other arguments specific to the method). To get it to learn, we call the *learn()* method. This implements *Displayable* so that we can display traces at multiple levels of detail (and perhaps with a GUI).

```

learnProblem.py — (continued)
344 from display import Displayable
345
346 class Learner(Displayable):
347     def __init__(self, dataset):
348         raise NotImplementedError("Learner.__init__") # abstract method
349
350     def learn(self):
351         """returns a predictor, a function from a tuple to a value for the target feature
352         """
353         raise NotImplementedError("learn") # abstract method

```

7.3 Learning With No Input Features

If we make the same prediction for each example, what prediction should we make? This can be used as a naive baseline; if a method does not do better than this, the the input features do not provide any useful information for the prediction. It is also the base case for some methods, such as decision-tree learning.

To run demo to compare different prediction methods on various evaluation criteria, in folder "aipython", load "learnNoInputs.py", using e.g., `ipython -i learnNoInputs.py`, and it prints some test results.

There are a few alternatives as to what could be allowed in a prediction:

- a point prediction, where we are only allowed to predict one of the values of the feature. For example, if the values of the feature are $\{0, 1\}$ we are only allowed to predict 0 or 1 or of the values are ratings in $\{1, 2, 3, 4, 5\}$, we can only predict one of these integers.
- a point prediction, where we are allowed to predict any value. For example, if the values of the feature are $\{0, 1\}$ we may be allowed to predict 0.3, 1, or even 1.7. For all of the criteria we can imagine, there is no point in

predicting a value greater than 1 or less than zero (but that doesn't mean we can't), but it is often useful to predict a value between 0 and 1. If the values are ratings in $\{1, 2, 3, 4, 5\}$, we may want to predict 3.4.

- a probability distribution over the values of the feature. For each value v , we predict a non-negative number p_v , such that the sum over all predictions is 1.

The following code assumes the second of these, where we can make a point prediction of any value (although median will only predict one of the actual values for the feature). The third can be implemented by having multiple indicator functions for the target.

Here are some prediction functions that take in a dataset of number and returns a prediction for the next case. Note that median will average the two middle values when there are an even number of examples. The mode will pick one of the values arbitrarily (here the larger) when more than one value has the maximum number of elements. So the median of $[0, 1]$ is 0.5, but the mode is 1.

```

learnNoInputs.py — Learning ignoring all input features
11 from learnProblem import squared_error, absolute_error, log_loss, mean
12 import math, random, statistics
13 import utilities # argmax for (element,value) pairs
14
15 class Predict(object):
16     """The class of prediction methods for a list of numbers
17     Please make the doc strings the same length, because they are used in tables.
18     Note that we don't need self argument, as we are creating Predict objects,
19     To use call Predict.laplace(data) etc."""
20
21     def mean(data):
22         "mean"
23         return mean(data)
24
25     def bounded_mean(data, bound=0.01):
26         "bounded mean"
27         return min(max(mean(data), bound), 1-bound)
28
29     def laplace(data):
30         "Laplace" # for Boolean (or 0/1 data only)
31         return mean(data, isum=1, icount=2)
32
33     def mode(data):
34         "mode"
35         counts = {}
36         for e in data:
37             if e in counts:
38                 counts[e] += 1
39         else:

```

```

40         counts[e] = 1
41     return utilities.argmaxx(counts.items())
42
43     def median(data):
44         "median"
45         return statistics.median(data)
46
47     all = [mean, bounded_mean, laplace, mode, median]

```

7.3.1 Evaluation

To evaluate a point prediction, we first generate some data from a simple (Bernoulli) distribution, where there are two possible values, 0 and 1 for the target feature. Given *prob*, a number in the range [0,1], this generate some training and test data where *prob* is the probability of each example being 1. To generate a 1 with probability *prob*, we generate a random number in range [0,1] and return 1 if that number is less than *prob*.

```

learnNoInputs.py — (continued)
49 def evaluate(train_size, predictor, error_measure, num_samples=10000, test_size=10 ):
50     """return the average error when
51     train_size is the number of training examples
52     predictor(training) -> [0,1]
53     error_measure(prediction,actual) -> non-negative reals
54     """
55     error = 0
56     for sample in range(num_samples):
57         prob = random.random()
58         training = [1 if random.random()<prob else 0 for i in range(train_size)]
59         prediction = predictor(training)
60         test = (1 if random.random()<prob else 0 for i in range(test_size))
61         error += sum( error_measure(prediction,actual) for actual in test)/test_size
62     return error/num_samples

```

Let's evaluate the predictions of the possible selections according to the different evaluation criteria, for various training sizes.

```

learnNoInputs.py — (continued)
64 def test_no_inputs(error_measures = [squared_error, absolute_error, log_loss]):
65     for train_size in [1,2,3,4,5,10,20,100,1000]:
66         print("For training size",train_size,":")
67         print("  Predictor","\t".join(error_measure.__doc__ for
68                                     error_measure in error_measures),sep="\t")
69         for predictor in Predict.all:
70             print(f"  {predictor.__doc__}",
71                   "\t".join("{:.7f}".format(evaluate(train_size, predictor, error_measure))
72                             for error_measure in error_measures),sep="\t")
73
74 if __name__ == "__main__":
75     test_no_inputs()

```


Exercise 7.4 Which predictor works best for low counts when the error is

- (a) Squared error
- (b) Absolute error
- (c) Log loss

You may need to try this a few times to make sure your answer is supported by the evidence. Does the difference from the other methods get more or less as the number of examples grow?

Exercise 7.5 Suggest some other predictions that only take the training data. Does your method do better than the given methods? A simple way to get other predictors is to vary the threshold of bounded average, or to change the pseudo-counts of the Laplace method (use other numbers instead of 1 and 2).

7.4 Decision Tree Learning

To run the decision tree learning demo, in folder "aipython", load "learnDT.py", using e.g., `ipython -i learnDT.py`, and it prints some test results. To try more examples, copy and paste the commented-out commands at the bottom of that file. This requires Python 3 with matplotlib.

The decision tree algorithm does binary splits, and assumes that all input features are binary functions of the examples. It stops splitting if there are no input features, the number of examples is less than a specified number of examples or all of the examples agree on the target feature.

```

learnDT.py — Learning a binary decision tree
11 from learnProblem import Learner, squared_error, absolute_error, log_loss, mean
12 from learnNoInputs import Predict
13 import math
14
15 class DT_learner(Learner):
16     def __init__(self,
17                   dataset,
18                   split_to_optimize=log_loss,      # to minimize for at each split
19                   leaf_prediction=Predict.mean,   # what to use for point prediction at leaves
20                   train=None,                     # used for cross validation
21                   min_number_examples=10):
22         self.dataset = dataset
23         self.target = dataset.target
24         self.split_to_optimize = split_to_optimize
25         self.leaf_prediction = leaf_prediction
26         self.min_number_examples = min_number_examples
27         if train is None:
28             self.train = self.dataset.train
29         else:
30             self.train = train

```

```

31
32     def learn(self):
33         return self.learn_tree(self.dataset.input_features, self.train)

```

The main recursive algorithm, takes in a set of input features and a set of training data. It first decides whether to split. If it doesn't split, it makes a point prediction, ignoring the input features.

It splits unless:

- there are no more input features
- there are fewer examples than *min_number_examples*,
- all the examples agree on the value of the target, or
- the best split makes all examples in the same partition.

If it splits, it selects the best split according to the evaluation critereon (assuming that is the only split it gets to do), and returns the condition to split on (in the variable *split*) and the corresponding partition of the examples.

```

learnDT.py — (continued)
35     def learn_tree(self, input_features, data_subset):
36         """returns a decision tree
37         for input_features is a set of possible conditions
38         data_subset is a subset of the data used to build this (sub)tree
39
40         where a decision tree is a function that takes an example and
41         makes a prediction on the target feature
42         """
43         if (input_features and len(data_subset) >= self.min_number_examples):
44             first_target_val = self.target(data_subset[0])
45             allagree = all(self.target(inst)==first_target_val for inst in data_subset)
46             if not allagree:
47                 split, partn = self.select_split(input_features, data_subset)
48                 if split: # the split succeeded in splitting the data
49                     false_examples, true_examples = partn
50                     rem_features = [fe for fe in input_features if fe != split]
51                     self.display(2,"Splitting on",split.__doc__,"with examples split",
52                                 len(true_examples),":",len(false_examples))
53                     true_tree = self.learn_tree(rem_features,true_examples)
54                     false_tree = self.learn_tree(rem_features,false_examples)
55                     def fun(e):
56                         if split(e):
57                             return true_tree(e)
58                         else:
59                             return false_tree(e)
60                     #fun = lambda e: true_tree(e) if split(e) else false_tree(e)
61                     fun.__doc__ = ("if "+split.__doc__+" then (" +true_tree.__doc__+"
62                                     ") else (" +false_tree.__doc__+"")
63                     return fun

```

```

64         # don't expand the trees but return a point prediction
65         prediction = self.leaf_prediction(self.target(e) for e in data_subset)
66         def leaf_fun(e):
67             return prediction
68         leaf_fun.__doc__ = "{:.7f}".format(prediction)
69         return leaf_fun

```

```

learnDT.py — (continued)
71     def select_split(self, input_features, data_subset):
72         """finds best feature to split on.
73
74         input_features is a non-empty list of features.
75         returns feature, partition
76         where feature is an input feature with the smallest error as
77             judged by split_to_optimize or
78             feature==None if there are no splits that improve the error
79         partition is a pair (false_examples, true_examples) if feature is not None
80         """
81         best_feat = None # best feature
82         # best_error = float("inf") # infinity - more than any error
83         best_error = training_error(self.dataset.target, data_subset,
84                                     self.split_to_optimize, self.leaf_prediction)
85         best_partition = None
86         for feat in input_features:
87             false_examples, true_examples = partition(data_subset, feat)
88             if false_examples and true_examples: #both partitons are non-empty
89                 err = (training_error(self.dataset.target, false_examples,
90                                     self.split_to_optimize, self.leaf_prediction)
91                       + training_error(self.dataset.target, true_examples,
92                                     self.split_to_optimize, self.leaf_prediction))
93                 self.display(3, " split on", feat.__doc__, "has error=", err,
94                             "splits into", len(true_examples), ":", len(false_examples))
95                 if err < best_error:
96                     best_feat = feat
97                     best_error=err
98                     best_partition = false_examples, true_examples
99                 self.display(3, "best split is on", best_feat.__doc__,
100                             "with err=", best_error)
101         return best_feat, best_partition
102
103     def partition(data_subset, feature):
104         """partitions the data_subset by the feature"""
105         true_examples = []
106         false_examples = []
107         for example in data_subset:
108             if feature(example):
109                 true_examples.append(example)
110             else:
111                 false_examples.append(example)
112         return false_examples, true_examples

```

```

113
114
115 def training_error(target, data_subset, eval_critereon, leaf_prediction):
116     """returns training error for dataset on the target (with no more splits)
117     We make a single prediction using leaf_prediction
118     It is evaluated using eval_critereon for each example
119     """
120     prediction = leaf_prediction(target(e) for e in data_subset)
121     error = sum(eval_critereon(prediction, target(e))
122                 for e in data_subset)
123     return error

```

Test cases:

```

learnDT.py — (continued)
125 from learnProblem import Data_set, Data_from_file
126
127 def testDT(data, print_tree=True, selections = Predict.all):
128     """Prints errors and the trees for various evaluation criteria and ways to select leaves.
129     """
130     evaluation_criteria = [squared_error, absolute_error, log_loss]
131     print("Split Choice", "Leaf Choice", '\t'.join(ecrit.__doc__
132                                                    for ecrit in evaluation_criteria), sep="\t")
133     for crit in evaluation_criteria:
134         for leaf in selections:
135             tree = DT_learner(data, split_to_optimize=crit, leaf_prediction=leaf).learn()
136             print(crit.__doc__, leaf.__doc__,
137                   "\t".join("{:.7f}".format(data.evaluate_dataset(data.test, tree, ecrit))
138                             for ecrit in evaluation_criteria), sep="\t")
139             if print_tree:
140                 print(tree.__doc__)
141
142 if __name__ == "__main__":
143     print("SPECT.csv"); testDT(data=Data_from_file('data/SPECT.csv', target_index=0), print_tree=False)
144     # print("carbool.csv"); testDT(data = Data_from_file('data/carbool.csv', target_index=-1))
145     # print("mail_reading.csv"); testDT(data = Data_from_file('data/mail_reading.csv', target_index=-1))
146     # print("holiday.csv"); testDT(data = Data_from_file('data/holiday.csv', num_train=19, target_index=-1))

```

Note that different runs may provide different values as they split the training and test sets differently. So if you have a hypothesis about what works better, make sure it is true for different runs.

Exercise 7.6 The current algorithm does not have a very sophisticated stopping criterion. What is the current stopping criterion? (Hint: you need to look at both *learn.tree* and *select.split*.)

Exercise 7.7 Extend the current algorithm to include in the stopping criterion

- (a) A minimum child size; don't use a split if one of the children has fewer elements than this.
- (b) A depth-bound on the depth of the tree.

- (c) An improvement bound such that a split is only carried out if error with the split is better than the error without the split by at least the improvement bound.

Which values for these parameters make the prediction errors on the test set the smallest? Try it on more than one dataset.

Exercise 7.8 Without any input features, it is often better to include a pseudo-count that is added to the counts from the training data. Modify the code so that it includes a pseudo-count for the predictions. When evaluating a split, including pseudo counts can make the split worse than no split. Does pruning with an improvement bound and pseudo-counts make the algorithm work better than with an improvement bound by itself?

Exercise 7.9 Some people have suggested using information gain (which is equivalent to greedy optimization of logloss) as the measure of improvement when building the tree, even in they want to have non-probabilistic predictions in the final tree. Does this work better than myopically choosing the split that is best for the evaluation criteria we will use to judge the final prediction?

7.5 Cross Validation and Parameter Tuning

To run the cross validation demo, in folder "aipython", load "learnCrossValidation.py", using e.g., `ipython -i learnCrossValidation.py`. Run the examples at the end to produce a graph like Figure 7.15. Note that different runs will produce different graphs, so your graph will not look like the one in the textbook. To try more examples, copy and paste the commented-out commands at the bottom of that file. This requires Python 3 with matplotlib.

The above decision tree overfits the data. One way to determine whether the prediction is overfitting is by cross validation. The code below implements k -fold cross validation, which can be used to choose the value of parameters to best fit the training data. If we want to use parameter tuning to improve predictions on a particular data set, we can only use the training data (and not the test data) to tune the parameter.

In k -fold cross validation, we partition the training set into k approximately equal-sized folds (each fold is an enumeration of examples). For each fold, we train on the other examples, and determine the error of the prediction on that fold. For example, if there are 10 folds, we train on 90% of the data, and then test on remaining 10% of the data. We do this 10 times, so that each example gets used as a test set once, and in the training set 9 times.

The code below creates one copy of the data, and multiple views of the data. For each fold, *fold* enumerates the examples in the fold, and *fold_complement* enumerates the examples not in the fold.

```

learnCrossValidation.py — Cross Validation for Parameter Tuning
11 from learnProblem import Data_set, Data_from_file, squared_error, absolute_error, log_loss
12 from learnDT import DT_learner
13 import matplotlib.pyplot as plt
14 import random
15
16 class K_fold_dataset(object):
17     def __init__(self, training_set, num_folds):
18         self.data = training_set.train.copy()
19         self.target = training_set.target
20         self.input_features = training_set.input_features
21         self.num_folds = num_folds
22         random.shuffle(self.data)
23         self.fold_boundaries = [(len(self.data)*i)//num_folds
24                                for i in range(0,num_folds+1)]
25
26     def fold(self, fold_num):
27         for i in range(self.fold_boundaries[fold_num],
28                        self.fold_boundaries[fold_num+1]):
29             yield self.data[i]
30
31     def fold_complement(self, fold_num):
32         for i in range(0,self.fold_boundaries[fold_num]):
33             yield self.data[i]
34         for i in range(self.fold_boundaries[fold_num+1],len(self.data)):
35             yield self.data[i]

```

The validation error is the average error for each example, where we test on each fold, and learn on the other folds.

```

learnCrossValidation.py — (continued)
37 def validation_error(self, learner, error_measure, **other_params):
38     error = 0
39     try:
40         for i in range(self.num_folds):
41             predictor = learner(self, train=list(self.fold_complement(i)),
42                                **other_params).learn()
43             error += sum( error_measure(predictor(e), self.target(e))
44                           for e in self.fold(i))
45     except ValueError:
46         return float("inf") #infinity
47     return error/len(self.data)

```

The *plot_error* method plots the average error as a function of a the minimum number of examples in decision-tree search, both for the validation set and for the test set. The error on the validation set can be used to tune the parameter — choose the value of the parameter that minimizes the error. The error on the test set cannot be used to tune the parameters; if it were to be used this way then it cannot be used to test.

```

learnCrossValidation.py — (continued)

```

```

49 def plot_error(data, criterion=squared_error, num_folds=5, xscale='linear'):
50     """Plots the error on the validation set and the test set
51     with respect to settings of the minimum number of examples.
52     xscale should be 'log' or 'linear'
53     """
54     plt.ion()
55     plt.xscale(xscale) # change between log and linear scale
56     plt.xlabel("minimum number of examples")
57     plt.ylabel("average "+criterion.__doc__)
58     folded_data = K_fold_dataset(data, num_folds)
59     verrors = [] # validation errors
60     terrors = [] # test set errors
61     for mne in range(1, len(data.train)+2):
62         verrors.append(folded_data.validation_error(DT_learner, criterion,
63                                                     min_number_examples=mne))
64         tree = DT_learner(data, criterion, min_number_examples=mne).learn()
65         terrors.append(data.evaluate_dataset(data.test, tree, criterion))
66     plt.plot(range(1, len(data.train)+2), verrors, ls='-', color='k',
67             label="validation for "+criterion.__doc__)
68     plt.plot(range(1, len(data.train)+2), terrors, ls='--', color='k',
69             label="test set for "+criterion.__doc__)
70     plt.legend()
71     plt.draw()
72
73 # The following produces Figure 7.15 of Poole and Mackworth [2017]
74 # Different runs produce different plots, because folds change.
75 # data = Data_from_file('data/SPECT.csv', target_index=0)
76 # plot_error(data) # warning, may take a long time depending on the dataset
77
78 #also try:
79 # data = Data_from_file('data/mail_reading.csv', target_index=-1)
80 # data = Data_from_file('data/carbool.csv', target_index=-1)

```

Note that different runs for the same data will have the same test error, but different validation error. If you rerun the `Data_from_file`, you will get the new test and training sets, and so the graph will change.

Exercise 7.10 Change the error plot so that it can evaluate the stopping criteria of the exercise of Section 7.6. Which criteria makes the most difference?

7.6 Linear Regression and Classification

Here we give a gradient descent searcher for linear regression and classification.

```

learnLinear.py — Linear Regression and Classification
11 from learnProblem import Learner
12 import random, math
13
14 class Linear_learner(Learner):

```

```

15     def __init__(self, dataset, train=None,
16                   learning_rate=0.1, max_init = 0.2,
17                   squashed=True):
18         """Creates a gradient descent searcher for a linear classifier.
19         The main learning is carried out by learn()
20
21         dataset provides the target and the input features
22         train provides a subset of the training data to use
23         number_iterations is the default number of steps of gradient descent
24         learning_rate is the gradient descent step size
25         max_init is the maximum absolute value of the initial weights
26         squashed specifies whether the output is a squashed linear function
27         """
28         self.dataset = dataset
29         self.target = dataset.target
30         if train==None:
31             self.train = self.dataset.train
32         else:
33             self.train = train
34         self.learning_rate = learning_rate
35         self.squashed = squashed
36         self.input_features = [one]+dataset.input_features # one is defined below
37         self.weights = {feat:random.uniform(-max_init,max_init)
38                         for feat in self.input_features}

```

predictor predicts the value of an example from the current parameter settings.
predictor_string gives a string representation of the predictor.

learnLinear.py — (continued)

```

40
41     def predictor(self,e):
42         """returns the prediction of the learner on example e"""
43         linpred = sum(w*f(e) for f,w in self.weights.items())
44         if self.squashed:
45             return sigmoid(linpred)
46         else:
47             return linpred
48
49     def predictor_string(self, sig_dig=3):
50         """returns the doc string for the current prediction function
51         sig_dig is the number of significant digits in the numbers"""
52         doc = "+".join(str(round(val,sig_dig))+"*"+feat.__doc__
53                        for feat,val in self.weights.items())
54         if self.squashed:
55             return "sigmoid("+ doc+")"
56         else:
57             return doc

```

learn is the main algorithm of the learner. It does *num_iter* steps of stochastic gradient descent with batch size = 1. The other parameters it gets from the class.


```

learnLinear.py — (continued)
59 def learn(self,num_iter=100):
60     for it in range(num_iter):
61         self.display(2,"prediction=",self.predictor_string())
62         for e in self.train:
63             predicted = self.predictor(e)
64             error = self.target(e) - predicted
65             update = self.learning_rate*error
66             for feat in self.weights:
67                 self.weights[feat] += update*feat(e)
68             #self.predictor.__doc__ = self.predictor_string()
69             #return self.predictor

```

one is a function that always returns 1. This is used for one of the input properties.

```

learnLinear.py — (continued)
71 def one(e):
72     "1"
73     return 1

```

sigmoid(*x*) is the function

$$\frac{1}{1 + e^{-x}}$$

```

learnLinear.py — (continued)
75 def sigmoid(x):
76     return 1/(1+math.exp(-x))

```

The following tests the learner on a data sets. Uncomment the other data sets for different examples.

```

learnLinear.py — (continued)
78 from learnProblem import Data_set, Data_from_file, evaluation_criteria
79 import matplotlib.pyplot as plt
80
81 def test(**args):
82     data = Data_from_file('data/SPECT.csv', target_index=0)
83     # data = Data_from_file('data/mail_reading.csv', target_index=-1)
84     # data = Data_from_file('data/carbool.csv', target_index=-1)
85     learner = Linear_learner(data,**args)
86     learner.learn()
87     print("function learned is", learner.predictor_string())
88     for ecrit in evaluation_criteria:
89         test_error = data.evaluate_dataset(data.test, learner.predictor, ecrit)
90         print("    Average", ecrit.__doc__, "error is", test_error)

```

The following plots the errors on the training and test sets as a function of the number of steps of gradient descent.

```

learnLinear.py — (continued)
92 def plot_steps(learner=None,
93               data = None,
94               criterion="sum-of-squares",
95               step=1,
96               num_steps=1000,
97               log_scale=True,
98               label=""):
99     """
100     plots the training and test error for a learner.
101     data is the
102     learner_class is the class of the learning algorithm
103     criterion gives the evaluation criterion plotted on the y-axis
104     step specifies how many steps are run for each point on the plot
105     num_steps is the number of points to plot
106
107     """
108     plt.ion()
109     plt.xlabel("step")
110     plt.ylabel("Average "+criterion+" error")
111     if log_scale:
112         plt.xscale('log') #plt.semilogx() #Makes a log scale
113     else:
114         plt.xscale('linear')
115     if data is None:
116         data = Data_from_file('data/holiday.csv', num_train=19, target_index=-1)
117         #data = Data_from_file('data/SPECT.csv', target_index=0)
118         # data = Data_from_file('data/mail_reading.csv', target_index=-1)
119         # data = Data_from_file('data/carbool.csv', target_index=-1)
120     random.seed(None) # reset seed
121     if learner is None:
122         learner = Linear_learner(data)
123     train_errors = []
124     test_errors = []
125     for i in range(1,num_steps+1,step):
126         test_errors.append(data.evaluate_dataset(data.test, learner.predictor, criterion))
127         train_errors.append(data.evaluate_dataset(data.train, learner.predictor, criterion))
128         learner.display(2, "Train error:",train_errors[-1],
129                        "Test error:",test_errors[-1])
130         learner.learn(num_iter=step)
131     plt.plot(range(1,num_steps+1,step),train_errors,ls='-',c='k',label="training errors")
132     plt.plot(range(1,num_steps+1,step),test_errors,ls='--',c='k',label="test errors")
133     plt.legend()
134     plt.draw()
135     learner.display(1, "Train error:",train_errors[-1],
136                   "Test error:",test_errors[-1])
137
138 if __name__ == "__main__":
139     test()
140

```

```

141 # This generates the figure
142 # from learnProblem import Data_set_augmented, prod_feat
143 # data = Data_from_file('data/SPECT.csv', prob_test=0.5, target_index=0)
144 # dataplus = Data_set_augmented(data, [], [prod_feat])
145 # plot_steps(data=data, num_steps=10000)
146 # plot_steps(data=dataplus, num_steps=10000) # warning very slow

```

Exercise 7.11 The squashed learner only makes predictions in the range (0,1). If the output values are {1,2,3,4} there is no use prediction less than 1 or greater than 4. Change the squashed learner so that it can learn values in the range (1,4). Test it on the file 'data/car.csv'.

The following plots the prediction as a function of the function of the number of steps of gradient descent. We first define a version of *range* that allows for real numbers (integers and floats).

```

learnLinear.py — (continued) —
147 def arange(start, stop, step):
148     """returns enumeration of values in the range [start, stop) separated by step.
149     like the built-in range(start, stop, step) but allows for integers and floats.
150     Note that rounding errors are expected with real numbers. (or use numpy.arange)
151     """
152     while start < stop:
153         yield start
154         start += step
155
156 def plot_prediction(learner=None,
157                    data = None,
158                    minx = 0,
159                    maxx = 5,
160                    step_size = 0.01, # for plotting
161                    label="function"):
162     plt.ion()
163     plt.xlabel("x")
164     plt.ylabel("y")
165     if data is None:
166         data = Data_from_file('data/simp_regr.csv', prob_test=0,
167                               boolean_features=False, target_index=-1)
168     if learner is None:
169         learner = Linear_learner(data, squashed=False)
170     learner.learning_rate=0.001
171     learner.learn(100)
172     learner.learning_rate=0.0001
173     learner.learn(1000)
174     learner.learning_rate=0.00001
175     learner.learn(10000)
176     learner.display(1, "function learned is", learner.predictor_string(),
177                    "error=", data.evaluate_dataset(data.train, learner.predictor, "sum-of-squares"))
178     plt.plot([e[0] for e in data.train], [e[-1] for e in data.train], "bo", label="data")
179     plt.plot(list(arange(minx, maxx, step_size)), [learner.predictor([x])
180                                                    for x in arange(minx, maxx, step_size)],

```

```

181 |                                     label=label)
182 |     plt.legend()
183 |     plt.draw()

```

learnLinear.py — (continued) —

```

185 | from learnProblem import Data_set_augmented, power_feat
186 | def plot_polynomials(data=None,
187 |                     learner_class = Linear_learner,
188 |                     max_degree=5,
189 |                     minx = 0,
190 |                     maxx = 5,
191 |                     num_iter = 100000,
192 |                     learning_rate = 0.0001,
193 |                     step_size = 0.01, # for plotting
194 |                     ):
195 |     plt.ion()
196 |     plt.xlabel("x")
197 |     plt.ylabel("y")
198 |     if data is None:
199 |         data = Data_from_file('data/simp_regr.csv', prob_test=0,
200 |                               boolean_features=False, target_index=-1)
201 |     plt.plot([e[0] for e in data.train],[e[-1] for e in data.train],"ko",label="data")
202 |     x_values = list(arange(minx,maxx,step_size))
203 |     line_styles = ['-','--','-.',':']
204 |     colors = ['0.5','k','k','k','k']
205 |     for degree in range(max_degree):
206 |         data_aug = Data_set_augmented(data,[power_feat(n) for n in range(1,degree+1)],
207 |                                       include_orig=False)
208 |         learner = learner_class(data_aug,squashed=False)
209 |         learner.learning_rate=learning_rate
210 |         learner.learn(num_iter)
211 |         learner.display(1,"For degree",degree,
212 |                         "function learned is", learner.predictor_string(),
213 |                         "error=",data.evaluate_dataset(data.train, learner.predictor, "sum-of-squares"))
214 |         ls = line_styles[degree % len(line_styles)]
215 |         col = colors[degree % len(colors)]
216 |         plt.plot(x_values,[learner.predictor([x]) for x in x_values], linestyle=ls, color=col,
217 |                  label="degree="+str(degree))
218 |         plt.legend(loc='upper left')
219 |         plt.draw()
220 |
221 | # Try:
222 | # plot_prediction()
223 | # plot_polynomials()
224 | #data = Data_from_file('data/mail_reading.csv', target_index=-1)
225 | #plot_prediction(data=data)

```

7.6.1 Batched Stochastic Gradient Descent

This implements batched stochastic gradient descent. If the batch size is 1, it can be simplified by not storing the differences in d , but applying them directly; this would be equivalent to the original code!

This overrides the learner *Linear_learner*. Note that the comparison with regular gradient descent is unfair as the number of updates per step is not the same. (How could it be made more fair?)

```

learnLinearBSGD.py — Linear Learner with Batched Stochastic Gradient Descent
11 from learnLinear import Linear_learner
12 import random, math
13
14 class Linear_learner_bsgd(Linear_learner):
15     def __init__(self, *args, batch_size=10, **kwargs):
16         Linear_learner.__init__(self, *args, **kwargs)
17         self.batch_size = batch_size
18
19     def learn(self, num_iter=None):
20         if num_iter is None:
21             num_iter = self.number_iterations
22         batch_size = min(self.batch_size, len(self.train))
23         d = {feat:0 for feat in self.weights}
24         for it in range(num_iter):
25             self.display(2, "prediction=", self.predictor_string())
26             for e in random.sample(self.train, batch_size):
27                 predicted = self.predictor(e)
28                 error = self.target(e) - predicted
29                 update = self.learning_rate*error
30                 for feat in self.weights:
31                     d[feat] += update*feat(e)
32             for feat in self.weights:
33                 self.weights[feat] += d[feat]
34                 d[feat]=0
35
36 # from learnLinear import plot_steps
37 # from learnProblem import Data_from_file
38 # data = Data_from_file('data/holiday.csv', target_index=-1)
39 # learner = Linear_learner_bsgd(data)
40 # plot_steps(learner = learner, data=data)
41
42 # to plot polynomials with batching (compare to SGD)
43 # from learnLinear import plot_polynomials
44 # plot_polynomials(learner_class = Linear_learner_bsgd)

```

7.7 Deep Neural Network Learning

This provides a modular implementation that implements the layers modularly. Layers can easily be configured in many configurations. A layer needs to

implement a function to compute the output values from the inputs and a way to back-propagate the error.

```

learnNN.py — Neural Network Learning
11 from learnProblem import Learner, Data_set, Data_from_file
12 from learnLinear import sigmoid, one
13 import random, math
14
15 class Layer(object):
16     def __init__(self, nn, num_outputs=None):
17         """Given a list of inputs, outputs will produce a list of length num_outputs.
18         nn is the neural network this is part of
19         num_outputs is the number of outputs for this layer.
20         """
21         self.nn = nn
22         self.num_inputs = nn.num_outputs # output of nn is the input to this layer
23         if num_outputs:
24             self.num_outputs = num_outputs
25         else:
26             self.num_outputs = nn.num_outputs # same as the inputs
27
28     def output_values(self, input_values):
29         """Return the outputs for this layer for the given input values.
30         input_values is a list of the inputs to this layer (of length num_inputs)
31         returns a list of length self.num_outputs
32         """
33         raise NotImplementedError("output_values") # abstract method
34
35     def backprop(self, errors):
36         """Backpropagate the errors on the outputs, return the errors on the inputs.
37         errors is a list of errors for the outputs (of length self.num_outputs).
38         Return the errors for the inputs to this layer (of length self.num_inputs).
39         You can assume that this is only called after corresponding output_values,
40         and it can remember information information required for the back-propagation.
41         """
42         raise NotImplementedError("backprop") # abstract method

```

A linear layer maintains an array of weights. *self.weights[o][i]* is the weight between input *i* and output *o*. A 1 is added to the inputs.

```

learnNN.py — (continued)
44 class Linear_complete_layer(Layer):
45     """a completely connected layer"""
46     def __init__(self, nn, num_outputs, max_init=0.2):
47         """A completely connected linear layer.
48         nn is a neural network that the inputs come from
49         num_outputs is the number of outputs
50         max_init is the maximum value for random initialization of parameters
51         """
52         Layer.__init__(self, nn, num_outputs)
53         # self.weights[o][i] is the weight between input i and output o

```

```

54         self.weights = [[random.uniform(-max_init, max_init)
55                             for inf in range(self.num_inputs+1)]
56                             for outf in range(self.num_outputs)]
57
58     def output_values(self, input_values):
59         """Returns the outputs for the input values.
60         It remembers the values for the backprop.
61
62         Note in self.weights there is a weight list for every output,
63         so wts in self.weights effectively loops over the outputs.
64         """
65         self.inputs = input_values + [1]
66         return [sum(w*val for (w,val) in zip(wts,self.inputs))
67                 for wts in self.weights]
68
69     def backprop(self, errors):
70         """Backpropagate the errors, updating the weights and returning the error in its inputs.
71         """
72         input_errors = [0]*(self.num_inputs+1)
73         for out in range(self.num_outputs):
74             for inp in range(self.num_inputs+1):
75                 input_errors[inp] += self.weights[out][inp] * errors[out]
76                 self.weights[out][inp] += self.nn.learning_rate * self.inputs[inp] * errors[out]
77         return input_errors[:-1] # remove the error for the "1"

```

learnNN.py — (continued)

```

79 class Sigmoid_layer(Layer):
80     """sigmoids of the inputs.
81     The number of outputs is equal to the number of inputs.
82     Each output is the sigmoid of its corresponding input.
83     """
84     def __init__(self, nn):
85         Layer.__init__(self, nn)
86
87     def output_values(self, input_values):
88         """Returns the outputs for the input values.
89         It remembers the output values for the backprop.
90         """
91         self.outputs= [sigmoid(inp) for inp in input_values]
92         return self.outputs
93
94     def backprop(self, errors):
95         """Returns the derivative of the errors"""
96         return [e*out*(1-out) for e,out in zip(errors, self.outputs)]

```

learnNN.py — (continued)

```

98 class ReLU_layer(Layer):
99     """Rectified linear unit (ReLU)  $f(z) = \max(0, z)$ .
100     The number of outputs is equal to the number of inputs.
101     """

```

```

102     def __init__(self, nn):
103         Layer.__init__(self, nn)
104
105     def output_values(self, input_values):
106         """Returns the outputs for the input values.
107         It remembers the input values for the backprop.
108         """
109         self.input_values = input_values
110         self.outputs = [max(0, inp) for inp in input_values]
111         return self.outputs
112
113     def backprop(self, errors):
114         """Returns the derivative of the errors"""
115         return [e if inp > 0 else 0 for e, inp in zip(errors, self.input_values)]

```

learnNN.py — (continued)

```

117 class NN(Learner):
118     def __init__(self, dataset, learning_rate=0.1):
119         self.dataset = dataset
120         self.learning_rate = learning_rate
121         self.input_features = dataset.input_features
122         self.num_outputs = len(self.input_features)
123         self.layers = []
124
125     def add_layer(self, layer):
126         """add a layer to the network.
127         Each layer gets values from the previous layer.
128         """
129         self.layers.append(layer)
130         self.num_outputs = layer.num_outputs
131
132     def predictor(self, ex):
133         """Predicts the value of the first output feature for example ex.
134         """
135         values = [f(ex) for f in self.input_features]
136         for layer in self.layers:
137             values = layer.output_values(values)
138         return values[0]
139
140     def predictor_string(self):
141         return "not implemented"

```

The *test* method learns a network and evaluates it according to various criteria.

learnNN.py — (continued)

```

143
144     def learn(self, num_iter):
145         """Learns parameters for a neural network using stochastic gradient decent.
146         num_iter is the number of iterations
147         """
148         for i in range(num_iter):

```



```

149         for e in random.sample(self.dataset.train, len(self.dataset.train)):
150             # compute all outputs
151             values = [f(e) for f in self.input_features]
152             for layer in self.layers:
153                 values = layer.output_values(values)
154             # backpropagate
155             errors = self.sum_squares_error([self.dataset.target(e)], values)
156             for layer in reversed(self.layers):
157                 errors = layer.backprop(errors)
158
159     def sum_squares_error(self, observed, predicted):
160         """Returns the errors for each of the target features.
161         """
162         return [obsd-pred for obsd, pred in zip(observed, predicted)]

```

This constructs a neural network consisting of neural network with one hidden layer. The hidden using used a ReLU activation function. The output layer used a sigmoid.

```

learnNN.py — (continued)
165 data = Data_from_file('data/mail_reading.csv', target_index=-1)
166 #data = Data_from_file('data/mail_reading_consis.csv', target_index=-1)
167 #data = Data_from_file('data/SPECT.csv', prob_test=0.5, target_index=0)
168 #data = Data_from_file('data/holiday.csv', target_index=-1) #, num_train=19)
169 nn1 = NN(data)
170 nn1.add_layer(Linear_complete_layer(nn1,3))
171 nn1.add_layer(Sigmoid_layer(nn1)) # comment this or the next
172 # nn1.add_layer(ReLU_layer(nn1))
173 nn1.add_layer(Linear_complete_layer(nn1,1))
174 nn1.add_layer(Sigmoid_layer(nn1))
175 nn1.learning_rate=0.1
176 #nn1.learn(100)
177
178 from learnLinear import plot_steps
179 import time
180 start_time = time.perf_counter()
181 plot_steps(learner = nn1, data = data, num_steps=10000)
182 for eg in data.train:
183     print(eg, nn1.predictor(eg))
184 end_time = time.perf_counter()
185 print("Time:", end_time - start_time)

```

Exercise 7.12 In the definition of *nn1* above, for each of the following, first hypothesize what will happen, then test your hypothesis, then explain whether you testing confirms your hypothesis or not. Test it for more than one data set, and use more than one run for each data set.

- Which fits the data better, having a sigmoid layer or a ReLU layer after the first linear layer?
- Which is faster, having a sigmoid layer or a ReLU layer after the first linear layer?

- (c) What happens if you have both the sigmoid layer and then a ReLU layer after the first linear layer and before the second linear layer?
- (d) What happens if you have neither the sigmoid layer nor a ReLU layer after the first linear layer?
- (e) What happens if you have a ReLU layer then a sigmoid layer after the first linear layer and before the second linear layer?

Exercise 7.13 Do some

It is even possible to define a perceptron layer. Warning: you may need to change the learning rate to make this work. Should I add it into the code? It doesn't follow the official line.

```
class PerceptronLayer(Layer):
    def __init__(self, nn):
        Layer.__init__(self, nn)

    def output_values(self, input_values):
        """Returns the outputs for the input values.
        """
        self.outputs= [1 if inp>0 else -1 for inp in input_values]
        return self.outputs

    def backprop(self, errors):
        """Pass the errors through"""
        return errors
```

7.8 Boosting

The following code implements functional gradient boosting for regression.

A Boosted dataset is created from a base dataset by subtracting the prediction of the offset function from each example. This does not save the new dataset, but generates it as needed. The amount of space used is constant, independent on the size of the data set.

```
_____learnBoosting.py — Functional Gradient Boosting_____
11 from learnProblem import Data_set, Learner
12
13 class Boosted_dataset(Data_set):
14     def __init__(self, base_dataset, offset_fun):
15         """new dataset which is like base_dataset,
16            but offset_fun(e) is subtracted from the target of each example e
17         """
18         self.base_dataset = base_dataset
19         self.offset_fun = offset_fun
20         Data_set.__init__(self, base_dataset.train, base_dataset.test,
21                           base_dataset.prob_test, base_dataset.target_index)
```

```

22
23     def create_features(self):
24         self.input_features = self.base_dataset.input_features
25         def newout(e):
26             return self.base_dataset.target(e) - self.offset_fun(e)
27         newout.frange = self.base_dataset.target.frange
28         self.target = newout

```

A boosting learner takes in a dataset and a base learner, and returns a new predictor. The base learner, takes a dataset, and returns a Learner object.

```

_____learnBoosting.py — (continued)_____
30 class Boosting_learner(Learner):
31     def __init__(self, dataset, base_learner_class):
32         self.dataset = dataset
33         self.base_learner_class = base_learner_class
34         mean = sum(self.dataset.target(e)
35                     for e in self.dataset.train)/len(self.dataset.train)
36         self.predictor = lambda e:mean # function that returns mean for each example
37         self.predictor.__doc__ = "lambda e:"+str(mean)
38         self.offsets = [self.predictor]
39         self.errors = [data.evaluate_dataset(data.test, self.predictor, "sum-of-squares")]
40         self.display(1,"Predict mean test set error=", self.errors[0] )
41
42
43     def learn(self, num_ensemble=10):
44         """adds num_ensemble learners to the ensemble.
45         returns a new predictor.
46         """
47         for i in range(num_ensemble):
48             train_subset = Boosted_dataset(self.dataset, self.predictor)
49             learner = self.base_learner_class(train_subset)
50             new_offset = learner.learn()
51             self.offsets.append(new_offset)
52             def new_pred(e, old_pred=self.predictor, off=new_offset):
53                 return old_pred(e)+off(e)
54             self.predictor = new_pred
55             self.errors.append(data.evaluate_dataset(data.test, self.predictor,"sum-of-squares"))
56             self.display(1,"After Iteration",len(self.offsets)-1,"test set error=", self.errors[-1])
57         return self.predictor

```

For testing, *sp_DT_learner* returns a function that constructs a decision tree learner where the minimum number of examples is a proportion of the number of training examples. The value of 0.9 tends to have one split, and a value of 0.5 tends to have two splits (but test it). Thus this can be used to construct small decision trees that can be used as weak learners.

```

_____learnBoosting.py — (continued)_____
59 # Testing
60
61 from learnDT import DT_learner

```

```

62 from learnProblem import Data_set, Data_from_file
63
64 def sp_DT_learner(min_prop=0.9):
65     def make_learner(dataset):
66         mne = len(dataset.train)*min_prop
67         return DT_learner(dataset,min_number_examples=mne)
68     return make_learner
69
70 data = Data_from_file('data/carbool.csv', target_index=-1)
71 #data = Data_from_file('data/SPECT.csv', target_index=0)
72 #data = Data_from_file('data/mail_reading.csv', target_index=-1)
73 #data = Data_from_file('data/holiday.csv', num_train=19, target_index=-1)
74 learner9 = Boosting_learner(data, sp_DT_learner(0.9))
75 #learner7 = Boosting_learner(data, sp_DT_learner(0.7))
76 #learner5 = Boosting_learner(data, sp_DT_learner(0.5))
77 predictor9 = learner9.learn(10)
78 for i in learner9.offsets: print(i.__doc__)
79 import matplotlib.pyplot as plt
80
81 def plot_boosting(data,steps=10, thresholds=[0.5,0.1,0.01,0.001], markers=['-','-.-','-.',':']):
82     learners = [Boosting_learner(data, sp_DT_learner(th)) for th in thresholds]
83     predictors = [learner.learn(steps) for learner in learners]
84     plt.ion()
85     plt.xscale('linear') # change between log and linear scale
86     plt.xlabel("number of trees")
87     plt.ylabel(" error")
88     for (learner,(threshold,marker)) in zip(learners,zip(thresholds,markers)):
89         plt.plot(range(len(learner.errors)), learner.errors, ls=marker,c='k',
90                  label=str(round(threshold*100))+ "% min example threshold")
91     plt.legend()
92     plt.draw()
93
94 # plot_boosting(data)

```

Reasoning Under Uncertainty

8.1 Representing Probabilistic Models

In the implementation of probabilistic models we will assume that variables are objects, rather than the strings we used for CSPs. (Note that in the CSP code, variables could be anything; we just used strings for the examples.) We use a class here because it is more amenable to extend to richer models, such as when we introduce time.

A variable consists of a name and a domain. The domain of a variable is a list or a tuple, as the ordering will matter in the representation of factors.

```
_____probVariables.py — Probabilistic Variables_____
11 import random
12
13 class Variable(object):
14     """A random variable.
15     name (string) - name of the variable
16     domain (list) - a list of the values for the variable.
17     Variables are ordered according to their name.
18     """
19
20     def __init__(self, name, domain, position=None):
21         """Variable
22         name a string
23         domain a list of printable values
24         position of form (x,y)
25         """
26         self.name = name # string
27         self.domain = domain # list of values
28         self.position = position if position else (random.random(), random.random())
29         self.size = len(domain)
```

```

30
31     def __str__(self):
32         return self.name
33
34     def __repr__(self):
35         return self.name # f"Variable({self.name})"

```

8.2 Representing Factors

A **factor** is, mathematically, a function from variables into a number; that is given a value for each of its variable, it gives a number. Factors are used for conditional probabilities, utilities in the next chapter, and are explicitly constructed by some algorithms (in particular variable elimination).

A variable assignment, or just **assignment**, is represented as a $\{variable : value\}$ dictionary. A factor can be evaluated when all of its variables are assigned. The method `get_value` evaluates the factor for an assignment. The assignment can include extra variables not in the factor. This method needs to be defined for every subclass.

```

_____probFactors.py — Factors for graphical models_____
11 from display import Displayable
12 import math
13
14 class Factor(Displayable):
15     nextid=0 # each factor has a unique identifier; for printing
16
17     def __init__(self,variables):
18         self.variables = variables # ordered list of variables
19         self.id = Factor.nextid
20         self.name = f"f{self.id}"
21         Factor.nextid += 1
22
23     def can_evaluate(self,assignment):
24         """True when the factor can be evaluated in the assignment
25         assignment is a {variable:value} dict
26         """
27         return all(v in assignment for v in self.variables)
28
29     def get_value(self,assignment):
30         """Returns the value of the factor given the assignment of values to variables.
31         Needs to be defined for each subclass.
32         """
33         assert self.can_evaluate(assignment)
34         raise NotImplementedError("get_value") # abstract method

```

The method `__str__` returns a brief definition (like `"f7(X,Y,Z)"`). The method `to_table` returns string representations of a table showing all of the assignments of values to variables, and the corresponding value.

```

probFactors.py — (continued)
36 def __str__(self):
37     """returns a string representing a summary of the factor"""
38     return f"{self.name}({','.join(str(var) for var in self.variables)})"
39
40 def to_table(self, variables=None):
41     """returns a string representation of the factor.
42     Allows for an arbitrary variable ordering.
43     variables is a list of the variables in the factor
44     (can contain other variables)"""
45     if variables==None:
46         variables = self.variables
47     else: #enforce ordering and allow for extra variables in ordering
48         variables = [v for v in variables if v in self.variables]
49     head = "\t".join(str(v) for v in variables)
50     return head+"\n"+self.ass_to_str(variables, {}, variables)
51
52 def ass_to_str(self, vars, asst, allvars):
53     #print(f"ass_to_str({vars}, {asst}, {allvars})")
54     if vars:
55         return "\n".join(self.ass_to_str(vars[1:], asst | {vars[0]:val}, allvars)
56                           for val in vars[0].domain)
57     else:
58         return ("\t".join(str(asst[var]) for var in allvars)
59               + "\t"+str(self.get_value(asst)) )
60
61 __repr__ = __str__

```

8.3 Conditional Probability Distributions

A **conditional probability distribution (CPD)** is a type of factor that represents a conditional probability. A CPD representing $P(X \mid Y_1 \dots Y_k)$ is a type of factor, where given values for X and each Y_i returns a number.

```

probFactors.py — (continued)
63 class CPD(Factor):
64     def __init__(self, child, parents):
65         """represents P(variable | parents)
66         """
67         self.parents = parents
68         self.child = child
69         Factor.__init__(self, parents+[child])
70
71     def __str__(self):
72         """A brief description of a factor using in tracing"""
73         if self.parents:
74             return f"P({self.child}|{' '.join(str(p) for p in self.parents)})"
75         else:
76             return f"P({self.child})"

```

```

77 |
78 |     __repr__ = __str__

```

8.3.1 Logistic Regression

A **logistic regression** CPD, for Boolean variable X represents $P(X=True \mid Y_1 \dots Y_k)$, using $k + 1$ real-values weights so

$$P(X=True \mid Y_1 \dots Y_k) = \text{sigmoid}(w_0 + \sum_i w_i Y_i)$$

where for Boolean Y_i , True is represented as 1 and False as 0.

```

_____probFactors.py — (continued) _____
80 | from learnLinear import sigmoid
81 |
82 | class LogisticRegression(CPD):
83 |     def __init__(self, child, parents, weights):
84 |         """A logistic regression representation of a conditional probability.
85 |         child is the Boolean (or 0/1) variable whose CPD is being defined
86 |         parents is the list of parents
87 |         weights is list of parameters, such that weights[i+1] is the weight for parents[i]
88 |         """
89 |         assert len(weights) == 1+len(parents)
90 |         CPD.__init__(self, child, parents)
91 |         self.weights = weights
92 |
93 |     def get_value(self, assignment):
94 |         assert self.can_evaluate(assignment)
95 |         prob = sigmoid(self.weights[0]
96 |                        + sum(self.weights[i+1]*assignment[self.parents[i]]
97 |                             for i in range(len(self.parents))))
98 |         if assignment[self.child]: #child is true
99 |             return prob
100 |         else:
101 |             return (1-prob)

```

8.3.2 Noisy-or

A **noisy-or**, for Boolean variable X with Boolean parents $Y_1 \dots Y_k$ is parametrized by $k + 1$ parameters p_0, p_1, \dots, p_k , where each $0 \leq p_i \leq 1$. The semantics is defined as though there are $k + 1$ hidden variables $Z_0, Z_1 \dots Z_k$, where $P(Z_0) = p_0$ and $P(Z_i \mid Y_i) = p_i$ for $i \geq 1$, and where X is true if and only if $Z_0 \vee Z_1 \vee \dots \vee Z_k$ (where \vee is “or”). Thus X is false if all of the Z_i are false. Intuitively, Z_0 is the probability of X when all Y_i are false and each Z_i is a noisy (probabilistic) measure that Y_i makes X true, and X only needs one to make it true.

```

_____probFactors.py — (continued) _____

```



```

103 class NoisyOR(CPD):
104     def __init__(self, child, parents, weights):
105         """A noisy representation of a conditional probability.
106         variable is the Boolean (or 0/1) child variable whose CPD is being defined
107         parents is the list of Boolean (or 0/1) parents
108         weights is list of parameters, such that weights[i+1] is the weight for parents[i]
109         """
110         assert len(weights) == 1+len(parents)
111         CPD.__init__(self, child, parents)
112         self.weights = weights
113
114     def get_value(self, assignment):
115         assert self.can_evaluate(assignment)
116         probfalse = (1-self.weights[0])*math.prod(1-self.weights[i+1]
117                                                    for i in range(len(self.parents))
118                                                    if assignment[self.parents[i]])
119         if assignment[self.child]:
120             return 1-probfalse
121         else:
122             return probfalse

```

8.3.3 Tabular Factors

A **tabular factor** is a factor that represents each assignment of values to variables separately. It is represented by a Python array (or python dict). If the variables are V_1, V_2, \dots, V_k , the value of $f(V_1 = v_1, V_2 = v_1, \dots, V_k = v_k)$ is stored in $f[v_1][v_2] \dots [v_k]$.

If the domain of V_i is $[0, \dots, n_i - 1]$ this can be represented as an array. Otherwise we can use a dictionary. Python is nice in that it doesn't care, whether an array or dict is used **except when enumerating the values**; enumerating a dict gives the keys (the variables) but enumerating an array gives the values. So we have to be careful not to do this.

```

_____probFactors.py — (continued)_____
124 from functools import reduce
125
126 class TabFactor(Factor):
127
128     def __init__(self, variables, values):
129         Factor.__init__(self, variables)
130         self.values = values
131
132     def get_value(self, assignment):
133         return self.get_val_rec(self.values, self.variables, assignment)
134
135     def get_val_rec(self, value, variables, assignment):
136         if variables == []:
137             return value
138         else:

```

```

139         return self.get_val_rec(value[assignment[variables[0]]],
140                                variables[1:],assignment)

```

Prob is a factor that represents a conditional probability by enumerating all of the values.

```

_____probFactors.py — (continued)_____
142 class Prob(CPD,TabFactor):
143     """A factor defined by a conditional probability table"""
144     def __init__(self,var,pars,cpt):
145         """Creates a factor from a conditional probability table, cpt
146         The cpt values are assumed to be for the ordering par+[var]
147         """
148         TabFactor.__init__(self,pars+[var],cpt)
149         self.child = var
150         self.parents = pars

```

8.4 Graphical Models

A graphical model consists of a set of variables and a set of factors. A belief network is a graphical model where all of the factors represent conditional probabilities. There are some operations (such as pruning variables) which are applicable to belief networks, but are not applicable to more general models. At the moment, we will treat them as the same.

```

_____probGraphicalModels.py — Graphical Models and Belief Networks_____
11 from display import Displayable
12 from probFactors import CPD
13 import matplotlib.pyplot as plt
14
15 class GraphicalModel(Displayable):
16     """The class of graphical models.
17     A graphical model consists of a title, a set of variables and a set of factors.
18
19     vars is a set of variables
20     factors is a set of factors
21     """
22     def __init__(self, title, variables=None, factors=None):
23         self.title = title
24         self.variables = variables
25         self.factors = factors

```

A **belief network** (also known as a **Bayesian network**) is a graphical model where all of the factors are conditional probabilities, and every variable has a conditional probability of it given its parents. This only checks the first condition, and builds some useful data structures.

```

_____probGraphicalModels.py — (continued)_____
27 class BeliefNetwork(GraphicalModel):

```

```

28 """The class of belief networks."""
29
30 def __init__(self, title, variables, factors):
31     """vars is a set of variables
32     factors is a set( of factors. All of the factors are instances of CPD (e.g., Prob).
33     """
34     GraphicalModel.__init__(self, title, variables, factors)
35     assert all(isinstance(f,CPD) for f in factors)
36     self.var2cpt = {f.child:f for f in factors}
37     self.var2parents = {f.child:f.parents for f in factors}
38     self.children = {n:[] for n in self.variables}
39     for v in self.var2parents:
40         for par in self.var2parents[v]:
41             self.children[par].append(v)
42     self.topological_sort_saved = None

```

The following creates a topological sort of the nodes, where the parents of a node come before the node in the resulting order. This is based on Kahn's algorithm from 1962.

```

_____probGraphicalModels.py — (continued)_____
44 def topological_sort(self):
45     """creates a topological ordering of variables such that the parents of
46     a node are before the node.
47     """
48     if self.topological_sort_saved:
49         return self.topological_sort_saved
50     next_vars = {n for n in self.var2parents if not self.var2parents[n] }
51     self.display(3,'topological_sort: next_vars',next_vars)
52     top_order=[]
53     while next_vars:
54         var = next_vars.pop()
55         self.display(3,'select variable',var)
56         top_order.append(var)
57         next_vars |= {ch for ch in self.children[var]
58                     if all(p in top_order for p in self.var2parents[ch])}
59         self.display(3,'var_with_no_parents_left',next_vars)
60     self.display(3,"top_order",top_order)
61     assert set(top_order)==set(self.var2parents),(top_order,self.var2parents)
62     self.topologicalsort_saved=top_order
63     return top_order

```

The **show** method uses matplotlib to show the graphical structure of a belief network.

```

_____probGraphicalModels.py — (continued)_____
65 def show(self):
66     plt.ion() # interactive
67     ax = plt.figure().gca()
68     ax.set_axis_off()
69     plt.title(self.title)

```

```

70     bbox = dict(boxstyle="round4,pad=1.0,rounding_size=0.5")
71     for var in reversed(self.topological_sort()):
72         if self.var2parents[var]:
73             for par in self.var2parents[var]:
74                 ax.annotate(var.name, par.position, xytext=var.position,
75                             arrowprops={'arrowstyle': '<-'}, bbox=bbox,
76                             ha='center')
77         else:
78             x,y = var.position
79             plt.text(x,y,var.name,bbox=bbox,ha='center')

```

8.4.1 Example Belief Networks

A Chain of 4 Variables

The first example belief network is a simple chain $A \rightarrow B \rightarrow C \rightarrow D$.

Please do not change this, as it is the example used for testing.

```

_____probGraphicalModels.py — (continued) _____
81 from probVariables import Variable
82 from probFactors import Prob, LogisticRegression, NoisyOR
83
84 boolean = [False, True]
85 A = Variable("A", boolean, position=(0,0.8))
86 B = Variable("B", boolean, position=(0.333,0.6))
87 C = Variable("C", boolean, position=(0.666,0.4))
88 D = Variable("D", boolean, position=(1,0.2))
89
90 f_a = Prob(A,[],[0.4,0.6])
91 f_b = Prob(B,[A],[[0.9,0.1],[0.2,0.8]])
92 f_c = Prob(C,[B],[[0.6,0.4],[0.3,0.7]])
93 f_d = Prob(D,[C],[[0.1,0.9],[0.75,0.25]])
94
95 bn_4ch = BeliefNetwork("4-chain", {A,B,C,D}, {f_a,f_b,f_c,f_d})

```

Report-of-Leaving Example

The second belief network, `bn_report`, is Example 8.15 of Poole and Mackworth [2017] (<http://artint.info>). The output of `bn_report.show()` is shown in Figure 8.1 of this document.

```

_____probGraphicalModels.py — (continued) _____
97 # Belief network report-of-leaving example (Example 8.15 shown in Figure 8.3) of
98 # Poole and Mackworth, Artificial Intelligence, 2017 http://artint.info
99
100 Alarm = Variable("Alarm", boolean, position=(0.366,0.633))
101 Fire = Variable("Fire", boolean, position=(0.633,0.9))
102 Leaving = Variable("Leaving", boolean, position=(0.366,0.366))
103 Report = Variable("Report", boolean, position=(0.366,0.1))

```




Figure 8.1: The report-of-leaving belief network

```

104 Smoke = Variable("Smoke", boolean, position=(0.9,0.633))
105 Tamper = Variable("Tamper", boolean, position=(0.1,0.9))
106
107 f_ta = Prob(Tamper,[],[0.98,0.02])
108 f-fi = Prob(Fire,[],[0.99,0.01])
109 f-sm = Prob(Smoke,[Fire],[[0.99,0.01],[0.1,0.9]])
110 f-al = Prob(Alarm,[Fire,Tamper],[[0.9999, 0.0001], [0.15, 0.85]], [[0.01, 0.99], [0.5, 0.5]])
111 f_lv = Prob(Leaving,[Alarm],[[0.999, 0.001], [0.12, 0.88]])
112 f-re = Prob(Report,[Leaving],[[0.99, 0.01], [0.25, 0.75]])
113
114 bn_report = BeliefNetwork("Report-of-leaving", {Tamper,Fire,Smoke,Alarm,Leaving,Report},
115                               {f_ta,f-fi,f-sm,f-al,f_lv,f-re})

```

Sprinkler Example

The third belief network is the sprinkler example from Pearl. The output of `bn_sprinkler.show()` is shown in Figure 8.2 of this document.

```

probGraphicalModels.py — (continued)
117 Season = Variable("Season", ["summer","winter"], position=(0.5,0.9))
118 Sprinkler = Variable("Sprinkler", ["on","off"], position=(0.9,0.6))
119 Rained = Variable("Rained", boolean, position=(0.1,0.6))
120 Grass_wet = Variable("Grass wet", boolean, position=(0.5,0.3))
121 Grass_shiny = Variable("Grass shiny", boolean, position=(0.1,0))

```




Figure 8.2: The sprinkler belief network

```

122 Shoes_wet = Variable("Shoes wet", boolean, position=(0.9,0))
123
124 f_season = Prob(Season,[],[0.5,0.5])
125 f_sprinkler = Prob(Sprinkler,[Season],[0.9,0.1,0.05,0.95])
126 f_rained = Prob(Rained,[Season],[0.7,0.3,0.2,0.8])
127 f_wet = Prob(Grass_wet,[Sprinkler,Rained], [1,0,0.1,0.9,0.2,0.8,0.02,0.98])
128 f_shiny = Prob(Grass_shiny, [Grass_wet], [0.95,0.05,0.3,0.7])
129 f_shoes = Prob(Shoes_wet, [Grass_wet], [0.92,0.08,0.35,0.65])
130
131 bn_sprinkler = BeliefNetwork("Pearl's Sprinkler Example",
132                             {Season, Sprinkler, Rained, Grass_wet, Grass_shiny, Shoes_wet},
133                             {f_season, f_sprinkler, f_rained, f_wet, f_shiny, f_shoes})

```

Bipartite diagnostic model with noisy-or

This belief network is a bipartite diagnostic model, with independent diseases, and the symptoms depend on the diseases, where the CPDs are defined using noisy-or. Bipartite means it is in two parts; the diseases are only connected to the symptoms and the symptoms are only connected to the diseases. The output of `bn_no1.show()` is shown in Figure 8.3 of this document.

```

135 Cough = Variable("Cough", boolean, (0.1,0.1))
136 Fever = Variable("Fever", boolean, (0.5,0.1))

```

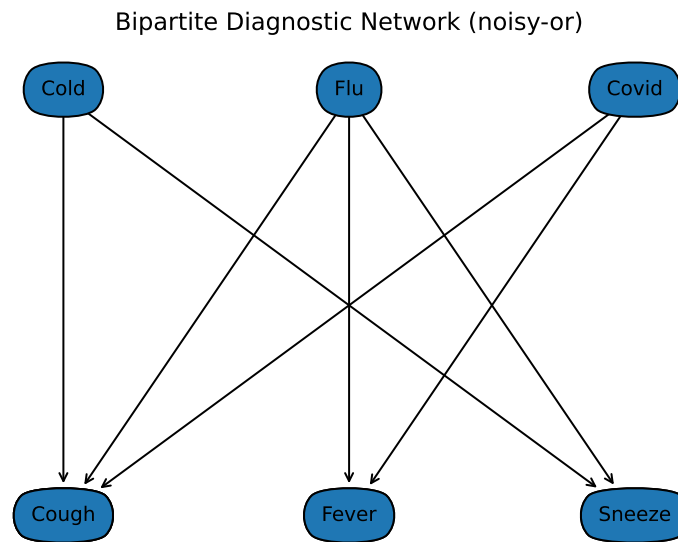


Figure 8.3: A partite diagnostic network

```

137 Sneeze = Variable("Sneeze", boolean, (0.9,0.1))
138 Cold = Variable("Cold",boolean, (0.1,0.9))
139 Flu = Variable("Flu",boolean, (0.5,0.9))
140 Covid = Variable("Covid",boolean, (0.9,0.9))
141
142 p_cold_no = Prob(Cold,[],[0.9,0.1])
143 p_flu_no = Prob(Flu,[],[0.95,0.05])
144 p_covid_no = Prob(Covid,[],[0.99,0.01])
145
146 p_cough_no = NoisyOR(Cough, [Cold,Flu,Covid], [0.1, 0.3, 0.2, 0.7])
147 p_fever_no = NoisyOR(Fever, [ Flu,Covid], [0.01, 0.6, 0.7])
148 p_sneeze_no = NoisyOR(Sneeze, [Cold,Flu ], [0.05, 0.5, 0.2 ])
149
150 bn_no1 = BeliefNetwork("Bipartite Diagnostic Network (noisy-or)",
151                        {Cough, Fever, Sneeze, Cold, Flu, Covid},
152                        {p_cold_no, p_flu_no, p_covid_no, p_cough_no, p_fever_no, p_sneeze_no})
153
154 # to see the conditional probability of Noisy-or do:
155 #print(p_cough_no.to_table())

```

Bipartite diagnostic model with noisy-or

The next belief network is a bipartite diagnostic model, with independent diseases, and the symptoms depend on the diseases, where the CPDs are defined

using logistic regression. It has the same graphical structure as the previous example (see Figure 8.3). This has the (approximately) the same conditional probabilities as the previous example when zero or one diseases are present. Note that $\text{sigmoid}(-2.2) \approx 0.1$

```

_____probGraphicalModels.py — (continued) _____
157
158 p_cold_lr = Prob(Cold,[],[0.9,0.1])
159 p_flu_lr = Prob(Flu,[],[0.95,0.05])
160 p_covid_lr = Prob(Covid,[],[0.99,0.01])
161
162 p_cough_lr = LogisticRegression(Cough, [Cold,Flu,Covid], [-2.2, 1.67, 1.26,
163 3.19])
164 p_fever_lr = LogisticRegression(Fever, [Flu,Covid], [-4.6, 5.02, 5.46])
165 p_sneeze_lr = LogisticRegression(Sneeze, [Cold,Flu ], [-2.94, 3.04, 1.79 ])
166
167 bn_lr1 = BeliefNetwork("Bipartite Diagnostic Network - logistic regression",
168 {Cough, Fever, Sneeze, Cold, Flu, Covid},
169 {p_cold_lr, p_flu_lr, p_covid_lr, p_cough_lr, p_fever_lr, p_sneeze_lr})
170
171 # to see the conditional probability of Noisy-or do:
172 #print(p_cough_lr.to_table())

```

8.5 Inference Methods

Each of the inference methods implements the query method that computes the posterior probability of a variable given a dictionary of $\{variable : value\}$ observations. The methods are Displayable because they implement the *display* method which is currently text-based.

```

_____probGraphicalModels.py — (continued) _____
173 from display import Displayable
174
175 class InferenceMethod(Displayable):
176     """The abstract class of graphical model inference methods"""
177     method_name = "unnamed" # each method should have a method name
178
179     def query(self, qvar, obs={}):
180         """returns a {value:prob} dictionary for the query variable"""
181         raise NotImplementedError("InferenceMethod query") # abstract method

```

We use bn_4ch as the test case, in particular $P(B \mid D = \text{true})$. This needs an error threshold, particularly for the approximate methods.

```

_____probGraphicalModels.py — (continued) _____
183 def testIM(self, threshold=0.000000001):
184     solver = self.bn_4ch
185     res = solver.query(B,{D:True})
186     correct_answer = 0.429632380245

```



```

187         assert correct_answer-threshold < res[True] < correct_answer+threshold, f"value {res[True]}
188         print(f"Unit test passed for {self.method_name}.")

```

8.6 Recursive Conditioning

An instance of a *RC* object takes in a graphical model. The query method uses recursive conditioning to compute the probability of a query variable given observations on other variables.

```

_____probRC.py — Recursive Conditioning for Graphical Models_____
11 import math
12 from probGraphicalModels import GraphicalModel, InferenceMethod
13 from probFactors import Factor
14 from utilities import dict_union
15
16 class RC(InferenceMethod):
17     """The class that queries graphical models using recursive conditioning
18
19     gm is graphical model to query
20     """
21     method_name = "recursive conditioning"
22
23     def __init__(self, gm=None):
24         self.gm = gm
25         self.cache = {(frozenset(), frozenset()):1}
26         ## self.max_display_level = 3
27
28     def query(self, var, obs={}, split_order=None):
29         """computes P(var|obs) where
30         var is a variable
31         obs is a variable:value dictionary
32         split_order is a list of the non-observed non-query variables in gm
33         """
34         if var in obs:
35             return {val:(1 if val == obs[var] else 0) for val in var.domain}
36         else:
37             if split_order == None:
38                 split_order = [v for v in self.gm.variables if (v not in obs) and v != var]
39                 unnorm = [self.rc(dict_union({var:val}, obs), self.gm.factors, split_order)
40                             for val in var.domain]
41                 p_obs = sum(unnorm)
42                 return {val:pr/p_obs for val, pr in zip(var.domain, unnorm)}

```

The following is the simplest search-based algorithm. It is exponential in the number of variables, so is not very useful. However, it is simple, and useful to understand before looking at the more complicated algorithm. Note that the above code does not call `rc0`; you will need to change the `self.rc` to `self.rc0` in above code to use it.

probRC.py — (continued)

```

44 def rc0(self, context, factors, split_order):
45     """simplest search algorithm"""
46     self.display(2,"calling rc0,", (context,factors))
47     if not factors:
48         return 1
49     elif to_eval := {fac for fac in factors if fac.can_evaluate(context)}:
50         self.display(3,"rc0 evaluating factors",to_eval)
51         val = math.prod(fac.get_value(context) for fac in to_eval)
52         return val * self.rc0(context, factors-to_eval, split_order)
53     else:
54         total = 0
55         var = split_order[0]
56         self.display(3, "rc0 branching on", var)
57         for val in var.domain:
58             total += self.rc0(dict_union({var:val},context), factors, split_order[1:])
59         self.display(3, "rc0 branching on", var,"returning", total)
60         return total

```

probRC.py — (continued)

```

62 def rc(self, context, factors, split_order):
63     """ returns the number \sum_{split_order} \prod_{factors} given assignments in context
64     context is a variable:value dictionary
65     factors is a set of factors
66     split_order is a list of variables in factors that are not in context
67     """
68     self.display(3,"calling rc,", (context,factors))
69     ce = (frozenset(context.items()), frozenset(factors)) # key for the cache entry
70     if ce in self.cache:
71         self.display(2,"rc cache lookup", (context,factors))
72         return self.cache[ce]
73 #     if not factors: # no factors; needed if you don't have forgetting and caching
74 #         return 1
75     elif vars_not_in_factors := {var for var in context
76                                 if not any(var in fac.variables for fac in factors)}:
77         # forget variables not in any factor
78         self.display(3,"rc forgetting variables", vars_not_in_factors)
79         return self.rc({key:val for (key,val) in context.items()
80                         if key not in vars_not_in_factors},
81                        factors, split_order)
82     elif to_eval := {fac for fac in factors if fac.can_evaluate(context)}:
83         # evaluate factors when all variables are assigned
84         self.display(3,"rc evaluating factors",to_eval)
85         val = math.prod(fac.get_value(context) for fac in to_eval)
86         if val == 0:
87             return 0
88         else:
89             return val * self.rc(context, {fac for fac in factors if fac not in to_eval}, split_order)
90     elif len(comp := connected_components(context, factors, split_order)) > 1:
91         # there are disconnected components

```

```

92         self.display(2,"splitting into connected components",comp)
93         return(math.prod(self.rc(context,f,eo) for (f,eo) in comp))
94     else:
95         assert split_order, "split_order should not be empty to get here"
96         total = 0
97         var = split_order[0]
98         self.display(3, "rc branching on", var)
99         for val in var.domain:
100             total += self.rc(dict_union({var:val},context), factors, split_order[1:])
101         self.cache[ce] = total
102         self.display(2, "rc branching on", var,"returning", total)
103         return total

```

connected_components returns a list of connected components, where a connected component is a set of factors and a set of variables, where the graph that connects variables and factors that involve them is connected. The connected components are built one at a time; with a current connected component. At all times factors is partitioned into 3 disjoint sets:

- component_factors containing factors in the current connected component where all factors that share a variable are already in the component
- factors_to_check containing factors in the current connected component where potentially some factors that share a variable are not in the component; these need to be checked
- other_factors the other factors that are not (yet) in the connected component

probRC.py — (continued)

```

105 def connected_components(context, factors, split_order):
106     """returns a list of (f,e) where f is a subset of factors and e is a subset of split_order
107     such that each element shares the same variables that are disjoint from other elements.
108     """
109     other_factors = set(factors) #copies factors
110     factors_to_check = {other_factors.pop()} # factors in connected component still to be checked
111     component_factors = set() # factors in first connected component already checked
112     component_variables = set() # variables in first connected component
113     while factors_to_check:
114         next_fac = factors_to_check.pop()
115         component_factors.add(next_fac)
116         new_vars = set(next_fac.variables) - component_variables - context.keys()
117         component_variables |= new_vars
118         for var in new_vars:
119             factors_to_check |= {f for f in other_factors if var in f.variables}
120             other_factors -= factors_to_check # set difference
121     if other_factors:
122         return ( [(component_factors,[e for e in split_order if e in component_variables])]
123                 + connected_components(context, other_factors, [e for e in split_order if e not in component_variables]) )
124     else:

```

```
125 |         return [(component_factors, split_order)]
```

Testing:

```

127 | from probGraphicalModels import bn_4ch, A,B,C,D,f_a,f_b,f_c,f_d
128 | bn_4chv = RC(bn_4ch)
129 | ## bn_4chv.query(A,{})
130 | ## bn_4chv.query(D,{})
131 | ## InferenceMethod.max_display_level = 3 # show more detail in displaying
132 | ## InferenceMethod.max_display_level = 1 # show less detail in displaying
133 | ## bn_4chv.query(A,{D:True},[C,B])
134 | ## bn_4chv.query(B,{A:True,D:False})
135 |
136 | from probGraphicalModels import bn_report,Alarm,Fire,Leaving,Report,Smoke,Tamper
137 | bn_reportRC = RC(bn_report) # answers queries using recursive conditioning
138 | ## bn_reportRC.query(Tamper,{})
139 | ## InferenceMethod.max_display_level = 0 # show no detail in displaying
140 | ## bn_reportRC.query(Leaving,{})
141 | ## bn_reportvRC.query(Tamper,{}, split_order=[Smoke,Fire,Alarm,Leaving,Report])
142 | ## bn_reportRC.query(Tamper,{Report:True})
143 | ## bn_reportvRC.query(Tamper,{Report:True,Smoke:False})
144 |
145 | from probGraphicalModels import bn_sprinkler, Season, Sprinkler, Rained, Grass_wet, Grass_shiny, Shoe
146 | bn_sprinklerv = RC(bn_sprinkler)
147 | ## bn_sprinklerv.query(Shoes_wet,{})
148 | ## bn_sprinklerv.query(Shoes_wet,{Rained:True})
149 | ## bn_sprinklerv.query(Shoes_wet,{Grass_shiny:True})
150 | ## bn_sprinklerv.query(Shoes_wet,{Grass_shiny:False,Rained:True})
151 |
152 | from probGraphicalModels import bn_no1, bn_lr1, Cough, Fever, Sneeze, Cold, Flu, Covid
153 | bn_no1v = RC(bn_no1)
154 | bn_lr1v = RC(bn_lr1)
155 | ## bn_lr1v.query(Flu, {Fever:1, Sneeze:1})
156 | ## bn_lr1v.query(Cough,{})
157 | ## bn_lr1v.query(Cold,{Cough:1,Sneeze:0,Fever:1})
158 | ## bn_lr1v.query(Flu,{Cough:0,Sneeze:1,Fever:1})
159 | ## bn_lr1v.query(Covid,{Cough:1,Sneeze:0,Fever:1})
160 | ## bn_lr1v.query(Covid,{Cough:1,Sneeze:0,Fever:1,Flu:0})
161 | ## bn_lr1v.query(Covid,{Cough:1,Sneeze:0,Fever:1,Flu:1})
162 |
163 | if __name__ == "__main__":
164 |     InferenceMethod.testIM(RC)

```

8.7 Variable Elimination

An instance of a *VE* object takes in a graphical model. The query method uses variable elimination to compute the probability of a variable given observations on some other variables.

```

_____probVE.py — Variable Elimination for Graphical Models_____
11 from probFactors import Factor, FactorObserved, FactorSum, factor_times
12 from probGraphicalModels import GraphicalModel, InferenceMethod
13
14 class VE(InferenceMethod):
15     """The class that queries Graphical Models using variable elimination.
16
17     gm is graphical model to query
18     """
19     method_name = "variable elimination"
20
21     def __init__(self, gm=None):
22         self.gm = gm
23
24     def query(self, var, obs={}, elim_order=None):
25         """computes P(var|obs) where
26         var is a variable
27         obs is a {variable:value} dictionary"""
28         if var in obs:
29             return {var:1 if val == obs[var] else 0 for val in var.domain}
30         else:
31             if elim_order == None:
32                 elim_order = self.gm.variables
33             projFactors = [self.project_observations(fact, obs)
34                           for fact in self.gm.factors]
35             for v in elim_order:
36                 if v != var and v not in obs:
37                     projFactors = self.eliminate_var(projFactors, v)
38             unnorm = factor_times(var, projFactors)
39             p_obs = sum(unnorm)
40             self.display(1, "Unnormalized probs:", unnorm, "Prob obs:", p_obs)
41             return {val:pr/p_obs for val, pr in zip(var.domain, unnorm)}

```

A *FactorObserved* is a factor that is the result of some observations on another factor. We don't store the values in a list; we just look them up as needed. The observations can include variables that are not in the list, but should have some intersection with the variables in the factor.

```

_____probFactors.py — (continued)_____
152 class FactorObserved(Factor):
153     def __init__(self, factor, obs):
154         Factor.__init__(self, [v for v in factor.variables if v not in obs])
155         self.observed = obs
156         self.orig_factor = factor
157
158     def get_value(self, assignment):
159         ass = assignment.copy()
160         for ob in self.observed:
161             ass[ob] = self.observed[ob]
162         return self.orig_factor.get_value(ass)

```

A *FactorSum* is a factor that is the result of summing out a variable from the product of other factors. I.e., it constructs a representation of:

$$\sum_{var} \prod_{f \in factors} f.$$

We store the values in a list in a lazy manner; if they are already computed, we used the stored values. If they are not already computed we can compute and store them.

```

164 class FactorSum(Factor):
165     def __init__(self, var, factors):
166         self.var_summed_out = var
167         self.factors = factors
168         vars = []
169         for fac in factors:
170             for v in fac.variables:
171                 if v is not var and v not in vars:
172                     vars.append(v)
173         Factor.__init__(self, vars)
174         self.values = {}
175
176     def get_value(self, assignment):
177         """lazy implementation: if not saved, compute it. Return saved value"""
178         asst = frozenset(assignment.items())
179         if asst in self.values:
180             return self.values[asst]
181         else:
182             total = 0
183             new_asst = assignment.copy()
184             for val in self.var_summed_out.domain:
185                 new_asst[self.var_summed_out] = val
186                 total += math.prod(fac.get_value(new_asst) for fac in self.factors)
187             self.values[asst] = total
188         return total

```

The method *factor_times* multiplies a set of factors that are all factors on the same variable (or on no variables). This is the last step in variable elimination before normalizing. It returns an array giving the product for each value of *variable*.

```

190 def factor_times(variable, factors):
191     """when factors are factors just on variable (or on no variables)"""
192     prods = []
193     facs = [f for f in factors if variable in f.variables]
194     for val in variable.domain:
195         ast = {variable: val}
196         prods.append(math.prod(f.get_value(ast) for f in facs))
197     return prods

```

To project observations onto a factor, for each variable that is observed in the factor, we construct a new factor that is the factor projected onto that variable. *Factor_observed* creates a new factor that is the result of assigning a value to a single variable.

```

43     def project_observations(self, factor, obs):
44         """Returns the resulting factor after observing obs
45
46         obs is a dictionary of {variable:value} pairs.
47         """
48         if any((var in obs) for var in factor.variables):
49             # a variable in factor is observed
50             return FactorObserved(factor, obs)
51         else:
52             return factor
53
54     def eliminate_var(self, factors, var):
55         """Eliminate a variable var from a list of factors.
56         Returns a new set of factors that has var summed out.
57         """
58         self.display(2, "eliminating ", str(var))
59         contains_var = []
60         not_contains_var = []
61         for fac in factors:
62             if var in fac.variables:
63                 contains_var.append(fac)
64             else:
65                 not_contains_var.append(fac)
66         if contains_var == []:
67             return factors
68         else:
69             newFactor = FactorSum(var, contains_var)
70             self.display(2, "Multiplying:", [str(f) for f in contains_var])
71             self.display(2, "Creating factor:", newFactor)
72             self.display(3, newFactor.to_table()) # factor in detail
73             not_contains_var.append(newFactor)
74             return not_contains_var
75
76 from probGraphicalModels import bn_4ch, A, B, C, D
77 bn_4chv = VE(bn_4ch)
78 ## bn_4chv.query(A, {})
79 ## bn_4chv.query(D, {})
80 ## InferenceMethod.max_display_level = 3 # show more detail in displaying
81 ## InferenceMethod.max_display_level = 1 # show less detail in displaying
82 ## bn_4chv.query(A, {D: True})
83 ## bn_4chv.query(B, {A: True, D: False})
84
85 from probGraphicalModels import bn_report, Alarm, Fire, Leaving, Report, Smoke, Tamper
86 bn_reportv = VE(bn_report) # answers queries using variable elimination

```

```

87 ## bn_reportv.query(Tamper,{})
88 ## InferenceMethod.max_display_level = 0 # show no detail in displaying
89 ## bn_reportv.query(Leaving,{})
90 ## bn_reportv.query(Tamper,{},elim_order=[Smoke,Report,Leaving,Alarm,Fire])
91 ## bn_reportv.query(Tamper,{Report:True})
92 ## bn_reportv.query(Tamper,{Report:True,Smoke:False})
93
94 from probGraphicalModels import bn_sprinkler, Season, Sprinkler, Rained, Grass_wet, Grass_shiny, Shoe
95 bn_sprinklerv = VE(bn_sprinkler)
96 ## bn_sprinklerv.query(Shoes_wet,{})
97 ## bn_sprinklerv.query(Shoes_wet,{Rained:True})
98 ## bn_sprinklerv.query(Shoes_wet,{Grass_shiny:True})
99 ## bn_sprinklerv.query(Shoes_wet,{Grass_shiny:False,Rained:True})
100
101 from probGraphicalModels import bn_lr1, Cough, Fever, Sneeze, Cold, Flu, Covid
102 vediag = VE(bn_lr1)
103 ## vediag.query(Cough,{})
104 ## vediag.query(Cold,{Cough:1,Sneeze:0,Fever:1})
105 ## vediag.query(Flu,{Cough:0,Sneeze:1,Fever:1})
106 ## vediag.query(Covid,{Cough:1,Sneeze:0,Fever:1})
107 ## vediag.query(Covid,{Cough:1,Sneeze:0,Fever:1,Flu:0})
108 ## vediag.query(Covid,{Cough:1,Sneeze:0,Fever:1,Flu:1})
109
110 if __name__ == "__main__":
111     InferenceMethod.testIM(VE)

```

8.8 Stochastic Simulation

8.8.1 Sampling from a discrete distribution

The method *sample_one* generates a single sample from a (possible unnormalized) distribution. *dist* is a $\{value : weight\}$ dictionary, where $weight \geq 0$. This returns a value with probability in proportion to its weight.

```

_____probStochSim.py — Probabilistic inference using stochastic simulation _____
11 import random
12 from probGraphicalModels import InferenceMethod
13
14 def sample_one(dist):
15     """returns the index of a single sample from normalized distribution dist."""
16     rand = random.random()*sum(dist.values())
17     cum = 0 # cumulative weights
18     for v in dist:
19         cum += dist[v]
20         if cum > rand:
21             return v

```

If we want to generate multiple samples, repeatedly calling *sample_one* may not be efficient. If we want to generate n samples, and the distribution is over

m values, *sample_one* takes time $O(mn)$. If m and n are of the same order of magnitude, we can do better.

The method *sample_multiple* generates multiple samples from a distribution defined by *dist*, where *dist* is a $\{value : weight\}$ dictionary, where $weight \geq 0$ and the weights cannot all be zero. This returns a list of values, of length *num_samples*, where each sample is selected with a probability proportional to its weight.

The method generates all of the random numbers, sorts them, and then goes through the distribution once, saving the selected samples.

```

_____probStochSim.py — (continued)_____
23 def sample_multiple(dist, num_samples):
24     """returns a list of num_samples values selected using distribution dist.
25     dist is a {value:weight} dictionary that does not need to be normalized
26     """
27     total = sum(dist.values())
28     rands = sorted(random.random()*total for i in range(num_samples))
29     result = []
30     dist_items = list(dist.items())
31     cum = dist_items[0][1] # cumulative sum
32     index = 0
33     for r in rands:
34         while r>cum:
35             index += 1
36             cum += dist_items[index][1]
37         result.append(dist_items[index][0])
38     return result

```

Exercise 8.1

What is the time and space complexity the following 4 methods to generate n samples, where m is the length of *dist*:

- n calls to *sample_one*
- sample_multiple*
- Create the cumulative distribution (choose how this is represented) and, for each random number, do a binary search to determine the sample associated with the random number.
- Choose a random number in the range $[i/n, (i+1)/n)$ for each $i \in \text{range}(n)$, where n is the number of samples. Use these as the random numbers to select the particles. (Does this give random samples?)

For each method suggest when it might be the best method.

The *test_sampling* method can be used to generate the statistics from a number of samples. It is useful to see the variability as a function of the number of samples. Try it for few samples and also for many samples.

```

_____probStochSim.py — (continued)_____
40 def test_sampling(dist, num_samples):

```

```

41     """Given a distribution, dist, draw num_samples samples
42     and return the resulting counts
43     """
44     result = {v:0 for v in dist}
45     for v in sample_multiple(dist, num_samples):
46         result[v] += 1
47     return result
48
49 # try the following queries a number of times each:
50 # test_sampling({1:1,2:2,3:3,4:4}, 100)
51 # test_sampling({1:1,2:2,3:3,4:4}, 100000)

```

8.8.2 Sampling Methods for Belief Network Inference

A *SamplingInferenceMethod* is an *InferenceMethod*, but the query method also takes arguments for the number of samples and the sample-order (which is an ordering of factors). The first methods assume a belief network (and not an undirected graphical model).

```

_____probStochSim.py — (continued)_____
53 class SamplingInferenceMethod(InferenceMethod):
54     """The abstract class of sampling-based belief network inference methods"""
55     def query(self,qvar,obs={},number_samples=1000,sample_order=None):
56         raise NotImplementedError("SamplingInferenceMethod query") # abstract

```

8.8.3 Rejection Sampling

```

_____probStochSim.py — (continued)_____
58 class RejectionSampling(SamplingInferenceMethod):
59     """The class that queries Graphical Models using Rejection Sampling.
60
61     bn is a belief network to query
62     """
63     method_name = "rejection sampling"
64
65     def __init__(self,bn=None):
66         self.bn = bn
67
68     def query(self,qvar,obs={},number_samples=1000,sample_order=None):
69         """computes P(qvar|obs) where
70         qvar is a variable.
71         obs is a {variable:value} dictionary.
72         sample_order is a list of variables where the parents
73         come before the variable.
74         """
75         if sample_order is None:
76             sample_order = self.bn.topological_sort()
77         self.display(2,*sample_order,sep="\t")

```

```

78     counts = {val:0 for val in qvar.domain}
79     for i in range(number_samples):
80         rejected = False
81         sample = {}
82         for nvar in sample_order:
83             fac = self.bn.var2cpt[nvar] #factor with nvar as child
84             val = sample_one({v:fac.get_value(sample|{nvar:v}) for v in nvar.domain})
85             self.display(2,val,end="\t")
86             if nvar in obs and obs[nvar] != val:
87                 rejected = True
88                 self.display(2,"Rejected")
89                 break
90             sample[nvar] = val
91         if not rejected:
92             counts[sample[qvar]] += 1
93             self.display(2,"Accepted")
94     tot = sum(counts.values())
95     # As well as the distribution we also include raw counts
96     return {c:v/tot if tot>0 else 1/len(qvar.domain) for (c,v) in counts.items()} | {"raw_count": tot}

```

8.8.4 Likelihood Weighting

Likelihood weighting includes a weight for each sample. Instead of rejecting samples based on observations, likelihood weighting changes the weights of the sample in proportion with the probability of the observation. The weight then becomes the probability that the variable would have been rejected.

```

_____probStochSim.py — (continued)_____
98 class LikelihoodWeighting(SamplingInferenceMethod):
99     """The class that queries Graphical Models using Likelihood weighting.
100
101     bn is a belief network to query
102     """
103     method_name = "likelihood weighting"
104
105     def __init__(self,bn=None):
106         self.bn = bn
107
108     def query(self,qvar,obs={},number_samples=1000,sample_order=None):
109         """computes P(qvar|obs) where
110         qvar is a variable.
111         obs is a {variable:value} dictionary.
112         sample_order is a list of factors where factors defining the parents
113         come before the factors for the child.
114         """
115         if sample_order is None:
116             sample_order = self.bn.topological_sort()
117         self.display(2,*[v for v in sample_order
118                        if v not in obs],sep="\t")
119         counts = {val:0 for val in qvar.domain}
120         for i in range(number_samples):

```

```

121     sample = {}
122     weight = 1.0
123     for nvar in sample_order:
124         fac = self.bn.var2cpt[nvar]
125         if nvar in obs:
126             sample[nvar] = obs[nvar]
127             weight *= fac.get_value(sample)
128         else:
129             val = sample_one({v:fac.get_value(sample|{nvar:v}) for v in nvar.domain})
130             self.display(2,val,end="\t")
131             sample[nvar] = val
132         counts[sample[qvar]] += weight
133         self.display(2,weight)
134     tot = sum(counts.values())
135     # as well as the distribution we also include the raw counts
136     return {c:v/tot for (c,v) in counts.items()} | {"raw_counts":counts}

```

Exercise 8.2 Change this algorithm so that it does **importance sampling** using a proposal distribution. It needs *sample_one* using a different distribution and then update the weight of the current sample. For testing, use a proposal distribution that only specifies probabilities for some of the variables (and the algorithm uses the probabilities for the network in other cases).

8.8.5 Particle Filtering

In this implementation, a particle is a $\{variable : value\}$ dictionary. Because adding a new value to dictionary involves a side effect, the dictionaries need to be copied during resampling.

```

_____probStochSim.py — (continued)_____
138 class ParticleFiltering(SamplingInferenceMethod):
139     """The class that queries Graphical Models using Particle Filtering.
140
141     bn is a belief network to query
142     """
143     method_name = "particle filtering"
144
145     def __init__(self,bn=None):
146         self.bn = bn
147
148     def query(self, qvar, obs={}, number_samples=1000, sample_order=None):
149         """computes P(qvar|obs) where
150         qvar is a variable.
151         obs is a {variable:value} dictionary.
152         sample_order is a list of factors where factors defining the parents
153         come before the factors for the child.
154         """
155         if sample_order is None:
156             sample_order = self.bn.topological_sort()
157         self.display(2,*[v for v in sample_order

```

```

158         if v not in obs], sep="\t")
159     particles = [{i} for i in range(number_samples)]
160     for nvar in sample_order:
161         fac = self.bn.var2cpt[nvar]
162         if nvar in obs:
163             weights = [fac.get_value(part|{nvar:obs[nvar]}) for part in particles]
164             particles = [p|{nvar:obs[nvar]} for p in resample(particles, weights, number_samples)]
165         else:
166             for part in particles:
167                 part[nvar] = sample_one({v:fac.get_value(part|{nvar:v}) for v in nvar.domain})
168                 self.display(2, part[nvar], end="\t")
169     counts = {val:0 for val in qvar.domain}
170     for part in particles:
171         counts[part[qvar]] += 1
172     tot = sum(counts.values())
173     # as well as the distribution we also include the raw counts
174     return {c:v/tot for (c,v) in counts.items()} | {"raw_counts":counts}

```

Resampling

Resample is based on *sample_multiple* but works with an array of particles. (Aside: Python doesn't let us use *sample_multiple* directly as it uses a dictionary, and particles, represented as dictionaries can't be the key of dictionaries).

```

_____probStochSim.py — (continued)_____
176 def resample(particles, weights, num_samples):
177     """returns num_samples copies of particles resampled according to weights.
178     particles is a list of particles
179     weights is a list of positive numbers, of same length as particles
180     num_samples is n integer
181     """
182     total = sum(weights)
183     rands = sorted(random.random()*total for i in range(num_samples))
184     result = []
185     cum = weights[0] # cumulative sum
186     index = 0
187     for r in rands:
188         while r>cum:
189             index += 1
190             cum += weights[index]
191         result.append(particles[index])
192     return result

```

8.8.6 Examples

```

_____probStochSim.py — (continued)_____
194 from probGraphicalModels import bn_4ch, A,B,C,D
195 bn_4chr = RejectionSampling(bn_4ch)

```

```

196 bn_4chL = LikelihoodWeighting(bn_4ch)
197 ## InferenceMethod.max_display_level = 2 # detailed tracing for all inference methods
198 ## bn_4chr.query(A,{})
199 ## bn_4chr.query(C,{})
200 ## bn_4chr.query(A,{C:True})
201 ## bn_4chr.query(B,{A:True,C:False})
202
203 from probGraphicalModels import bn_report,Alarm,Fire,Leaving,Report,Smoke,Tamper
204 bn_reportr = RejectionSampling(bn_report) # answers queries using rejection sampling
205 bn_reportL = LikelihoodWeighting(bn_report) # answers queries using rejection sampling
206 bn_reportp = ParticleFiltering(bn_report) # answers queries using particle filtering
207 ## bn_reportr.query(Tamper,{})
208 ## bn_reportr.query(Tamper,{})
209 ## bn_reportr.query(Tamper,{Report:True})
210 ## InferenceMethod.max_display_level = 0 # no detailed tracing for all inference methods
211 ## bn_reportr.query(Tamper,{Report:True},number_samples=100000)
212 ## bn_reportr.query(Tamper,{Report:True,Smoke:False})
213 ## bn_reportr.query(Tamper,{Report:True,Smoke:False},number_samples=100)
214
215 ## bn_reportL.query(Tamper,{Report:True,Smoke:False},number_samples=100)
216 ## bn_reportL.query(Tamper,{Report:True,Smoke:False},number_samples=100)
217
218
219 from probGraphicalModels import bn_sprinkler,Season, Sprinkler
220 from probGraphicalModels import Rained, Grass_wet, Grass_shiny, Shoes_wet
221 bn_sprinklerr = RejectionSampling(bn_sprinkler) # answers queries using rejection sampling
222 bn_sprinklerL = LikelihoodWeighting(bn_sprinkler) # answers queries using rejection sampling
223 bn_sprinklerp = ParticleFiltering(bn_sprinkler) # answers queries using particle filtering
224 #bn_sprinklerr.query(Shoes_wet,{Grass_shiny:True,Rained:True})
225 #bn_sprinklerL.query(Shoes_wet,{Grass_shiny:True,Rained:True})
226 #bn_sprinklerp.query(Shoes_wet,{Grass_shiny:True,Rained:True})
227
228 if __name__ == "__main__":
229     InferenceMethod.testIM(RejectionSampling, threshold=0.1)
230     InferenceMethod.testIM(LikelihoodWeighting, threshold=0.1)
231     InferenceMethod.testIM(LikelihoodWeighting, threshold=0.1)

```

Exercise 8.3 This code keeps regenerating the distribution of a variable given its parents. Implement one or both of the following, and compare them to the original. Make *cond_dist* return a slice that corresponds to the distribution, and then use the slice instead of the dictionary (a list slice does not generate new data structures). Make *cond_dist* remember values it has already computed, and only return these.

8.8.7 Gibbs Sampling

The following implements **Gibbs sampling**, a form of **Markov Chain Monte Carlo MCMC**.

probStochSim.py — (continued) —

```

233 #import random
234 #from probGraphicalModels import InferenceMethod
235
236 #from probStochSim import sample_one, SamplingInferenceMethod
237
238 class GibbsSampling(SamplingInferenceMethod):
239     """The class that queries Graphical Models using Gibbs Sampling.
240
241     bn is a graphical model (e.g., a belief network) to query
242     """
243     method_name = "Gibbs sampling"
244
245     def __init__(self, bn=None):
246         self.bn = bn
247
248     def query(self, qvar, obs={}, number_samples=1000, burn_in=100, sample_order=None):
249         """computes P(qvar|obs) where
250         qvar is a variable.
251         obs is a {variable:value} dictionary.
252         sample_order is a list of non-observed variables in order, or
253         if sample_order None, the variables are shuffled at each iteration.
254         """
255         counts = {val:0 for val in qvar.domain}
256         if sample_order is not None:
257             variables = sample_order
258         else:
259             variables = [v for v in self.bn.variables if v not in obs]
260         var_to_factors = {v:set() for v in self.bn.variables}
261         for fac in self.bn.factors:
262             for var in fac.variables:
263                 var_to_factors[var].add(fac)
264         sample = {var:random.choice(var.domain) for var in variables}
265         self.display(2, "Sample:", sample)
266         sample.update(obs)
267         for i in range(burn_in + number_samples):
268             if sample_order == None:
269                 random.shuffle(variables)
270             for var in variables:
271                 # get unnormalized probability distribution of var given its neighbours
272                 vardist = {val:1 for val in var.domain}
273                 for val in var.domain:
274                     sample[var] = val
275                     for fac in var_to_factors[var]: # Markov blanket
276                         vardist[val] *= fac.get_value(sample)
277                     sample[var] = sample_one(vardist)
278             if i >= burn_in:
279                 counts[sample[qvar]] +=1
280         tot = sum(counts.values())
281         # as well as the computed distribution, we also include raw counts
282         return {c:v/tot for (c,v) in counts.items()} | {"raw_counts":counts}

```

```

283 |
284 | #from probGraphicalModels import bn_4ch, A,B,C,D
285 | bn_4chg = GibbsSampling(bn_4ch)
286 | ## InferenceMethod.max_display_level = 2 # detailed tracing for all inference methods
287 | bn_4chg.query(A,{})
288 | ## bn_4chg.query(D,{})
289 | ## bn_4chg.query(B,{D:True})
290 | ## bn_4chg.query(B,{A:True,C:False})
291 |
292 | from probGraphicalModels import bn_report,Alarm,Fire,Leaving,Report,Smoke,Tamper
293 | bn_reportg = GibbsSampling(bn_report)
294 | ## bn_reportg.query(Tamper,{Report:True},number_samples=1000)
295 |
296 | if __name__ == "__main__":
297 |     InferenceMethod.testIM(GibbsSampling, threshold=0.1)

```

Exercise 8.4 Change the code so that it can have multiple query variables. Make the list of query variable be an input to the algorithm, so that the default value is the list of all non-observed variables.

Exercise 8.5 In this algorithm, explain where it computes the probability of a variable given its Markov blanket. Instead of returning the average of the samples for the query variable, it is possible to return the average estimate of the probability of the query variable given its Markov blanket. Does this converge to the same answer as the given code? Does it converge faster, slower, or the same?

8.8.8 Plotting Behaviour of Stochastic Simulators

The stochastic simulation runs can give different answers each time they are run. For the algorithms that give the same answer in the limit as the number of samples approaches infinity (as do all of these algorithms), the algorithms can be compared by comparing the accuracy for multiple runs. Summary statistics like the variance may provide some information, but the assumptions behind the variance being appropriate (namely that the distribution is approximately Gaussian) may not hold for cases where the predictions are bounded and often skewed.

It is more appropriate to plot the distribution of predictions over multiple runs. The *plot_stats* method plots the prediction of a particular variable (or for the partition function) for a number of runs of the same algorithm. On the *x*-axis, is the prediction of the algorithm. On the *y*-axis is the number of runs with prediction less than or equal to the *x* value. Thus this is like a cumulative distribution over the predictions, but with counts on the *y*-axis.

Note that for runs where there are no samples that are consistent with the observations (as can happen with rejection sampling), the prediction of probability is 1.0 (as a convention for 0/0).

That variable *what* contains the query variable, or *what* is “*prob_ev*”, the probability of evidence.

probStochSim.py — (continued)

```

299 import matplotlib.pyplot as plt
300
301 def plot_stats(method, qvar, qval, obs, number_runs=1000, **queryargs):
302     """Plots a cumulative distribution of the prediction of the model.
303     method is a InferenceMethod (that implements appropriate query(.))
304     plots P(qvar=qval | obs)
305     qvar is the query variable, qval is corresponding value
306     obs is the {variable:value} dictionary representing the observations
307     number_iterations is the number of runs that are plotted
308     **queryargs is the arguments to query (often number_samples for sampling methods)
309     """
310     plt.ion()
311     plt.xlabel("value")
312     plt.ylabel("Cumulative Number")
313     method.max_display_level, prev_mdl = 0, method.max_display_level #no display
314     answers = [method.query(qvar, obs, **queryargs)
315                 for i in range(number_runs)]
316     values = [ans[qval] for ans in answers]
317     label = f"{method.method_name} P({qvar}={qval})|{'', '.'}.join(f'{var}={val}' for (var, val) in obs.items())"
318     values.sort()
319     plt.plot(values, range(number_runs), label=label)
320     plt.legend() #loc="upper left")
321     plt.draw()
322     method.max_display_level = prev_mdl # restore display level
323
324 # Try:
325 # plot_stats(bn_reportr, Tamper, True, {Report: True, Smoke: True}, number_samples=1000, number_runs=1000)
326 # plot_stats(bn_reportL, Tamper, True, {Report: True, Smoke: True}, number_samples=1000, number_runs=1000)
327 # plot_stats(bn_reportp, Tamper, True, {Report: True, Smoke: True}, number_samples=1000, number_runs=1000)
328 # plot_stats(bn_reportr, Tamper, True, {Report: True, Smoke: True}, number_samples=100, number_runs=1000)
329 # plot_stats(bn_reportL, Tamper, True, {Report: True, Smoke: True}, number_samples=100, number_runs=1000)
330 # plot_stats(bn_reportg, Tamper, True, {Report: True, Smoke: True}, number_samples=1000, number_runs=1000)
331
332 def plot_mult(methods, example, qvar, qval, obs, number_samples=1000, number_runs=1000):
333     for method in methods:
334         solver = method(example)
335         if isinstance(method, SamplingInferenceMethod):
336             plot_stats(solver, qvar, qval, obs, number_samples, number_runs)
337         else:
338             plot_stats(solver, qvar, qval, obs, number_runs)
339
340 from probRC import RC
341 # Try following (but it takes a while..)
342 methods = [RC, RejectionSampling, LikelihoodWeighting, ParticleFiltering, GibbsSampling]
343 # plot_mult(methods, bn_report, Tamper, True, {Report: True, Smoke: False}, number_samples=100, number_runs=1000)
344 # plot_mult(methods, bn_report, Tamper, True, {Report: True, Smoke: True}, number_samples=100, number_runs=1000)
345
346 # Sprinkler Example:
347 # plot_stats(bn_sprinklerr, Shoes_wet, True, {Grass_shiny: True, Rained: True}, number_samples=1000)

```

```
348 | # plot_stats(bn_sprinklerL, Shoes_wet, True, {Grass_shiny: True, Rained: True}, number_samples=1000)
```

8.9 Hidden Markov Models

This code for hidden Markov models is independent of the graphical models code, to keep it simple. Section 8.10 gives code that models hidden Markov models, and more generally, dynamic belief networks, using the graphical models code.

This HMM code assumes there are multiple Boolean observation variables that depend on the current state and are independent of each other given the state.

```

_____probHMM.py — Hidden Markov Model_____
11 import random
12 from probStochSim import sample_one, sample_multiple
13
14 class HMM(object):
15     def __init__(self, states, obsvars, pobs, trans, indist):
16         """A hidden Markov model.
17         states - set of states
18         obsvars - set of observation variables
19         pobs - probability of observations, pobs[i][s] is P(Obs_i=True | State=s)
20         trans - transition probability - trans[i][j] gives P(State=j | State=i)
21         indist - initial distribution - indist[s] is P(State_0 = s)
22         """
23         self.states = states
24         self.obsvars = obsvars
25         self.pobs = pobs
26         self.trans = trans
27         self.indist = indist

```

Consider the following example. Suppose you want to unobtrusively keep track of an animal in a triangular enclosure using sound. Suppose you have 3 microphones that provide unreliable (noisy) binary information at each time step. The animal is either close to one of the 3 points of the triangle or in the middle of the triangle.

```

_____probHMM.py — (continued)_____
29 # state
30 #     0=middle, 1,2,3 are corners
31 states1 = {'middle', 'c1', 'c2', 'c3'} # states
32 obs1 = {'m1', 'm2', 'm3'} # microphones

```

The observation model is as follows. If the animal is in a corner, it will be detected by the microphone at that corner with probability 0.6, and will be independently detected by each of the other microphones with a probability of 0.1. If the animal is in the middle, it will be detected by each microphone with a probability of 0.4.

```

probHMM.py — (continued)
34 # pobs gives the observation model:
35 #pobs[mi][state] is P(mi=on | state)
36 closeMic=0.6; farMic=0.1; midMic=0.4
37 pobs1 = {'m1':{'middle':midMic, 'c1':closeMic, 'c2':farMic, 'c3':farMic}, # mic 1
38         'm2':{'middle':midMic, 'c1':farMic, 'c2':closeMic, 'c3':farMic}, # mic 2
39         'm3':{'middle':midMic, 'c1':farMic, 'c2':farMic, 'c3':closeMic}} # mic 3

```

The transition model is as follows: If the animal is in a corner it stays in the same corner with probability 0.80, goes to the middle with probability 0.1 or goes to one of the other corners with probability 0.05 each. If it is in the middle, it stays in the middle with probability 0.7, otherwise it moves to one the corners, each with probability 0.1.

```

probHMM.py — (continued)
41 # trans specifies the dynamics
42 # trans[i] is the distribution over states resulting from state i
43 # trans[i][j] gives P(S=j | S=i)
44 sm=0.7; mmc=0.1 # transition probabilities when in middle
45 sc=0.8; mcm=0.1; mcc=0.05 # transition probabilities when in a corner
46 trans1 = {'middle':{'middle':sm, 'c1':mmc, 'c2':mmc, 'c3':mmc}, # was in middle
47          'c1':{'middle':mcm, 'c1':sc, 'c2':mcc, 'c3':mcc}, # was in corner 1
48          'c2':{'middle':mcm, 'c1':mcc, 'c2':sc, 'c3':mcc}, # was in corner 2
49          'c3':{'middle':mcm, 'c1':mcc, 'c2':mcc, 'c3':sc}} # was in corner 3

```

Initially the animal is in one of the four states, with equal probability.

```

probHMM.py — (continued)
51 # initially we have a uniform distribution over the animal's state
52 indist1 = {st:1.0/len(states1) for st in states1}
53
54 hmm1 = HMM(states1, obs1, pobs1, trans1, indist1)

```

8.9.1 Exact Filtering for HMMs

A *HMMVEfilter* has a current state distribution which can be updated by observing or by advancing to the next time.

```

probHMM.py — (continued)
56 from display import Displayable
57
58 class HMMVEfilter(Displayable):
59     def __init__(self, hmm):
60         self.hmm = hmm
61         self.state_dist = hmm.indist
62
63     def filter(self, obsseq):
64         """updates and returns the state distribution following the sequence of
65         observations in obsseq using variable elimination.
66

```

```

67     Note that it first advances time.
68     This is what is required if it is called sequentially.
69     If that is not what is wanted initially, do an observe first.
70     """
71     for obs in obsseq:
72         self.advance()    # advance time
73         self.observe(obs) # observe
74     return self.state_dist
75
76     def observe(self, obs):
77         """updates state conditioned on observations.
78         obs is a list of values for each observation variable"""
79         for i in self.hmm.obsvars:
80             self.state_dist = {st:self.state_dist[st]*(self.hmm.pobs[i][st]
81                                                         if obs[i] else (1-self.hmm.pobs[i][st]))
82                               for st in self.hmm.states}
83         norm = sum(self.state_dist.values()) # normalizing constant
84         self.state_dist = {st:self.state_dist[st]/norm for st in self.hmm.states}
85         self.display(2,"After observing",obs,"state distribution:",self.state_dist)
86
87     def advance(self):
88         """advance to the next time"""
89         nextstate = {st:0.0 for st in self.hmm.states} # distribution over next states
90         for j in self.hmm.states:    # j ranges over next states
91             for i in self.hmm.states: # i ranges over previous states
92                 nextstate[j] += self.hmm.trans[i][j]*self.state_dist[i]
93         self.state_dist = nextstate

```

The following are some queries for *hmm1*.

```

_____probHMM.py — (continued)_____
95 hmm1f1 = HMMVEfilter(hmm1)
96 # hmm1f1.filter([{'m1':0, 'm2':1, 'm3':1}, {'m1':1, 'm2':0, 'm3':1}])
97 ## HMMVEfilter.max_display_level = 2 # show more detail in displaying
98 # hmm1f2 = HMMVEfilter(hmm1)
99 # hmm1f2.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':1, 'm3':0}, {'m1':1, 'm2':0, 'm3':0},
100 #                {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0},
101 #                {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':1}, {'m1':0, 'm2':0, 'm3':1},
102 #                {'m1':0, 'm2':0, 'm3':1}])
103 # hmm1f3 = HMMVEfilter(hmm1)
104 # hmm1f3.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':1, 'm2':0, 'm3':0}, {'m1':
105 #
106 # How do the following differ in the resulting state distribution?
107 # Note they start the same, but have different initial observations.
108 ## HMMVEfilter.max_display_level = 1 # show less detail in displaying
109 # for i in range(100): hmm1f1.advance()
110 # hmm1f1.state_dist
111 # for i in range(100): hmm1f3.advance()
112 # hmm1f3.state_dist

```

Exercise 8.6 The localization example in the book is a controlled HMM, where

there is a given action at each time and the transition depends on the action. Change the code to allow for controlled HMMs. Hint: the action only influences the state transition.

Exercise 8.7 The representation assumes that there are a list of Boolean observations. Extend the representation so that the each observation variable can have multiple discrete values. You need to choose a representation for the model, and change the algorithm.

8.9.2 Particle Filtering for HMMs

In this implementation a particle is just a state. If you want to do some form of smoothing, a particle should probably be a history of states. This maintains, *particles*, an array of states, *weights* an array of (non-negative) real numbers, such that *weights*[*i*] is the weight of *particles*[*i*].

```

probHMM.py — (continued)
113 from display import Displayable
114 from probStochSim import resample
115
116 class HMMparticleFilter(Displayable):
117     def __init__(self, hmm, number_particles=1000):
118         self.hmm = hmm
119         self.particles = [sample_one(hmm.indist)
120                           for i in range(number_particles)]
121         self.weights = [1 for i in range(number_particles)]
122
123     def filter(self, obsseq):
124         """returns the state distribution following the sequence of
125         observations in obsseq using particle filtering.
126
127         Note that it first advances time.
128         This is what is required if it is called after previous filtering.
129         If that is not what is wanted initially, do an observe first.
130         """
131         for obs in obsseq:
132             self.advance() # advance time
133             self.observe(obs) # observe
134             self.resample_particles()
135             self.display(2, "After observing", str(obs),
136                         "state distribution:", self.histogram(self.particles))
137         self.display(1, "Final state distribution:", self.histogram(self.particles))
138         return self.histogram(self.particles)
139
140     def advance(self):
141         """advance to the next time.
142         This assumes that all of the weights are 1."""
143         self.particles = [sample_one(self.hmm.trans[st])
144                           for st in self.particles]
145

```

```

146 def observe(self, obs):
147     """reweighs the particles to incorporate observations obs"""
148     for i in range(len(self.particles)):
149         for obv in obs:
150             if obs[obv]:
151                 self.weights[i] *= self.hmm.pobs[obv][self.particles[i]]
152             else:
153                 self.weights[i] *= 1-self.hmm.pobs[obv][self.particles[i]]
154
155 def histogram(self, particles):
156     """returns list of the probability of each state as represented by
157     the particles"""
158     tot=0
159     hist = {st: 0.0 for st in self.hmm.states}
160     for (st,wt) in zip(self.particles,self.weights):
161         hist[st]+=wt
162         tot += wt
163     return {st:hist[st]/tot for st in hist}
164
165 def resample_particles(self):
166     """resamples to give a new set of particles."""
167     self.particles = resample(self.particles, self.weights, len(self.particles))
168     self.weights = [1] * len(self.particles)

```

The following are some queries for *hmm1*.

```

probHMM.py — (continued)
170 hmm1pf1 = HMMparticleFilter(hmm1)
171 # HMMparticleFilter.max_display_level = 2 # show each step
172 # hmm1pf1.filter([{'m1':0, 'm2':1, 'm3':1}, {'m1':1, 'm2':0, 'm3':1}])
173 # hmm1pf2 = HMMparticleFilter(hmm1)
174 # hmm1pf2.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':1, 'm3':0}, {'m1':1, 'm2':0, 'm3':0},
175 #                 {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0},
176 #                 {'m1':0, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':1}, {'m1':0, 'm2':0, 'm3':1},
177 #                 {'m1':0, 'm2':0, 'm3':1}])
178 # hmm1pf3 = HMMparticleFilter(hmm1)
179 # hmm1pf3.filter([{'m1':1, 'm2':0, 'm3':0}, {'m1':0, 'm2':0, 'm3':0}, {'m1':1, 'm2':0, 'm3':0}, {'m1':

```

Exercise 8.8 A form of importance sampling can be obtained by not resampling. Is it better or worse than particle filtering? Hint: you need to think about how they can be compared. Is the comparison different if there are more states than particles?

Exercise 8.9 Extend the particle filtering code to continuous variables and observations. In particular, suppose the state transition is a linear function with Gaussian noise of the previous state, and the observations are linear functions with Gaussian noise of the state. You may need to research how to sample from a Gaussian distribution.

8.9.3 Generating Examples

The following code is useful for generating examples.

```

181 def simulate(hmm,horizon):
182     """returns a pair of (state sequence, observation sequence) of length horizon.
183     for each time t, the agent is in state_sequence[t] and
184     observes observation_sequence[t]
185     """
186     state = sample_one(hmm.indist)
187     obsseq=[]
188     stateseq=[]
189     for time in range(horizon):
190         stateseq.append(state)
191         newobs = {obs:sample_one({0:1-hmm.pobs[obs][state],1:hmm.pobs[obs][state]})
192                   for obs in hmm.obsvars}
193         obsseq.append(newobs)
194         state = sample_one(hmm.trans[state])
195     return stateseq,obsseq
196
197 def simobs(hmm,stateseq):
198     """returns observation sequence for the state sequence"""
199     obsseq=[]
200     for state in stateseq:
201         newobs = {obs:sample_one({0:1-hmm.pobs[obs][state],1:hmm.pobs[obs][state]})
202                   for obs in hmm.obsvars}
203         obsseq.append(newobs)
204     return obsseq
205
206 def create_eg(hmm,n):
207     """Create an annotated example for horizon n"""
208     seq,obs = simulate(hmm,n)
209     print("True state sequence:",seq)
210     print("Sequence of observations:\n",obs)
211     hmmfilter = HMMVEfilter(hmm)
212     dist = hmmfilter.filter(obs)
213     print("Resulting distribution over states:\n",dist)

```

8.10 Dynamic Belief Networks

A **dynamic belief network (DBN)** is a belief network that extends in time.

There are a number of ways that reasoning can be carried out in a DBN, including:

- Rolling out the DBN for some time period, and using standard belief network inference. The latest time that needs to be in the rolled out network is the time of the latest observation or the time of a query (whichever is

later). This allows us to observe any variables at any time and query any variables at any time. This is covered in Section 8.10.2.

- An unrolled belief network may be very large, and we might only be interested in asking about “now”. In this case we can just representing the variables “now”. In this approach we can observe and query the current variables. We can then move to the next time. This does not allow for arbitrary historical queries (about the past or the future), but can be much simpler. This is covered in Section 8.10.3.

8.10.1 Representing Dynamic Belief Networks

To specify a DBN, think about the distribution *now*. *Now* will be represented as time 1. Each variable will have a corresponding previous variable; these will be created together.

A dynamic belief network consists of:

- A set of features. A variable is a feature-time pair.
- An initial distribution over the features “now” (time 1). This is a belief network with all variables being time 1 variables.
- A specification of the dynamics. We define the how the variables *now* (time 1) depend on variables *now* and the previous time (time 0), in such a way that the graph is acyclic.

```

_____probDBN.py — Dynamic belief networks_____
11 from probVariables import Variable
12 from probGraphicalModels import GraphicalModel, BeliefNetwork
13 from probFactors import Prob
14 from probVE import VE
15 from display import Displayable
16
17 class DBNvariable(Variable):
18     """A random variable that incorporates the stage (time)
19
20     A variable can have both a name and an index. The index defaults to 1.
21     """
22     def __init__(self, name, domain=[False, True], index=1):
23         Variable.__init__(self, f"{name}_{index}", domain)
24         self.basename = name
25         self.domain = domain
26         self.index = index
27         self.previous = None
28
29     def __lt__(self, other):
30         if self.name != other.name:
31             return self.name < other.name

```



```

32         else:
33             return self.index < other.index
34
35     def __gt__(self, other):
36         return other < self
37
38 def variable_pair(name, domain=[False, True]):
39     """returns a variable and its predecessor. This is used to define 2-stage DBNs
40
41     If the name is X, it returns the pair of variables X_prev, X_now"""
42     var_now = DBNvariable(name, domain, index='now')
43     var_prev = DBNvariable(name, domain, index='prev')
44     var_now.previous = var_prev
45     return var_prev, var_now

```

A *FactorRename* is a factor that is the result renaming the variables in the factor. It takes a factor, *fac*, and a *{new : old}* dictionary, where *new* is the name of a variable in the resulting factor and *old* is the corresponding name in *fac*. This assumes that the all variables are renamed.

```

_____probDBN.py — (continued)_____
47 class FactorRename(Factor):
48     def __init__(self, fac, renaming):
49         """A renamed factor.
50         fac is a factor
51         renaming is a dictionary of the form {new:old} where old and new var variables,
52         where the variables in fac appear exactly once in the renaming
53         """
54         Factor.__init__(self, [n for (n,o) in renaming.items() if o in fac.variables])
55         self.orig_fac = fac
56         self.renaming = renaming
57
58     def get_value(self, assignment):
59         return self.orig_fac.get_value({self.renaming[var]:val
60                                         for (var,val) in assignment.items()
61                                         if var in self.variables})

```

The following class renames the variables of a conditional probability distribution. It is used for template models (e.g., dynamic decision networks or relational models)

```

_____probDBN.py — (continued)_____
63 class CPDRename(FactorRename, CPD):
64     def __init__(self, cpd, renaming):
65         renaming_inverse = {old:new for (new,old) in renaming.items()}
66         CPD.__init__(self, renaming_inverse[cpd.child], [renaming_inverse[p] for p in cpd.parents])
67         self.orig_fac = cpd
68         self.renaming = renaming

```

```

_____probDBN.py — (continued)_____
70 class DBN(Displayable):

```

```

71     """The class of stationary Dynamic Belief networks.
72     * name is the DBN name
73     * vars_now is a list of current variables (each must have
74     previous variable).
75     * transition_factors is a list of factors for P(X|parents) where X
76     is a current variable and parents is a list of current or previous variables.
77     * init_factors is a list of factors for P(X|parents) where X is a
78     current variable and parents can only include current variables
79     The graph of transition factors + init factors must be acyclic.
80
81     """
82     def __init__(self, title, vars_now, transition_factors=None, init_factors=None):
83         self.title = title
84         self.vars_now = vars_now
85         self.vars_prev = [v.previous for v in vars_now]
86         self.transition_factors = transition_factors
87         self.init_factors = init_factors
88         self.var_index = {} # var_index[v] is the index of variable v
89         for i,v in enumerate(vars_now):
90             self.var_index[v]=i

```

Here is a 3 variable DBN:

```

probDBN.py — (continued)
92 A0,A1 = variable_pair("A", domain=[False,True])
93 B0,B1 = variable_pair("B", domain=[False,True])
94 C0,C1 = variable_pair("C", domain=[False,True])
95
96 # dynamics
97 pc = Prob(C1,[B1,C0],[[0.03,0.97],[0.38,0.62],[0.23,0.77],[0.78,0.22]])
98 pb = Prob(B1,[A0,A1],[[0.5,0.5],[0.77,0.23],[0.4,0.6],[0.83,0.17]])
99 pa = Prob(A1,[A0,B0],[[0.1,0.9],[0.65,0.35],[0.3,0.7],[0.8,0.2]])
100
101 # initial distribution
102 pa0 = Prob(A1,[],[0.9,0.1])
103 pb0 = Prob(B1,[A1],[0.3,0.7],[0.8,0.2])
104 pc0 = Prob(C1,[],[0.2,0.8])
105
106 db1 = DBN("Simple DBN",[A1,B1,C1],[pa,pb,pc],[pa0,pb0,pc0])

```

Here is the animal example

```

probDBN.py — (continued)
108 from probHMM import closeMic, farMic, midMic, sm, mmc, sc, mcm, mcc
109
110 Pos_0,Pos_1 = variable_pair("Position",domain=[0,1,2,3])
111 Mic1_0,Mic1_1 = variable_pair("Mic1")
112 Mic2_0,Mic2_1 = variable_pair("Mic2")
113 Mic3_0,Mic3_1 = variable_pair("Mic3")
114
115 # conditional probabilities - see hmm for the values of sm,mmc, etc
116 ppos = Prob(Pos_1, [Pos_0],

```

```

117         [[sm, mmc, mmc, mmc], #was in middle
118         [mcm, sc, mcc, mcc], #was in corner 1
119         [mcm, mcc, sc, mcc], #was in corner 2
120         [mcm, mcc, mcc, sc]]) #was in corner 3
121 pm1 = Prob(Mic1_1, [Pos_1], [[1-midMic, midMic], [1-closeMic, closeMic],
122                             [1-farMic, farMic], [1-farMic, farMic]])
123 pm2 = Prob(Mic2_1, [Pos_1], [[1-midMic, midMic], [1-farMic, farMic],
124                             [1-closeMic, closeMic], [1-farMic, farMic]])
125 pm3 = Prob(Mic3_1, [Pos_1], [[1-midMic, midMic], [1-farMic, farMic],
126                             [1-farMic, farMic], [1-closeMic, closeMic]])
127 ipos = Prob(Pos_1, [], [0.25, 0.25, 0.25, 0.25])
128 dbn_an =DBN("Animal DBN",[Pos_1,Mic1_1,Mic2_1,Mic3_1],
129            [ppos, pm1, pm2, pm3],
130            [ipos, pm1, pm2, pm3])

```

8.10.2 Unrolling DBNs

```

probDBN.py — (continued)
132 class BNfromDBN(BeliefNetwork):
133     """Belief Network unrolled from a dynamic belief network
134     """
135
136     def __init__(self,dbn,horizon):
137         """dbn is the dynamic belief network being unrolled
138         horizon>0 is the number of steps (so there will be horizon+1 variables for each DBN variable)
139         """
140         self.name2var = {var.basename: [DBNvariable(var.basename,var.domain,index) for index in range(horizon+1)]
141                         for var in dbn.vars_now}
142         self.display(1,f"name2var={self.name2var}")
143         variables = {v for vs in self.name2var.values() for v in vs}
144         self.display(1,f"variables={variables}")
145         bnfactors = {CPDrename(fac,{self.name2var[var.basename][0]:var
146                                for var in fac.variables})
147                     for fac in dbn.init_factors}
148         bnfactors |= {CPDrename(fac,{self.name2var[var.basename][i]:var
149                                for var in fac.variables if var.index=='prev'}
150                        | {self.name2var[var.basename][i+1]:var
151                           for var in fac.variables if var.index=='now'})
152                     for fac in dbn.transition_factors
153                     for i in range(horizon)}
154         self.display(1,f"bnfactors={bnfactors}")
155         BeliefNetwork.__init__(self, dbn.title, variables, bnfactors)

```

Here are two examples. Note that we need to use `bn.name2var['B'][2]` to get the variable B2 (B at time 2).

```

157 # Try
158 #from probRC import RC

```

```

159 |bn = BNfromDBN(dbn1,2) # construct belief network
160 |drc = RC(bn)           # initialize recursive conditioning
161 |B2 = bn.name2var['B'][2]
162 |drc.query(B2) #P(B2)
163 |drc.query(bn.name2var['B'][1],{bn.name2var['B'][0]:True,bn.name2var['C'][1]:False}) #P(B1|B0,C1)

```

8.10.3 DBN Filtering

If we only wanted to ask questions about the current state, we can save space by forgetting the history variables.

```

_____probDBN.py — (continued) _____
164 |class DBNVEfilter(VE):
165 |    def __init__(self,dbn):
166 |        self.dbn = dbn
167 |        self.current_factors = dbn.init_factors
168 |        self.current_obs = {}
169 |
170 |    def observe(self, obs):
171 |        """updates the current observations with obs.
172 |        obs is a variable:value dictionary where variable is a current
173 |        variable.
174 |        """
175 |        assert all(self.current_obs[var]==obs[var] for var in obs
176 |                   if var in self.current_obs),"inconsistent current observations"
177 |        self.current_obs.update(obs) # note 'update' is a dict method
178 |
179 |    def query(self,var):
180 |        """returns the posterior probability of current variable var"""
181 |        return VE(GraphicalModel(self.dbn.title,self.dbn.vars_now,self.current_factors)).query(var,self)
182 |
183 |    def advance(self):
184 |        """advance to the next time"""
185 |        prev_factors = [self.make_previous(fac) for fac in self.current_factors]
186 |        prev_obs = {var.previous:val for var,val in self.current_obs.items()}
187 |        two_stage_factors = prev_factors + self.dbn.transition_factors
188 |        self.current_factors = self.elim_vars(two_stage_factors,self.dbn.vars_prev,prev_obs)
189 |        self.current_obs = {}
190 |
191 |    def make_previous(self,fac):
192 |        """Creates new factor from fac where the current variables in fac
193 |        are renamed to previous variables.
194 |        """
195 |        return FactorRename(fac, {var.previous:var for var in fac.variables})
196 |
197 |    def elim_vars(self,factors, vars, obs):
198 |        for var in vars:
199 |            if var in obs:
200 |                factors = [self.project_observations(fac,obs) for fac in factors]
201 |            else:
202 |                factors = self.eliminate_var(factors, var)

```

203 | **return** factors

Example queries:

```
_____probDBN.py — (continued) _____  
205 | #df = DBNVEfilter(dbn1)  
206 | #df.observe({B1:True}); df.advance(); df.observe({C1:False})  
207 | #df.query(B1) #P(B1|B0,C1)  
208 | #df.advance(); df.query(B1)  
209 | #dfa = DBNVEfilter(dbn_an)  
210 | # dfa.observe({Mic1_1:0, Mic2_1:1, Mic3_1:1})  
211 | # dfa.advance()  
212 | # dfa.observe({Mic1_1:1, Mic2_1:0, Mic3_1:1})  
213 | # dfa.query(Pos_1)
```


Planning with Uncertainty

9.1 Decision Networks

The decision network code builds on the representation for belief networks of Chapter 8.

We first allow for factors that define the utility. Here the **utility** is a function of the variables in *vars*. In a **utility table** the utility is defined in terms of a, a list that enumerates the values as in Section 8.3.3.

```
-----decnNetworks.py — Representations for Decision Networks-----
11 from probGraphicalModels import GraphicalModel, BeliefNetwork
12 from probFactors import Factor, CPD, TabFactor, factor_times, Prob
13 from probVariables import Variable
14 import matplotlib.pyplot as plt
15
16 class Utility(Factor):
17     """A factor defining a utility"""
18     pass
19
20 class UtilityTable(TabFactor, Utility):
21     """A factor defining a utility using a table"""
22     def __init__(self, vars, table, position=None):
23         """Creates a factor on vars from the table.
24         The table is ordered according to vars.
25         """
26         TabFactor.__init__(self, vars, table)
27         self.position = position
```

A **decision variable** is like a random variable with a string name, and a domain, which is a list of possible values. The decision variable also includes the parents, a list of the variables whose value will be known when the decision is made. It also includes a position, which is only used for plotting.

```

decnNetworks.py — (continued)
29 class DecisionVariable(Variable):
30     def __init__(self, name, domain, parents, position=None):
31         Variable.__init__(self, name, domain, position)
32         self.parents = parents
33         self.all_vars = set(parents) | {self}

```

A decision network is a graphical model where the variables can be random variables or decision variables. Among the factors we assume there is one utility factor.

```

decnNetworks.py — (continued)
35 class DecisionNetwork(BeliefNetwork):
36     def __init__(self, title, vars, factors):
37         """vars is a list of variables
38         factors is a list of factors (instances of CPD and Utility)
39         """
40         GraphicalModel.__init__(self, title, vars, factors) # don't call init for BeliefNetwork
41         self.var2parents = ({v : v.parents for v in vars if isinstance(v,DecisionVariable)}
42                             | {f.child:f.parents for f in factors if isinstance(f,CPD)})
43         self.children = {n:[] for n in self.variables}
44         for v in self.var2parents:
45             for par in self.var2parents[v]:
46                 self.children[par].append(v)
47         self.utility_factor = [f for f in factors if isinstance(f,Utility)][0]
48         self.topological_sort_saved = None

```

The split order ensures that the parents of a decision node are split before the decision node, and no other variables (if that is possible).

```

decnNetworks.py — (continued)
50 def split_order(self):
51     so = []
52     tops = self.topological_sort()
53     for v in tops:
54         if isinstance(v,DecisionVariable):
55             so += [p for p in v.parents if p not in so]
56             so.append(v)
57     so += [v for v in tops if v not in so]
58     return so

```

```

decnNetworks.py — (continued)
60 def show(self):
61     plt.ion() # interactive
62     ax = plt.figure().gca()
63     ax.set_axis_off()
64     plt.title(self.title)
65     for par in self.utility_factor.variables:
66         ax.annotate("Utility", par.position, xytext=self.utility_factor.position,
67                    arrowprops={'arrowstyle': '<-'}, bbox=dict(boxstyle="sawtooth,pad=1.0",
68                    ha='center')

```





Figure 9.1: The umbrella decision network

```

69     for var in reversed(self.topological_sort()):
70         if isinstance(var, DecisionVariable):
71             bbox = dict(boxstyle="square,pad=1.0",color="green")
72         else:
73             bbox = dict(boxstyle="round4,pad=1.0,rounding_size=0.5")
74         if self.var2parents[var]:
75             for par in self.var2parents[var]:
76                 ax.annotate(var.name, par.position, xytext=var.position,
77                             arrowprops={'arrowstyle': '<-'},bbox=bbox,
78                             ha='center')
79         else:
80             x,y = var.position
81             plt.text(x,y,var.name,bbox=bbox,ha='center')

```

9.1.1 Example Decision Networks

Umbrella Decision Network

Here is a simple “umbrella” decision network. The output of `umbrella_dn.show()` is shown in Figure 9.1.

```

decnNetworks.py — (continued)
83 Weather = Variable("Weather", ["NoRain", "Rain"], position=(0.5,0.8))
84 Forecast = Variable("Forecast", ["Sunny", "Cloudy", "Rainy"], position=(0,0.4))
85 # Each variant uses one of the following:
86 Umbrella = DecisionVariable("Umbrella", ["Take", "Leave"], {Forecast}, position=(0.5,0))
87
88 p_weather = Prob(Weather, [], [0.7, 0.3])
89 p_forecast = Prob(Forecast, [Weather], [[0.7, 0.2, 0.1], [0.15, 0.25, 0.6]])
90 umb_utility = UtilityTable([Weather, Umbrella], [[20, 100], [70, 0]], position=(1,0.4))
91

```

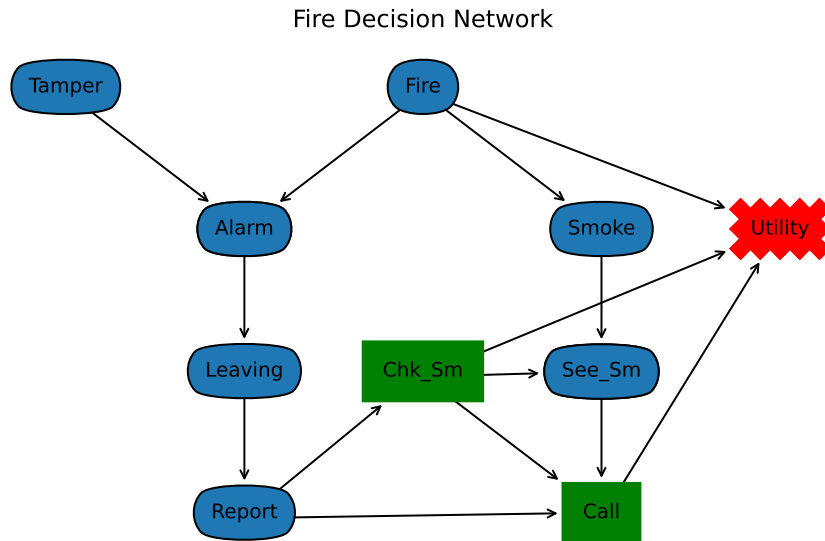


Figure 9.2: Fire Decision Network

```

92 | umbrella_dn = DecisionNetwork("Umbrella Decision Network",
93 |                               {Weather, Forecast, Umbrella},
94 |                               {p_weather, p_forecast, umb_utility})

```

The following is a variant with the umbrella decision having 2 parents; nothing else has changed. This is interesting because one of the parents is not needed; if the agent knows the weather, it can ignore the forecast.

```

_____decnNetworks.py — (continued) _____
96 | Umbrella2p = DecisionVariable("Umbrella", ["Take", "Leave"], {Forecast, Weather}, position=(0.5,0))
97 | umb_utility2p = UtilityTable([Weather, Umbrella2p], [[20, 100], [70, 0]], position=(1,0.4))
98 | umbrella_dn2p = DecisionNetwork("Umbrella Decision Network (extra arc)",
99 |                               {Weather, Forecast, Umbrella2p},
100 |                               {p_weather, p_forecast, umb_utility2p})

```

Fire Decision Network

The fire decision network of Figure 9.2 (showing the result of `fire_dn.show()`) is represented as:

```

_____decnNetworks.py — (continued) _____
102 | boolean = [False, True]
103 | Alarm = Variable("Alarm", boolean, position=(0.25,0.633))
104 | Fire = Variable("Fire", boolean, position=(0.5,0.9))

```

```

105 Leaving = Variable("Leaving", boolean, position=(0.25,0.366))
106 Report = Variable("Report", boolean, position=(0.25,0.1))
107 Smoke = Variable("Smoke", boolean, position=(0.75,0.633))
108 Tamper = Variable("Tamper", boolean, position=(0,0.9))
109
110 See_Sm = Variable("See_Sm", boolean, position=(0.75,0.366) )
111 Chk_Sm = DecisionVariable("Chk_Sm", boolean, {Report}, position=(0.5, 0.366))
112 Call = DecisionVariable("Call", boolean,{See_Sm,Chk_Sm,Report}, position=(0.75,0.1))
113
114 f_ta = Prob(Tamper,[],[0.98,0.02])
115 f-fi = Prob(Fire,[],[0.99,0.01])
116 f-sm = Prob(Smoke,[Fire],[[0.99,0.01],[0.1,0.9]])
117 f-al = Prob(Alarm,[Fire,Tamper],[[0.9999, 0.0001], [0.15, 0.85]], [[0.01, 0.99], [0.5, 0.5]])
118 f-lv = Prob(Leaving,[Alarm],[[0.999, 0.001], [0.12, 0.88]])
119 f-re = Prob(Report,[Leaving],[[0.99, 0.01], [0.25, 0.75]])
120 f-ss = Prob(See_Sm,[Chk_Sm,Smoke],[[1,0],[1,0]],[[1,0],[0,1]])
121
122 ut = UtilityTable([Chk_Sm,Fire,Call],[[0,-200],[-5000,-200]],[[-20,-220],[-5020,-220]]], position=
123
124 fire_dn = DecisionNetwork("Fire Decision Network",
125                           {Tamper,Fire,Alarm,Leaving,Smoke,Call,See_Sm,Chk_Sm,Report},
126                           {f_ta,f-fi,f-sm,f-al,f-lv,f-re,f-ss,ut})

```

Cheating Decision Network

The following is the representation of the cheating decision of Figure 9.3. Note that we keep the names of the variables short (less than 8 characters) so that the tables look good when printed.

```

-----decnNetworks.py --- (continued) -----
128 grades = ['A','B','C','F']
129 Watched = Variable("Watched", boolean, position=(0,0.9))
130 Caught1 = Variable("Caught1", boolean, position=(0.2,0.7))
131 Caught2 = Variable("Caught2", boolean, position=(0.6,0.7))
132 Punish = Variable("Punish", ["None","Suspension","Recorded"], position=(0.8,0.9))
133 Grade_1 = Variable("Grade_1", grades, position=(0.2,0.3))
134 Grade_2 = Variable("Grade_2", grades, position=(0.6,0.3))
135 Fin_Grd = Variable("Fin_Grd", grades, position=(0.8,0.1))
136 Cheat_1 = DecisionVariable("Cheat_1", boolean, set(), position=(0,0.5)) #no parents
137 Cheat_2 = DecisionVariable("Cheat_2", boolean, {Cheat_1,Caught1}, position=(0.4,0.5))
138
139 p-wa = Prob(Watched,[],[0.7, 0.3])
140 p-cc1 = Prob(Caught1,[Watched,Cheat_1],[[1.0, 0.0], [0.9, 0.1]], [[1.0, 0.0], [0.5, 0.5]])
141 p-cc2 = Prob(Caught2,[Watched,Cheat_2],[[1.0, 0.0], [0.9, 0.1]], [[1.0, 0.0], [0.5, 0.5]])
142 p-pun = Prob(Punish,[Caught1,Caught2],[[1.0, 0.0, 0.0], [0.5, 0.4, 0.1]], [[0.6, 0.2, 0.2], [0.2, 0.5, 0.3]])
143 p-gr1 = Prob(Grade_1,[Cheat_1], [{'A':0.2, 'B':0.3, 'C':0.3, 'D': 0.2}, {'A':0.5, 'B':0.3, 'C':0.2, 'D':0.0}])
144 p-gr2 = Prob(Grade_2,[Cheat_2], [{'A':0.2, 'B':0.3, 'C':0.3, 'D': 0.2}, {'A':0.5, 'B':0.3, 'C':0.2, 'D':0.0}])
145 p-fg = Prob(Fin_Grd,[Grade_1,Grade_2],
146             {'A':{'A':1.0, 'B':0.0, 'C': 0.0, 'D':0.0},
147              'B':{'A':0.5, 'B':0.5, 'C': 0.0, 'D':0.0},

```





Figure 9.4: A decision network that is a chain of 3 decisions

Chain of 3 decisions

The following example is a finite-stage fully-observable Markov decision process with a single reward (utility) at the end. It is interesting because the parents do not include all predecessors. The methods we use will work without change on this, even though the agent does not condition on all of its previous observations and actions. The output of `ch3.show()` is shown in Figure 9.4.

```

-----decnNetworks.py ----- (continued) -----
171 S0 = Variable('S0', boolean, position=(0,0.5))
172 D0 = DecisionVariable('D0', boolean, {S0}, position=(1/7,0.1))
173 S1 = Variable('S1', boolean, position=(2/7,0.5))
174 D1 = DecisionVariable('D1', boolean, {S1}, position=(3/7,0.1))
175 S2 = Variable('S2', boolean, position=(4/7,0.5))
176 D2 = DecisionVariable('D2', boolean, {S2}, position=(5/7,0.1))
177 S3 = Variable('S3', boolean, position=(6/7,0.5))
178
179 p_s0 = Prob(S0, [], [0.5,0.5])
180 tr = [[[0.1, 0.9], [0.9, 0.1]], [[0.2, 0.8], [0.8, 0.2]]] # 0 is flip, 1 is keep value
181 p_s1 = Prob(S1, [D0,S0], tr)
182 p_s2 = Prob(S2, [D1,S1], tr)
183 p_s3 = Prob(S3, [D2,S2], tr)
184
185 ch3U = UtilityTable([S3],[0,1], position=(7/7,0.9))

```

```

186 |
187 | ch3 = DecisionNetwork("3-chain", {S0,D0,S1,D1,S2,D2,S3},{p_s0,p_s1,p_s2,p_s3,ch3U})
188 | #rc3 = RC_DN(ch3)
189 | #rc3.optimize()
190 | #rc3.opt_policy

```

9.1.2 Recursive Conditioning for decision networks

An instance of a RC_DN object takes in a decision network. The query method uses recursive conditioning to compute the expected utility of the optimal policy. `self.opt_policy` becomes the optimal policy.

```

_____decnNetworks.py — (continued)_____
192 | import math
193 | from probGraphicalModels import GraphicalModel, InferenceMethod
194 | from probFactors import Factor
195 | from utilities import dict_union
196 | from probRC import connected_components
197 |
198 | class RC_DN(InferenceMethod):
199 |     """The class that queries graphical models using recursive conditioning
200 |
201 |     gm is graphical model to query
202 |     """
203 |
204 |     def __init__(self, gm=None):
205 |         self.gm = gm
206 |         self.cache = {(frozenset(), frozenset()):1}
207 |         ## self.max_display_level = 3
208 |
209 |     def optimize(self, split_order=None):
210 |         """computes expected utility, and creates optimal decision functions, where
211 |         elim_order is a list of the non-observed non-query variables in gm
212 |         """
213 |         if split_order == None:
214 |             split_order = self.gm.split_order()
215 |         self.opt_policy = {}
216 |         return self.rc({}, self.gm.factors, split_order)

```

The following is the simplest search-based algorithm. It is exponential in the number of variables, so is not very useful. However, it is simple, and useful to understand before looking at the more complicated algorithm. Note that the above code does not call `rc0`; you will need to change the `self.rc` to `self.rc0` in above code to use it.

```

_____decnNetworks.py — (continued)_____
218 |     def rc0(self, context, factors, split_order):
219 |         """simplest search algorithm"""
220 |         self.display(2, "calling rc0, ", (context, factors), "with S0", split_order)

```

```

221         if not factors:
222             return 1
223         elif to_eval := {fac for fac in factors if fac.can_evaluate(context)}:
224             self.display(3, "rc0 evaluating factors", to_eval)
225             val = math.prod(fac.get_value(context) for fac in to_eval)
226             return val * self.rc0(context, factors-to_eval, split_order)
227         else:
228             var = split_order[0]
229             self.display(3, "rc0 branching on", var)
230             if isinstance(var, DecisionVariable):
231                 assert set(context) <= set(var.parents), f"cannot optimize {var} in context {context}"
232                 maxres = -math.inf
233                 for val in var.domain:
234                     self.display(3, "In rc0, branching on", var, "=", val)
235                     newres = self.rc0(dict_union({var:val}, context), factors, split_order[1:])
236                     if newres > maxres:
237                         maxres = newres
238                         theval = val
239                 self.opt_policy[frozenset(context.items())] = (var, theval)
240                 return maxres
241             else:
242                 total = 0
243                 for val in var.domain:
244                     total += self.rc0(dict_union({var:val}, context), factors, split_order[1:])
245                 self.display(3, "rc0 branching on", var, "returning", total)
246                 return total

```

We can combine the optimization for decision networks above, with the improvements of recursive conditioning used for graphical models (Section 8.6, page 153).

decnNetworks.py — (continued)

```

248 def rc(self, context, factors, split_order):
249     """ returns the number  $\sum_{split\_order} \prod_{factors}$  given assignments in context
250     context is a variable:value dictionary
251     factors is a set of factors
252     split_order is a list of variables in factors that are not in context
253     """
254     self.display(3, "calling rc,", (context, factors))
255     ce = (frozenset(context.items()), frozenset(factors)) # key for the cache entry
256     if ce in self.cache:
257         self.display(2, "rc cache lookup", (context, factors))
258         return self.cache[ce]
259     # if not factors: # no factors; needed if you don't have forgetting and caching
260     #     return 1
261     elif vars_not_in_factors := {var for var in context
262                                if not any(var in fac.variables for fac in factors)}:
263         # forget variables not in any factor
264         self.display(3, "rc forgetting variables", vars_not_in_factors)
265         return self.rc({key:val for (key,val) in context.items()
266                        if key not in vars_not_in_factors},

```

```

267         factors, split_order)
268     elif to_eval := {fac for fac in factors if fac.can_evaluate(context)}:
269         # evaluate factors when all variables are assigned
270         self.display(3, "rc evaluating factors", to_eval)
271         val = math.prod(fac.get_value(context) for fac in to_eval)
272         if val == 0:
273             return 0
274         else:
275             return val * self.rc(context, {fac for fac in factors if fac not in to_eval}, split_order)
276     elif len(comp := connected_components(context, factors, split_order)) > 1:
277         # there are disconnected components
278         self.display(2, "splitting into connected components", comp)
279         return(math.prod(self.rc(context, f, eo) for (f, eo) in comp))
280     else:
281         assert split_order, f"split_order empty rc({context},{factors})"
282         var = split_order[0]
283         self.display(3, "rc branching on", var)
284         if isinstance(var, DecisionVariable):
285             assert set(context) <= set(var.parents), f"cannot optimize {var} in context {context}"
286             maxres = -math.inf
287             for val in var.domain:
288                 self.display(3, "In rc, branching on", var, "=", val)
289                 newres = self.rc(dict_union({var:val}, context), factors, split_order[1:])
290                 if newres > maxres:
291                     maxres = newres
292                     theval = val
293                 self.opt_policy[frozenset(context.items())] = (var, theval)
294                 self.cache[ce] = maxres
295             return maxres
296         else:
297             total = 0
298             for val in var.domain:
299                 total += self.rc(dict_union({var:val}, context), factors, split_order[1:])
300             self.display(3, "rc branching on", var, "returning", total)
301             self.cache[ce] = total
302             return total

```

Here is how to run the optimize the example decision networks:

```

decnNetworks.py — (continued)
304 # Umbrella decision network
305 #urc = RC_DN(umberella_dn)
306 #urc.optimize()
307 #urc.opt_policy
308
309 #rc_fire = RC_DN(fire_dn)
310 #rc_fire.optimize()
311 #rc_fire.opt_policy
312
313 #rc_cheat = RC_DN(cheat_dn)
314 #rc_cheat.optimize()

```



```

315 #rc_cheat.opt_policy
316
317 #rc_ch3 = RC_DN(ch3)
318 #rc_ch3.optimize()
319 #rc_ch3.opt_policy

```

9.1.3 Variable elimination for decision networks

VE_DN is variable elimination for decision networks. The method *optimize* is used to optimize all the decisions. Note that *optimize* requires a legal elimination ordering of the random and decision variables, otherwise it will give an exception. (A decision node can only be maximized if the variables that are not its parents have already been eliminated.)

```

decnNetworks.py — (continued)
321 from probVE import VE
322
323 class VE_DN(VE):
324     """Variable Elimination for Decision Networks"""
325     def __init__(self,dn=None):
326         """dn is a decision network"""
327         VE.__init__(self,dn)
328         self.dn = dn
329
330     def optimize(self,elim_order=None,obs={}):
331         if elim_order == None:
332             elim_order = reversed(self.gm.split_order())
333         policy = []
334         proj_factors = [self.project_observations(fact,obs)
335                         for fact in self.dn.factors]
336         for v in elim_order:
337             if isinstance(v,DecisionVariable):
338                 to_max = [fac for fac in proj_factors
339                           if v in fac.variables and set(fac.variables) <= v.all_vars]
340                 assert len(to_max)==1, "illegal variable order "+str(elim_order)+" at "+str(v)
341                 newFac = FactorMax(v, to_max[0])
342                 policy.append(newFac.decision_fun)
343                 proj_factors = [fac for fac in proj_factors if fac is not to_max[0]]+[newFac]
344                 self.display(2,"maximizing",v,"resulting factor",newFac.brief() )
345                 self.display(3,newFac)
346             else:
347                 proj_factors = self.eliminate_var(proj_factors, v)
348             assert len(proj_factors)==1,"Should there be only one element of proj_factors?"
349             value = proj_factors[0].get_value({})
350             return value,policy

decnNetworks.py — (continued)
352 class FactorMax(Factor):
353     """A factor obtained by maximizing a variable in a factor.

```

```

354     Also builds a decision_function. This is based on FactorSum.
355     """
356
357     def __init__(self, dvar, factor):
358         """dvar is a decision variable.
359         factor is a factor that contains dvar and only parents of dvar
360         """
361         self.dvar = dvar
362         self.factor = factor
363         vars = [v for v in factor.variables if v is not dvar]
364         Factor.__init__(self, vars)
365         self.values = [None]*self.size
366         self.decision_fun = FactorDF(dvar, vars, [None]*self.size)
367
368     def get_value(self, assignment):
369         """lazy implementation: if saved, return saved value, else compute it"""
370         index = self.assignment_to_index(assignment)
371         if self.values[index]:
372             return self.values[index]
373         else:
374             max_val = float("-inf") # -infinity
375             new_asst = assignment.copy()
376             for elt in self.dvar.domain:
377                 new_asst[self.dvar] = elt
378                 fac_val = self.factor.get_value(new_asst)
379                 if fac_val > max_val:
380                     max_val = fac_val
381                     best_elt = elt
382             self.values[index] = max_val
383             self.decision_fun.values[index] = best_elt
384             return max_val

```

A decision function is a stored factor.

```

decnNetworks.py — (continued)
386 class FactorDF(TabFactor):
387     """A decision function"""
388     def __init__(self, dvar, vars, values):
389         TabStored.__init__(self, vars, values)
390         self.dvar = dvar
391         self.name = str(dvar) # Used in printing

```

Here are some example queries:

```

decnNetworks.py — (continued)
393 # Example queries:
394 # v,p = VE_DN(fire_dn).optimize(); print(v)
395 # for df in p: print(df, "\n")
396
397 # VE_DN.max_display_level = 3 # if you want to show lots of detail
398 # v,p = VE_DN(cheat_dn).optimize(); print(v)
399 # for df in p: print(df, "\n") # print decision functions

```

9.2 Markov Decision Processes

We will represent a **Markov decision process (MDP)** directly, rather than using the recursive conditioning or variable elimination code, as we did for decision networks.

```

mdpProblem.py — Representations for Markov Decision Processes
11 from utilities import argmaxd
12 import random
13 import matplotlib.pyplot as plt
14 from matplotlib.widgets import Button, CheckButtons
15
16 class MDP(object):
17     """A Markov Decision Process. Must define:
18     self.states the set (or list) of states
19     self.actions the set (or list) of actions
20     self.discount a real-valued discount
21     """
22
23     def __init__(self, states, actions, discount, init=0):
24         self.states = states
25         self.actions = actions
26         self.discount = discount
27         self.initv = self.v = {s: init for s in self.states}
28         self.initq = self.q = {s: {a: init for a in self.actions} for s in self.states}
29
30     def P(self, s, a):
31         """Transition probability function
32         returns a dictionary of {s1:p1} such that P(s1 | s,a)=p1. Other probabilities are zero.
33         """
34         raise NotImplementedError("P") # abstract method
35
36     def R(self, s, a):
37         """Reward function R(s,a)
38         returns the expected reward for doing a in state s.
39         """
40         raise NotImplementedError("R") # abstract method

```

Two state partying example (Example 9.27 in Poole and Mackworth [2017]):

```

mdpExamples.py — MDP Examples
11 from mdpProblem import MDP, GridMDP
12
13 class party(MDP):
14     """Simple 2-state, 2-Action Partying MDP Example"""
15     def __init__(self, discount=0.9):
16         states = {'healthy', 'sick'}
17         actions = {'relax', 'party'}
18         MDP.__init__(self, states, actions, discount)
19
20     def R(self, s, a):

```

```

21     "R(s,a)"
22     return { 'healthy': {'relax': 7, 'party': 10},
23             'sick': {'relax': 0, 'party': 2 }}[s][a]
24
25     def P(self,s,a):
26         "returns a dictionary of {s1:p1} such that P(s1 | s,a)=p1. Other probabilities are zero."
27         phealthy = { # P('healthy' | s, a)
28                     'healthy': {'relax': 0.95, 'party': 0.7},
29                     'sick': {'relax': 0.5, 'party': 0.1 }}[s][a]
30         return {'healthy':phealthy, 'sick':1-phealthy}

```

The next example is the tiny game from Example 12.1 and Figure 12.1 of Poole and Mackworth [2017]. The state is represented as (x,y) where x counts from zero from the left, and y counts from zero upwards, so the state $(0,0)$ is on the bottom-left state. The actions are *upC* for up-careful, and *upR* for up-risky. (Note that GridMDP is just a type of MDP for which we have methods to show; you can assume it is just MDP here).

```

mdpExamples.py — (continued)
33 class tiny(GridMDP):
34     def __init__(self, discount=0.9):
35         actions = ['right', 'upC', 'left', 'upR']
36         self.x_dim = 2 # x-dimension
37         self.y_dim = 3
38         states = [(x,y) for x in range(self.x_dim) for y in range(self.y_dim)]
39         # for GridMDP
40         self.xoff = {'right':0.25, 'upC':0, 'left':-0.25, 'upR':0}
41         self.yoff = {'right':0, 'upC':-0.25, 'left':0, 'upR':0.25}
42         GridMDP.__init__(self, states, actions, discount)
43
44     def P(self,s,a):
45         """return a dictionary of {s1:p1} if P(s1 | s,a)=p1. Other probabilities are zero.
46         """
47         (x,y) = s
48         if a == 'right':
49             return {(1,y):1}
50         elif a == 'upC':
51             return {(x,min(y+1,2)):1}
52         elif a == 'left':
53             if (x,y) == (0,2): return {(0,0):1}
54             else: return {(0,y): 1}
55         elif a == 'upR':
56             if x==0:
57                 if y<2: return {(x,y):0.1, (x+1,y):0.1, (x,y+1):0.8}
58                 else: # at (0,2)
59                     return {(0,0):0.1, (1,2): 0.1, (0,2): 0.8}
60             elif y < 2: # x==1
61                 return {(0,y):0.1, (1,y):0.1, (1,y+1):0.8}
62             else: # at (1,2)
63                 return {(0,2):0.1, (1,2): 0.9}
64

```

```

65     def R(self,s,a):
66         (x,y) = s
67         if a == 'right':
68             return [0,-1][x]
69         elif a == 'upC':
70             return [-1,-1,-2][y]
71         elif a == 'left':
72             if x==0:
73                 return [-1, -100, 10][y]
74             else: return 0
75         elif a == 'upR':
76             return [[-0.1, -10, 0.2],[-0.1, -0.1, -0.9]][x][y]
77             # at (0,2) reward is 0.1*10+0.8*-1=0.2

```

Here is the domain of Example 9.28 of Poole and Mackworth [2017]. Here the state is represented as (x,y) where x counts from zero from the left, and y counts from zero upwards, so the state $(0,0)$ is on the bottom-left state.

mdpExamples.py — (continued)

```

79 class grid(GridMDP):
80     """ x_dim * y_dim grid with rewarding states"""
81     def __init__(self, discount= 0.9, x_dim=10, y_dim=10):
82         self.x_dim = x_dim # size in x-direction
83         self.y_dim = y_dim # size in y-direction
84         actions = ['up', 'down', 'right', 'left']
85         states = [(x,y) for x in range(y_dim) for y in range(y_dim)]
86         self.rewarding_states = {(3,2):-10, (3,5):-5, (8,2):10, (7,7):3}
87         self.fling_states = {(8,2), (7,7)}
88         self.xoff = {'right':0.25, 'up':0, 'left':-0.25, 'down':0}
89         self.yoff = {'right':0, 'up':0.25, 'left':0, 'down':-0.25}
90         GridMDP.__init__(self, states, actions, discount)
91
92     def intended_next(self,s,a):
93         """returns the next state in the direction a.
94         This is where the agent will end up if to goes in its intended_direction
95         (which it does with probability 0.7).
96         """
97         (x,y) = s
98         if a=='up':
99             return (x, y+1 if y+1 < self.y_dim else y)
100         if a=='down':
101             return (x, y-1 if y > 0 else y)
102         if a=='right':
103             return (x+1 if x+1 < self.x_dim else x,y)
104         if a=='left':
105             return (x-1 if x > 0 else x,y)
106
107     def P(self,s,a):
108         """return a dictionary of {s1:p1} if P(s1 | s,a)=p1. Other probabilities are zero.
109         Corners are tricky because different actions result in same state.
110         """

```

```

111     if s in self.fling_states:
112         return {(0,0): 0.25, (self.x_dim-1,0):0.25, (0,self.y_dim-1):0.25, (self.x_dim-1,self.y_dim-1):0.25}
113     res = dict()
114     for ai in self.actions:
115         s1 = self.intended_next(s,ai)
116         ps1 = 0.7 if ai==a else 0.1
117         if s1 in res: # occurs in corners
118             res[s1] += ps1
119         else:
120             res[s1] = ps1
121     return res
122
123 def R(self,s,a):
124     if s in self.rewarding_states:
125         return self.rewarding_states[s]
126     else:
127         (x,y) = s
128         rew = 0
129         # rewards from crashing:
130         if y==0: ## on bottom.
131             rew += -0.7 if a == 'down' else -0.1
132         if y==self.y_dim-1: ## on top.
133             rew += -0.7 if a == 'up' else -0.1
134         if x==0: ## on left
135             rew += -0.7 if a == 'left' else -0.1
136         if x==self.x_dim-1: ## on right.
137             rew += -0.7 if a == 'right' else -0.1
138     return rew

```

9.2.1 Value Iteration

This implements value iteration.

This uses indexes of the states and actions (not the names). The value function is represented so $v[s]$ is the value of state with index s . A Q function is represented so $q[s][a]$ is the value for doing action with index a state with index s . Similarly a policy π is represented as a list where $pi[s]$, where s is the index of a state, returns the index of the action.

```

_____mdpProblem.py — (continued) _____
42 def vi(self, n):
43     """carries out n iterations of value iteration, updating value function self.v
44     Returns a Q-function, value function, policy
45     """
46     print("calling vi")
47     assert n>0,"You must carry out at least one iteration of vi. n="+str(n)
48     #v = v0 if v0 is not None else {s:0 for s in self.states}
49     for i in range(n):
50         self.q = {s: {a: self.R(s,a)+self.discount*sum(p1*self.v[s1]
51                                     for (s1,p1) in self.P(s,a).items())

```

```

52         for a in self.actions}
53         for s in self.states}
54         self.v = {s: max(self.q[s][a] for a in self.actions)
55                 for s in self.states}
56         self.pi = {s: argmaxd(self.q[s])
57                  for s in self.states}
58         return self.q, self.v, self.pi

```

The following shows how this can be used.

```

_____mdpExamples.py — (continued)_____
140 ## Testing value iteration
141 # Try the following:
142 # pt = party(discount=0.9)
143 # pt.vi(1)
144 # pt.vi(100)
145 # party(discount=0.99).vi(100)
146 # party(discount=0.4).vi(100)
147
148 # gr = grid()
149 # gr.show()
150 # q,v,pi = gr.vi(100)
151 # q[(7,2)]

```

9.2.2 Showing Grid MDPs

A GridMDP is a type of MDP where the states are (x,y) positions. It is a special sort of MDP only because we have methods to show it.

```

_____mdpProblem.py — (continued)_____
60 class GridMDP(MDP):
61     def __init__(self, states, actions, discount):
62         MDP.__init__(self, states, actions, discount)
63
64     def show(self):
65         plt.ion() # interactive
66         fig,(self.ax) = plt.subplots()
67         plt.subplots_adjust(bottom=0.2)
68         stepB = Button(plt.axes([0.8,0.05,0.1,0.075]), "step")
69         stepB.on_clicked(self.on_step)
70         resetB = Button(plt.axes([0.6,0.05,0.1,0.075]), "reset")
71         resetB.on_clicked(self.on_reset)
72         self.qcheck = CheckButtons(plt.axes([0.2,0.05,0.35,0.075]),
73                                   ["show q-values","show policy"])
74         self.qcheck.on_clicked(self.show_vals)
75         self.show_vals(None)
76         plt.show()
77
78     def show_vals(self,event):
79         self.ax.cla()

```

```

80     array = [[self.v[(x,y)] for x in range(self.x_dim)]
81               for y in range(self.y_dim)]
82     self.ax.pcolormesh([x-0.5 for x in range(self.x_dim+1)],
83                       [x-0.5 for x in range(self.y_dim+1)],
84                       array, edgecolors='black', cmap='summer')
85     # for cmap see https://matplotlib.org/stable/tutorials/colors/colormaps.html
86     if self.qcheck.get_status()[1]: # "show policy"
87         for (x,y) in self.q:
88             maxv = max(self.q[(x,y)][a] for a in self.actions)
89             for a in self.actions:
90                 if self.q[(x,y)][a] == maxv:
91                     # draw arrow in appropriate direction
92                     self.ax.arrow(x,y,self.xoff[a]*2,self.yoff[a]*2,
93                                  color='red',width=0.05, head_width=0.2, length_includes_head=True)
94     if self.qcheck.get_status()[0]: # "show q-values"
95         self.show_q(event)
96     else:
97         self.show_v(event)
98     self.ax.set_xticks(range(self.x_dim))
99     self.ax.set_xticklabels(range(self.x_dim))
100    self.ax.set_yticks(range(self.y_dim))
101    self.ax.set_yticklabels(range(self.y_dim))
102    plt.draw()
103
104    def on_step(self,event):
105        self.vi(1)
106        self.show_vals(event)
107
108    def show_v(self,event):
109        """show values"""
110        for (x,y) in self.v:
111            self.ax.text(x,y,"{val: .2f}".format(val=self.v[(x,y)]),ha='center')
112
113    def show_q(self,event):
114        """show q-values"""
115        for (x,y) in self.q:
116            for a in self.actions:
117                self.ax.text(x+self.xoff[a],y+self.yoff[a],
118                             "{val: .2f}".format(val=self.q[(x,y)][a]),ha='center')
119
120    def on_reset(self,event):
121        self.v = self.initv
122        self.q = self.initq
123        self.show_vals(event)

```

Figure 9.6 shows the user interface, which can be obtained using `tiny().show()`, resizing it, checking “show q-values” and “show policy”, and clicking “step” a few times.

Figure ?? shows the user interface, which can be obtained using `grid().show()`, resizing it, checking “show q-values” and “show policy”, and clicking “step” a




Figure 9.5: Interface for tiny example, after a number of steps. Each rectangle represents a state. In each rectangle are the 4 Q-values for the state. The leftmost number is the for the left action; the rightmost number is for the right action; the upper most is for the *upR* (up-risky) action and the lowest number is for the *upC* action. The arrow points to the action(s) with the maximum Q-value.

few times.

Exercise 9.1 Computing q before v may seem like a waste of space because we don't need to store q in order to compute value function or the policy. Change the algorithm so that it loops through the states and actions once per iteration, and only stores the value function and the policy. Note that to get the same results as before, you would need to make sure that you use the previous value of v in the computation not the current value of v . Does using the current value of v hurt the algorithm or make it better (in approaching the actual value function)?

9.2.3 Asynchronous Value Iteration

This implements asynchronous value iteration, storing Q .




Figure 9.6: Interface for grid example, after a number of steps. Each rectangle represents a state. In each rectangle are the 4 Q-values for the state. The leftmost number is the for the left action; the rightmost number is for the right action; the upper most is for the up action and the lowest number is for the down action. The arrow points to the action(s) with the maximum Q-value.

A Q function is represented so $q[s][a]$ is the value for doing action with index a state with index s .

```

_____mdpProblem.py — (continued) _____
125     def avi(self,n):
126         states = list(self.states)
127         actions = list(self.actions)
128         for i in range(n):
129             s = random.choice(states)
130             a = random.choice(actions)
131             self.q[s][a] = (self.R(s,a) + self.discount *
132                             sum(p1 * max(self.q[s1][a1]
133                                           for a1 in self.actions)
134                                for (s1,p1) in self.P(s,a).items()))
135         return Q

```

The following shows how avi can be used.

```

_____mdpExamples.py — (continued) _____
154 ## Testing asynchronous value iteration
155 # Try the following:
156 # pt = party(discount=0.9)
157 # pt.avi(10)
158 # pt.vi(1000)
159
160 # gr = grid()
161 # q = gr.avi(100000)
162 # q[(7,2)]

```

Exercise 9.2 Implement value iteration that stores the V -values rather than the Q -values. Does it work better than storing Q ? (What might better mean?)

Exercise 9.3 In asynchronous value iteration, try a number of different ways to choose the states and actions to update (e.g., sweeping through the state-action pairs, choosing them at random). Note that the best way may be to determine which states have had their Q -values change the most, and then update the previous ones, but that is not so straightforward to implement, because you need to find those previous states.

Learning with Uncertainty

10.1 K-means

The k-means learner maintains two lists that suffice as sufficient statistics to classify examples, and to learn the classification:

- *class_counts* is a list such that *class_counts*[*c*] is the number of examples in the training set with *class* = *c*.
- *feature_sum* is a list such that *feature_sum*[*i*][*c*] is sum of the values for the *i*'th feature *i* for members of class *c*. The average value of the *i*th feature in class *i* is

$$\frac{\text{feature_sum}[i][c]}{\text{class_counts}[c]}$$

The class is initialized by randomly assigning examples to classes, and updating the statistics for *class_counts* and *feature_sum*.

```
learnKMeans.py — k-means learning
11 from learnProblem import Data_set, Learner, Data_from_file
12 import random
13 import matplotlib.pyplot as plt
14
15 class K_means_learner(Learner):
16     def __init__(self, dataset, num_classes):
17         self.dataset = dataset
18         self.num_classes = num_classes
19         self.random_initialize()
20
21     def random_initialize(self):
```

```

22     # class_counts[c] is the number of examples with class=c
23     self.class_counts = [0]*self.num_classes
24     # feature_sum[i][c] is the sum of the values of feature i for class c
25     self.feature_sum = [[0]*self.num_classes
26                         for feat in self.dataset.input_features]
27     for eg in self.dataset.train:
28         cl = random.randrange(self.num_classes) # assign eg to random class
29         self.class_counts[cl] += 1
30         for (ind,feat) in enumerate(self.dataset.input_features):
31             self.feature_sum[ind][cl] += feat(eg)
32     self.num_iterations = 0
33     self.display(1,"Initial class counts: ",self.class_counts)

```

The distance from (the mean of) a class to an example is the sum, over all features, of the sum-of-squares differences of the class mean and the example value.

```

learnKMeans.py — (continued) —
35     def distance(self,cl,eg):
36         """distance of the eg from the mean of the class"""
37         return sum( (self.class_prediction(ind,cl)-feat(eg))**2
38                   for (ind,feat) in enumerate(self.dataset.input_features))
39
40     def class_prediction(self,feat_ind,cl):
41         """prediction of the class cl on the feature with index feat_ind"""
42         if self.class_counts[cl] == 0:
43             return 0 # there are no examples so we can choose any value
44         else:
45             return self.feature_sum[feat_ind][cl]/self.class_counts[cl]
46
47     def class_of_eg(self,eg):
48         """class to which eg is assigned"""
49         return (min((self.distance(cl,eg),cl)
50                   for cl in range(self.num_classes)))[1]
51         # second element of tuple, which is a class with minimum distance

```

One step of k-means updates the *class_counts* and *feature_sum*. It uses the old values to determine the classes, and so the new values for *class_counts* and *feature_sum*. At the end it determines whether the values of these have changes, and then replaces the old ones with the new ones. It returns an indicator of whether the values are stable (have not changed).

```

learnKMeans.py — (continued) —
53     def k_means_step(self):
54         """Updates the model with one step of k-means.
55         Returns whether the assignment is stable.
56         """
57         new_class_counts = [0]*self.num_classes
58         # feature_sum[i][c] is the sum of the values of feature i for class c
59         new_feature_sum = [[0]*self.num_classes
60                           for feat in self.dataset.input_features]

```

```

61         for eg in self.dataset.train:
62             cl = self.class_of_eg(eg)
63             new_class_counts[cl] += 1
64             for (ind,feat) in enumerate(self.dataset.input_features):
65                 new_feature_sum[ind][cl] += feat(eg)
66         stable = (new_class_counts == self.class_counts) and (self.feature_sum == new_feature_sum)
67         self.class_counts = new_class_counts
68         self.feature_sum = new_feature_sum
69         self.num_iterations += 1
70         return stable
71
72
73     def learn(self,n=100):
74         """do n steps of k-means, or until convergence"""
75         i=0
76         stable = False
77         while i<n and not stable:
78             stable = self.k_means_step()
79             i += 1
80             self.display(1,"Iteration",self.num_iterations,
81                         "class counts: ",self.class_counts," Stable=",stable)
82         return stable
83
84     def show_classes(self):
85         """sorts the data by the class and prints in order.
86         For visualizing small data sets
87         """
88         class_examples = [[] for i in range(self.num_classes)]
89         for eg in self.dataset.train:
90             class_examples[self.class_of_eg(eg)].append(eg)
91         print("Class","Example",sep='\t')
92         for cl in range(self.num_classes):
93             for eg in class_examples[cl]:
94                 print(cl,*eg,sep='\t')
95
96     def plot_error(self, maxstep=20):
97         """Plots the sum-of-squares error as a function of the number of steps"""
98         plt.ion()
99         plt.xlabel("step")
100        plt.ylabel("Ave sum-of-squares error")
101        train_errors = []
102        if self.dataset.test:
103            test_errors = []
104        for i in range(maxstep):
105            self.learn(1)
106            train_errors.append( sum(self.distance(self.class_of_eg(eg),eg)
107                                   for eg in self.dataset.train)
108                               /len(self.dataset.train))
109            if self.dataset.test:
110                test_errors.append( sum(self.distance(self.class_of_eg(eg),eg)

```

```

111         for eg in self.dataset.test)
112             /len(self.dataset.test))
113     plt.plot(range(1,maxstep+1),train_errors,
114             label=str(self.num_classes)+" classes. Training set")
115     if self.dataset.test:
116         plt.plot(range(1,maxstep+1),test_errors,
117                 label=str(self.num_classes)+" classes. Test set")
118     plt.legend()
119     plt.draw()
120
121 %data = Data_from_file('data/emdata1.csv', num_train=10, target_index=2000) % trivial example
122 data = Data_from_file('data/emdata2.csv', num_train=10, target_index=2000)
123 %data = Data_from_file('data/emdata0.csv', num_train=14, target_index=2000) % example from textbook
124 kml = K_means_learner(data,2)
125 num_iter=4
126 print("Class assignment after",num_iter,"iterations:")
127 kml.learn(num_iter); kml.show_classes()
128
129 # Plot the error
130 # km2=K_means_learner(data,2); km2.plot_error(20) # 2 classes
131 # km3=K_means_learner(data,3); km3.plot_error(20) # 3 classes
132 # km13=K_means_learner(data,13); km13.plot_error(20) # 13 classes
133
134 # data = Data_from_file('data/carbool.csv', target_index=2000,boolean_features=True)
135 # kml = K_means_learner(data,3)
136 # kml.learn(20); kml.show_classes()
137 # km3=K_means_learner(data,3); km3.plot_error(20) # 3 classes
138 # km3=K_means_learner(data,30); km3.plot_error(20) # 30 classes

```

Exercise 10.1 Change *boolean_features = True* flag to allow for numerical features. K-means assumes the features are numerical, so we want to make non-numerical features into numerical features (using characteristic functions) but we probably don't want to change numerical features into Boolean.

Exercise 10.2 If there are many classes, some of the classes can become empty (e.g., try 100 classes with carbool.csv). Implement a way to put some examples into a class, if possible. Two ideas are:

- (a) Initialize the classes with actual examples, so that the classes will not start empty. (Do the classes become empty?)
- (b) In *class prediction*, we test whether the code is empty, and make a prediction of 0 for an empty class. It is possible to make a different prediction to "steal" an example (but you should make sure that a class has a consistent value for each feature in a loop).

Make your own suggestions, and compare it with the original, and whichever of these you think may work better.

10.2 EM

In the following definition, a class, c , is a integer in range $[0, \text{num_classes})$. i is an index of a feature, so $\text{feat}[i]$ is the i th feature, and a feature is a function from tuples to values. val is a value of a feature.

A model consists of 2 lists, which form the sufficient statistics:

- *class_counts* is a list such that $\text{class_counts}[c]$ is the number of tuples with $\text{class} = c$, where each tuple is weighted by its probability, i.e.,

$$\text{class_counts}[c] = \sum_{t:\text{class}(t)=c} P(t)$$

- *feature_counts* is a list such that $\text{feature_counts}[i][\text{val}][c]$ is the weighted count of the number of tuples t with $\text{feat}[i](t) = \text{val}$ and $\text{class}(t) = c$, each tuple is weighted by its probability, i.e.,

$$\text{feature_counts}[i][\text{val}][c] = \sum_{t:\text{feat}[i](t)=\text{val} \text{ and } \text{class}(t)=c} P(t)$$

```

learnEM.py — EM Learning
11 from learnProblem import Data_set, Learner, Data_from_file
12 import random
13 import math
14 import matplotlib.pyplot as plt
15
16 class EM_learner(Learner):
17     def __init__(self, dataset, num_classes):
18         self.dataset = dataset
19         self.num_classes = num_classes
20         self.class_counts = None
21         self.feature_counts = None

```

The function *em_step* goes through the training examples, and updates these counts. The first time it is run, when there is no model, it uses random distributions.

```

learnEM.py — (continued)
23 def em_step(self, orig_class_counts, orig_feature_counts):
24     """updates the model."""
25     class_counts = [0]*self.num_classes
26     feature_counts = [{val:[0]*self.num_classes
27                         for val in feat.frange}
28                       for feat in self.dataset.input_features]
29     for tple in self.dataset.train:
30         if orig_class_counts: # a model exists
31             tpl_class_dist = self.prob(tple, orig_class_counts, orig_feature_counts)
32         else:                  # initially, with no model, return a random distribution

```

```

33         tpl_class_dist = random_dist(self.num_classes)
34         for cl in range(self.num_classes):
35             class_counts[cl] += tpl_class_dist[cl]
36             for (ind, feat) in enumerate(self.dataset.input_features):
37                 feature_counts[ind][feat(tple)][cl] += tpl_class_dist[cl]
38         return class_counts, feature_counts

```

prob computes the probability of a class c for a tuple $tple$, given the current statistics.

$$\begin{aligned}
 P(c \mid tple) &\propto P(c) * \prod_i P(X_i = tple(i) \mid c) \\
 &= \frac{\text{class_counts}[c]}{\text{len}(\text{self.dataset})} * \prod_i \frac{\text{feature_counts}[i][\text{feat}_i(tple)][c]}{\text{class_counts}[c]} \\
 &\propto \frac{\prod_i \text{feature_counts}[i][\text{feat}_i(tple)][c]}{\text{class_counts}[c]^{|\text{feats}|-1}}
 \end{aligned}$$

The last step is because $\text{len}(\text{self.dataset})$ is a constant (independent of c). $\text{class_counts}[c]$ can be taken out of the product, but needs to be raised to the power of the number of features, and one of them cancels.

```

learnEM.py — (continued)
40 def prob(self, tple, class_counts, feature_counts):
41     """returns a distribution over the classes for tuple tple in the model defined by the counts
42     """
43     feats = self.dataset.input_features
44     unnorm = [prod(feature_counts[i][feat(tple)][c]
45                   for (i, feat) in enumerate(feats))
46              /(class_counts[c]**(len(feats)-1))
47              for c in range(self.num_classes)]
48     thesum = sum(unnorm)
49     return [un/thesum for un in unnorm]

```

learn does n steps of EM:

```

learnEM.py — (continued)
51 def learn(self, n):
52     """do n steps of em"""
53     for i in range(n):
54         self.class_counts, self.feature_counts = self.em_step(self.class_counts,
55                                                                self.feature_counts)

```

The following is for visualizing the classes. It prints the dataset ordered by the probability of class c .

```

learnEM.py — (continued)
57 def show_class(self, c):
58     """sorts the data by the class and prints in order.
59     For visualizing small data sets
60     """
61     sorted_data = sorted((self.prob(tple, self.class_counts, self.feature_counts)[c],

```

```

62         ind, # preserve ordering for equal probabilities
63         tpl)
64         for (ind,tpl) in enumerate(self.dataset.train))
65     for cc,r,tpl in sorted_data:
66         print(cc,*tpl,sep='\t')

```

The following are for evaluating the classes.

The probability of a tuple can be evaluated by marginalizing over the classes:

$$\begin{aligned}
 P(tple) &= \sum_c P(c) * \prod_i P(X_i = tple(i) \mid c) \\
 &= \sum_c \frac{cc[c]}{\text{len}(\text{self.dataset})} * \prod_i \frac{fc[i][\text{feat}_i(tple)][c]}{cc[c]}
 \end{aligned}$$

where cc is the class count and fc is feature count. $\text{len}(\text{self.dataset})$ can be distributed out of the sum, and $cc[c]$ can be taken out of the product:

$$= \frac{1}{\text{len}(\text{self.dataset})} \sum_c \frac{1}{cc[c]^{\#feats-1}} * \prod_i fc[i][\text{feat}_i(tple)][c]$$

Given the probability of each tuple, we can evaluate the logloss, as the negative of the log probability:

```

learnEM.py — (continued)
68 def logloss(self, tple):
69     """returns the logloss of the prediction on tple, which is -log(P(tple))
70     based on the current class counts and feature counts
71     """
72     feats = self.dataset.input_features
73     res = 0
74     cc = self.class_counts
75     fc = self.feature_counts
76     for c in range(self.num_classes):
77         res += prod(fc[i][feat(tple)][c]
78                     for (i, feat) in enumerate(feats)) / (cc[c]**(len(feats)-1))
79     if res > 0:
80         return -math.log2(res/len(self.dataset.train))
81     else:
82         return float("inf") #infinity
83
84 def plot_error(self, maxstep=20):
85     """Plots the logloss error as a function of the number of steps"""
86     plt.ion()
87     plt.xlabel("step")
88     plt.ylabel("Ave Logloss (bits)")
89     train_errors = []
90     if self.dataset.test:
91         test_errors = []
92     for i in range(maxstep):
93         self.learn(1)

```

```

94         train_errors.append( sum(self.logloss(tple) for tple in self.dataset.train)
95                               /len(self.dataset.train))
96         if self.dataset.test:
97             test_errors.append( sum(self.logloss(tple) for tple in self.dataset.test)
98                                 /len(self.dataset.test))
99         plt.plot(range(1,maxstep+1),train_errors,
100                  label=str(self.num_classes)+" classes. Training set")
101         if self.dataset.test:
102             plt.plot(range(1,maxstep+1),test_errors,
103                      label=str(self.num_classes)+" classes. Test set")
104         plt.legend()
105         plt.draw()
106
107     def prod(L):
108         """returns the product of the elements of L"""
109         res = 1
110         for e in L:
111             res *= e
112         return res
113
114     def random_dist(k):
115         """generate k random numbers that sum to 1"""
116         res = [random.random() for i in range(k)]
117         s = sum(res)
118         return [v/s for v in res]
119
120     data = Data_from_file('data/emdata2.csv', num_train=10, target_index=2000)
121     eml = EM_learner(data,2)
122     num_iter=2
123     print("Class assignment after",num_iter,"iterations:")
124     eml.learn(num_iter); eml.show_class(0)
125
126     # Plot the error
127     # em2=EM_learner(data,2); em2.plot_error(40) # 2 classes
128     # em3=EM_learner(data,3); em3.plot_error(40) # 3 classes
129     # em13=EM_learner(data,13); em13.plot_error(40) # 13 classes
130
131     # data = Data_from_file('data/carbool.csv', target_index=2000,boolean_features=False)
132     # [f.frange for f in data.input_features]
133     # eml = EM_learner(data,3)
134     # eml.learn(20); eml.show_class(0)
135     # em3=EM_learner(data,3); em3.plot_error(60) # 3 classes
136     # em30=EM_learner(data,30); em30.plot_error(60) # 30 classes

```

Exercise 10.3 For the EM data, where there are naturally 2 classes, 3 classes does better on the training set after a while than 2 classes, but worse on the test set. Explain why. Hint: look what the 3 classes are. Use "em3.show_class(i)" for each of the classes $i \in [0,3)$.

Exercise 10.4 Write code to plot the logloss as a function of the number of classes (from 1 to say 15) for a fixed number of iterations. (From the experience with the

existing code, think about how many iterations is appropriate.)

Multiagent Systems

11.1 Minimax

Here we consider two-player zero-sum games. Here a player only wins when another player loses. This can be modeled as where there is a single utility which one agent (the maximizing agent) is trying minimize and the other agent (the minimizing agent) is trying to minimize.

11.1.1 Creating a two-player game

```
masProblem.py — A Multiagent Problem
11 from display import Displayable
12
13 class Node(Displayable):
14     """A node in a search tree. It has a
15     name a string
16     isMax is True if it is a maximizing node, otherwise it is minimizing node
17     children is the list of children
18     value is what it evaluates to if it is a leaf.
19     """
20     def __init__(self, name, isMax, value, children):
21         self.name = name
22         self.isMax = isMax
23         self.value = value
24         self.allchildren = children
25
26     def isLeaf(self):
27         """returns true if this is a leaf node"""
28         return self.allchildren is None
29
```

```

30     def children(self):
31         """returns the list of all children."""
32         return self.allchildren
33
34     def evaluate(self):
35         """returns the evaluation for this node if it is a leaf"""
36         return self.value

```

The following gives the tree from Figure 11.5 of the book. Note how 888 is used as a value here, but never appears in the trace.

```

_____masProblem.py — (continued) _____
38 fig10_5 = Node("a", True, None, [
39     Node("b", False, None, [
40         Node("d", True, None, [
41             Node("h", False, None, [
42                 Node("h1", True, 7, None),
43                 Node("h2", True, 9, None)]],
44             Node("i", False, None, [
45                 Node("i1", True, 6, None),
46                 Node("i2", True, 888, None)]])],
47         Node("e", True, None, [
48             Node("j", False, None, [
49                 Node("j1", True, 11, None),
50                 Node("j2", True, 12, None)]],
51             Node("k", False, None, [
52                 Node("k1", True, 888, None),
53                 Node("k2", True, 888, None)]])]),
54     Node("c", False, None, [
55         Node("f", True, None, [
56             Node("l", False, None, [
57                 Node("l1", True, 5, None),
58                 Node("l2", True, 888, None)]],
59             Node("m", False, None, [
60                 Node("m1", True, 4, None),
61                 Node("m2", True, 888, None)]])],
62     Node("g", True, None, [
63         Node("n", False, None, [
64             Node("n1", True, 888, None),
65             Node("n2", True, 888, None)]],
66         Node("o", False, None, [
67             Node("o1", True, 888, None),
68             Node("o2", True, 888, None)]])])])])

```

The following is a representation of a **magic-sum game**, where players take turns picking a number in the range [1,9], and the first player to have 3 numbers that sum to 15 wins. Note that this is a syntactic variant of **tic-tac-toe** or **naughts and crosses**. To see this, consider the numbers on a **magic square** (Figure 11.1); 3 numbers that add to 15 correspond exactly to the winning positions of tic-tac-toe played on the magic square.

Note that we do not remove symmetries. (What are the symmetries? How

6	1	8
7	5	3
2	9	4

Figure 11.1: Magic Square

do the symmetries of tic-tac-toe translate here?)

masProblem.py — (continued)

```

70
71 class Magic_sum(Node):
72     def __init__(self, xmove=True, last_move=None,
73                 available=[1,2,3,4,5,6,7,8,9], x=[], o=[]):
74         """This is a node in the search for the magic-sum game.
75         xmove is True if the next move belongs to X.
76         last_move is the number selected in the last move
77         available is the list of numbers that are available to be chosen
78         x is the list of numbers already chosen by x
79         o is the list of numbers already chosen by o
80         """
81         self.isMax = self.xmove = xmove
82         self.last_move = last_move
83         self.available = available
84         self.x = x
85         self.o = o
86         self.allchildren = None #computed on demand
87         lm = str(last_move)
88         self.name = "start" if not last_move else "o="+lm if xmove else "x="+lm
89
90     def children(self):
91         if self.allchildren is None:
92             if self.xmove:
93                 self.allchildren = [
94                     Magic_sum(xmove = not self.xmove,
95                             last_move = sel,
96                             available = [e for e in self.available if e is not sel],
97                             x = self.x+[sel],
98                             o = self.o)
99                     for sel in self.available]
100             else:
101                 self.allchildren = [
102                     Magic_sum(xmove = not self.xmove,
103                             last_move = sel,
104                             available = [e for e in self.available if e is not sel],
105                             x = self.x,
106                             o = self.o+[sel])
107                     for sel in self.available]
108         return self.allchildren
109

```

```
110 def isLeaf(self):
111     """A leaf has no numbers available or is a win for one of the players.
112     We only need to check for a win for o if it is currently x's turn,
113     and only check for a win for x if it is o's turn (otherwise it would
114     have been a win earlier).
115     """
116     return (self.available == [] or
117             (sum_to_15(self.last_move,self.o)
118              if self.xmove
119              else sum_to_15(self.last_move,self.x)))
120
121 def evaluate(self):
122     if self.xmove and sum_to_15(self.last_move,self.o):
123         return -1
124     elif not self.xmove and sum_to_15(self.last_move,self.x):
125         return 1
126     else:
127         return 0
128
129 def sum_to_15(last,selected):
130     """is true if last, together with two other elements of selected sum to 15.
131     """
132     return any(last+a+b == 15
133               for a in selected if a != last
134               for b in selected if b != last and b != a)
```

11.1.2 Minimax and α - β Pruning

This is a naive depth-first **minimax algorithm**:

```

11  def minimax(node,depth):
12      """returns the value of node, and a best path for the agents
13      """
14      if node.isLeaf():
15          return node.evaluate(),None
16      elif node.isMax:
17          max_score = float("-inf")
18          max_path = None
19          for C in node.children():
20              score,path = minimax(C,depth+1)
21              if score > max_score:
22                  max_score = score
23                  max_path = C.name,path
24          return max_score,max_path
25      else:
26          min_score = float("inf")
27          min_path = None
28          for C in node.children():
29              score,path = minimax(C,depth+1)
30              if score < min_score:
31                  min_score = score
32                  min_path = C.name,path
33          return min_score,min_path

```

The following is a depth-first minimax with α - β **pruning**. It returns the value for a node as well as a best path for the agents.

```

masMiniMax.py — (continued)
35 def minimax_alpha_beta(node,alpha,beta,depth=0):
36     """node is a Node, alpha and beta are cutoffs, depth is the depth
37     returns value, path
38     where path is a sequence of nodes that results in the value
39     """
40     node.display(2," "*depth,"minimax_alpha_beta(",node.name,", ",alpha, ", ",beta,")")
41     best=None # only used if it will be pruned
42     if node.isLeaf():
43         node.display(2," "*depth,"returning leaf value",node.evaluate())
44         return node.evaluate(),None
45     elif node.isMax:
46         for C in node.children():
47             score,path = minimax_alpha_beta(C,alpha,beta,depth+1)
48             if score >= beta: # beta pruning
49                 node.display(2," "*depth,"pruned due to beta=",beta,"C=",C.name)
50                 return score, None
51             if score > alpha:
52                 alpha = score
53                 best = C.name, path
54             node.display(2," "*depth,"returning max alpha",alpha,"best",best)
55             return alpha,best
56     else:
57         for C in node.children():
58             score,path = minimax_alpha_beta(C,alpha,beta,depth+1)
59             if score <= alpha: # alpha pruning
60                 node.display(2," "*depth,"pruned due to alpha=",alpha,"C=",C.name)
61                 return score, None
62             if score < beta:
63                 beta=score
64                 best = C.name,path
65             node.display(2," "*depth,"returning min beta",beta,"best=",best)
66             return beta,best

```

Testing:

```

masMiniMax.py — (continued)
68 from masProblem import fig10_5, Magic_sum, Node
69
70 # Node.max_display_level=2 # print detailed trace
71 # minimax_alpha_beta(fig10_5, -9999, 9999,0)
72 # minimax_alpha_beta(Magic_sum(), -9999, 9999,0)
73
74 #To see how much time alpha-beta pruning can save over minimax, uncomment the following:
75 ## import timeit
76 ## timeit.Timer("minimax(Magic_sum(),0)",setup="from __main__ import minimax, Magic_sum"
77 ##             ).timeit(number=1)
78 ## trace=False

```

```
79 | ## timeit.Timer("minimax_alpha_beta(Magic_sum(), -9999, 9999,0)",
80 | ##              setup="from __main__ import minimax_alpha_beta, Magic_sum"
81 | ##              ).timeit(number=1)
```


Reinforcement Learning

12.1 Representing Agents and Environments

When the learning agent does an action in the environment, it observes a $(state, reward)$ pair from the environment. The *state* is the world state; this is the fully observable assumption.

An RL environment implements a *do(action)* method that returns a $(state, reward)$ pair.

```
rlProblem.py — Representations for Reinforcement Learning
11 import random
12 from display import Displayable
13 from utilities import flip
14
15 class RL_env(Displayable):
16     def __init__(self, actions, state):
17         self.actions = actions # set of actions
18         self.state = state    # initial state
19
20     def do(self, action):
21         """do action
22         returns state, reward
23         """
24         raise NotImplementedError("RL_env.do") # abstract method
```

Here is the definition of the simple 2-state, 2-action party/relax decision.

```
rlProblem.py — (continued)
26 class Healthy_env(RL_env):
27     def __init__(self):
28         RL_env.__init__(self, ["party", "relax"], "healthy")
29
```

```

30 def do(self, action):
31     """updates the state based on the agent doing action.
32     returns state,reward
33     """
34     if self.state=="healthy":
35         if action=="party":
36             self.state = "healthy" if flip(0.7) else "sick"
37             reward = 10
38         else: # action=="relax"
39             self.state = "healthy" if flip(0.95) else "sick"
40             reward = 7
41     else: # self.state=="sick"
42         if action=="party":
43             self.state = "healthy" if flip(0.1) else "sick"
44             reward = 2
45         else:
46             self.state = "healthy" if flip(0.5) else "sick"
47             reward = 0
48     return self.state,reward

```

12.1.1 Simulating an environment from an MDP

Given the definition for an MDP (page 195), *Env_from_MDP* takes in an MDP and simulates the environment with those dynamics.

Note that the MDP does not contain enough information to simulate a system, because it loses any dependency between the rewards and the resulting state; here we assume the agent always received the average reward for the state and action.

```

rlProblem.py — (continued)
50 class Env_from_MDP(RL_env):
51     def __init__(self, mdp):
52         initial_state = mdp.states[0]
53         RL_env.__init__(self,mdp.actions, initial_state)
54         self.mdp = mdp
55         self.action_index = {action:index for (index,action) in enumerate(mdp.actions)}
56         self.state_index = {state:index for (index,state) in enumerate(mdp.states)}
57
58     def do(self, action):
59         """updates the state based on the agent doing action.
60         returns state,reward
61         """
62         action_ind = self.action_index[action]
63         state_ind = self.state_index[self.state]
64         self.state = pick_from_dist(self.mdp.trans[state_ind][action_ind], self.mdp.states)
65         reward = self.mdp.reward[state_ind][action_ind]
66         return self.state, reward
67
68     def pick_from_dist(dist,values):

```





Figure 12.1: Monster game

```

69     """
70     e.g. pick_from_dist([0.3,0.5,0.2],['a','b','c']) should pick 'a' with probability 0.3, etc.
71     """
72     ran = random.random()
73     i=0
74     while ran>dist[i]:
75         ran -= dist[i]
76         i += 1
77     return values[i]

```

12.1.2 Simple Game

This is for the game depicted in Figure 12.1.

```

_____rlSimpleEnv.py — Simple game_____
11 import random
12 from utilities import flip
13 from rlProblem import RL_env
14
15 class Simple_game_env(RL_env):
16     xdim = 5
17     ydim = 5
18
19     vwalls = [(0,3), (0,4), (1,4)] # vertical walls right of these locations
20     hwalls = [] # not implemented
21     crashed_reward = -1
22
23     prize_locs = [(0,0), (0,4), (4,0), (4,4)]
24     prize_appears_prob = 0.3
25     prize_reward = 10

```

```

26
27 monster_locs = [(0,1), (1,1), (2,3), (3,1), (4,2)]
28 monster_appears_prob = 0.4
29 monster_reward_when_damaged = -10
30 repair_stations = [(1,4)]
31
32 actions = ["up","down","left","right"]
33
34 def __init__(self):
35     # State:
36     self.x = 2
37     self.y = 2
38     self.damaged = False
39     self.prize = None
40     # Statistics
41     self.number_steps = 0
42     self.total_reward = 0
43     self.min_reward = 0
44     self.min_step = 0
45     self.zero_crossing = 0
46     RL_env.__init__(self, Simple_game_env.actions,
47                     (self.x, self.y, self.damaged, self.prize))
48     self.display(2,"","Step","Tot Rew","Ave Rew",sep="\t")
49
50 def do(self,action):
51     """updates the state based on the agent doing action.
52     returns state,reward
53     """
54     reward = 0.0
55     # A prize can appear:
56     if self.prize is None and flip(self.prize_appears_prob):
57         self.prize = random.choice(self.prize_locs)
58     # Actions can be noisy
59     if flip(0.4):
60         actual_direction = random.choice(self.actions)
61     else:
62         actual_direction = action
63     # Modeling the actions given the actual direction
64     if actual_direction == "right":
65         if self.x==self.xdim-1 or (self.x,self.y) in self.vwalls:
66             reward += self.crashed_reward
67         else:
68             self.x += 1
69     elif actual_direction == "left":
70         if self.x==0 or (self.x-1,self.y) in self.vwalls:
71             reward += self.crashed_reward
72         else:
73             self.x += -1
74     elif actual_direction == "up":
75         if self.y==self.ydim-1:

```

```

76         reward += self.crashed_reward
77     else:
78         self.y += 1
79     elif actual_direction == "down":
80         if self.y==0:
81             reward += self.crashed_reward
82         else:
83             self.y += -1
84     else:
85         raise RuntimeError("unknown_direction "+str(direction))
86
87     # Monsters
88     if (self.x,self.y) in self.monster_locs and flip(self.monster_appears_prob):
89         if self.damaged:
90             reward += self.monster_reward_when_damaged
91         else:
92             self.damaged = True
93     if (self.x,self.y) in self.repair_stations:
94         self.damaged = False
95
96     # Prizes
97     if (self.x,self.y) == self.prize:
98         reward += self.prize_reward
99         self.prize = None
100
101     # Statistics
102     self.number_steps += 1
103     self.total_reward += reward
104     if self.total_reward < self.min_reward:
105         self.min_reward = self.total_reward
106         self.min_step = self.number_steps
107     if self.total_reward>0 and reward>self.total_reward:
108         self.zero_crossing = self.number_steps
109     self.display(2,"",self.number_steps,self.total_reward,
110                 self.total_reward/self.number_steps,sep="\t")
111
112     return (self.x, self.y, self.damaged, self.prize), reward

```

12.1.3 Evaluation and Plotting

```

r1Plot.py — RL Plotter
11 import matplotlib.pyplot as plt
12
13 def plot_rl(ag, label=None, yplot='Total', step_size=None,
14            steps_explore=1000, steps_exploit=1000, xscale='linear'):
15     """
16     plots the agent ag
17     label is the label for the plot
18     yplot is 'Average' or 'Total'

```

```

19     step_size is the number of steps between each point plotted
20     steps_explore is the number of steps the agent spends exploring
21     steps_exploit is the number of steps the agent spends exploiting
22     xscale is 'log' or 'linear'
23
24     returns total reward when exploring, total reward when exploiting
25     """
26     assert yplot in ['Average', 'Total']
27     if step_size is None:
28         step_size = max(1, (steps_explore+steps_exploit)//500)
29     if label is None:
30         label = ag.label
31     ag.max_display_level, old_md1 = 1, ag.max_display_level
32     plt.ion()
33     plt.xscale(xscale)
34     plt.xlabel("step")
35     plt.ylabel(yplot+" reward")
36     steps = []      # steps
37     rewards = []    # return
38     ag.restart()
39     step = 0
40     while step < steps_explore:
41         ag.do(step_size)
42         step += step_size
43         steps.append(step)
44         if yplot == "Average":
45             rewards.append(ag.acc_rewards/step)
46         else:
47             rewards.append(ag.acc_rewards)
48     acc_rewards_exploring = ag.acc_rewards
49     ag.explore, explore_save = 0, ag.explore
50     while step < steps_explore+steps_exploit:
51         ag.do(step_size)
52         step += step_size
53         steps.append(step)
54         if yplot == "Average":
55             rewards.append(ag.acc_rewards/step)
56         else:
57             rewards.append(ag.acc_rewards)
58     plt.plot(steps, rewards, label=label)
59     plt.legend(loc="upper left")
60     plt.draw()
61     ag.max_display_level = old_md1
62     ag.explore=explore_save
63     return acc_rewards_exploring, ag.acc_rewards-acc_rewards_exploring

```

12.2 Q Learning

To run the Q-learning demo, in folder “aipython”, load “rlQTest.py”, and copy and paste the example queries at the bottom of that file. This assumes Python 3.

```

rlQLearner.py — Q Learning
11 import random
12 from display import Displayable
13 from utilities import argmaxe, flip
14
15 class RL_agent(Displayable):
16     """An RL_Agent
17     has percepts (s, r) for some state s and real reward r
18     """
19
20 class Q_learner(RL_agent):
21     """A Q-learning agent has
22     belief-state consisting of
23     state is the previous state
24     q is a {(state,action):value} dict
25     visits is a {(state,action):n} dict. n is how many times action was done in state
26     acc_rewards is the accumulated reward
27
28     it observes (s, r) for some world-state s and real reward r
29     """
30
31 def __init__(self, env, discount, explore=0.1, fixed_alpha=True, alpha=0.2,
32             alpha_fun=lambda k:1/k,
33             qinit=0, label="Q_learner"):
34     """env is the environment to interact with.
35     discount is the discount factor
36     explore is the proportion of time the agent will explore
37     fixed_alpha specifies whether alpha is fixed or varies with the number of visits
38     alpha is the weight of new experiences compared to old experiences
39     alpha_fun is a function that computes alpha from the number of visits
40     qinit is the initial value of the Q's
41     label is the label for plotting
42     """
43     RL_agent.__init__(self)
44     self.env = env
45     self.actions = env.actions
46     self.discount = discount
47     self.explore = explore
48     self.fixed_alpha = fixed_alpha
49     self.alpha = alpha

```

```

50     self.alpha_fun = alpha_fun
51     self.qinit = qinit
52     self.label = label
53     self.restart()

```

`restart` is used to make the learner relearn everything. This is used by the plotter to create new plots.

```

_____rlQLearner.py — (continued)_____
55     def restart(self):
56         """make the agent relearn, and reset the accumulated rewards
57         """
58         self.acc_rewards = 0
59         self.state = self.env.state
60         self.q = {}
61         self.visits = {}

```

`do` takes in the number of steps.

```

_____rlQLearner.py — (continued)_____
63     def do(self,num_steps=100):
64         """do num_steps of interaction with the environment"""
65         self.display(2,"s\ta\tr\ts'\tQ")
66         alpha = self.alpha
67         for i in range(num_steps):
68             action = self.select_action(self.state)
69             next_state,reward = self.env.do(action)
70             if not self.fixed_alpha:
71                 k = self.visits[(self.state, action)] = self.visits.get((self.state, action),0)+1
72                 alpha = self.alpha_fun(k)
73             self.q[(self.state, action)] = (
74                 (1-alpha) * self.q.get((self.state, action),self.qinit)
75                 + alpha * (reward + self.discount
76                     * max(self.q.get((next_state, next_act),self.qinit)
77                         for next_act in self.actions)))
78             self.display(2,self.state, action, reward, next_state,
79                 self.q[(self.state, action)], sep='\t')
80             self.state = next_state
81             self.acc_rewards += reward

```

`select_action` is used to select the next action to perform. This can be reimplemented to give a different exploration strategy.

```

_____rlQLearner.py — (continued)_____
83     def select_action(self, state):
84         """returns an action to carry out for the current agent
85         given the state, and the q-function
86         """
87         if flip(self.explore):
88             return random.choice(self.actions)
89         else:
90             return argmaxe((next_act, self.q.get((state, next_act),self.qinit))

```

```
91 | for next_act in self.actions)
```

Exercise 12.1 Implement a soft-max action selection. Choose a temperature that works well for the domain. Explain how you picked this temperature. Compare the epsilon-greedy, soft-max and optimism in the face of uncertainty.

Exercise 12.2 Implement SARSA. Hint: it does not do a *max* in *do*. Instead it needs to choose *next_act* before it does the update.

12.2.1 Testing Q-learning

The first tests are for the 2-action 2-state

```

rlQTest.py — RL Q Tester
11 from rlProblem import Healthy_env
12 from rlLearner import Q_learner
13 from rlPlot import plot_rl
14
15 env = Healthy_env()
16 ag = Q_learner(env, 0.7)
17 ag_opt = Q_learner(env, 0.7, qinit=100, label="optimistic" ) # optimistic agent
18 ag_exp_l = Q_learner(env, 0.7, explore=0.01, label="less explore")
19 ag_exp_m = Q_learner(env, 0.7, explore=0.5, label="more explore")
20 ag_disc = Q_learner(env, 0.9, qinit=100, label="disc 0.9")
21 ag_va = Q_learner(env, 0.7, qinit=100, fixed_alpha=False, alpha_fun=lambda k:10/(9+k), label="alpha=1/k")
22
23 # ag.max_display_level = 2
24 # ag.do(20)
25 # ag.q # get the learned q-values
26 # ag.max_display_level = 1
27 # ag.do(1000)
28 # ag.q # get the learned q-values
29 # plot_rl(ag, yplot="Average")
30 # plot_rl(ag_opt, yplot="Average")
31 # plot_rl(ag_exp_l, yplot="Average")
32 # plot_rl(ag_exp_m, yplot="Average")
33 # plot_rl(ag_disc, yplot="Average")
34 # plot_rl(ag_va, yplot="Average")
35
36 from mdpExamples import mdpt
37 from rlProblem import Env_from_MDP
38 envt = Env_from_MDP(mdpt)
39 agt = Q_learner(envt, 0.8)
40 # agt.do(20)
41
42 from rlSimpleEnv import Simple_game_env
43 senv = Simple_game_env()
44 sag1 = Q_learner(senv, 0.9, explore=0.2, fixed_alpha=True, alpha=0.1)
45 # plot_rl(sag1, steps_explore=100000, steps_exploit=100000, label="alpha="+str(sag1.alpha))
46 sag2 = Q_learner(senv, 0.9, explore=0.2, fixed_alpha=False)
47 # plot_rl(sag2, steps_explore=100000, steps_exploit=100000, label="alpha=1/k")

```

```

48 | sag3 = Q_learner(senv,0.9,explore=0.2,fixed_alpha=False,alpha_fun=lambda k:10/(9+k))
49 | # plot_rl(sag3,steps_explore=100000,steps_exploit=100000,label="alpha=10/(9+k)")

```

12.3 Model-based Reinforcement Learner

To run the demo, in folder “aipython”, load “rlModelLearner.py”, and copy and paste the example queries at the bottom of that file. This assumes Python 3.

A model-based reinforcement learner builds a Markov decision process model of the domain, simultaneously learns the model and plans with that model.

The model-based reinforcement learner used the following data structures:

- $q[s,a]$ is dictionary that, given a (s,a) pair returns the Q -value, the estimate of the future (discounted) value of being in state s and doing action a .
- $r[s,a]$ is dictionary that, given a (s,a) pair returns the average reward from doing a in state s .
- $t[s,a,s']$ is dictionary that, given a (s,a,s') tuple returns the number of times a was done in state s , with the result being state s' .
- $visits[s,a]$ is dictionary that, given a (s,a) pair returns the number of times action a was carried out in state s .
- $res_states[s,a]$ is dictionary that, given a (s,a) pair returns the list of resulting states that have occurred when action a was carried out in state s . This is used in the asynchronous value iteration to determine the s' states to sum over.
- $visits_list$ is a list of (s,a) pair that have been carried out. This is used to ensure there is no divide-by zero in the asynchronous value iteration. Note that this could be constructed from r , $visits$ or res_states by enumerating the keys, but needs to be a list for *random.choice*, and we don't want to keep recreating it.

```

_____rlModelLearner.py — Model-based Reinforcement Learner_____
11 | import random
12 | from rlQLearner import RL_agent
13 | from display import Displayable
14 | from utilities import argmaxe, flip
15 |
16 | class Model_based_reinforcement_learner(RL_agent):
17 |     """A Model-based reinforcement learner
18 |     """
19 |

```



```

20 def __init__(self, env, discount, explore=0.1, qinit=0,
21             updates_per_step=10, label="MBR_learner"):
22     """env is the environment to interact with.
23     discount is the discount factor
24     explore is the proportion of time the agent will explore
25     qinit is the initial value of the Q's
26     updates_per_step is the number of AVI updates per action
27     label is the label for plotting
28     """
29     RL_agent.__init__(self)
30     self.env = env
31     self.actions = env.actions
32     self.discount = discount
33     self.explore = explore
34     self.qinit = qinit
35     self.updates_per_step = updates_per_step
36     self.label = label
37     self.restart()

```

rlModelLearner.py — (continued)

```

39 def restart(self):
40     """make the agent relearn, and reset the accumulated rewards
41     """
42     self.acc_rewards = 0
43     self.state = self.env.state
44     self.q = {} # {(st,action):q_value} map
45     self.r = {} # {(st,action):reward} map
46     self.t = {} # {(st,action,st_next):count} map
47     self.visits = {} # {(st,action):count} map
48     self.res_states = {} # {(st,action):set_of_states} map
49     self.visits_list = [] # list of (st,action)
50     self.previous_action = None

```

rlModelLearner.py — (continued)

```

52 def do(self, num_steps=100):
53     """do num_steps of interaction with the environment
54     for each action, do updates_per_step iterations of asynchronous value iteration
55     """
56     for step in range(num_steps):
57         pst = self.state # previous state
58         action = self.select_action(pst)
59         self.state, reward = self.env.do(action)
60         self.acc_rewards += reward
61         self.t[(pst, action, self.state)] = self.t.get((pst, action, self.state), 0) + 1
62         if (pst, action) in self.visits:
63             self.visits[(pst, action)] += 1
64             self.r[(pst, action)] += (reward - self.r[(pst, action)]) / self.visits[(pst, action)]
65             self.res_states[(pst, action)].add(self.state)
66         else:
67             self.visits[(pst, action)] = 1

```

```

68         self.r[(pst,action)] = reward
69         self.res_states[(pst,action)] = {self.state}
70         self.visits_list.append((pst,action))
71         st,act = pst,action    #initial state-action pair for AVI
72         for update in range(self.updates_per_step):
73             self.q[(st,act)] = self.r[(st,act)]+self.discount*(
74                 sum(self.t[st,act,rst]/self.visits[st,act]*
75                     max(self.q.get((rst,nact),self.qinit) for nact in self.actions)
76                     for rst in self.res_states[(st,act)]))
77             st,act = random.choice(self.visits_list)

```

rlModelLearner.py — (continued)

```

79     def select_action(self, state):
80         """returns an action to carry out for the current agent
81         given the state, and the q-function
82         """
83         if flip(self.explore):
84             return random.choice(self.actions)
85         else:
86             return argmaxe((next_act, self.q.get((state, next_act),self.qinit))
87                             for next_act in self.actions)

```

rlModelLearner.py — (continued)

```

89 from rlQTest import senv # simple game environment
90 mbl1 = Model_based_reinforcement_learner(senv,0.9,updates_per_step=10)
91 # plot_rl(mbl1,steps_explore=100000,steps_exploit=100000,label="model-based(10)")
92 mbl2 = Model_based_reinforcement_learner(senv,0.9,updates_per_step=1)
93 # plot_rl(mbl2,steps_explore=100000,steps_exploit=100000,label="model-based(1)")

```

Exercise 12.3 If there was only one update per step, the algorithm can be made simpler and use less space. Explain how. Does it make it more efficient? Is it worthwhile having more than one update per step for the games implemented here?

Exercise 12.4 It is possible to implement the model-based reinforcement learner by replacing q , r , $visits$, res_states with a single dictionary that returns a tuple (q, r, v, tm) where q , r and v are numbers, and tm is a map from resulting states into counts. Does this make the algorithm easier to understand? Does this make the algorithm more efficient?

Exercise 12.5 If the states and the actions were mapped into integers, the dictionaries could be implemented more efficiently as arrays. This entails an extra step in specifying problems. Implement this for the simple game. Is it more efficient?

12.4 Reinforcement Learning with Features

To run the demo, in folder “aipython”, load “rlFeatures.py”, and copy and paste the example queries at the bottom of that file. This assumes Python 3.

12.4.1 Representing Features

A feature is a function from state and action. To construct the features for a domain, we construct a function that takes a state and an action and returns the list of all feature values for that state and action. This feature set is redesigned for each problem.

`get_features(state, action)` returns the feature values appropriate for the simple game.

```

11 from rlSimpleEnv import Simple_game_env
12 from rlProblem import RL_env
13
14 def get_features(state, action):
15     """returns the list of feature values for the state-action pair
16     """
17     assert action in Simple_game_env.actions
18     (x,y,d,p) = state
19     # f1: would go to a monster
20     f1 = monster_ahead(x,y,action)
21     # f2: would crash into wall
22     f2 = wall_ahead(x,y,action)
23     # f3: action is towards a prize
24     f3 = towards_prize(x,y,action,p)
25     # f4: damaged and action is toward repair station
26     f4 = towards_repair(x,y,action) if d else 0
27     # f5: damaged and towards monster
28     f5 = 1 if d and f1 else 0
29     # f6: damaged
30     f6 = 1 if d else 0
31     # f7: not damaged
32     f7 = 1-f6
33     # f8: damaged and prize ahead
34     f8 = 1 if d and f3 else 0
35     # f9: not damaged and prize ahead
36     f9 = 1 if not d and f3 else 0
37     features = [1,f1,f2,f3,f4,f5,f6,f7,f8,f9]
38     for pr in Simple_game_env.prize_locs+[None]:
39         if p==pr:
40             features += [x, 4-x, y, 4-y]
41         else:
42             features += [0, 0, 0, 0]
43     # fp04 feature for y when prize is at 0,4
44     # this knows about the wall to the right of the prize
45     if p==(0,4):
46         if x==0:
47             fp04 = y
48         elif y<3:
49             fp04 = y
50     else:

```

```

51         fp04 = 4-y
52     else:
53         fp04 = 0
54     features.append(fp04)
55     return features
56
57 def monster_ahead(x,y,action):
58     """returns 1 if the location expected to get to by doing
59     action from (x,y) can contain a monster.
60     """
61     if action == "right" and (x+1,y) in Simple_game_env.monster_locs:
62         return 1
63     elif action == "left" and (x-1,y) in Simple_game_env.monster_locs:
64         return 1
65     elif action == "up" and (x,y+1) in Simple_game_env.monster_locs:
66         return 1
67     elif action == "down" and (x,y-1) in Simple_game_env.monster_locs:
68         return 1
69     else:
70         return 0
71
72 def wall_ahead(x,y,action):
73     """returns 1 if there is a wall in the direction of action from (x,y).
74     This is complicated by the internal walls.
75     """
76     if action == "right" and (x==Simple_game_env.xdim-1 or (x,y) in Simple_game_env.vwalls):
77         return 1
78     elif action == "left" and (x==0 or (x-1,y) in Simple_game_env.vwalls):
79         return 1
80     elif action == "up" and y==Simple_game_env.ydim-1:
81         return 1
82     elif action == "down" and y==0:
83         return 1
84     else:
85         return 0
86
87 def towards_prize(x,y,action,p):
88     """action goes in the direction of the prize from (x,y)"""
89     if p is None:
90         return 0
91     elif p==(0,4): # take into account the wall near the top-left prize
92         if action == "left" and (x>1 or x==1 and y<3):
93             return 1
94         elif action == "down" and (x>0 and y>2):
95             return 1
96         elif action == "up" and (x==0 or y<2):
97             return 1
98         else:
99             return 0
100    else:

```

```

101     px,py = p
102     if p==(4,4) and x==0:
103         if (action=="right" and y<3) or (action=="down" and y>2) or (action=="up" and y<2):
104             return 1
105         else:
106             return 0
107     if (action == "up" and y<py) or (action == "down" and py<y):
108         return 1
109     elif (action == "left" and px<x) or (action == "right" and x<px):
110         return 1
111     else:
112         return 0
113
114 def towards_repair(x,y,action):
115     """returns 1 if action is towards the repair station.
116     """
117     if action == "up" and (x>0 and y<4 or x==0 and y<2):
118         return 1
119     elif action == "left" and x>1:
120         return 1
121     elif action == "right" and x==0 and y<3:
122         return 1
123     elif action == "down" and x==0 and y>2:
124         return 1
125     else:
126         return 0
127
128 def simp_features(state,action):
129     """returns a list of feature values for the state-action pair
130     """
131     assert action in Simple_game_env.actions
132     (x,y,d,p) = state
133     # f1: would go to a monster
134     f1 = monster_ahead(x,y,action)
135     # f2: would crash into wall
136     f2 = wall_ahead(x,y,action)
137     # f3: action is towards a prize
138     f3 = towards_prize(x,y,action,p)
139     return [1,f1,f2,f3]

```

12.4.2 Feature-based RL learner

This learns a linear function approximation of the Q-values. It requires the function *get_features* that given a state and an action returns a list of values for all of the features. Each environment requires this function to be provided.

rlFeatures.py — Feature-based Reinforcement Learner

```

11 import random
12 from rlQLearner import RL_agent
13 from display import Displayable

```

```

14 | from utilities import argmaxe, flip
15 |
16 | class SARSA_LFA_learner(RL_agent):
17 |     """A SARSA_LFA learning agent has
18 |     belief-state consisting of
19 |         state is the previous state
20 |         q is a {(state,action):value} dict
21 |         visits is a {(state,action):n} dict. n is how many times action was done in state
22 |         acc_rewards is the accumulated reward
23 |
24 |     it observes (s, r) for some world-state s and real reward r
25 |     """
26 |     def __init__(self, env, get_features, discount, explore=0.2, step_size=0.01,
27 |                 winit=0, label="SARSA_LFA"):
28 |         """env is the feature environment to interact with
29 |         get_features is a function get_features(state,action) that returns the list of feature values
30 |         discount is the discount factor
31 |         explore is the proportion of time the agent will explore
32 |         step_size is gradient descent step size
33 |         winit is the initial value of the weights
34 |         label is the label for plotting
35 |         """
36 |         RL_agent.__init__(self)
37 |         self.env = env
38 |         self.get_features = get_features
39 |         self.actions = env.actions
40 |         self.discount = discount
41 |         self.explore = explore
42 |         self.step_size = step_size
43 |         self.winit = winit
44 |         self.label = label
45 |         self.restart()

```

`restart()` is used to make the learner relearn everything. This is used by the plotter to create new plots.

```

47 |     def restart(self):
48 |         """make the agent relearn, and reset the accumulated rewards
49 |         """
50 |         self.acc_rewards = 0
51 |         self.state = self.env.state
52 |         self.features = self.get_features(self.state, list(self.env.actions)[0])
53 |         self.weights = [self.winit for f in self.features]
54 |         self.action = self.select_action(self.state)

```

`do` takes in the number of steps.

```

56 |     def do(self, num_steps=100):
57 |         """do num_steps of interaction with the environment"""
58 |         self.display(2, "s\t a\t r\t ts\t tQ\t ddelta")

```

```

59     for i in range(num_steps):
60         next_state, reward = self.env.do(self.action)
61         self.acc_rewards += reward
62         next_action = self.select_action(next_state)
63         feature_values = self.get_features(self.state, self.action)
64         oldQ = dot_product(self.weights, feature_values)
65         nextQ = dot_product(self.weights, self.get_features(next_state, next_action))
66         delta = reward + self.discount * nextQ - oldQ
67         for i in range(len(self.weights)):
68             self.weights[i] += self.step_size * delta * feature_values[i]
69         self.display(2, self.state, self.action, reward, next_state,
70                     dot_product(self.weights, feature_values), delta, sep='\t')
71         self.state = next_state
72         self.action = next_action
73
74     def select_action(self, state):
75         """returns an action to carry out for the current agent
76         given the state, and the q-function.
77         This implements an epsilon-greedy approach
78         where self.explore is the probability of exploring.
79         """
80         if flip(self.explore):
81             return random.choice(self.actions)
82         else:
83             return argmaxe((next_act, dot_product(self.weights,
84                                                     self.get_features(state, next_act)))
85                             for next_act in self.actions)
86
87     def show_actions(self, state=None):
88         """prints the value for each action in a state.
89         This may be useful for debugging.
90         """
91         if state is None:
92             state = self.state
93         for next_act in self.actions:
94             print(next_act, dot_product(self.weights, self.get_features(state, next_act)))
95
96     def dot_product(l1, l2):
97         return sum(e1*e2 for (e1, e2) in zip(l1, l2))

```

Test code:

```

rlFeatures.py — (continued)
100 from rlQTest import env # simple game environment
101 from rlSimpleGameFeatures import get_features, simp_features
102 from rlPlot import plot_rl
103
104 fa1 = SARSA_LFA_learner(env, get_features, 0.9, step_size=0.01)
105 #fa1.max_display_level = 2
106 #fa1.do(20)
107 #plot_rl(fa1, steps_explore=10000, steps_exploit=10000, label="SARSA_LFA(0.01)")

```

```

108 | fas1 = SARSA_LFA_learner(senv, simp_features, 0.9, step_size=0.01)
109 | #plot_rl(fas1, steps_explore=10000, steps_exploit=10000, label="SARSA_LFA(simp)")

```

Exercise 12.6 How does the step-size affect performance? Try different step sizes (e.g., 0.1, 0.001, other sizes in between). Explain the behaviour you observe. Which step size works best for this example. Explain what evidence you are basing your prediction on.

Exercise 12.7 Does having extra features always help? Does it sometime help? Does whether it helps depend on the step size? Give evidence for your claims.

Exercise 12.8 For each of the following first predict, then plot, then explain the behaviour you observed:

- (a) SARSA_LFA, Model-based learning (with 1 update per step) and Q-learning for 10,000 steps 20% exploring followed by 10,000 steps 100% exploiting
- (b) SARSA_LFA, model-based learning and Q-learning for
 - i) 100,000 steps 20% exploring followed by 100,000 steps 100% exploit
 - ii) 10,000 steps 20% exploring followed by 190,000 steps 100% exploit
- (c) Suppose your goal was to have the best accumulated reward after 200,000 steps. You are allowed to change the exploration rate at a fixed number of steps. For each of the methods, which is the best position to start exploiting more? Which method is better? What if you wanted to have the best reward after 10,000 or 1,000 steps?

Based on this evidence, explain when it is preferable to use SARSA_LFA, Model-based learner, or Q-learning.

Important: you need to run each algorithm more than once. Your explanation should include the variability as well as the typical behavior.

12.5 Multiagent Learning

The next code of for multiple agents that learn when interacting with other agents. This code is designed to be extended, and as such is restricted to being two agents, a single state, and the only observation is the reward. Coordinating agents can't easily implement that agent architecture. However, in that architecture, an agent calls the environment. That architecture was chosen because it was simple. However, it does not really work when there are multiple agents, instead we have a controller that tells the agents the percepts (here the percepts are just the reward).

```

masLearn.py — Simulations of agents learning
11 | from display import Displayable
12 | import utilities # argmaxall for (element,value) pairs
13 | import matplotlib.pyplot as plt
14 | import random
15 |
16 | class GameAgent(Displayable):

```



```

17     next_id=0
18     def __init__(self, actions):
19         """
20         Actions is the set of actions the agent can do. It needs to be told that!
21         """
22         self.actions = actions
23         self.id = GameAgent.next_id
24         GameAgent.next_id += 1
25         self.display(2,f"Agent {self.id} has actions {actions}")
26         self.dist = {act:1 for act in actions} # unnormalized distribution
27         self.total_score = 0
28
29     def init_action(self):
30         """ The initial action.
31         Act randomly initially
32         Could be overridden (but I'm not sure why you would).
33         """
34         self.act = random.choice(self.actions)
35         return self.act
36
37     def select_action(self, reward):
38         """
39         Select the action given the reward.
40         This implements "Act randomly" and should be overridden!
41         """
42         self.total_score += reward
43         self.act = random.choice(self.actions)
44         return self.act

```

masLearn.py — (continued)

```

46 class SimpleCountingAgent(GameAgent):
47     """This agent just counts the number of times (it thinks) it has won and does the
48     actions it thinks is most likely to win.
49     """
50     def __init__(self, actions, prior_count=1):
51         """
52         Actions is the set of actions the agent can do. It needs to be told that!
53         """
54         GameAgent.__init__(self, actions)
55         self.prior_count = prior_count
56         self.dist = {a: prior_count for a in self.actions} # unnormalized distribution
57         self.averew = 0
58         self.num_steps = 0
59
60     def select_action(self, reward):
61         self.total_score += reward
62         self.num_steps += 1
63         self.display(2,f"The reward for agent {self.id} was {reward}")
64         self.averew = self.averew+(reward-self.averew)/self.num_steps
65         if reward>self.averew:

```

```

66         self.dist[self.act] += 1
67     else:
68         for otheract in self.actions:
69             if otheract != self.act:
70                 self.dist[otheract] += 1/(len(self.actions))
71     self.display(2,f"Distribution for agent {self.id} is {normalize(self.dist)}")
72     self.act = select_from_dist(self.dist)
73     self.display(2,f"Agent {self.id} did {self.act}")
74     return self.act

```

masLearn.py — (continued)

```

76 class SimpleQAgent(GameAgent):
77     """This agent maintains the Q-function for each state.
78     (Or just the average reward as the future state is all the same).
79     Chooses the best action using
80     """
81     def __init__(self, actions, q_init=100, alpha=0.1, prob_step_size=0.001, min_prob=0.01):
82         """
83         Actions is the set of actions the agent can do. It needs to be told that!
84         q_init is the initial q-values
85         alpha is the step size for action estimate
86         prob_step_size is the step size for probability change
87         min_prob is the minimum a probability should become
88         """
89         GameAgent.__init__(self, actions)
90         self.Q = {a:q_init for a in self.actions}
91         self.dist = normalize({a:0.7+random.random() for a in self.actions}) # start with random dist
92         zero
93         self.alpha = alpha
94         self.prob_step_size = prob_step_size
95         self.min_prob = min_prob
96         self.num_steps = 1 # (1 because it is only used after initial step)
97     def select_action(self, reward):
98         self.total_score += reward
99         self.display(2,f"The reward for agent {self.id} was {reward}")
100        self.Q[self.act] += self.alpha*(reward-self.Q[self.act])
101        a_best = utilities.argmaxall(self.Q.items())
102        for a in self.actions:
103            if a in a_best:
104                self.dist[a] += self.prob_step_size
105            else:
106                self.dist[a] -= min(self.dist[a], self.prob_step_size)
107                self.dist[a] = max(self.dist[a],self.min_prob)
108        self.dist = normalize(self.dist)
109        self.display(2,f"Distribution for agent {self.id} is {self.dist}")
110        self.act = select_from_dist(self.dist)
111        self.display(2,f"Agent {self.id} did {self.act}")
112        return self.act
113

```

```

114 def normalize(dist):
115     """unnorm dict is a {value:number} dictionary, where the numbers are all non-negative
116     returns dict where the numbers sum to one
117     """
118     tot = sum(dist.values())
119     return {var:val/tot for (var,val) in dist.items()}
120
121 def select_from_dist(dist):
122     rand = random.random()
123     for (act,prob) in normalize(dist).items():
124         rand -= prob
125         if rand < 0:
126             return act

```

The simulator takes a game and simulates the game:

```

masLearn.py — (continued)
128 class SimulateGame(Displayable):
129     def __init__(self, game, agents):
130         self.game = game
131         self.agents = agents # list of agents
132         self.action_history = []
133         self.reward_history = []
134         self.dist_history = []
135         self.actions = tuple(ag.init_action() for ag in self.agents)
136         self.num_steps = 0
137
138     def go(self, steps):
139         for i in range(steps):
140             self.num_steps += 1
141             self.rewards = self.game.play(self.actions)
142             self.reward_history.append(self.rewards)
143             self.actions = tuple(self.agents[i].select_action(self.rewards[i])
144                                 for i in range(self.game.num_agents))
145             self.action_history.append(self.actions)
146             self.dist_history.append([normalize(ag.dist) for ag in self.agents])
147             print("Scores:", ' '.join(f"Agent {ag.id} average reward={ag.total_score/self.num_steps}"
148                                     #return self.reward_history, self.action_history
149
150     def action_dist(self, which_actions=[1,1]):
151         """ which actions is [a0,a1]
152         returns the empirical distribution of actions for agents,
153         where ai specifies the index of the actions for agent i
154         """
155         return [sum(1 for a in sim.action_history
156                     if a[i]==gm.actions[i][which_actions[i]])/len(sim.action_history)
157                 for i in range(2)]

```

```

masLearn.py — (continued)
159
160 def plot_dynamics(self, x_action=0, y_action=0):

```

```

161 plt.ion() # make it interactive
162 agents = self.agents
163 x_act = self.game.actions[0][x_action]
164 y_act = self.game.actions[1][y_action]
165 plt.xlabel(f"Action {self.agents[0].actions[x_action]} for Agent {agents[0].id}")
166 plt.ylabel(f"Action {self.agents[1].actions[y_action]} for Agent {agents[1].id}")
167 plt.plot([self.dist_history[t][0][x_act] for t in range(len(self.dist_history))],
168          [self.dist_history[t][1][y_act] for t in range(len(self.dist_history))])
169 plt.legend()

```

The following are some games from Poole and Mackworth [2017].

```

masLearn.py — (continued)
172 class ShoppingGame(Displayable):
173     def __init__(self):
174         self.num_agents = 2
175         self.actions = [['shopping', 'football']]*2
176
177     def play(self, actions):
178         return {('football', 'football'): (2,1),
179                ('football', 'shopping'): (0,0),
180                ('shopping', 'football'): (0,0),
181                ('shopping', 'shopping'): (1,2)}[actions]
182
183
184 class SoccerGame(Displayable):
185     def __init__(self):
186         self.num_agents = 2
187         self.actions = [['left', 'right']]*2
188
189     def play(self, actions):
190         return {('left', 'left'): (0.6, 0.4),
191                ('left', 'right'): (0.2, 0.8),
192                ('right', 'left'): (0.3, 0.7),
193                ('right', 'right'): (0.9,0.1)
194                }[actions]
195
196 class GameShow(Displayable):
197     def __init__(self):
198         self.num_agents = 2
199         self.actions = [['take', 'give']]*2
200
201     def play(self, actions):
202         return {('take', 'take'): (100, 100),
203                ('take', 'give'): (1100, 0),
204                ('give', 'take'): (0, 1100),
205                ('give', 'give'): (1000,1000)
206                }[actions]
207
208
209 class UniqueNEGGameExample(Displayable):

```

```

210     def __init__(self):
211         self.num_agents = 2
212         self.actions = [['a1', 'b1', 'c1'], ['d2', 'e2', 'f2']]
213
214     def play(self, actions):
215         return {('a1', 'd2'): (3, 5),
216                 ('a1', 'e2'): (5, 1),
217                 ('a1', 'f2'): (1, 2),
218                 ('b1', 'd2'): (1, 1),
219                 ('b1', 'e2'): (2, 9),
220                 ('b1', 'f2'): (6, 4),
221                 ('c1', 'd2'): (2, 6),
222                 ('c1', 'e2'): (4, 7),
223                 ('c1', 'f2'): (0, 8)
224                 }[actions]
225
226 # Choose one:
227 # gm = ShoppingGame()
228 # gm = SoccerGame()
229 # gm = GameShow()
230 # gm = UniqueNEGameExample()
231
232 # Choose one:
233 # sim=SimulateGame(gm,[SimpleQAgent(gm.actions[0]), SimpleQAgent(gm.actions[1])]); sim.go(10000)
234 # sim= SimulateGame(gm,[SimpleCountingAgent(gm.actions[0]), SimpleCountingAgent(gm.actions[1])]);
235 # sim=SimulateGame(gm,[SimpleCountingAgent(gm.actions[0]), SimpleQAgent(gm.actions[1])]); sim.go(1
236
237
238 # sim.plot_dynamics()
239
240 # empirical proportion that agents did their action at index 1:
241 # sim.action_dist([1,1])
242
243 # learned distribution for agent 0
244 # sim.agents[0].dist

```


Relational Learning

13.1 Collaborative Filtering

Based on gradient descent algorithm of Koren, Y., Bell, R. and Volinsky, C., Matrix Factorization Techniques for Recommender Systems, IEEE Computer 2009.

This assumes the form of the dataset from movielens (<http://grouplens.org/datasets/movielens/>). The rating are a set of (*user, item, rating, timestamp*) tuples.

```
_____relnCollFilt.py — Latent Property-based Collaborative Filtering _____
11 import random
12 import matplotlib.pyplot as plt
13 import urllib.request
14 from learnProblem import Learner
15 from display import Displayable
16
17 class CF_learner(Learner):
18     def __init__(self,
19                 rating_set,          # a Rating_set object
20                 rating_subset = None, # subset of ratings to be used as training ratings
21                 test_subset = None,   # subset of ratings to be used as test ratings
22                 step_size = 0.01,    # gradient descent step size
23                 reglz = 1.0,         # the weight for the regularization terms
24                 num_properties = 10,  # number of hidden properties
25                 property_range = 0.02 # properties are initialized to be between
26                                     # -property_range and property_range
27             ):
28         self.rating_set = rating_set
29         self.ratings = rating_subset or rating_set.training_ratings # whichever is not empty
30         if test_subset is None:
```

```

31         self.test_ratings = self.rating_set.test_ratings
32     else:
33         self.test_ratings = test_subset
34         self.step_size = step_size
35         self.reglz = reglz
36         self.num_properties = num_properties
37         self.num_ratings = len(self.ratings)
38         self.ave_rating = (sum(r for (u,i,r,t) in self.ratings)
39                             /self.num_ratings)
40         self.users = {u for (u,i,r,t) in self.ratings}
41         self.items = {i for (u,i,r,t) in self.ratings}
42         self.user_bias = {u:0 for u in self.users}
43         self.item_bias = {i:0 for i in self.items}
44         self.user_prop = {u:[random.uniform(-property_range,property_range)
45                               for p in range(num_properties)]
46                           for u in self.users}
47         self.item_prop = {i:[random.uniform(-property_range,property_range)
48                               for p in range(num_properties)]
49                             for i in self.items}
50         self.zeros = [0 for p in range(num_properties)]
51         self.iter=0
52
53     def stats(self):
54         self.display(1,"ave sumsq error of mean for training=",
55                     sum((self.ave_rating-rating)**2 for (user,item,rating,timestamp)
56                         in self.ratings)/len(self.ratings))
57         self.display(1,"ave sumsq error of mean for test=",
58                     sum((self.ave_rating-rating)**2 for (user,item,rating,timestamp)
59                         in self.test_ratings)/len(self.test_ratings))
60         self.display(1,"error on training set",
61                     self.evaluate(self.ratings))
62         self.display(1,"error on test set",
63                     self.evaluate(self.test_ratings))

```

learn carries out *num_iter* steps of gradient descent.

relnCollFilt.py — (continued)

```

65     def prediction(self,user,item):
66         """Returns prediction for this user on this item.
67         The use of .get() is to handle users or items not in the training set.
68         """
69         return (self.ave_rating
70                 + self.user_bias.get(user,0) #self.user_bias[user]
71                 + self.item_bias.get(item,0) #self.item_bias[item]
72                 + sum([self.user_prop.get(user,self.zeros)[p]*self.item_prop.get(item,self.zeros)[p]
73                       for p in range(self.num_properties)]))
74
75     def learn(self, num_iter = 50):
76         """ do num_iter iterations of gradient descent."""
77         for i in range(num_iter):
78             self.iter += 1

```



```

79         abs_error=0
80         sumsq_error=0
81         for (user,item,rating,timestamp) in random.sample(self.ratings,len(self.ratings)):
82             error = self.prediction(user,item) - rating
83             abs_error += abs(error)
84             sumsq_error += error * error
85             self.user_bias[user] -= self.step_size*error
86             self.item_bias[item] -= self.step_size*error
87             for p in range(self.num_properties):
88                 self.user_prop[user][p] -= self.step_size*error*self.item_prop[item][p]
89                 self.item_prop[item][p] -= self.step_size*error*self.user_prop[user][p]
90         for user in self.users:
91             self.user_bias[user] -= self.step_size*self.reglz* self.user_bias[user]
92             for p in range(self.num_properties):
93                 self.user_prop[user][p] -= self.step_size*self.reglz*self.user_prop[user][p]
94         for item in self.items:
95             self.item_bias[item] -= self.step_size*self.reglz*self.item_bias[item]
96             for p in range(self.num_properties):
97                 self.item_prop[item][p] -= self.step_size*self.reglz*self.item_prop[item][p]
98         self.display(1,"Iteration",self.iter,
99                     "(Ave Abs,AveSumSq) training =",self.evaluate(self.ratings),
100                    "test =",self.evaluate(self.test_ratings))

```

evaluate evaluates current predictions on the rating set:

```

relnCollFilt.py — (continued)
102 def evaluate(self,ratings):
103     """returns (average_absolute_error, average_sum_squares_error) for ratings
104     """
105     abs_error = 0
106     sumsq_error = 0
107     if not ratings: return (0,0)
108     for (user,item,rating,timestamp) in ratings:
109         error = self.prediction(user,item) - rating
110         abs_error += abs(error)
111         sumsq_error += error * error
112     return abs_error/len(ratings), sumsq_error/len(ratings)

```

13.1.1 Alternative Formulation

An alternative formulation is to regularize after each update.

13.1.2 Plotting

```

relnCollFilt.py — (continued)
114 def plot_predictions(self, examples="test"):
115     """
116     examples is either "test" or "training" or the actual examples
117     """

```

```

118     if examples == "test":
119         examples = self.test_ratings
120     elif examples == "training":
121         examples = self.ratings
122     plt.ion()
123     plt.xlabel("prediction")
124     plt.ylabel("cumulative proportion")
125     self.actuals = [[] for r in range(0,6)]
126     for (user,item,rating,timestamp) in examples:
127         self.actuals[rating].append(self.prediction(user,item))
128     for rating in range(1,6):
129         self.actuals[rating].sort()
130         numrat=len(self.actuals[rating])
131         yvals = [i/numrat for i in range(numrat)]
132         plt.plot(self.actuals[rating], yvals, label="rating="+str(rating))
133     plt.legend()
134     plt.draw()

```

This plots a single property. Each $(user, item, rating)$ is plotted where the x -value is the value of the property for the user, the y -value is the value of the property for the item, and the rating is plotted at this (x, y) position. That is, $rating$ is plotted at the (x, y) position $(p(user), p(item))$.

```

relnCollFilt.py — (continued)
136     def plot_property(self,
137         p,                # property
138         plot_all=False, # true if all points should be plotted
139         num_points=200 # number of random points plotted if not all
140     ):
141         """plot some of the user-movie ratings,
142         if plot_all is true
143         num_points is the number of points selected at random plotted.
144
145         the plot has the users on the x-axis sorted by their value on property p and
146         with the items on the y-axis sorted by their value on property p and
147         the ratings plotted at the corresponding x-y position.
148         """
149         plt.ion()
150         plt.xlabel("users")
151         plt.ylabel("items")
152         user_vals = [self.user_prop[u][p]
153             for u in self.users]
154         item_vals = [self.item_prop[i][p]
155             for i in self.items]
156         plt.axis([min(user_vals)-0.02,
157             max(user_vals)+0.05,
158             min(item_vals)-0.02,
159             max(item_vals)+0.05])
160         if plot_all:
161             for (u,i,r,t) in self.ratings:
162                 plt.text(self.user_prop[u][p],

```

```

163         self.item_prop[i][p],
164         str(r))
165     else:
166         for i in range(num_points):
167             (u,i,r,t) = random.choice(self.ratings)
168             plt.text(self.user_prop[u][p],
169                     self.item_prop[i][p],
170                     str(r))
171     plt.show()

```

13.1.3 Creating Rating Sets

A rating set can be read from the Internet or read from a local file. The default is to read the Movielens 100K dataset from the Internet. It would be more efficient to save the dataset as a local file, and then set *local_file = True*, as then it will not need to download the dataset every time the program is run.

```

relnCollFilt.py — (continued)
173 class Rating_set(Displayable):
174     def __init__(self,
175                 date_split=892000000,
176                 local_file=False,
177                 url="http://files.grouplens.org/datasets/movielens/ml-100k/u.data",
178                 file_name="u.data"):
179         self.display(1,"reading...")
180         if local_file:
181             lines = open(file_name,'r')
182         else:
183             lines = (line.decode('utf-8') for line in urllib.request.urlopen(url))
184         all_ratings = (tuple(int(e) for e in line.strip().split('\t'))
185                        for line in lines)
186         self.training_ratings = []
187         self.training_stats = {1:0, 2:0, 3:0, 4:0 ,5:0}
188         self.test_ratings = []
189         self.test_stats = {1:0, 2:0, 3:0, 4:0 ,5:0}
190         for rate in all_ratings:
191             if rate[3] < date_split: # rate[3] is timestamp
192                 self.training_ratings.append(rate)
193                 self.training_stats[rate[2]] += 1
194             else:
195                 self.test_ratings.append(rate)
196                 self.test_stats[rate[2]] += 1
197         self.display(1,"...read:", len(self.training_ratings),"training ratings and",
198                     len(self.test_ratings),"test ratings")
199         tr_users = {user for (user,item,rating,timestamp) in self.training_ratings}
200         test_users = {user for (user,item,rating,timestamp) in self.test_ratings}
201         self.display(1,"users:",len(tr_users),"training,",len(test_users),"test,",
202                     len(tr_users & test_users),"in common")
203         tr_items = {item for (user,item,rating,timestamp) in self.training_ratings}
204         test_items = {item for (user,item,rating,timestamp) in self.test_ratings}
205         self.display(1,"items:",len(tr_items),"training,",len(test_items),"test,",

```

```

206         len(tr_items & test_items),"in common")
207     self.display(1,"Rating statistics for training set: ",self.training_stats)
208     self.display(1,"Rating statistics for test set: ",self.test_stats)

```

Sometimes it is useful to plot a property for all (*user, item, rating*) triples. There are too many such triples in the data set. The method *create_top_subset* creates a much smaller dataset where this makes sense. It picks the most rated items, then picks the users who have the most ratings on these items. It is designed for depicting the meaning of properties, and may not be useful for other purposes.

```

relnCollFilt.py — (continued)
210     def create_top_subset(self, num_items = 30, num_users = 30):
211         """Returns a subset of the ratings by picking the most rated items,
212         and then the users that have most ratings on these, and then all of the
213         ratings that involve these users and items.
214         """
215         items = {item for (user,item,rating,timestamp) in self.training_ratings}
216
217         item_counts = {i:0 for i in items}
218         for (user,item,rating,timestamp) in self.training_ratings:
219             item_counts[item] += 1
220
221         items_sorted = sorted((item_counts[i],i) for i in items)
222         top_items = items_sorted[-num_items:]
223         set_top_items = set(item for (count, item) in top_items)
224
225         users = {user for (user,item,rating,timestamp) in self.training_ratings}
226         user_counts = {u:0 for u in users}
227         for (user,item,rating,timestamp) in self.training_ratings:
228             if item in set_top_items:
229                 user_counts[user] += 1
230
231         users_sorted = sorted((user_counts[u],u)
232                                for u in users)
233         top_users = users_sorted[-num_users:]
234         set_top_users = set(user for (count, user) in top_users)
235         used_ratings = [ (user,item,rating,timestamp)
236                          for (user,item,rating,timestamp) in self.training_ratings
237                          if user in set_top_users and item in set_top_items]
238         return used_ratings
239
240     movielens = Rating_set()
241     learner1 = CF_learner(movielens, num_properties = 1)
242     #learner1.learn(50)
243     # learner1.plot_predictions(examples = "training")
244     # learner1.plot_predictions(examples = "test")
245     #learner1.plot_property(0)
246     #movielens_subset = movielens.create_top_subset(num_items = 20, num_users = 20)
247     #learner_s = CF_learner(movielens, rating_subset=movielens_subset, test_subset=[], num_properties=1)
248     #learner_s.learn(1000)

```

```
249 | #learner_s.plot_property(0,plot_all=True)
```


Version History

- 2021-05-13 Version 0.9.0 Major revisions to chapters 8 and 9. Introduced recursive conditioning, simplified much code. New section on multi-agent reinforcement learning.
- 2020-11-04 Version 0.8.6 simplified value iteration for MDPs.
- 2020-10-20 Version 0.8.4 planning simplified, and gives error if goal not part of state (by design). Fixed arc costs.
- 2020-07-21 Version 0.8.2 added positions and string to constraints
- 2019-09-17 Version 0.8.0 rerepresented blocks world (Section 6.1.2) due to bug found by Donato Meoli.

Bibliography

- Lichman, M. (2013), UCI machine learning repository. URL <http://archive.ics.uci.edu/ml>. 107
- Poole, D. L. and Mackworth, A. K. (2017), *Artificial Intelligence: foundations of computational agents*. Cambridge University Press, 2nd edition, URL <https://artint.info>. 148

Index

- α - β pruning, 220
- A* search, 41
- A* Search, 44
- action, 85
- agent, 21, 223
- argmax, 18
- assignment, 52, 142
- assumable, 82
- asynchronous value iteration, 201
- augmented feature, 116
- batched stochastic gradient descent, 133
- Bayesian network, 146
- belief network, 146
- blocks world, 88
- Boolean feature, 108
- bottom-up proof, 79
- branch-and-bound search, 46
- class
 - Action_instance*, 101
 - Agent*, 21
 - Arc*, 34
 - Askable*, 77
 - Assumable*, 82
 - BNfromDBN*, 179
 - BeliefNetwork*, 146
 - Boosted_dataset*, 138
 - Boosting_learner*, 139
 - Branch_and_bound*, 47
 - CF_learner*, 241
 - CPDrename*, 177
 - CSP*, 52
 - CSP_from_STRIPS*, 98
 - Clause*, 77
 - Con_solver*, 61
 - Constraint*, 51
 - DBN*, 177
 - DBNVEfilter*, 180
 - DBNvariable*, 176
 - DT_learner*, 121
 - Data_from_file*, 112
 - Data_set*, 108
 - Data_set_augmented*, 116
 - DecisionNetwork*, 184
 - DecisionVariable*, 183
 - Displayable*, 17
 - EM_learner*, 209
 - Env_from_MDP*, 224
 - Environment*, 22
 - FactorDF*, 194

- FactorMax*, 193
- FactorObserved*, 157
- FactorRename*, 177
- FactorSum*, 158
- Forward_STRIPS*, 91
- FrontierPQ*, 42
- GibbsSampling*, 166
- GraphicalModel*, 146
- GridMDP*, 196, 199
- HMM*, 170
- HMMVEfilter*, 171
- HMMparticleFilter*, 173
- Healthyenv*, 223
- InferenceMethod*, 152
- KB*, 78
- KBA*, 82
- K_fold_dataset*, 125
- K_means_learner*, 205
- Layer*, 134
- Learner*, 118
- LikelihoodWeighting*, 163
- Linear_complete_layer*, 134
- Linear_learner*, 127
- Linear_learner_bsgd*, 133
- MDP*, 195
- Magic_sum*, 217
- Model_based_reinforcement_learner*, 232
- NN*, 136
- Node*, 215
- POP_node*, 102
- POP_search_from_STRIPS*, 103
- ParticleFiltering*, 164
- Path*, 36
- Planning_problem*, 86
- Plot_env*, 30
- Plot_prices*, 24
- Prob*, 146
- Q_learner*, 229
- RC*, 153
- RC_DN*, 190
- RL_agent*, 229
- RL_env*, 223
- Rating_set*, 245
- ReLU_layer*, 135
- Regression_STRIPS*, 95
- RejectionSampling*, 162
- Rob_body*, 26
- Rob_env*, 25
- Rob_middle_layer*, 28
- Rob_top_layer*, 29
- Runtime_distribution*, 74
- SARSA_LFA_learner*, 237
- SLSearcher*, 67
- STRIPS_domain*, 86
- SamplingInferenceMethod*, 162
- Search_from_CSP*, 59
- Search_problem*, 33
- Search_problem_from_explicit_graph*, 35
- Search_with_AC_from_CSP*, 65
- Searcher*, 41
- SearcherMPP*, 45
- Sigmoid_layer*, 135
- Simple_game_env*, 225
- State*, 91
- Strips*, 85
- Subgoal*, 95
- TP_agent*, 24
- TP_env*, 22
- TabFactor*, 145
- Updatable_priority_queue*, 72
- Utility*, 183
- UtilityTable*, 183
- VE*, 156
- VE_DN*, 193
- Variable*, 141
- clause, 77
- collaborative filtering, 241
- condition, 51
- conditional probability distribution (CPD), 143
- consistency algorithms, 61
- constraint, 51
- constraint satisfaction problem, 51
- copy_with_assign, 64
- CPD (conditional probability distribution), 143
- cross validation, 125
- CSP, 51

- consistency, 61
- domain splitting, 63, 65
- search, 59
- stochastic local search, 67
- currying, 53
- data set, 108
- DBN
 - filtering, 180
 - unrolling, 179
- DBN (dynamic belief network), 175
- decision network, 183
- decision tree learning, 121
- decision variable, 183
- deep learning, 133
- dict_union, 19
- display, 17
- Displayable, 17
- domain splitting, 63, 65
- dynamic belief network (DBN), 175
 - representation, 176
- EM, 209
- environment, 21, 22, 223
- example, 108
- explicit graph, 34
- factor, 142, 145
- factor_times, 158
- feature, 108
- feature engineering, 107
- file
 - agentEnv.py*, 25
 - agentMiddle.py*, 28
 - agentTop.py*, 29
 - agents.py*, 21
 - cspConsistency.py*, 61
 - cspExamples.py*, 53
 - cspProblem.py*, 51
 - cspSLS.py*, 67
 - cspSearch.py*, 59
 - decnNetworks.py*, 183
 - display.py*, 17
 - learnBoosting.py*, 138
 - learnCoordinate.py*, 240
 - learnCrossValidation.py*, 125
 - learnDT.py*, 121
 - learnEM.py*, 209
 - learnKMeans.py*, 205
 - learnLinear.py*, 127
 - learnLinearBSGD.py*, 133
 - learnNN.py*, 134
 - learnNoInputs.py*, 119
 - learnProblem.py*, 108
 - logicAssumables.py*, 82
 - logicBottomUp.py*, 79
 - logicProblem.py*, 77
 - logicTopDown.py*, 81
 - masMiniMax.py*, 219
 - masProblem.py*, 215
 - mdpExamples.py*, 195
 - mdpProblem.py*, 195
 - probDBN.py*, 176
 - probFactors.py*, 142
 - probGraphicalModels.py*, 146
 - probHMM.py*, 170
 - probRC.py*, 153
 - probStochSim.py*, 160
 - probVE.py*, 156
 - probVariables.py*, 141
 - pythonDemo.py*, 13
 - relnCollFilt.py*, 241
 - rlFeatures.py*, 237
 - rlModelLearner.py*, 232
 - rlPlot.py*, 227
 - rlProblem.py*, 223
 - rlQLearner.py*, 229
 - rlQTest.py*, 231
 - rlSimpleEnv.py*, 225
 - rlSimpleGameFeatures.py*, 235
 - searchBranchAndBound.py*, 47
 - searchGeneric.py*, 41
 - searchMPP.py*, 45
 - searchProblem.py*, 33
 - searchTest.py*, 48
 - stripsCSPPlanner.py*, 98
 - stripsForwardPlanner.py*, 91
 - stripsHeuristic.py*, 93
 - stripsPOP.py*, 101
 - stripsProblem.py*, 85
 - stripsRegressionPlanner.py*, 95

- utilities.py*, 18
- filtering, 171, 173
 - DBN, 180
- forward planning, 90
- game, 215
- Gibbs sampling, 166
- graphical model, 146
- heuristic planning, 92, 97
- hidden Markov model, 170
- hierarchical controller, 25
- HMM
 - exact filtering, 171
 - particle filtering, 173
- HMM (hidden Markov models), 170
- importance sampling, 164
- ipython, 10
- k-means, 205
- kernel, 116
- knowledge base, 78
- learner, 118
- learning, 107–140, 205–213, 223–247
 - batched stochastic gradient descent, 133
 - cross validation, 125
 - decision tree, 121
 - deep learning, 133
 - EM, 209
 - k-means, 205
 - linear regression, 127
 - linear classification, 127
 - neural network, 133
 - no inputs, 118
 - reinforcement, 223–240
 - relational, 241
 - supervised, 107–140
 - with uncertainty, 205–213
- likelihood weighting, 163
- linear regression, 127
- linear classification, 127
- logistic regression, 144
- magic square, 216
- magic-sum game, 216
- Markov Chain Monte Carlo, 166
- Markov decision process, 195
- max_display_level, 17
- MCMC, 166
- MDP, 195, 224
- method
 - consistent*, 53
 - holds*, 52
 - maxh*, 93
 - zero*, 91
- minimax, 215
- minimax algorithm, 219
- minsets, 83
- model-based reinforcement learner, 232
- multiagent system, 215
- multiple path pruning, 45
- n-queens problem, 58
- naughts and crosses, 216
- neural network, 133
- noisy-or, 144
- NotImplementedError, 21
- partial-order planner, 101
- particle filtering, 164
 - HMMs, 173
- planning, 85–106, 183–203
 - CSP, 98
 - decision network, 183
 - forward, 90
 - MDP, 195
 - partial order, 101
 - regression, 94
 - with certainty, 85–106
 - with learning, 232
 - with uncertainty, 183–203
- plotting
 - agents in time, 24
 - reinforcement learning, 227
 - robot environment, 30
 - runtime distribution, 74
 - stochastic simulation, 168
- predictor, 109

- Prob, 146
- probability, 141
- proof
 - bottom-up, 79
 - top-down, 81
- proposition, 77
- Python, 9
- Q learning, 229
- RC, 153, 190
- recursive conditioning (RC), 153
- recursive conditioning for decision networks, 190
- regression planning, 94
- reinforcement learning, 223–240
 - environment, 223
 - feature-based, 234
 - model-based, 232
 - Q-learning, 229
- rejection sampling, 162
- relational learning, 241
- resampling, 165
- robot
 - body, 26
 - environment, 25
 - middle layer, 28
 - plotting, 30
 - top layer, 29
- robot delivery domain, 86
- runtime, 15
- runtime distribution, 74
- sampling, 160
 - importance sampling, 164
 - belief networks, 162
 - likelihood weighting, 163
 - particle filtering, 164
 - rejection, 162
- scope, 51
- search, 33
 - A*, 41
 - branch-and-bound, 46
 - multiple path pruning, 45
- search_with_any_conflict, 69
- search_with_var_pq, 70
- show, 147
- sigmoid, 129
- stochastic local search, 67
 - any-conflict, 69
 - two-stage choice, 70
- stochastic simulation, 160
- tabular factor, 145
- test
 - SLS, 74
- tic-tac-toe, 216
- top-down proof, 81
- uncertainty, 141
- unit test, 19, 44, 58, 80, 81
- unrolling
 - DBN, 179
- updatable priority queue, 72
- utility, 183
- utility table, 183
- value iteration, 198
- variable, 51, 141
- variable elimination (VE), 156
- variable elimination for decision networks, 193
- VE, 156
- visualize, 17
- yield, 14